



Flaris Language Guide

The Flaris programming language · flaris-lang.org · github.com/flaris-lang/flaris · flaris-lang.org/discord

Version 1.0.0.9 **Generated** 2026-07-27 **License** see repository

© 2026 SolidInvention AB — created and maintained by SolidInvention AB

Table of Contents

1. [Introduction](#)
 2. [Getting Started - CLI](#)
 3. [Lexical Structure](#)
 4. [Variables & Types](#)
 5. [Expressions](#)
 6. [Control Flow](#)
 7. [Functions](#)
 8. [Classes](#)
 9. [Fibers & Async](#)
 10. [Analyzer & Diagnostics](#)
 11. [Debugging](#)
 12. [10-Minute Tours](#)
 13. [Package Manager \(flarispmp\)](#)
-

1. Introduction

Flaris is a lightweight, embedded scripting language designed for high performance, portability, and deep integration with host applications. It combines the expressiveness of modern scripting languages with the control of a virtual machine built around **fibers**, **async/await**, and **cooperative scheduling**.

Flaris is:

- **Deterministic** - execution order is always clear and reproducible.
- **Fast** - written in C with a compact bytecode interpreter.
- **Embeddable** - ideal for applications, tools, servers, and games.
- **Safe by default** - unsafe features (FFI, raw pointers/memory) require explicit enabling (`--unsafe`). Optional static analysis helps catch mistakes early. Runtime-version only without compiler and eval/compile functionality provided (`flaristrt`)
- **Concurrent** - fibers and async functions allow concurrency without threads.

Three core principles:

1. **Explicit over implicit** - nothing runs unless you call or resume it.
2. **Minimal magic** - the runtime is predictable, transparent, and debuggable.
3. **Simplicity without compromise** - features are easy to use but designed to scale.

Typical use cases: game scripting, tool automation, embedded device logic, interactive consoles, server-side utilities, scripting in C/C++ software, simulation and testing frameworks.

Supported Targets

Flaris targets 64-bit platforms: x86-64 and ARM64 on Linux, macOS, and Windows. Compiled `.flx` bytecode is portable between all of them - you can compile on Apple-silicon macOS and deploy on an x86-64 Linux server without recompiling.

The one exception is code that works with raw memory addresses (`FFI`, `Memory.*`, `Buffer.GetAddress`). Addresses are process-local and should never be serialized or shared across machines. Use `Os.Arch()` and `Os.Name()` to guard any platform-specific paths.

2. Getting Started - CLI

The `flarism` executable is how you run Flaris programs.

Invocation Modes

Compile and run (most common):

```
flarism script.flx
```

Compile to bytecode (for distribution):

```
flarism --compile script.flx output.flx
```

For performance:

```
flarism --compile script.flx output.flx --strip --small
```

Run precompiled bytecode:

```
flarism --exec script.flx
```

Read source from stdin - use `-` as the filename:

```
echo 'fn Main() { Console.WriteLine("hi"); }' | flarism -  
cat script.flx | flarism - --verbose
```

Compile from stdin to bytecode:

```
cat script.flx | flarism --compile - output.flx
```

Execute bytecode from stdin:

```
cat output.flx | flarism --exec -
```

The stdin sentinel `-` works wherever a source or bytecode filename is accepted. This enables pipelines where the program never touches the filesystem:

```
generate_code | flarism - # generate and run on the fly
```

Disassemble bytecode (debug):

```
flarism --disasm script.flx
```

Check JIT eligibility - compile without running and print, per function, whether it runs as native code under `-jit` and (if not) why:

```
flarism --check script.flx
# ✓ dot (line 20)
# · label (line 34) - operator on a non-numeric operand (e.g. string concatenation)
```

With `--jit`, every eligible function is JIT-compiled automatically (no annotation needed) - `--check` shows you which ones qualify. See Reference R11.

Run inline code (quick test):

```
flarism --eval 'Console.WriteLine("Hello");'
```

Get SHA-256 fingerprint (for imports):

```
flarism --sha256 script.flx
```

Self-contained binaries

Bundle a program and the runtime into a single native executable - one file to distribute, with nothing to install on the target machine (similar to Go). `--embed` compiles your program (or takes precompiled `.flx`) and appends it to the small compiler-free `flaris` runtime; the result runs your program directly when launched.

From source, in one step:

```
flarism --embed app.flx myapp
./myapp arg1 arg2 # runs standalone - no flaris on PATH, no --libs
```

Statically link library dependencies with `--bundle`, so the shipped binary needs no `~/flaris/libs`:

```
flarism --embed app.flx myapp --bundle='./mylib.flx'
```

Embed precompiled bytecode directly (detected by content, not extension):

```
flarism --embed app.flx myapp
```

`--embed <input> <output>` accepts the same `--bundle`, `--sign`, `--no-opt`, and `--strip` options as `--compile`. Runtime behavior is also decided here: pass `--jit`, `--unsafe`, or VM limits (`--stack=`, `--frames=`, ...) alongside `--embed` and they are baked into the binary and applied on every start - a standalone binary hands all its CLI arguments to the script, so embedding is the only place to set them (nothing is enabled by default). It needs the `flaris` runtime stub next to `flarism` or on `PATH` (the official installer places both). The stub has no compiler and no `eval`, so a bundled binary ships only the bytecode you built. Runs on Linux, macOS (including Apple Silicon), and Windows - the payload rides in trailing space every executable format tolerates; see **Reference R1** for platform and signing notes.

Package signing (Ed25519)

`.flx` bytecode can carry an embedded Ed25519 signature. The official standard library is signed with the `flaris-lang.org` key, which is compiled into the runtime as a built-in trust anchor - so signed code verifies with nothing to set up. A tampered file always fails to load; `--require-signed` additionally refuses anything unsigned or signed by an untrusted key. (Full details in `reference.md`.)

Sign when compiling (the secret key inline or as a keyfile path):

```
flarism --compile mylib.flx mylib.flx --sign=publisher.sk
```

Inspect a file's signature (machine-readable - used by `flarism`):

```
flarism --sig-info mylib.flx
# unsigned
# signed|<trusted|valid|invalid>|<pubkey-hex>|<signer>
```

Trust a publisher (append a key to `~/.flaris/trusted_keys`):

```
flarism --import-key "acme corp" <ed25519-pubkey-hex>
```

Enforce signatures at load:

```
flarism --require-signed --exec app.flx
```

Common Flags

Flag	Effect
<code>--verbose</code>	Verbose output (timing, info)
<code>--trace</code>	Very verbose (opcode-level trace)
<code>--time</code>	Timing report
<code>--jit</code>	Enable JIT: auto-compiles every eligible function to native code (use <code>--check</code> to see which)
<code>--mem</code>	Memory report at shutdown (peak RSS, slab footprint, leak detection)
<code>--stats</code>	Statistics (allocations, counters)
<code>--unsafe</code>	Enable FFI / unsafe code
<code>--small</code>	Enable strip of symbols

Flag	Effect
<code>--strip</code>	Strip debug symbols (production)
<code>--no-opt</code>	Disable optimizations
<code>--no-verify</code>	Disable FFI library SHA-256 checks (not <code>library()</code> pins)
<code>--require-signed</code>	Refuse to load any <code>.flx</code> that is unsigned, tampered, or signed by an untrusted key

Development workflow:

Show idle time (VM is waiting for fibers or other) and total running time.

```
flarismv app.flx --verbose --time
```

Production:

```
flarismv --exec app.flx
```

See **Reference R1** for the complete CLI specification.

3. Lexical Structure

Flaris source files are UTF-8 encoded. Identifiers use ASCII letters, digits, and underscores: `[A-Za-z_][A-Za-z0-9_]*`.

Keywords

| Keyword | Purpose | |-----|-----| | `fn` | Declare function | | `async` | Mark function as asynchronous | | `static` | Declare static method in class | | `inline` | Suggest function inlining (hint) | | `let` / `var` | Declare mutable variable (identical) | | `global` | Declare module-level global variable | | `const` | Declare compile-time constant | | `if` / `else` | Conditional | | `for` | C-style loop | | `foreach` | Iterate over collection elements | | `iter` / `from` / `to` | Range-based integer loop | | `while` | While loop | | `switch` / `case` / `default` | Multi-branch conditional | | `break` / `continue` | Loop control | | `return` | Return from function | | `yield` | Yield control from fiber | | `await` | Wait for async result | | `nil` / `true` / `false` | Literals | | `class` / `this` / `super` / `new` | Object-oriented | | `import` / `export` / `as` | Modules | | `try` / `catch` / `finally` / `throw` | Exceptions | | `guard` | Assert non-nil (throws if nil) | | `enum` | Define enumeration | | `breakpoint` | Debug pause | | `in` | Membership test / foreach iterator |

Literals

Nil and Booleans:

```
nil
true
false
```

Integers (64-bit signed):

```
42
0xFF
0b1001
1_000_000    // underscores allowed for readability
```

Floats (IEEE-754 double):

```
3.14
1.0e6
-2.5E-3
3.14_15_92
```

Strings:

```
"hello\nworld 🐱"    // single-line with escape sequences
"""
multi-line string
preserves formatting
"""
```

Escape sequences: `\n`, `\r`, `\t`, `\"`, `\\`, `\xNN`, `\uNNNN`

Character literals (single Unicode codepoint):

```
'a'
'ä'
'🐱'
'\n'
'\u00E5'
```

Comments

```
// single-line comment

/* multi-line
   comment */

/* nested /* block */ comments are supported */
```

Operators

```
// Arithmetic
+ - * / % ^^    // ^^ is power
```

```
// Bitwise
& | ^ ~ << >>

// Logical
&& || !
and or // aliases for && ||

// Comparison
== != < <= > >= is

// Assignment + compound
= += -= *= /= %= ^=
&= |= ^= <<= >>= ??=

// Special
?? // null coalescing
=> // arrow function
```

Punctuation

Semicolons are **required** at the end of every statement. Newlines do not end statements.

```
let x = 10; // required semicolon
let y = 20;; // extra semicolons are tolerated
```

4. Variables & Types

Declaring Variables

Use `let` or `var` for **local variables** inside functions and blocks (they are identical):

```
let name = "Alice";
var count = 0;
```

Rules:

- Must always have an initializer - `let x;` is illegal, use `let x = nil;`
- Block-scoped - variables are only visible inside their `{ }` block
- Can be reassigned freely
- `let` and `var` are **not valid at module (top) level** - use `global` there

Use `const` for immutable bindings:

```
const PI = 3.1415926535;
const MAX = 128;
```

`const` prevents rebinding, but not mutation of the object itself:

```
const cfg = { port: 8080 };
cfg.port = 9090;    // ✓ allowed (mutating the object)
cfg = {};           // ✗ illegal (reassigning the binding)
```

Globals

Use `global` to declare a module-level global, from the top-level code, or from inside a function or nested scope:

```
global shared_int: int = 2;

fn setup() {
  global config = { debug: true }; // visible everywhere in the module
}
```

- `global` always creates (or warns if already exists) a top-level variable
- Access and mutation of globals is slower than locals - prefer locals in hot paths

Types at a Glance

Flaris is dynamically typed. Every value has a type at runtime.

Primitives: `nil`, `bool`, `int`, `float`, `char`, `string`

Containers: `array`, `object`

Structured: `class`, `instance`, `exception`

Execution: `fiber`, `fn`, `module`

Binary: `block`, `pointer`

Type annotations on function signatures are **optional** and produce warnings, not errors. They don't change runtime behavior. See **Reference R3** for the full type annotation specification.

Quick example:

```
fn add(a: int, b: int): int {
  return a + b;
}
```

Scoping

Variables are block-scoped. Every `{ }` creates a new scope:

```
let x = 1;
{
```

```
let y = 2;
  Console.WriteLine(x, y); // 1 2
}
// y is not accessible here
```

Inner scopes can **shadow** outer variables:

```
let x = 1;
{
  let x = 2;           // shadows outer x
  Console.WriteLine(x); // 2
}
Console.WriteLine(x);  // 1
```

Globals

Top-level declarations are module globals - visible anywhere in the file and shared across all fibers:

```
global counter = 0;

fn tick() {
  counter += 1;
}
```

Use `global` to explicitly promote a variable to module scope from inside a function:

```
fn init() {
  global appName = "MyApp";
}
```

Closures

Flaris supports closures: inner functions capture variables from their enclosing scope by value at the point the inner function is created. Both parameters and body-local variables of the outer function are captured.

Single capture - parameter:

```
fn makeAdder(x) {
  return fn(n) { return x + n; };
}

let add5 = makeAdder(5);
let add10 = makeAdder(10);
Console.WriteLine(add5(3)); // 8
Console.WriteLine(add10(3)); // 13
```

Each call to `makeAdder` produces an independent closure with its own snapshot of `x`.

Multi-capture - parameter + body local:

```
fn makeGreeter(prefix) {  
    let sep = ": ";  
    return fn(name) { return prefix + sep + name; };  
}  
  
let hello = makeGreeter("Hello");  
Console.WriteLine(hello("World")); // Hello: World  
Console.WriteLine(hello("Flaris")); // Hello: Flaris
```

Capture semantics - by value at bind time:

Integers, booleans, strings, and other scalar values are snapshotted when the inner function is created. Mutating the outer variable after the inner function is created does not affect the captured copy.

Reference types (arrays, objects) share the same heap object - mutations are visible through the closure:

```
fn makeCounter() {  
    let state = [0];  
    return fn() {  
        state[0] = state[0] + 1;  
        return state[0];  
    };  
}  
  
let c = makeCounter();  
Console.WriteLine(c()); // 1  
Console.WriteLine(c()); // 2  
Console.WriteLine(c()); // 3
```

Loop snapshot capture:

When a closure is created inside a loop, the current value of the loop variable is captured at that iteration. Each closure gets its own independent snapshot:

```
fn makeFns() {  
    let funcs = [];  
    iter(i from 0 to 5) {  
        Array.Append(funcs, fn() { return i; });  
    }  
    return funcs;  
}  
  
let fs = makeFns();  
Console.WriteLine(fs[0]()); // 0  
Console.WriteLine(fs[3]()); // 3
```

Globals are still accessible:

Module-level globals are always accessible from inner functions, as before:

```
global multiplier = 3;

fn makeTransform() {
  return fn(x) {
    return x * multiplier;
  };
}

let transform = makeTransform();
Console.WriteLine(transform(5)); // 15
multiplier = 2;
Console.WriteLine(transform(5)); // 10 (reads live global)
```

Variables and Fibers

Each fiber has its own call stack and local variables. Globals are shared:

```
global shared = 0;

fn worker(id) {
  let local = id * 10; // local to this fiber's stack
  shared += 1;         // shared across all fibers
  yield local;
}
```

Common Mistakes

- `let x;` is illegal - always provide an initializer - forces developers to set a reasonable initial value
- `let` and `var` at the top level are illegal - use `global name = value;` at module scope
- `const` prevents rebinding the name, not mutating the value
- Declaring `let x` inside a function does not modify a global `x` - use `global x = value;` to declare a new global from within a function
- Reusing the same name in nested scopes shadows the outer variable
- `global x = value;` on an already-declared global produces a warning and redefines it - not a silent no-op

5. Expressions

An **expression** produces a value. Expressions become **statements** when terminated with `;`.

Important: **Assignments do not produce a value at runtime.**

In the grammar, `a = b` is parsed as an assignment expression node, but the compiler emits assignment bytecode that **does not leave a value on the stack**. Treat assignments as **statement-only**: you cannot write `let x = (y = 3)` or use `x = y` as a condition.

Literals

```
nil           // nil
true          // bool
42            // int
3.14          // float
"hello"       // string
'A'           // char - single Unicode codepoint, single quotes
[1, 2, 3]     // array
{ a: 1, b: 2 } // object (keys are strings)
```

Char literals use single quotes. The `Char` module provides classification and conversion - see **Reference R8 - Char Module**.

Instance method syntax on strings, chars, and arrays

`String`, `Char`, and `Array` values support calling their built-in module methods directly on the value using dot notation. This is purely syntactic sugar - both forms are equivalent and can be mixed freely. When the variable is typed, the compiler resolves the call at compile time and emits identical bytecode to the module-style form.

```
// Char instance methods
let c = String.CharAt("hello", 0);
c.IsAlpha()      // true   - same as Char.IsAlpha(c)
c.ToUpper()      // 'H'    - same as Char.ToUpper(c)
c.Code()         // 104    - same as Char.Code(c)

// String instance methods
let s = "Hello, World!";
s.Length()       // 13     - same as String.Length(s)
s.ToLower()      // "hello, world!"
s.Contains("World") // true
s.Substr(7, 5)    // "World"
s.Replace("World", "Flaris") // "Hello, Flaris!"
s.Split(",")      // ["Hello", " World!"]
" hi ".Trim()     // "hi"

// Array instance methods (functions that take the array as first arg)
var a: array = [];
a = a.Append(1);
a = a.Append(2); // same as Array.Append(a, 2)
a.First()        // 1
a.Last()         // 2
a.Contains(1)     // true
```

```
a.Sum() // 3.0
a = a.Reverse()
let evens = a.Where(fn(x) { return x % 2 == 0; });
let doubled = a.Select(fn(x) { return x * 2; });
```

Note: `Array.Create` and similar factory functions that do not take an array as first argument are not available as instance methods (`a.Create(5)` will not dispatch to `Array.Create`).

Operators

Arithmetic: `+` `-` `*` `/` `%` `^^`

- `/` yields float if any term is float, else int
- `^^` is exponentiation: `2 ^^ 3` - 8

Bitwise: `&` `|` `^` `~` `<<` `>>` (and, or, xor, not, shl, shr)

Comparison: `==` `!=` `<` `<=` `>` `>=` `is`

Logical: `&&` `||` `!` - aliases: `and` `or`

- Short-circuiting: `x && y` only evaluates `y` if `x` is truthy
- Falsy: `nil`, `false`, `0`, `0.0`, `""`, `'\0'`, empty arrays/objects; everything else is truthy
- `x && y` yields the value of `y` when `x` is truthy, otherwise `false`; `x || y` yields `true` when `x` is truthy, otherwise the value of `y`
- `and` / `or` are exact tokenizer aliases - identical precedence and behaviour to `&&` / `||`

Increment / Decrement: `++` `--` (postfix)

`++` and `--` only work on plain variables (locals and globals). They mutate in place and **produce no value** - they cannot be used inside expressions.

```
let i = 0;
let s = "";
i++; // ok - i is now 1
let x = i++; // ❌ compile error - ++ produces no value
obj.count++; // ❌ compile error - use obj.count += 1 instead
arr[0]++; // ❌ compile error - use arr[0] += 1 instead
s++; // ❌ compile error
```

For member properties and array elements, use compound assignment instead:

```
obj.count += 1;
arr[i] -= 1;
```

Null coalescing: `x ?? y` - returns `y` only if `x` is `nil`

```
let name = user.name ?? "Anonymous";
```

Ternary: `condition ? exprIfTrue : exprIfFalse`

```
let label = age >= 18 ? "adult" : "minor";
```

Member Access

```
obj.key          // reads a property
obj.key = val;   // sets a property (statement)
```

Null-safe: accessing `.key` on `nil` returns `nil` - it does not throw.

```
let user = nil;
user.name;    // nil (safe)
user.items;   // nil (safe)
```

Indexing

```
array[i]        // 0-based integer index
object["key"]   // string key (equivalent to object.key)
string[i]       // byte at index i
```

Null-safe: `nil[x]` returns `nil`.

guard

`guard(expr)` asserts that `expr` is not `nil`. If it is, throws `Exception.GuardCheck`:

```
let addr = guard(user.address); // throws if nil
```

Use `guard` when a `nil` value means something has gone wrong. For optional values, just use `??` or check directly.

Function Calls

```
fn add(a, b) { return a + b; }
let x = add(1, 2); // 3
```

await and yield

```
let result = await asyncFn(); // wait for async function result
yield value;                  // suspend fiber, produce value
```

Note:

- `await expr`; used as a **statement** discards the resulting value.

- `let x = await expr;` used in an **expression** keeps the value.

Operator Precedence (high to low)

Level	Operators	Associativity
1	Literals, variables, <code>()</code>	-
2	<code>()</code> call, <code>[]</code> index, <code>.</code> member	Left
3	<code>-</code> <code>!</code> <code>~</code> <code>await</code> <code>yield</code> (prefix)	Right
4	<code>*</code> <code>/</code> <code>%</code> <code>^^</code> <code>&</code> <code> </code> <code>^</code> <code><<</code> <code>>></code>	Left
5	<code>+</code> <code>-</code>	Left
6	<code><</code> <code><=</code> <code>></code> <code>>=</code>	Left
7	<code>==</code> <code>!=</code> <code>is</code> <code>in</code>	Left
8	<code>&&</code>	Left
9	<code> </code>	Left
10	<code>??</code>	Left
11	<code>?:</code> (ternary)	Right
12	<code>=</code> <code>+=</code> etc. (assignment)	Right

When in doubt, use parentheses.

6. Control Flow

Every statement ends with `;`. Newlines are not statement terminators.

If / Else

```
if (condition) {  
    ...  
} else if (other) {  
    ...  
} else {  
    ...  
}
```

Truthiness: `nil`, `false`, `0`, `0.0`, `""`, `'\0'` and empty arrays/objects are falsy; everything else is truthy.

While

```
let i = 0;  
while (i < 10) {  
    Console.WriteLine(i);  
    i += 1;  
}
```

Do / While

A `do { ... } while (cond);` loop tests its condition **after** the body, so the body always runs at least once. The trailing `;` is required.

```
let line = "";
do {
    line = Console.ReadLine();
    process(line);
} while (line != "quit");
```

For

```
for (let i = 0; i < 10; i += 1) {
    Console.WriteLine(i);
}
```

Any of the three clauses may be omitted. An empty condition means "always true", so `for (;;)` is a bare infinite loop (exit it with `break`, `return` or a `throw`):

```
for (;;) {
    let cmd = next();
    if (cmd == nil) { break; }
    run(cmd);
}

for (let i = 0; ; i += 1) {    // no condition - step and break inside the body
    if (i >= limit) { break; }
}
```

Foreach

Iterates arrays, objects, and strings. Null-safe - iterating `nil` runs zero times.

```
foreach (item in [1, 2, 3]) {
    Console.WriteLine(item);
}

// With key
foreach (value, key in obj) {
    Console.WriteLine(key, "=", value);
}

// String characters
foreach (ch in "hello") {
```

```
Console.WriteLine(ch);  
}
```

Iterable	Key	Value
Array	integer index	element value
Object	property name (string)	property value
String	character index	UTF-8 character
nil	-	no iterations

Iter

High-performance integer range loop - no allocation, optimized bytecode:

```
iter (i from 0 to 10) {  
    sum += i;  
}
```

- Only ascending (`from < to`)
- Bounds evaluated once before the loop

Use `iter` for tight numeric loops; use `foreach` for container iteration.

Break & Continue

```
for (let i = 0; i < 10; i += 1) {  
    if (i == 5) break;      // exit loop  
    if (i % 2 == 0) continue; // skip to next iteration  
    Console.WriteLine(i);  
}
```

Switch

```
switch (value) {  
    case 1:  
        doOne();  
        break;  
    case 2:  
    case 3:  
        doTwoOrThree();  
        break;  
    default:  
        doOther();  
        break;  
}
```

Cases fall through unless you `break`. `break` inside a switch does not affect outer loops.

Return

```
fn add(a, b) {  
    return a + b;  
}  
  
fn log(msg) {  
    Console.WriteLine(msg);  
    return;    // returns nil  
}
```

Exceptions

`Exception` is a built-in base **class**. Anything thrown is an `Exception` instance (or a subclass instance). Instances carry `Code` (int), `Error` (message string) and `StackTrace`, plus a `ToString()` method.

Throw:

```
throw new Exception(Exception.DivByZero, "Cannot divide by zero");  
  
// Shorthand – desugars to `new Exception(code, msg)`:  
throw(Exception.DivByZero, "Cannot divide by zero");
```

`throw` requires an `Exception` (or subclass) instance; throwing anything else is a type error.

Custom exception types - extend `Exception`:

```
class NetworkError : Exception {}  
class TimeoutError : NetworkError {}  
  
throw new NetworkError(503, "service unavailable");
```

Try / Catch / Finally:

```
try {  
    let v = 1 / 0;  
} catch (err) {  
    Console.WriteLine("Error:", err.Error);  
    Console.WriteLine("Code:", err.Code);  
} finally {  
    cleanup();    // always runs: normal exit, exception, return, break, continue  
}
```

There is a single `catch` clause that binds the caught instance. To handle specific exception types, dispatch on the value with a type `switch - case <Class>: is an instanceof test` (subclass-aware), so order most-specific-first and use a bare `Exception` (or `default`) as the catch-all:

```
try {
    doRequest();
} catch (e) {
    switch (e) {
        case TimeoutError: retry();          break;
        case NetworkError: reconnect();       break;
        case Exception:    log(e.ToString()); break; // any other exception
    }
}
```

Cases fall through like in C - end each case with `break` (or `return`) unless fallthrough is intended.

Re-raise from a catch with `throw e; .` The same `instanceof` test is available as the `is` operator: `if (e is NetworkError) { ... }.`

`finally` runs on every exit path - normal completion, a caught (or re-raised) exception, and `return` / `break` / `continue` that leave the `try` or `catch` (the leave runs every enclosing `finally` first, even across loop boundaries). If a `finally` itself throws, that exception supersedes any in flight.

Control may **not** leave a `finally` body via `return`, `break`, or `continue` (compile error) - a `finally` is cleanup that always runs to completion. A `throw` from a `finally` is allowed (it propagates to an upstream handler), and a `break` / `continue` targeting a loop opened *inside* the `finally` is fine (it does not leave the `finally`).

If no handler exists, the exception propagates up through fibers. An uncaught exception terminates the VM.

Note: `case <ImportedClass>:` across module boundaries is matched by value, not type, for now (cross-module type resolution is a later addition); the compiler warns when a switch label's type is unknown.

Import

Flaris imports require **compiled bytecode** (`.flx`) and a **version requirement**:

```
import { sqrt } from library("./mathlib", "1.0");
import { sqrt, pow } from library("./mathlib", "1.0");
import { Foo as Bar } from library("./classes", "1");
```

Library names are case-sensitive

The first argument resolves to a file: `library("Signals")` loads `Signals.flx` from the `--libs` path, `library("./util/Math")` loads `./util/Math.flx`, and a URL loads that exact path. Resolution is an **exact, case-sensitive filename match** - the name is never lowercased. Case-insensitive filesystems (macOS/Windows) forgive a mismatch, but on **Linux** `library("signals")` will not find `Signals.flx`. Keep

the import string, the on-disk filename, and any published/pinned name identical; the convention is a capitalized first letter (`Signals` , `Http` , `BigInt`).

Imported **symbol names are case-sensitive too** and must match the library's `export` exactly - `import { Jwt }` matches `export { Jwt }`, not `JWT` or `jwt`.

Where libraries are found

A `library("Name")` import is resolved by trying, in order:

1. the name as a path - `library("./mathlib")` → `./mathlib.flx`, relative to the working directory;
2. each `;`-separated entry of `--libs=<paths>`;
3. the **per-user install directory** - `~/flaris/libs` (POSIX) or `%USERPROFILE%\flaris\libs` (Windows), overridable with the `FLARIS_LIBS` environment variable.

The first candidate whose embedded version satisfies the requirement wins, so an explicit `--libs` entry always takes precedence over a globally-installed copy. Step 3 means a package manager can install a library once into the home directory and have every script resolve it with **no `--libs` flag**.

The second argument is a version string matched against the version embedded in the `.flx` file at compile time. Set the version during compilation with `--version=` - the default is `1.0.0.0`.

Version syntax

String	Meaning
<code>"1"</code>	Major 1, any minor/patch/build
<code>"1.2"</code>	Major 1, minor 2, any patch/build
<code>"1.2.3"</code>	Major 1, minor 2, patch 3, any build
<code>"1.2.3.4"</code>	Exact match on all four components
<code>"=1.2"</code>	Exactly 1.2.x.x (explicit exact)
<code>">=1.2"</code>	Any version $\geq 1.2.0.0$
<code>"<2.0"</code>	Any version $< 2.0.0.0$
<code>">=1.2 <2.0"</code>	Range: $1.2.0.0 \leq v < 2.0.0.0$

The **major version always requires an exact match** regardless of operator - a module compiled as `2.x` will never satisfy a `"1.x"` requirement. This enforces ABI stability.

Optional fingerprint verification

A third argument adds a tamper check: the **whole-file SHA-256** of the `.flx`, identical to `sha256sum file.flx`. Obtain it with `flarism --sha256 file.flx`:

```
import { Jwt } from library("./Jwt", "1.0", "a3f1...64hexchars...");
```

The hash is recomputed over the actual bytes loaded, so it detects any tampering. An optional `sha256:` prefix is accepted (the form `flarism` records in `flaris.json`). If it does not match the loaded file the VM halts immediately, regardless of the global `--no-verify` flag. This is the recommended pattern for security-sensitive dependencies.

Compiling a library with a version

```
# Compile with version
flarism --compile mylib.flx mylib.flx --min-version=1.2.0

# Print fingerprint
flarism --sha256 mylib.flx
```

Only symbols explicitly `export` ed from a module are importable.

Imported types

Type-aware features - the `is` operator, a type `switch` (`case <Class>:`), and typed `catch` dispatch - need the compiler to know an imported symbol's *type*. Flaris resolves that two ways:

1. Automatically from the dependency's `.flx`. Each `export` records the exported symbol's type in the compiled module (functions, classes, constants, ...). When you `import` from it, the compiler reads those types back, so an imported class is recognised as a class:

```
import { ApiError } from library("./errors", "1.0"); // ApiError is known to be a class

try { request(); }
catch (e) {
  switch (e) {
    case ApiError: handle(e); // instanceof dispatch – resolved across modules
    case Exception: log(e);
  }
}
```

This requires the dependency's `.flx` to be **compiled before** the module that imports it.

2. Explicitly, with a `:type` annotation. You can also annotate the import, which always takes precedence (useful when the dependency isn't compiled yet, or to be explicit):

```
import { ApiError:class, MAX_RETRIES:int, connect:fn } from library("./net", "1.0");
```

Recognised type names: `nil`, `bool`, `int`, `float`, `char`, `string`, `array`, `object`, `class`, `instance`, `fn`, `module`, `fiber`, `stream`, `block`, `exception`, `any`, `number`.

If a symbol's type is unknown (no annotation and the dependency wasn't available at compile time), it is treated as `any` - code still runs correctly at runtime (type tests fall back to dynamic checks), it's just not statically resolved.

Static linking (self-contained bundles)

Pass `--bundle='<file1.flx>;<file2.flx>;...'` when compiling to embed all imported `.flx` files directly into the output. The result is a single `.flx` that runs anywhere without the dependency files present:

```
flarism --compile app.flx app_bundle.flx --bundle='mathlib.flx' --libs='.'
flarism --exec app_bundle.flx          # mathlib.flx not needed on disk
```

Inspect the bundle:

```
flarism --disasm app_bundle.flx
```

Shows the TOC listing each bundled dep (name, version, offset, size) followed by a full disassembly of each chunk.

A bundled `.flx` still needs a VM to run it (`flarism/flaris`). To ship a single **native executable** with the runtime included, embed the same bundle with `--embed` (see *Self-contained binaries* under §2):

```
flarism --embed app.flx app --bundle='mathlib.flx'
./app          # standalone – no VM or deps needed
```

7. Functions

Functions are first-class values created with `fn`. Modifiers `static` or `async` is placed between `fn` and `name`.

```
fn <static> <async> <inline> name(<args:?type> ? :type)
```

Argument limit - functions accept up to 16. For larger configurations, pass an object:

```
// preferred for many params
fn draw(opts) {
  let x = opts.x;
  let color = opts.color;
  ...
}
```

Basic Declaration

```
fn greet(name) {
  return "Hello, " + name + "!";
}

greet("Alice"); // Hello, Alice!
```

Functions as Values

```
let f = greet;          // store in variable
f("Bob");               // call through variable
```

```
fn apply(func, val) {  
    return func(val);  
}  
  
apply(greet, "Carol"); // pass as argument
```

Type Annotations (optional)

Type annotations on parameters and return values produce compile-time warnings, not errors:

```
fn add(a: int, b: int): int {  
    return a + b;  
}  
  
add(5, 10);           // ✓ OK  
add("x", "y");       // ⚠ Warning at compile time, still runs
```

Use `any` to explicitly accept all types for a parameter - this disables type checking for that parameter while still annotating the others:

```
fn log(label: str, value: any) {  
    print(label, value);  
}  
  
log("result", 42);           // ✓ OK  
log("result", [1,2,3]);     // ✓ OK – any type accepted
```

Type annotations are **completely optional** - all existing code works without them. See **Reference R3** for the full type keyword list and union type syntax.

Parameters

By default all parameters are required. Append `?` to a parameter name to mark a parameter as optional - the VM will pass `nil` for any omitted trailing arguments:

```
fn add(a, b) { return a + b; }  
  
add(1, 2);    // OK  
add(1);       // Error: too few arguments  
  
fn greet(name?: string) {  
    Console.WriteLine("Hello " + (name ?? "stranger"));  
}  
  
greet("World"); // Hello World  
greet();        // Hello stranger
```

Optional parameters must come after all required ones. They can be combined with union types:

```
fn log(msg: string, level?: string|nil) {  
    // level is nil when omitted  
}
```

For variable numbers of arguments, use an array:

```
fn sum(numbers) {  
    let total = 0;  
    foreach (n in numbers) { total += n; }  
    return total;  
}  
  
sum([1, 2, 3, 4, 5]);
```

Return Values

Functions return `nil` if no `return` is given:

```
fn classify(x) {  
    if (x < 0) return "negative";  
    if (x > 0) return "positive";  
    return "zero";  
}
```

Async Functions

Async functions run as fibers and return immediately:

```
import { Http } from library("Http", "1.0");  
  
fn async fetchData(url) {  
    let response = await Http.Get(url);  
    return response.body;  
}  
  
let fiber = fetchData("http://example.com"); // returns fiber  
let data = await fiber;                      // wait for result
```

Error handling in async:

```
fn async safeLoad(url): object|nil {  
    try {  
        let resp = await Http.Get(url);  
        return resp.body;  
    } catch (e) {
```

```
        Console.WriteLine("Error:", e.Error);
        return nil;
    }
}
```

Anonymous Functions

```
let square = fn(x) {
    return x * x;
};

Array.Select([1, 2, 3], fn(x) { return x * 2; });
```

Note: type annotations **are supported** on anonymous functions as well.

Arrow Functions

Single-expression shorthand - the expression is automatically returned:

```
let double = fn(x) => x * 2;
let add = fn(a, b) => a + b;

fn square(x) => x * x;    // also works as named function
```

Recursion

```
fn factorial(n) {
    if (n <= 1) return 1;
    return n * factorial(n - 1);
}
```

Function Modifiers

Modifiers come after `fn`:

```
fn async load() { ... }    // async fiber
fn static helper() { ... } // class-level (no this)
fn inline calc(x) => x * 2; // hint to inline

fn static async fetch(url) { ... } // combine modifiers
```

`inline` is a hint: the compiler expands the body at each call site when it can do so without changing behavior, and otherwise emits a normal call (e.g. when a name in the body would collide with a caller local, on deep inline recursion, or on arity mismatch). Restrictions: the body may contain at most one `return`, and only

as its last statement - a `return` inside a nested block (`if`, `loop`, `try`) is a compile error. Arguments are always evaluated exactly once, left to right. Note that `VM.PatchFunction` does not affect call sites that were inlined.

Common Patterns

Factory function:

```
fn createUser(name, age) {  
  return {  
    name: name,  
    age: age,  
    greet: fn() { Console.WriteLine("Hi, I'm " + name); }  
  };  
}  
  
let user = createUser("Bob", 25);  
user.greet();
```

Higher-order functions:

```
let nums = [1, 2, 3, 4, 5];  
let doubled = Array.Select(nums, fn(x) => x * 2);  
let evens = Array.Where(nums, fn(x) => x % 2 == 0);  
let total = Array.Reduce(nums, fn(acc, x) => acc + x, 0);
```

Error handling with union return:

```
fn divide(a, b): number|nil {  
  if (b == 0) return nil;  
  return a / b;  
}  
  
let result = divide(10, 2);  
if (result == nil) {  
  Console.WriteLine("Cannot divide by zero");  
}
```

8. Classes

Classes provide lightweight object-oriented structure. Flaris classes support instance methods, static methods, single inheritance, async methods, and reflection.

Basic Class

```
class Point {  
    let x = 0;    // instance fields MUST be declared in the class body  
    let y = 0;  
  
    fn Constructor(x, y) {  
        this.x = x;  
        this.y = y;  
    }  
  
    fn move(dx, dy) {  
        this.x += dx;  
        this.y += dy;  
    }  
  
    fn ToString() {  
        return "Point(" + str(this.x) + ", " + str(this.y) + ")";  
    }  
  
    fn static origin() {  
        return new Point(0, 0);  
    }  
}  
  
let p = new Point(10, 20);  
p.move(5, -3);  
Console.WriteLine(p.ToString());           // Point(15, 17)  
Console.WriteLine(Point.origin().ToString()); // Point(0, 0)
```

Inside the class body you may declare:

- `fn Constructor(...)` - optional constructor
- `fn method(...)` - instance methods
- `fn static method(...)` - static methods
- `fn async method(...)` - async instance methods
- `let field = <const expr>` - **instance field default** (initializer must be a compile-time constant)
- `const NAME = <const expr>` - **static class constant** (stored on the class, not the instance)

Class body initializers

- Instance field defaults declared with `let` in the class body must be **compile-time constant expressions**.
- `const` declarations inside a class define **static members** (class-level constants).

Sealed instance fields

A class instance has a **fixed set of fields**: every instance field must be declared in the class body with `let` or `var` (own fields) or inherited from a base class. You **cannot** introduce a new field by assigning to it - not from the constructor, not from a method, and not from outside the instance:

```
class Point {
  let x = 0;
  let y = 0;

  fn Constructor(x, y) {
    this.x = x;      // OK - declared field
    this.y = y;      // OK - declared field
    this.z = 0;      // ERROR - 'z' is not a declared field
  }
}

let p = new Point(1, 2);
p.color = "red";    // ERROR - 'color' is not a declared field
```

This is enforced at three levels:

- **Compiler** and **analyzer** reject `this.X = ...` writes to an undeclared field at compile time (error 1014). If the class extends a base that lives in an imported module (whose fields are not visible during this compile), the static check is relaxed for that class and the runtime guard below applies.
- **VM** raises an exception at runtime if any write (`this.X = ...`, `instance.X = ...`, or a compound assignment) would create a field that is not declared on the instance's class or one of its bases.

To hold optional/late-bound data, declare the field up front with a `nil` default (`let handle = nil;`) and assign it later. If you genuinely need an open-ended bag of dynamic keys, use a plain **object literal** (`{}`) instead of a class instance - object literals remain fully dynamic.

Creating Instances

```
let p = new Point(10, 20);
```

1. Allocate a new object with class `Point`
2. Call `Constructor(10, 20)` if defined
3. Return the constructed instance

If `Constructor` is missing, the instance is created without initialization.

The Constructor

```
class User {
  let name = nil;
  let age = 0;
```

```
fn Constructor(name, age) {  
    this.name = name;  
    this.age = age;  
}  
}
```

- Instance fields **must be declared** in the class body with `let / var`. The constructor assigns them; it cannot introduce new fields. See **Sealed instance fields** below.
- Any explicit `return` value is ignored - `new` always returns the instance
- If `Constructor` throws, construction fails

The Destructor

`Destructor` is called automatically when the last reference to an instance is released (ref-count drops to zero). Use it to release external resources such as file handles, network connections, or native memory blocks.

```
class FileHandle {  
    let file = nil;  
    fn Constructor(path) {  
        this.file = File.Open(path, "w");  
    }  
  
    fn Write(text) {  
        File.Write(this.file, text);  
    }  
  
    fn Destructor() {  
        File.Close(this.file);  
    }  
}  
  
fn Main() {  
    let f = new FileHandle("out.txt");  
    f.Write("hello");  
    // f goes out of scope here - Destructor is called automatically  
}
```

Rules and caveats:

- `Destructor` receives no arguments and its return value is ignored
- `this` refers to the instance being freed, exactly as in any other method
- If `Destructor` throws, the exception is swallowed and a warning is printed - it does **not** propagate
- `Destructor` is **not** inherited; each class must declare its own if needed
- Avoid storing `this` anywhere inside `Destructor` - resurrection (re-referencing the dying object) is not supported and results in a dangling pointer

Reference cycles

Flaris frees objects by reference counting, and nothing traces the heap afterwards. That makes destruction predictable, but it means **a cycle is never freed**: if two objects point at each other, neither count ever reaches zero.

The common shape is a child that keeps a reference back to its parent:

```
class Parser {
    let _sub = nil;                // parent holds the child
    fn AddSub() {
        this._sub = new SubParser(this);
        return this._sub;
    }
    fn ProgName() { return "tool"; }
}

class SubParser {
    let _parent = nil;            // ...and the child holds the parent
    fn Constructor(parent) {
        this._parent = parent;    // cycle - neither is ever freed
    }
    fn Name() { return this._parent.ProgName() + " sub"; }
}
```

Nothing here is reported at compile time and the program behaves correctly - it just never releases those objects.

Usually the back-reference exists only to read a value or two, so the fix is to copy what you need and drop the reference:

```
class SubParser {
    let _parentProg = "";
    fn Constructor(parent) {
        this._parentProg = parent.ProgName(); // copy the value, not the object
    }
    fn Name() { return this._parentProg + " sub"; }
}
```

When the child genuinely needs the live parent, pass it as an argument so the reference lasts only for the call, or clear the field in a teardown method before the structure is dropped. The same rule covers any loop - two objects referring to each other, or a closure that captures an object being stored back on it:

```
let self = this;
Fiber.Run(fn() { self.Loop(); }); // fine - Fiber.Run returns an int id
this._fiber = Fiber.New(fn() { self.Loop(); }); // cycle - the field holds a fiber
// whose closure captures `self`
```

`Fiber.New` returns a fiber **object**, so keeping it in a field of the same object the closure captures closes the loop. `Fiber.Run` hands back an integer id instead, which holds nothing and is always safe to store.

Check with `--mem`, which prints a leak dump at shutdown and exits non-zero when anything is left over:

```
flarism --mem myprogram.flr
```

A non-zero leak count that stays the same across runs is almost always a cycle. See Reference R7 "Memory Management" for the underlying model.

Inheritance

Single inheritance with `::`:

```
class Animal {
  let name = nil;
  fn Constructor(name) {
    this.name = name;
  }
  fn speak() {
    return this.name + " makes a sound";
  }
}

class Dog : Animal {
  fn speak() {
    return this.name + " barks"; // 'name' inherited from Animal
  }
}

let d = new Dog("Rex");
d.speak(); // "Rex barks"
```

Method lookup walks: instance - derived class - base class - ...

Calling the base constructor:

```
class ColorPoint : Point {
  let color = nil;
  fn Constructor(x, y, color) {
    super(x, y); // calls Point.Constructor(x, y)
    this.color = color;
  }
}
```

Calling an overridden base method:

`super.method(args)` calls the base-class implementation from inside an override. The lookup is static: it resolves on the base of the class that *defines* the current method (not the receiver's dynamic class), so a multi-level chain of overrides each reaches its own base without looping:

```
class Animal {
  fn describe() { return "animal"; }
}

class Dog : Animal {
  fn describe() { return super.describe() + " that barks"; }
}

let d = new Dog();
d.describe(); // "animal that barks"
```

Calling a method that does not exist on the base chain raises a runtime error (a super call never falls back to the receiver's own class - that would re-enter the override). `super.method` is only available inside instance methods of a class with a base class.

Static Methods

Static methods belong to the class, not instances:

```
class MathUtils {
  fn static abs(n) {
    if (n < 0) return -n;
    return n;
  }
}

MathUtils.abs(-42); // 42
```

Static methods are inherited and can be overridden.

Async Methods

```
import { Http } from library("Http", "1.0");

class Loader {
  let data = nil;
  fn async load(url) {
    let data = await Http.Get(url);
    this.data = data;
    return true;
  }
}
```

```
let l = new Loader();  
let ok = await l.load("http://example.com");
```

`this` is preserved across `await` because it lives in the fiber's call frame.

Bound Methods

When you read a method from an instance, you get a bound method that remembers its receiver:

```
let m = p.move;    // bound to p  
m(5, 10);          // equivalent to p.move(5, 10)
```

This makes callbacks natural - you can pass `obj.method` directly.

Classes as Values

Classes are objects and can be stored and passed around:

```
let C = MyClass;  
let obj = new C();  
C.Version = "1.0";    // add static fields dynamically
```

Because the class in a `new` expression is just a value, it may be reached through a dotted path - a class held in an object field or a namespace-like container:

```
let registry = { Widget: MyClass };  
let w = new registry.Widget(1, 2);    // new on a member path
```

Reflection

```
Class.IsClass(val)           // true if val is a class  
Class.IsInstance(val)        // true if val is an instance  
Class.Name(MyClass)          // "MyClass"  
Class.InstanceOf(obj, MyClass) // inheritance check  
Class.SubclassOf(Child, Parent) // class hierarchy check  
Class.GetBase(MyClass)        // parent class  
Class.InstanceMethods(MyClass) // ["Constructor", "move", ...]  
Class.StaticMethods(MyClass)  // ["origin", "create", ...]  
Class.Invoke(target, "method", args...) // dynamic method call  
Class.Fields(instance)        // ["x", "y", "color"]
```

Object helpers:

```
Object.Clone(value)           // deep clone  
Object.HasKey(obj, key)       // own-property check  
Object.Count(obj)             // number of own fields  
Object.Entries(obj)           // [[key, value], ...]
```

```
Object.TryGet(obj, key, default) // safe lookup
Object.Clear(obj)                // remove all fields
```

Common Mistakes

- Forgetting `new` - `MyClass()` is not the same as `new MyClass()`
- Expecting multiple inheritance - only single inheritance is supported
- `return` in Constructor is ignored - `new` always returns the instance
- `Class.InstanceOf(instance, class)` - argument order matters
- `Class.SubclassOf(child_class, parent_class)` - argument order matters
- Storing `this` inside `Destructor` - resurrection is not supported; the object is freed immediately after `Destructor` returns
- Expecting `Destructor` to be inherited - each class must declare its own

9. Fibers & Async

Fibers are lightweight cooperatively-scheduled execution contexts. They allow concurrency, coroutines, and `async/await` without OS threads.

Key properties:

- No preemption - a fiber runs until it yields, awaits, returns, or sleeps
- No data races by default
- Cooperative scheduling

Concurrency model

Flaris supports two complementary models that run on the same scheduler:

Fibers - manual, controllable coroutines:

- `yield` produces values
- `Fiber.Resume` advances a fiber one step at a time

Async/await - automatic, high-level tasks:

- `fn async` creates a fiber-based task
- `await` runs it to completion and returns the final result
- `yield` inside `async` functions is a scheduler hint, not visible to the caller

Creating Fibers

Explicitly:

```
let f = Fiber.New(myFunction);
```

Implicitly via `async` functions:

```
fn async load() { return 42; }  
let f = load();    // creates a fiber, does not run yet
```

A newly created fiber does not run until resumed.

Resuming

`Fiber.Resume(f)` - runs the fiber until it yields or returns, then resumes the caller:

```
fn counter() {  
    yield 1;  
    yield 2;  
    return 3;  
}  
  
let f = Fiber.New(counter);  
Fiber.Resume(f);    // 1  
Fiber.Resume(f);    // 2  
Fiber.Resume(f);    // 3
```

`Fiber.Resume(f, true)` - **detached** mode, runs fiber in background:

```
Fiber.Resume(f, true); // caller continues immediately, fiber runs independently
```

Fiber States

State	Meaning
READY	Queued to run
RUNNING	Currently executing
YIELD	Yielded a value
WAIT_FIBER	Waiting for another fiber
WAIT_IO	Sleeping or waiting for IO/timer
FINISHED	Completed

Yield

`yield value` suspends the fiber and returns the value to whoever called `Fiber.Resume`:

```
fn producer() {  
    yield "first";  
    yield "second";  
    return "done";  
}
```

Inside async functions, `yield` acts as a scheduler hint - it does not affect the `await` return value.

Await

`await f` waits for fiber `f` to complete and returns its final `return` value (not yield values):

```
fn async compute(x) {  
    yield x;          // ignored by await  
    return x * 2;  
}  
  
let v = await compute(5); // 10
```

Awaiting an async function:

```
fn async load() { return 123; }  
let v = await load();    // 123
```

Non-blocking built-ins

Several standard-library functions have an `Async` variant that offloads blocking work to a background thread so other fibers keep running:

Sync	Async	Module
<code>File.ReadText(path)</code>	<code>File.ReadTextAsync(path)</code>	File
<code>File.ReadAllBytes(path)</code>	<code>File.ReadAllBytesAsync(path)</code>	File
<code>File.WriteText(path, text)</code>	<code>File.WriteTextAsync(path, text)</code>	File
<code>Net.ResolveDNS(host)</code>	<code>Net.ResolveDNSAsync(host)</code>	Net
<code>Os.Execute(cmd)</code>	<code>Os.ExecuteAsync(cmd)</code>	Os

Usage - call with `await` inside an async function:

```
fn async loadFiles(a, b) {  
    let text = await File.ReadTextAsync(a);  
    let bytes = await File.ReadAllBytesAsync(b);  
    return text;  
}
```

The calling fiber suspends at each `await` and is resumed automatically when the operation completes. Other fibers continue to run in the meantime.

Sleep

```
Fiber.Sleep(500); // sleep 500ms, non-blocking for other fibers
```

Other fibers continue to run while one is sleeping.

Error Handling

Exceptions from fibers propagate to the awaiting caller:

```
fn async fail() {
    throw(42, "something went wrong");
}

try {
    await fail();
} catch (e) {
    Console.WriteLine(e.Error); // "something went wrong"
}
```

Example: Interleaving

```
fn async worker(name, n) {
    let i = 0;
    while (i < n) {
        Console.WriteLine(name, i);
        yield;
        i += 1;
    }
}

let f1 = worker("A", 2);
let f2 = worker("B", 2);

// Scheduler interleaves them:
// A 0 | B 0 | A 1 | B 1
let r1 = await f1;
let r2 = await f2;
```

Example: Async Chain

```
fn async inner(x) {
    yield x;
    return x * 2;
}

fn async outer(x) {
    let z = await inner(x);
    return z + 1;
}

Console.WriteLine(await outer(5)); // 11
```

Common Mistakes

- `await` returns the final `return` value - not `yield` values

- Fibers don't run until resumed - create + schedule explicitly
- Detached fiber exceptions abort that fiber (and then the VM if uncaught)
- There is no preemption - `yield` in tight loops affects all fibers
- `yield` outside a fiber context is an error

10. Analyzer & Diagnostics

After parsing, Flaris runs a semantic analyzer. It resolves names, builds scopes, validates language rules, and emits helpful errors and warnings before bytecode is generated.

Diagnostics are reported like:

```
✖ Error: file:line:column: 1003: 'x' is not defined.  
⚠ Warning: file:line:column: 2000: Local 'x' is never read (prefix with '_' to suppress).
```

CODE is a stable number - errors are 1000+, warnings 2000+. The full table is in the reference under *R12 - Static Analyzer*. Diagnostics go to stderr, sorted by source position.

- **Errors** stop compilation (after a limit).
- **Warnings** allow compilation to continue.

Common Errors and Fixes

Undefined symbol / function

Symptoms

- `'x'` is not defined.
- Call to undefined function `'foo'`.

Fix

- Declare it (`let x = ...;`) in an accessible scope.
- Check spelling and casing.
- If it is a method, call it on the receiver: `obj.method()`.

Redeclaration

Symptom

- Redeclaration of `'name'`.

Fix

- Rename the inner declaration, or remove `let` and reassign instead.

Invalid `this`

Symptom

- ``this`` is not available here.

Fix

- Use `this` only inside **instance** methods (not static functions).

Invalid `super`

Symptom

- ``super`` is not available here.

Fix

- Use `super(...)` only inside an instance constructor, and `super.method(...)` only inside instance methods, of a class that has a base class.

`await` outside `async`

Symptom

- ``await`` used outside `async` function.

Fix

- Mark the containing function as `fn async ...`, or remove `await`.

Assigning to `const`

Symptom

- Cannot assign to `const 'X'`.

Fix

- Use `let` if rebinding is intended.
- Keep `const` if you only mutate object contents (e.g., `cfg.port = 1;` is OK) but do not rebind `cfg`.

Missing return on some paths

Symptom

- Function `'name'` has a declared return type but can finish without returning a value.

Fix

- Add `return` in every branch, or allow `nil` return if that is intended.

Assignment used as a condition

Symptom

- Invalid assignment (produces no output) in `if`-statement (also in `while`-statement, in `for` condition-statement)

Fix Assignments do not produce values in compiled bytecode. Rewrite:

```
// bad
if (x = get()) { ... }

// good
let tmp = get();
x = tmp;
if (tmp) { ... }
```

Common Warnings and Fixes

Unused variable / parameter

Symptoms

- Local 'x' is never read (prefix with '_' to suppress).
- Global 'x' is never read.
- Parameter 'p' is never used (prefix with '_' to suppress).

Fix

- Use it, remove it, or prefix the name with `_` to document that it is intentionally unused. The `_` prefix works for locals, parameters and globals alike. Loop and `catch` variables are exempt automatically.
- To silence one specific warning instead, add `// flaris-ignore[unused-symbol]` on the line (or on the line above it), or pass `-Wno-unused-symbol`. Errors cannot be silenced - see *R12 - Static Analyzer* in the reference.

Unreachable code

Symptom

- Unreachable code.

Fix

- Remove code after `return`, `throw`, or other guaranteed exits, or restructure with a condition.

Shadowing

Symptom

- 'x' shadows declaration at line N.

Fix

- Rename the inner variable to avoid confusion.

Class member or constructor does not exist

Symptoms

- Class 'Counter' has no method 'Nope'.
- Class 'Counter' has no member 'missing'.
- Constructor of 'Counter' expects 0–1 arguments, got 3.

Fix

- Instances are sealed: declare the field with `let` in the class body, or fix the name. These are only reported when the receiver's class is provable - `this`, a local from `new C()`, or a parameter annotated `v: C`.

Circular inheritance

Symptom

- Class 'A' is part of a circular inheritance chain.

Fix

- Break the cycle. Base classes may be declared in any order, so this is a real cycle, not a forward reference. A cycle cannot be left to the runtime: method lookup would walk the base chain forever.

Operator does not accept an operand type

Symptom

- Operator '-' does not accept string as its left operand.
- Operator '<<' does not accept float as its left operand.

Fix

- Convert explicitly - `(int)x`, `Convert.ToInt(s)`. Note that `<<`, `>>`, `&`, `|` and `^` run on integers only and never widen to float. `+` is never reported: it falls back to string concatenation for any operand pair.

11. Debugging

Flaris debugging is built directly into the language and VM. No external debugger required - and no separate adapter process either: the VM speaks the Debug Adapter Protocol itself, so an editor talks straight to it.

In VS Code

Install the extension from `vscode-flaris/` and press `F5` on any `.fls` file. Breakpoints go in the gutter, `F10`/`F11` step, and the Variables pane expands objects, class instances and arrays.

Two things are worth knowing because they have no equivalent in most debuggers:

- **Each fiber is a thread.** The editor's thread picker is a fiber picker - you can read the stack of any parked fiber while the program is stopped.
- **Right-click a frame in the Call Stack → Open Disassembly View** to see the bytecode around it.

Breakpoints can be added and removed while the program runs, and the pause button interrupts it at the next source line. A `launch.json` is only needed to change defaults:

```
{
  "type": "flaris",
  "request": "launch",
  "name": "Debug current Flaris file",
  "program": "${file}",
  "cwd": "${workspaceFolder}",
  "vmArgs": ["--libs=./libs"]
}
```

Any other DAP-capable editor can attach with `flarism --dap=<port> app.fl`, which listens on loopback and waits for a client before running anything.

Everything below works the same from a terminal, which is what you want over SSH or in a container.

Debug Symbols

Debug symbols map VM instructions back to source code (file, line, function). Enabled by default.

```
flarism script.fl          # symbols enabled (default)
flarism script.fl --strip  # strip symbols for production
```

With symbols: error messages show file, line, and function name.

Without symbols: errors show only memory addresses.

Overhead is minimal (~5-10% bytecode size, ~2-3% runtime).

Breakpoints

Run with `--debug` and the program stops at the start of `Main`, where you can set breakpoints by function or line and step through from there - no source edits needed:

```
$ flarism --debug app.fl
[Entry] Main (app.fl:40) [fiber 1]
(fdb) b processData          # or: b 42, or b app.fl:42, or b Order.Describe
breakpoint 1 at app.fl:12 (processData)
(fdb) c

[Breakpoint 1] processData (app.fl:12) [fiber 1]
(fdb) n                      # step over; s steps in, finish runs to the return
[Step] processData (app.fl:13) [fiber 1]
(fdb)                        # an empty line repeats the last step
```

`--debug` disables the JIT: a JIT-compiled function runs natively and never reports its lines, so breakpoints inside one could not fire.

If an exception is never caught, `--debug` opens the prompt at the throw site with every frame still standing, so you can read the values that caused it:

```
[Unhandled exception: code 42] level3 (app.flr:3) [fiber 1]
(fdb) locals
  [0] x                = (int) 11
  [1] doomed            = (int) 22
```

Without `--debug` the same failure is reported with its own code and message and the process exits with that code. `fibers` lists every live fiber and what it is waiting on, which is usually what you need when a cooperative program stops making progress rather than crashing.

You can also stop from the source itself. `breakpoint <code>;` pauses execution at that point:

```
fn processData(items) {
  breakpoint 100;    // entry

  for (let i = 0; i < len(items); i += 1) {
    if (items[i] < 0) {
      breakpoint 200; // triggered for negative values
    }
  }

  breakpoint 300;    // exit
}
```

Reaching one opens a prompt where you can walk the call stack and inspect values, then continue:

```
[Breakpoint:100] processData (app.flr:42) [fiber 1]
(fdb) locals
  [0] items          = (array) [4, -7, 12]
  [1] i               = (int) 1
(fdb) p items[1]
items[1] = (int) -7
(fdb) bt
-> #0  processData (app.flr:42)
     #1  Main (app.flr:9)
(fdb) c
```

`bt` walks the stack, `up/down` pick a frame, `locals` and `p <path>` inspect it, `list` shows the surrounding source, `dis` shows the bytecode around the instruction the frame is stopped at (and `dis ir <fn>` the JIT IR of a compiled function), and `help` lists everything. Inspection never changes the program: the prompt allocates nothing and raises nothing.

The prompt opens when stdin is a terminal, so a `breakpoint` left in committed code prints its location and continues under a pipe or in CI instead of hanging. Pass `--debug` to open it anyway. With `--verbose` the surrounding bytecode is shown as well. Use meaningful codes to track investigation stages.

Conditional breakpoints: no special syntax is needed - the condition is ordinary code, so any expression works:

```
if (data[i] == nil) {  
    breakpoint 200; // only breaks when data is nil  
}
```

The Debug Module

```
Debug.Value(obj)           // deep inspection: type, refs, address, value  
Debug.Assert(actual, expected) // halt with diff if not equal  
Debug.Pool()               // memory allocator statistics  
Debug.StackPtr()           // current stack depth (int)  
Debug.Refs(obj)            // reference count of obj  
Debug.GuardAddress(addr)   // software memory watchpoint  
Debug.Here(msg?)           // print current file:line:function
```

Debug.Value example:

```
let user = { name: "Alice", age: 30 };  
Debug.Value(user);  
// DBG-VAL: [0x7f8a4c002a40] type object. Refs 1. Value: { name: "Alice", age: 30 }
```

Debug.Assert example:

```
Debug.Assert(result, 42);  
// On failure: x Assertion failed on line 56: Left: 43 Right: 42
```

Debug.GuardAddress - software memory watchpoint:

```
Debug.GuardAddress(data); // any read/write to this address triggers a pause  
process(data);
```

Debug.Here - quick location tracing:

```
fn complexLogic() {  
    Debug.Here();           // prints file:line:function  
    Debug.Here("inside if"); // with optional message  
}
```

Verbose Modes

```
flarism script.flx          # normal (errors + output only)  
flarism script.flx --verbose # verbose (timing, info, function entry)  
flarism script.flx --trace   # very verbose (every opcode)  
flarism script.flx --verbose --time # verbose + timing report  
flarism script.flx --verbose --stats # verbose + allocation statistics
```

Typical Workflows

Development debugging:

```
fn investigate() {
    let data = loadData();
    Debug.Value(data);           // inspect loaded state
    breakpoint 100;              // pause here
    Debug.GuardAddress(data);     // watch for corruption
    process(data);
    Debug.Assert(data.valid, true); // verify post-condition
}
```

Memory leak detection:

```
flarism script.flx --stats
```

```
let before = VM.CurrentAllocations();
repeat1000Times();
VM.CompactMemory();
let after = VM.CurrentAllocations();
if (after > before + 10) {
    Console.WriteLine("Leak! Delta:", after - before);
}
```

Performance profiling:

```
flarism script.flx --verbose --time
# Output: [Time] Execution took 34.943 ms. Idle time 0ms
```

Development vs Production

```
# Development - all debug features
flarism script.flx --verbose --time

# Production - stripped, fast (compiled with --small and --strip )
flarism --exec app.flx
```

Best Practices

- Always keep debug symbols during development (default, don't use `--strip`)
- Use meaningful breakpoint codes - treat them as documentation
- Combine `--verbose` + `breakpoint` for step-through investigation
- Use `Debug.Assert` liberally in test code
- Use `Debug.GuardAddress` for hard-to-find memory corruption

See **Reference R11** for the complete Debug module API and opcode reference.

12. 10-Minute Tours

These tours help you get productive quickly if you're coming from another language.

Tour: Coming from C#

Mental Model

Flaris $\approx$ C# syntax + scripting flexibility + VM predictability

What feels familiar: classes, methods, `try/catch`, loops, `async/await`

What is different: no static typing, no GC, fibers instead of Tasks, `nil` instead of null references

Side-by-Side Basics

Variables:

C#	Flaris
<code>int x = 10;</code>	<code>let x = 10;</code>
<code>var x = 10;</code>	<code>var x = 10;</code>
<code>const int MAX = 100;</code>	<code>const MAX = 100;</code>
<code>x++</code>	<code>x += 1;</code>

Functions:

C#:

```
int Add(int a, int b) { return a + b; }  
Func<int, int> f = x => x * 2;
```

Flaris:

```
fn add(a: int, b: int): int { return a + b; }  
var f = fn(x) => x * 2;
```

Arrays vs List<T> + LINQ:

C#:

```
var xs = new List<int> { 1, 2, 3 };  
var doubled = xs.Select(x => x * 2).ToList();
```

Flaris:

```
var xs = [1, 2, 3];
var doubled = Array.Select(xs, fn(x) => x * 2);
```

null vs nil:

C# (dangerous):

```
User u = null;
Console.WriteLine(u.Name); // NullReferenceException!
```

Flaris (safe):

```
var u = nil;
Console.WriteLine(u.name); // nil (no exception)
```

All nil access is safe by default. Use `guard()` when you want to assert non-nil.

Classes:

C#:

```
class Point {
    public int X, Y;
    public Point(int x, int y) { X = x; Y = y; }
    public void Move(int dx, int dy) { X += dx; Y += dy; }
    public static Point Origin() => new Point(0, 0);
}
```

Flaris:

```
class Point {
    var x:int = 0;
    var y:int = 0;

    fn Constructor(x, y) {
        this.x = x;
        this.y = y;
    }
    fn move(dx, dy) {
        this.x += dx;
        this.y += dy;
    }
    fn static origin() {
        return new Point(0, 0);
    }
}
```

```
}  
}
```

Key differences: no access modifiers, dynamic dispatch by default, no type declarations.

Async/Await:

C#:

```
async Task<int> Compute() { return 42; }  
var v = await Compute();
```

Flaris:

```
fn async compute() { return 42; }  
var v = await compute(); // compute() returns a fiber
```

Mental model: a **fiber** $\approx$ a `Task`, but cooperatively scheduled. No OS threads, no preemption, no race conditions.

Errors:

```
// Nearly identical to C#  
try {  
    doWork();  
} catch (e) {  
    Console.WriteLine(e.Error);  
} finally {  
    cleanup();  
}
```

Mental Summary for C# Developers

- ✓ Think like C# for **structure**
- ✓ Think like a scripting language for **flexibility**
- ✓ Think like a VM for **performance**
- ✗ Don't rely on static typing
- ✗ Don't rely on GC behavior - use reference counting patterns

Tour: Coming from JavaScript

Mental Model

Flaris $\approx$ JavaScript syntax + explicit semantics + VM predictability

What feels familiar: dynamic typing, objects and arrays, closures, `async/await`

What is different: no `undefined`, safe `nil`, explicit globals, fibers not event loop

Concept	JavaScript	Flaris
undefined	Yes	No - just <code>nil</code>
null	Dangerous	Safe <code>nil</code>
Globals	Implicit	Explicit <code>let</code> at top level
Prototypes	Implicit	Classes
Event loop	Mandatory	Fiber scheduler
Implicit coercion	Yes	No

Side-by-Side

Variables:

```
// JavaScript
let x = 10;
x++;

// Flaris
let x = 10;
x += 1;    // or: x++ (postfix only, no prefix ++)
```

Arrays:

```
// JavaScript
xs.map(x => x * x);

// Flaris
Array.Select(xs, fn(x) => x * x);
```

Objects:

```
// JavaScript
let user = { name: "Ada" };
console.log(user.age); // undefined

// Flaris
let user = { name: "Ada" };
Console.WriteLine(user.age); // nil
```

Functions:

```
// JavaScript
const f = x => x * 2;
```

```
// Flaris
let f = fn(x) => x * 2;
```

Async/Await:

```
// JavaScript (event loop)
async function f() { return 42; }
let v = await f();

// Flaris (fiber scheduler)
fn async f() { return 42; }
let v = await f();
```

Classes:

```
// JavaScript
class Point {
  constructor(x, y) { this.x = x; this.y = y; }
}

// Flaris (instance fields must be declared)
class Point {
  let x = 0; let y = 0;
  fn Constructor(x, y) { this.x = x; this.y = y; }
}
```

Mental Summary for JavaScript Developers

- ✓ Familiar syntax
- ✓ Predictable semantics - no coercion surprises
- ✓ No event-loop surprises
- ✗ No `undefined`
- ✗ No implicit coercion

Tour: Coming from Rust

Mental Model

Rust's safety goals applied to a dynamic, embeddable VM

What feels familiar: predictable execution, no undefined behavior, explicit control flow, strong safety defaults

What is different: dynamic typing, reference counting (not borrowing), runtime errors (not `Result<T,E>`)

Concept	Rust	Flaris
Typing	Static	Dynamic

Concept	Rust	Flaris
Memory	Ownership + borrowing	Reference counting
Nulls	<code>Option<T></code>	Safe <code>nil</code>
Errors	<code>Result<T, E></code>	Runtime exceptions
Concurrency	Threads + async	Fibers
Abstraction cost	Zero-cost	Explicit, visible

Side-by-Side

Variables:

```
let mut x = 10;
x += 1;
const MAX: i32 = 100;
```

```
let x = 10;
x += 1;      // or: x++ (postfix only, no prefix ++)
const MAX = 100;
// All variables are mutable unless const. No type annotations required.
```

Functions:

```
fn add(a: i32, b: i32) -> i32 { a + b }
```

```
fn add(a, b) { return a + b; }
// or with optional annotations:
fn add(a: int, b: int): int { return a + b; }
```

Arrays vs `Vec<T>`:

```
let xs = vec![1, 2, 3];
let sum: i32 = xs.iter().sum();
```

```
let xs = [1, 2, 3];
let sum = Array.Reduce(xs, fn(acc, x) => acc + x, 0);
```

`Option<T>` vs `nil`:

```
let u: Option<User> = None;
// u.name -> compile error
```

```
let u = nil;
Console.WriteLine(u.name); // nil - safe, no panic
```

Error handling:

```
fn div(a: i32, b: i32) -> Result<i32, String> {
    if b == 0 { Err("div by zero".into()) }
    else { Ok(a / b) }
}
```

```
fn div(a, b) {
    if (b == 0) throw(100, "div by zero");
    return a / b;
}
try { div(10, 0); } catch (e) { Console.WriteLine(e.Error); }
// Rust: explicit propagation. Flaris: explicit error boundaries.
```

Concurrency:

```
std::thread::spawn(|| { /* work */ });
```

```
let f = Fiber.New(fn() { /* work */ });
Fiber.Resume(f);
// Or async:
fn async task() { return 42; }
let v = await task();
```

Mental model: no data races by default, cooperative scheduling.

Mental Summary for Rust Developers

- ✓ Predictable execution
- ✓ No undefined behavior
- ✓ Explicit performance costs
- ✓ Strong safety defaults
- ✗ No compile-time type guarantees
- ✗ No zero-cost abstractions
- ✗ Errors are runtime, not typed

Think of Flaris as **a safe, dynamic VM - not a systems language**.

13. Package Manager (flarispmm)

flarispmm is the official Flaris package manager - written entirely in Flaris itself. It installs third-party libraries from any public or private git repository (or from a direct `.flx` URL), compiles them to `.flx` bytecode, and places them in the per-user library directory (`~/.flaris/libs`, or `$FLARIS_LIBS`) where flarispmm resolves them automatically - no `--libs` flag needed.

Installation

`flarisp` ships as two files alongside `flarism`:

File	Purpose
<code>flarisp.flx</code>	Compiled package manager bytecode
<code>flarisp</code>	Shell wrapper - runs <code>flarism --exec flarisp.flx "\$@"</code>

Copy both to the same directory as `flarism` (e.g. `/usr/local/bin/`). The wrapper finds `flarisp.flx` next to itself automatically.

```
# Verify installation
flarisp version
# flarisp 0.4.0
```

`flarisp.flx` itself is signed with the official `flaris-lang.org` key - verify any copy with `flarism --sig-info flarisp.flx`.

Prerequisites: `git` in `PATH` for git-based packages; `curl` (ships with macOS, most Linux distros, and Windows 10+) or `wget` for direct `.flx` downloads.

Quick Start

```
# 1. Create a manifest in your project directory
flarisp init

# 2. Add a package from GitHub (or any git host)
flarisp add https://github.com/user/mylib

# 3. Run your program - installed packages resolve automatically
flarism myapp.flx
```

That's the entire workflow. `flarisp add` clones, compiles, verifies, and registers the package in one step.

The `flaris.json` Manifest

Every project managed by `flarisp` has a `flaris.json` at its root. `flarisp init` creates one:

```
{
  "name": "my-app",
  "version": "0.1.0",
  "description": "My Flaris application",
  "author": "Your Name",
  "dependencies": {
    "https://github.com/user/mylib": "main",
    "https://github.com/user/other": "v1.2.0",
    "https://example.com/tool.flx": "direct"
```

```
{,
  "integrity": {
    "https://example.com/tool.flx": { "tool.flx": "sha256:..." }
  },
  "publishers": {
    "https://example.com/tool.flx": { "key": "ed25519:...", "signer": "example.com" }
  }
}
```

Dependency keys are git clone URLs or direct `.flx` download URLs. **Values** are the branch name, tag, or `"main"` for the default branch (`"direct"` for `.flx` URLs). Version pinning with a tag (`"v1.2.0"`) is recommended for production - it gives you a reproducible build.

The `integrity` and `publishers` sections are maintained by `flarisp` itself - they record SHA-256 pins and signed packages' publisher keys (see [Security](#) below). Commit `flaris.json` including these sections: anyone who runs `flarisp install` from your checkout then gets exactly the bytes you pinned, or a hard refusal.

URLs are restricted to `https://` (plus `ssh://git@` for git) - plain-http and exotic git transports are refused, both from the command line and from a cloned manifest.

Commands

`flarisp init`

Creates `flaris.json` in the current directory. Does nothing if one already exists.

```
flarisp init
# [✓] Created flaris.json
```

`flarisp add <url> [version] [--key <pubkey>]`

Clones the package, compiles all `.fls` files in its root, installs the resulting `.flx` files into `~/.flaris/libs`, records their SHA-256 pins, and registers the dependency in `flaris.json`. For a direct `.flx` URL it downloads instead of cloning, and additionally pins the publisher's Ed25519 key if the file is signed.

```
# Default branch
flarisp add https://github.com/user/mylib

# Specific tag
flarisp add https://github.com/user/mylib v2.0.0

# Private / self-hosted git
flarisp add https://git.mycompany.com/internal/utils
```

```
# Pre-compiled, signed .flx - verify the publisher key while adding
flarisp add https://example.com/tool.flx --key <ed25519-pubkey-hex>
```

`add` is the only command allowed to *move* an existing integrity pin - use it deliberately when you intend to accept a new release of a direct `.flx` package.

After `add`, import the package by the name of its exported `.fls` file:

```
import { MyLib } from library("MyLib", "1.0");
```

Source is cached in `~/.flaris/libs/.src/` - re-running `add` on an already-installed package skips the clone and only recompiles.

`flarisp install`

Reads `flaris.json` and installs every listed dependency. Use this after cloning a project someone else created:

```
git clone https://github.com/user/my-project
cd my-project
flarisp install
flarism main.fl
```

`install` is strictly reproducing what was pinned: every download and every locally compiled module is verified against the recorded `sha256: pin` (and, for signed `.flx` packages, the pinned publisher key) **before** it is moved into `~/.flaris/libs`. A mismatch is a hard refusal - `install` never re-pins. If upstream legitimately released something new, accept it explicitly with `flarisp add` (direct `.flx`) or `flarisp update` (git packages).

`flarisp list`

Shows all registered packages and whether their source is present locally:

```
flarisp list
# Packages in flaris.json:
#   mylib @ v2.0.0   [installed]
#     https://github.com/user/mylib
#   utils @ main    [installed]
#     https://github.com/user/utils
```

`flarisp update [url]`

Pulls the latest commits for all packages (or one specific package), recompiles, and re-pins the integrity hashes to the new build:

```
flarisp update # update everything
flarisp update https://github.com/user/mylib # update one package
```

Note: If the package is pinned to a tag ("v1.2.0"), `git pull` will have nothing new to fetch. To upgrade to a newer tag, use `flarisp remove` then `flarisp add` with the new version.

`flarisp remove <url>`

Removes the compiled `.flx` files produced by a package and unregisters it - including its integrity and publisher pins - from `flaris.json`. The cloned source in `~/.flaris/libs/.src/` is kept as a cache - a subsequent `add` reuses it and skips the network round-trip.

```
flarisp remove https://github.com/user/mylib
# [.] Deleted MyLib.flx
# [✓] Removed mylib from flaris.json
```

Using Installed Packages

Installed packages live in `~/.flaris/libs` (override with `$FLARIS_LIBS`) as `.flx` files - the same per-user directory `flarism` searches automatically, so no flag is needed.

Inside your code, import exactly as you would any `stdlib` module - the name comes from the `.fls` filename in the package:

```
// Package installed from https://github.com/stefansolid/flaris-test-package
// which contains Greet.fls → compiled to ~/.flaris/libs/Greet.flx

import { Greet } from library("Greet", "1.0");

fn Main() {
    Greet.Hello("Coder"); // Hello, Coder!
    Greet.Shout("it works"); // IT WORKS!!!
}
```

Run:

```
flarism myapp.fls
```

Directory Layout

The manifest stays in your project; the installed packages are per-user:

```

my-project/
├─ flaris.json      ← manifest + pins (commit this)
└─ main.flx

~/.flaris/libs/    ← or $FLARIS_LIBS
├─ MyLib.flx       ← installed packages (build artifacts)
├─ Utils.flx
├─ .staging/       ← downloads/compiles awaiting verification
│                  (never resolvable by the module loader)
└─ .src/           ← cloned git sources (cache)
    ├─ github.com_user_mylib/
    └─ github.com_user_utils/

```

Nothing package-related needs a `.gitignore` entry - only `flaris.json` lives in the project, and it *should* be committed. Anyone who checks out your project runs `flarisp install` to reproduce the exact pinned dependencies.

Creating a Package

Any git repository containing `.fls` files is a valid Flaris package. The minimal structure:

```

my-package/
├─ flaris.json    ← required: declares name and version
└─ MyLib.flx     ← one or more library files

```

`flaris.json` for a package:

```

{
  "name": "MyLib",
  "version": "1.0.0",
  "description": "What this library does",
  "author": "Your Name",
  "homepage": "https://github.com/user/my-package"
}

```

`MyLib.flx` - define and export your API:

```

// MyLib.flx - Example library
//
// Usage:
//   import { MyLib } from library("MyLib", "1.0");
//   MyLib.Greet("world");

class MyLib {

  fn static Greet(name?: string) {

```

```
    let n = name != nil ? name : "World";
    Console.WriteLine("Hello, " + n + "!");
}

fn static Version() : string {
    return "1.0.0";
}
}

export { MyLib };
```

Publish by pushing to any git host - no registry needed. Users install with:

```
flarisp add https://github.com/you/my-package
```

Version Tags

Tag your releases so users can pin to a stable version:

```
git tag v1.0.0
git push origin v1.0.0
```

Users can then install the exact release:

```
flarisp add https://github.com/you/my-package v1.0.0
```

Version Matching

When `flarisp` compiles a package it embeds the `version` field from the package's `flaris.json` into the bytecode using `--version=<ver>`. The `library()` call in your code specifies a version requirement that must match:

```
// Matches any 1.x release
import { MyLib } from library("MyLib", "1.0");

// Exact match
import { MyLib } from library("MyLib", "1.2.3");

// Range syntax
import { MyLib } from library("MyLib", ">=1.0 <2.0");
```

If the version in the bytecode does not satisfy the requirement, the VM halts before execution. This prevents accidentally running an incompatible version of a library.

Security - Integrity and Publisher Pinning

`flarisp` enforces three layers, all recorded in `flaris.json` and all verified while a file is still staged - a package that fails any check never reaches the import path:

- Exact-bytes integrity (all packages).** Every installed `.flx` gets a `sha256:` pin under `integrity` (trust-on-first-use at `add` time). On every later `install`, the downloaded file - or the locally recompiled module, for git packages - must hash to exactly the pinned value or it is refused. `install` can never move a pin; only an explicit `add` (direct `.flx`) or `update` (git) accepts new bytes.
- Publisher continuity (signed `.flx` packages).** When a signed `.flx` is added, the publisher's Ed25519 key is pinned under `publishers`. Every later download must be signed by the *same* key - a tampered file, a stripped signature, or a silently swapped key all block the install. Accept an intentional key rotation by re-adding with `--key <new-pubkey>`.
- Publisher verification at add (`--key`).** `flarisp add <url> --key <pubkey>` refuses the package unless it is signed by exactly that key - use it when the publisher advertises their key out-of-band (README, website).

Git/source packages are compiled locally and left unsigned on purpose: a signature applied at compile time would attest to the *installer*, not the publisher. They are trusted via the git ref plus the sha256 pin of the compiled output instead. Note that the compiled hash embeds debug info (including the local source path), so git-package pins are machine-specific; a pin mismatch on another machine means "recompile differs - review and run `flarisp update`", not necessarily an attack.

Signature inspection is delegated to the VM: `flarism --sig-info <file.flx>` prints `unsigned or signed|<trusted|valid|invalid>|<pubkey>|<signer>`.

VM-side Fingerprint Pinning

For security-sensitive dependencies you can additionally pin the fingerprint at the *import site*, enforced by the VM itself on every run. Get the fingerprint after installing:

```
flarism --sha256 ~/.flaris/libs/MyLib.flx
# Fingerprint(SHA256): [a3f1...64hexchars...]
```

Then lock the import to that exact binary:

```
import { MyLib } from library("MyLib", "1.0", "a3f1...64hexchars...");
```

If the bytecode does not match - even if the version is correct - the VM halts immediately. Use this for libraries that handle credentials, crypto, or authentication.

The fingerprint is the whole-file SHA-256 (`flarism --sha256 == sha256sum`), recomputed from the bytes actually loaded - there is no self-certifying hash inside the file. A `sha256:` prefix (the form `flarisp` stores under `integrity` in `flaris.json`) is accepted, so you can paste that value directly. The pin is always enforced and is not affected by `--no-verify`.

Publisher signatures

Where the fingerprint pins the *exact bytes*, an Ed25519 **publisher signature** pins the *author* - and survives version bumps. When you add a signed `.flx` package, `flarisp` records the publisher's public key under `publishers` in `flaris.json`, and every later `install/update` must be signed by that same key, or it is refused (catching a tampered file or a silently swapped key):

```
# Pin the publisher to the key you expect (cross-checked before recording)
flarisp add https://example.com/MyLib.flx --key <ed25519-pubkey-hex>
```

Git/source packages are compiled locally and left unsigned on purpose - a signature applied at install time would attest to *you*, not the publisher - so they rely on the git ref plus fingerprint instead. To enforce signatures at runtime, trust the publisher's key (`flarism --import-key <name> <pubkey>`) and run with `--require-signed`.

A `.flx` is executable code, not a sandbox. It can touch the filesystem, spawn processes, and (with `--unsafe`) call FFI and raw memory - so only run bytecode you trust, and pin third-party libraries to a known hash. Note that `--jit` (JIT) widens the trust surface: JIT-compiled code omits the interpreter's runtime type checks, so never enable `--jit` on untrusted bytecode. See Reference R1 "Security model".

Environment Variables

Variable	Effect
FLARISM	Override the path to <code>flarism</code> used for compiling packages. Useful when multiple versions are installed.

*End of Flaris Guide. For complete API reference, type system details, `stdlib` documentation, performance tuning, and VM internals, see the **Flaris Reference**.*