



Flaris Language Reference

The Flaris programming language · flaris-lang.org · github.com/flaris-lang/flaris · flaris-lang.org/discord

Version 1.0.0.9 **Generated** 2026-07-25 **License** see repository

© 2026 SolidInvention AB — created and maintained by SolidInvention AB

Table of Contents

- [R1 - CLI Reference](#)
- [R2 - Type System](#)
- [R3 - Built-in Functions](#)
- [R4 - The Exception class](#)
- [R5 - Operator Precedence](#)
- [R6 - VM Limits](#)
- [R7 - Performance Model](#)
- [R8 - Standard Library](#)
- [R9 - Debug Reference](#)
- [R10 - Writing Fast Flaris Code](#)
- [R11 - JIT Compilation](#)
- [R12 - Static Analyzer](#)

R1 - CLI Reference

The Flaris VM binary is `flarismvm`.

Invocation Modes

Command	Description
<code>flarismvm <file.flx> [options]</code>	Compile and run source file
<code>flarismvm - [options]</code>	Compile and run source from stdin
<code>flarismvm --compile <input.flx> <output.flx> [options]</code>	Compile source to bytecode
<code>flarismvm --compile - <output.flx> [options]</code>	Compile source from stdin to bytecode
<code>flarismvm --exec <file.flx> [options]</code>	Execute compiled bytecode
<code>flarismvm --exec - [options]</code>	Execute bytecode from stdin
<code>flarismvm --disasm <file.flx></code>	Disassemble bytecode (debug view)
<code>flarismvm --eval '<code>' [options]</code>	Run code directly from string
<code>flarismvm --sha256 <file.flx></code>	Print the whole-file SHA-256 (identical to <code>sha256sum</code>) - the value to pin via <code>library(file, version, hash)</code>
<code>flarismvm --sig-info <file.flx></code>	Print the embedded signature status, machine-readable: <code>unsigned</code> , or <code>signed <trusted valid invalid> <pubkey-hex> <signer></code> (see <i>Package signing</i>)
<code>flarismvm --import-key <name> <pubkey></code>	Add a 64-hex Ed25519 public key (with a label) to <code>~/.flaris/trusted_keys</code>

Command	Description
<code>flarism --check <file.fl></code>	Compile without running: full analyzer diagnostics plus a JIT eligibility report - see R11 and R12
<code>flarism -Wno-<name> ...</code>	Silence one analyzer warning, e.g. <code>-Wno-unused-symbol</code> - see R12
<code>flarism -Werror ...</code>	Treat any analyzer warning as an error
<code>flarism -w ...</code>	Silence all analyzer warnings
<code>flarism --diagnostics json ...</code>	Emit diagnostics as a JSON array on stdout
<code>flarism --version</code>	Print version string
<code>flarism --compile <input.fl></code> <code><output.fl> --bundle='<file1.fl>;...'</code> <code>[options]</code>	Compile and bundle named modules into a self-contained <code>.fl</code>
<code>flarism --embed <input.fl> .fl></code> <code><output> [options]</code>	Build a self-contained native executable (program + runtime in one file). See <i>Self-contained executables</i>

Flags - Disablers (default ON)

Flag	Effect
<code>--no-opt</code>	Disable compiler optimizations
<code>--strip</code>	Disable debug symbols (smaller, faster bytecode)
<code>--small</code>	Extra size optimizations for compiled <code>.fl</code> ; also removes string symbol names
<code>--no-verify</code>	Disable FFI library (Ffi.Load) SHA-256 checks

Flags - Enablers (default OFF)

Flag	Effect
<code>--unsafe</code>	Enable unsafe code and FFI
<code>--mem</code>	Memory report at shutdown: peak RSS, slab footprint, peak/total object counts, heap bytes, block regions, and leak count. Exits non-zero if leaks are detected.
<code>--stats</code>	Print VM statistics after execution (implies <code>--time</code>)
<code>--time</code>	Print timing report
<code>--verbose</code>	Verbose output (also enables the timing report)
<code>--list-keys</code>	Print the built-in and trusted signing keys, then exit
<code>--trace</code>	Very verbose / print disassembled bytecode before execution
<code>--jit</code>	Enable JIT: automatically compiles every eligible function - embeds JIT code at compile time and runs it natively at run time (ARM64 and x86-64)
<code>--debug</code>	Stop at the start of <code>Main</code> and open the debugger prompt; disables <code>--jit</code> (see R9)
<code>--dap=<port></code>	Serve the Debug Adapter Protocol on loopback for an editor; implies <code>--debug</code>

Diagnostics are colored only when stderr is a VT-capable TTY; setting the `NO_COLOR` environment variable (non-empty, no-color.org convention) forces plain output.

Exit Codes

Code	Meaning
0	Success
<code>fn Main</code> return	A numeric return value from <code>Main</code> becomes the process exit code
exception code	An uncaught exception exits with the exception's numeric code
1	Tokenizer error
2	Compile error (bad source on the CLI or in a compiled unit)
3	Assert failure (<code>Assert</code> / <code>AssertEq</code> or an internal invariant)

Code	Meaning
4	File error (missing or unreadable source/bytecode)
5	Debugger quit (q)
6	Reserved (analyzer failures currently surface as code 2)
7	Fingerprint/pin mismatch on a signed or pinned import
8	Reserved

VM Tuning Flags

These override the compiled-in defaults. Values must be within min/max ranges listed in [R6](#).

Flag	Effect
<code>--stack=<int></code>	Override evaluation stack size
<code>--fifo=<int></code>	Override object cache (FIFO) size
<code>--slabs=<int></code>	Override slab allocator count
<code>--fibers=<int></code>	Override maximum fiber count
<code>--frames=<int></code>	Override maximum call-frames
<code>--io-threads=<int></code>	Worker threads for the I/O pool (1-16). Default: half the logical CPUs, minimum 4. Raise it if <code>--stats</code> reports jobs run inline.

Imports

Flag	Effect
<code>--min-version=<M.m.p.b></code>	Set the minimum runtime version stamped into compiled <code>.flx</code> output. Compile mode only; silently ignored at runtime. Examples: <code>--min-version=1.0</code> , <code>--min-version=2.1.3</code> . Default: <code>1.0.0.0</code>
<code>--libs=<path></code>	Set default library loading-path (otherwise <code>.</code>) for VM to load libraries for later resolving imports
<code>--sign=<hex keyfile></code>	Compile mode only. Embed an Ed25519 signature in the <code>.flx</code> . Accepts the 128-hex-char secret key inline or a path to a file containing it. See <i>Package signing</i> .
<code>--require-signed</code>	Runtime. Refuse to load any <code>.flx</code> that is unsigned, has an invalid signature, or is signed by a key not in <code>~/.flaris/trusted_keys</code> .

Module name resolution (case-sensitive)

`library(name, ...)` (and `VM.Import`) resolve `name` to a file by **exact, case-sensitive match**:
`library("Signals")` → `Signals.flx` (on the `--libs` path), `library("./x/Y")` → `./x/Y.flx`, a URL → that exact path. The name is never lowercased, so `library("signals")` finds `Signals.flx` only on case-insensitive filesystems (macOS/Windows) and **fails on Linux**. Keep the import string, the on-disk filename, and any pinned name byte-identical (convention: capitalized first letter). Imported symbol names must also match the library's `export` exactly - `{ Jwt } ≠ { JWT }`.

Search order for a non-path name: (1) the name as a path relative to the working directory, (2) each `;-` separated entry of `--libs`, (3) the per-user install dir - `~/.flaris/libs` (POSIX) / `%USERPROFILE%\flaris\libs` (Windows), or `$FLARIS_LIBS` if set. The first version-satisfying candidate wins (`--libs` beats the global dir). Step 3 lets a package manager install once into the home directory and have `flarism script.flx` resolve imports with no `--libs` flag.

Library integrity (pinning)

`library(name, version, hash)` (and `VM.Import(name, version, hash)`) verify the loaded module against `hash` - the **whole-file SHA-256** of the `.flx`, identical to `sha256sum file.flx` and to `flarism --sha256`

`file.flx`. The hash is recomputed over the actual bytes loaded, so it detects any tampering of a downloaded library. Supply it from a trusted source (a registry, a README, or `flaris.json`).

- An optional `sha256:` prefix is accepted, so a value recorded by flarisp in `flaris.json` (`"sha256:<hex>"`) can be pasted verbatim.
- A mismatch halts the VM (fail-closed) and is **not** affected by `--no-verify` (`--no-verify` only disables the FFI shared-library check in `Ffi.Load`).
- Omit the hash to load without pinning.
- No hash is embedded in the `.flx` itself: a self-recomputed embedded hash would be self-certifying and offer no protection against tampering. The trust anchor is the hash you pass in, delivered out-of-band.

Package signing (Ed25519)

Where the pinned hash above answers *"are these the same bytes I vetted?"* (trust on first use), an **embedded Ed25519 signature** answers *"who produced this?"* - authenticity that survives an untrusted transport, registry, or mirror.

A signature is carried inside the `.flx` header (192-byte header; algorithm + 32-byte public key + 64-byte signature). It covers the SHA-256 of the **entire file with the signature field masked to zero**, so it authenticates the header metadata (name, version, entry), the constants, and the code together.

Producing a signed `.flx`:

```
# One-time: generate a key pair (keep the secret key offline).
flarism --eval 'let kp = Crypto.Ed25519KeyPair();
               File.WriteAllText("publisher.sk", kp.SecretKey);    // 128 hex chars
               File.WriteAllText("publisher.pk", kp.PublicKey)'    // 64 hex chars

# Compile and sign (signing is deterministic - the .flx stays reproducible).
flarism --compile mylib.flx mylib.flx --sign=publisher.sk
```

The public key is printed at compile time and embedded in the file; verify it later with `Crypto.Ed25519Verify`.

Verifying at load:

Situation	Default (soft)	With <code>--require-signed</code>
Unsigned <code>.flx</code>	runs (no warning)	rejected
Valid signature, signer in <code>trusted_keys</code>	runs	runs
Valid signature, signer not in <code>trusted_keys</code>	runs	rejected (untrusted signer)
Invalid / tampered signature	rejected (fail-closed)	rejected

- A **tampered** signed file always fails to load, in both modes - the signature no longer matches the recomputed digest.
- **The official Flaris signing key is compiled into the runtime** as a built-in trust anchor, so the bundled standard library verifies under `--require-signed` with no file to download or configure. The published key is also on the website as a cross-check.
- Additional publishers are trusted via `~/.flaris/trusted_keys` (or `$FLARIS_TRUSTED_KEYS`): one 64-hex public key per line, optionally followed by whitespace and a human-readable label, e.g. `b6e2...79be`

`acme corp`. Blank lines and `#` comments are ignored. This file is **additive** on top of the built-in anchor - add a third-party publisher's key here to trust their packages too. Use `flarism --import-key <name> <pubkey>` to append one safely (it dedupes). The label (and the built-in `flaris-lang.org` name) is shown as the *signer* by `flarism --disasm` and `flarism --sig-info`.

- Signature verification is independent of `--no-verify` and of the `library(..., hash)` pin; use signing for *authenticity* and the hash pin for *exact-artifact* checks.

Security model (trust boundary)

A `.flx` is **executable code, not a sandboxed format**. Running one is equivalent to running a native program: it can read and write files, spawn processes (`0s.Run`), open sockets, and - with `--unsafe` - call arbitrary C via FFI and access raw memory. On load the VM validates bytecode structure (validate chunks, JIT-IR validation, size/count caps) to contain *malformed* input safely, but it is **not** a security sandbox for *malicious* input.

- **Only run `.flx` files you trust.** For third-party libraries, pin the whole-file hash via `library(name, version, hash)` (above) so you execute exactly the artifact you vetted, and/or require a trusted Ed25519 signature (`--require-signed`, see *Package signing*) to authenticate the publisher. These are the primary defenses.
- **`--jit` (JIT) widens the trust surface.** JIT-compiled functions omit the runtime type checks the interpreter performs, so a *tampered* `.flx` run with `--jit` can crash or corrupt memory where the interpreter would raise a clean error. Do not enable `--jit` on untrusted bytecode.
- `--unsafe` grants FFI and raw-memory access; never combine it with untrusted code.

Bytecode verification

Before any chunk runs, the loader structurally validates it. This *contains malformed input* - a corrupt or truncated `.flx` is rejected with a clean error instead of crashing - but it is **not** type verification and **not** a sandbox for hostile input (see the trust-boundary note above). Validation runs in three passes over each function chunk and recurses into nested functions (nesting depth is bounded):

1. **Decode & operand checks.** Every instruction is decoded and its operands are range-checked: constant-pool indices must be in range (and property/method keys must actually be strings), local slots must be `< 128` (`MAX_LOCALS`), call and invoke argument counts must be `<= 16` (`MAX_CALL_ARGUMENTS`), and built-in module/function references must exist. Each instruction's decoded length must also match its canonical length, so an opcode cannot be mis-sized. Instruction boundaries are recorded for pass 2.
2. **Branch-target check.** Every jump, loop, branch, `try`, `foreach`/iterator and jump-table target must land exactly on a recorded instruction boundary - control can never jump into the middle of an instruction.
3. **Stack-depth check.** An abstract simulation computes the maximum value-stack depth the chunk can reach, which the VM reserves up front per call frame so the value stack cannot overflow at run time. Any opcode the simulation does not model is rejected (fail-closed) rather than assumed harmless.

Size and count caps also apply: `10 MB` of bytecode and `16,385` constants per function (see *Limits*). A chunk that fails any check is refused at load time; an invalid instruction encountered at run time raises `Exception.IllegalInstruction` (code `16`), a hard failure. These guarantees hold for the interpreter; `--jit` omits the interpreter's runtime type checks and so widens the trust surface as noted above.

Bundle files

A bundle is a self-contained `.flx` file produced with `--bundle=<file.flx>;...`. It embeds named dependency chunks in a single file with a TOC at the start. At runtime the VM pre-loads all bundled deps before execution begins - no filesystem access occurs for those modules.

Limitations:

- Only named modules will be loaded, and they have to be readable from direct path or via `--libs` folder

Inspecting a bundle:

```
flarism --disasm my_bundle.flx
```

The disassembler shows the bundle TOC, each dep chunk, and the main chunk.

Self-contained executables

`--embed <input> <output>` produces a single native executable that carries both your program and the runtime - one file to distribute, nothing to install on the target machine (comparable to Go's static binaries).

```
# From source, in one step
flarism --embed app.flx myapp
./myapp arg1 arg2          # runs standalone - no flaris on PATH, no --libs

# Statically link library dependencies into the same binary
flarism --embed app.flx myapp --bundle='./mylib.flx'

# Embed precompiled bytecode directly
flarism --embed app.flx myapp
```

How it works. The full `flarism` compiles the input (or reads the `.flx` directly), then appends the bytecode to the compiler-free `flaris` runtime, followed by a fixed 24-byte trailer recording the payload's offset and length. On startup any Flaris binary reads that trailer out of its own image; if a payload is present it runs the embedded program and passes all CLI arguments to the script, otherwise it behaves as the normal VM. The payload rides *after* the executable image, so the OS loader ignores it and a plain file copy still works.

Details:

- **Input** is detected by content (the `FLS2` magic), not by extension: a `.flx` is compiled first, a `.flx` is embedded as-is.
- The same compile-time options as `--compile` apply to a `.flx` input: `--bundle`, `--sign`, `--no-opt`, `--strip`.
- **Requires the `flaris` runtime stub.** It is resolved next to `flarism` first, then on `PATH` (the official installer places both). Only the full `flarism` can build embeds; `flaris` itself has no compiler and no `--embed`.
- The embedded runtime has no compiler and no `eval` - a bundled binary ships exactly the bytecode you built.
- Combine with `--bundle` so the shipped binary needs no `~/flaris/libs` and no `--libs`.

Platform notes. The mechanism is platform-agnostic: ELF, Mach-O, and PE all tolerate the payload as trailing "overlay" bytes past the mapped image, so self-contained binaries run on Linux, macOS (including Apple Silicon, where the runtime's ad-hoc signature covers the stub and the trailing payload is tolerated), and Windows. The only per-platform gap is code-signing for distribution: appending bytes invalidates strict macOS codesign /notarization and Windows Authenticode (and an unsigned `.exe` may draw SmartScreen friction, though it still runs).

Common Recipes

```
# Development – run with timing and verbose
flarism app.flx --verbose --time

# Production – execute pre-compiled bytecode
flarism --exec app.flx

# Compile for production (no debug symbols, optimized)
flarism --compile app.flx app.flx --strip

# Enable FFI, skip the FFI shared-library (.so) SHA-256 check
flarism app.flx --unsafe --no-verify

# Compile-only: analyzer diagnostics + JIT eligibility report, no execution
flarism --check app.flx

# Compile library with explicit version
flarism --compile mylib.flx mylib.flx --min-version=1.2.0

# Print the whole-file SHA-256 (== sha256sum) to pin as the third arg to library()
flarism --sha256 mylib.flx

# Compile and embed an Ed25519 signature, then run requiring a trusted signer
flarism --compile mylib.flx mylib.flx --sign=publisher.sk
flarism --exec mylib.flx --require-signed

# Inspect a package's signature (machine-readable) and trust a publisher
flarism --sig-info mylib.flx
flarism --import-key "acme corp" <ed25519-pubkey-hex>

# Compile with all imports bundled into a single .flx
flarism --compile app.flx app.flx --bundle='./data.flx'

# Bundle with explicit version
flarism --compile app.flx app.flx --bundle='./data.flx' --min-version=2.0.0

# Inspect bundle contents
flarism --disasm app.flx
```

```
# Build a single self-contained executable (program + runtime), libs bundled in
flarism --embed app.flx app --bundle='./data.flx'
./app

# Compile with JIT code embedded (for every eligible function)
flarism --compile --jit app.flx app.flx

# Run with JIT active (runs eligible functions as native code)
flarism --jit --exec app.flx

# See which functions are JIT-eligible during compilation
flarism --verbose app.flx

# Read source from stdin, compile and run
echo 'fn Main() { Console.WriteLine("hello"); }' | flarism -

# Compile from stdin to a .flx file
cat app.flx | flarism --compile - app.flx

# Execute pre-compiled bytecode piped via stdin
cat app.flx | flarism --exec -

# End-to-end pipe: compile and execute without touching the filesystem
cat app.flx | flarism --compile - /tmp/app.flx && flarism --exec - < /tmp/app.flx
```

R2 - Type System

Lexical Structure

How source text is broken into tokens (the numeric/float literal forms are detailed under [Type Limits](#)).

Encoding. Source is UTF-8. A leading UTF-8 BOM (EF BB BF) is skipped. Identifiers are ASCII: a leading `_`/letter followed by `_`/letters/digits.

Comments.

```
// line comment - to end of line
/* block comment */
/* block /* comments */ nest */    // fully balanced
```

Numeric literals. Decimal, `0x/0X` hex, `0b/0B` binary, `0o/0O` octal, and floats with a fractional part and/or `e/E` exponent (optional `+/-`). A digit is required on both sides of the decimal point (`0.5/5.0` , never `.5` or `5.`). `_` digit separators are allowed in every base and in the fraction and exponent (`0xFF_FF` , `0b1010_1010` , `1_000.000_5`); they are removed before the value is parsed.

String literals. `"..."` (single line) and `"""..."""` (triple-quoted, spans newlines, verbatim - no escape processing). Inside `"..."` a backslash at end of line is a line continuation (the newline is dropped). A raw

newline inside `"..."` is an error - use `\n`, a line continuation, or `"""..."""`. Recognised escapes:

- `\n \r \t \b \f` - control characters (newline, return, tab, backspace, form-feed)
- `\\ \" \'` - literal backslash, double quote, single quote
- `\xHH` - one byte from two hex digits
- `\uXXXX` - Unicode code point from exactly four hex digits (max `U+FFFF`)
- `\u{H..H}` - Unicode code point from 1-6 hex digits, up to `U+10FFFF` (e.g. `\u{1F600}` for 🤪)

Both `\u` forms map a surrogate value to `U+FFFD`. Use `\u{...}` for code points above `U+FFFF` (emoji, etc.); the four-digit `\uXXXX` cannot reach them. An unknown escape keeps the character literally (`"\q"` is `q`). Because Flaris strings are NUL-terminated at runtime, a `\0` or `\x00` **in a string literal is a compile error** - put binary data in a `Buffer`. String literals are capped at 64 KB of content.

Char literals. `'...'` holds exactly one Unicode code point: a raw UTF-8 character (`'a'`, `'å'`, `'🤪'`) or an escape (`'\n'`, `'\x41'`, `'å'`, `'\u{1F600}'`). Unlike strings, `'\0'` is valid - it is the code point 0, not an embedded NUL.

Keywords. `fn async static inline let var const global if else for foreach iter from to while switch case default break continue return yield await guard nil true false class this super new enum import export library try catch finally throw breakpoint`. The words `and` or `in` `is` `as` are operators. (`var` is an alias for `let`.)

Runtime Types

These are the actual types values carry at runtime.

Primitive Types

Type	Description
<code>nil</code>	Absence of value. Only one instance exists.
<code>bool</code>	Boolean <code>true</code> or <code>false</code> .
<code>char</code>	Single Unicode codepoint (unsigned 32-bit).
<code>int</code>	Signed 64-bit integer. Values within $\pm 2^{60}$ are tagged immediates (no allocation); larger values are heap-allocated. Full <code>int64</code> range supported.
<code>float</code>	64-bit IEEE-754 double-precision.
<code>string</code>	Immutable UTF-8 byte sequence.

Composite Types

Type	Description
<code>array</code>	Dynamic, zero-indexed array. Max 10,000,000 elements.
<code>object</code>	Hash-based dictionary with string keys (property names). Indexing with a non-string key (<code>o[42] = x</code>) raises. For <code>int/char/float</code> or other value keys use <code>Collections.HashMap</code> / <code>Collections.OrderedMap</code> . Max 10,000,000 properties. Property iteration order is not stable across processes (see <code>Json.Canonicalize</code> for a deterministic ordering).
<code>exception</code>	Exception with code and message.

Structured Types

Type	Description
<code>class</code>	Class definition object.
<code>instance</code>	Instance of a class.

Execution Types

Type	Description
<code>fiber</code>	Lightweight cooperative coroutine.
<code>fn</code>	Callable function value. May be a named function, an anonymous function, or a closure (anonymous function with captured outer variables).
<code>module</code>	Built-in module namespace.

Binary Types

Type	Description
<code>block</code>	Raw memory buffer. Element count × element size bytes.
<code>pointer</code>	Runtime allocated pointer (FFI use).
<code>stream</code>	File/IO stream handle.

Type Limits

Integer

Flaris integers are full signed 64-bit. Values within $\pm 2^{60}$ are stored as **tagged pointer immediates** (zero heap allocation, fast path). Values outside that range fall back to heap-allocated objects - same type, same semantics, just slower to create.

Property	Value
Representation	tagged immediate for $ x < 2^{60}$, heap-allocated otherwise
Minimum	$-9,223,372,036,854,775,808$ (-2^{63})
Maximum	$9,223,372,036,854,775,807$ ($2^{63}-1$)
Fast (no-alloc) range	$-1,152,921,504,606,846,976$ to $1,152,921,504,606,846,975$ ($\sim \pm 2^{60}$)
Overflow behavior	Wraps (no trap)
Division	<code>int / int</code> → <code>int</code> (truncates toward zero); <code>float</code> if either operand is <code>float</code>
Power	<code>int ^^ int</code> → <code>int</code> (negative exponent truncates toward zero: $2 ^^ -1 = 0$); <code>float</code> if either operand is <code>float</code>

Integer literals:

```
42          // decimal
-10         // negative
0x1234      // hexadecimal
0b1001001   // binary
0o755       // octal
1_000_000   // underscore separator (any base)
0xFF_FF     // separators work in hex / binary / octal too
```

`_` digit separators are accepted in every base (and in a float's fraction and exponent); they are stripped before the value is parsed.

Float

Property	Value
Standard	IEEE-754 double precision (64-bit)
Precision	~15–17 significant decimal digits
Minimum positive	2.2250738585072014e−308
Maximum	1.7976931348623157e+308
Special values	Infinity, −Infinity, NaN

Display vs. round-trip: `str()` and `Console.WriteLine/Print` use a **compact** form - an integral value prints with no decimal point or exponent (`5000.0` shows as `5000`), and a fractional value uses `%g` (up to ~6 significant digits, so `0.1 + 0.2` shows as `0.3`). For a **precise, round-trippable** string (full `%.17g`) use `Convert.ToString` / `String.ToString` instead.

Float literals:

```
3.14
-0.5
1.0e6    // scientific notation
6.02e23  // exponent, optional +/- sign
1_000.5  // underscore separators (integer, fraction and exponent)
```

A digit is required on both sides of the decimal point: write `0.5` , not `.5` (a leading `.` is the member-access operator), and `5.0` , not `5.` (`5.foo` is a method/property access on `5`). Mixed int + float operations always yield `float` .

String

Property	Value
Encoding	UTF-8
Immutable	Yes - mutation creates a new string
Indexing	Byte-level (returns 1-byte string)
Max size	256 MB

UTF-8 Indexing Model

Strings store raw UTF-8 bytes. Characters may occupy 1–4 bytes. The language distinguishes:

- **Byte index** - position in the raw byte buffer. Used by `foreach` , `String.UTF8CpAt` , `String.UTF8CpNext` . Range `0..string.byteLength` .
- **Code-point index** - the Nth Unicode scalar value. Used by `String.UTF8GetCharAt` , `String.UTF8Substr` , `String.UTF8ByteIndexOf` . Range `0..String.UTF8Length(s)` .

Use `String.*` UTF-8 helpers for correct Unicode handling. Standard string indexing (`s[i]`) returns bytes, not characters.

Char

`char` is an unsigned 32-bit Unicode codepoint value. The `char()` built-in creates one from an integer codepoint.

Character escapes in a `'...'` literal denote **one code point**: `'\n'`, `'\t'`, `'\0'`, `'\xHH'` (U+00HH), `'\uXXXX'`, and `'\u{H...H}'` (up to U+10FFFF). `'\xHH'` is a code point, not a raw byte - `'\xE0'` is U+00E0 (`'à'`), not a lone byte.

Arithmetic with `char` :

- `char + char` and `char + string` **concatenate** into a `string` (each `char` contributes its UTF-8 encoding): `'a' + 'b'` is `"ab"`, and `'a' + 1` is **not** used for concatenation.
- `char + int` (and other numeric operators) treat the `char` as its integer code point: `'a' + 1` is `98`, `'z' - 'a'` is `25`.

Bool

Falsy values: `nil`, `false`, `0`, `0.0`, `""`, `'\0'`, empty arrays and empty objects. Everything else is truthy.

Type Annotations

Type annotations are optional. They produce **warnings**, not errors. All existing unannotated code continues to work.

Annotation Syntax

```
// Parameter types only
fn greet(name: string) { ... }

// Return type only
fn add(a, b): int { ... }

// Full signature
fn calculate(x: int, y: int): int { return x * y; }

// Mixed – some typed, some not
fn format(template: string, value) { return template + str(value); }
```

Type Keywords

Primitive:

Keyword	Runtime Type	Notes
<code>int</code>	Integer	Full int64; tagged immediate for $ x < 2^{60}$
<code>float</code>	Float	64-bit IEEE-754
<code>number</code>	int or float	Shorthand for <code>int float</code>
<code>string</code>	String	UTF-8 text
<code>bool</code>	Boolean	true / false
<code>char</code>	Char	32-bit codepoint

Keyword	Runtime Type	Notes
<code>nil</code>	Nil	Null / absent

Collection:

Keyword	Runtime Type
<code>array</code>	Dynamic array
<code>object</code>	Key-value hashmap

Execution:

Keyword	Runtime Type
<code>fn</code>	Callable function
<code>fiber</code>	Cooperative coroutine
<code>class</code>	Class definition
<code>instance</code>	Class instance

Special:

Keyword	Meaning
<code>any</code>	Accepts any type - disables checking for that parameter
<code>stream</code>	File / stream handle
<code>block</code>	Binary data buffer
<code>pointer</code>	Runtime pointer (FFI)

Union Types

Multiple acceptable types separated by `|`:

```
fn findUser(id: int): object|nil { ... }
fn abs(n: int|float): int|float { ... }

// 'number' is shorthand for int|float
fn abs(n: number): number { ... }
```

The checker validates that provided values match **at least one** type in the union.

Array Type Annotations

Wrap an element type in `[...]` to annotate an array variable or parameter:

```
let scores:[int] = [10, 20, 30];
let names:[string] = ["alice", "bob"];

fn sum(values:[int]): int { ... }
fn process(rows:[object]): nil { ... } // nested arrays
```

The element type is checked by the analyzer but is not tracked at runtime - the value's runtime type remains array. The `[T]` form can appear anywhere a plain type name is valid, including unions:

```
let data:[int]|nil = nil;
```

Optional Parameters

Append `?` directly after the parameter name to mark it as optional. The caller may omit trailing optional arguments; the VM automatically passes `nil` for each one not provided.

```
fn greet(name?: string)      { ... } // name is nil if omitted
fn log(msg: string, level?: int) { ... } // level is nil if omitted
fn tag(a: string, b?, c?: string) { ... } // b is any?, c is string?
```

Rules:

- `?` goes on the **name**, before the `:` - `x?`, `x?: int`, both valid. `x: int?` is not valid.
- Optional parameters must trail all required ones.
- Omitted arguments always arrive as `nil` - use `??` to substitute a default inside the function body.
- Type checking still applies for arguments that *are* provided.

```
fn greet(name?: string) {
    Console.WriteLine("Hello " + (name ?? "stranger"));
}

greet("World"); // Hello World
greet();        // Hello stranger
```

Variable Type Annotations

Form	Scope	Notes
<code>let name = expr</code>	Local (block / function)	Cannot be used at module top level
<code>var name = expr</code>	Local (block / function)	Identical to <code>let</code>
<code>const name = expr</code>	Local or global	Immutable binding
<code>global name = expr</code>	Module (global)	Accessible from all scopes in the file. Redeclaring in the same file is a compile error; redefining from another compile unit (<code>VM.Eval</code> , another module) warns and rebinds
<code>global name:type = expr</code>	Module (global)	Type-annotated global

Variables can be optionally annotated with a type at declaration time using `let name:type = value`. Once annotated, only values of that type may be assigned.

```
// Inside a function - use let/var:
let y:int = 2;
y = "test";           // Error - y declared as int

let i:instance = nil;
```

```
i = new MyClass(); // OK - instance, string, object, array may be nil

let scores:[int] = [1, 2, 3]; // array of int

// At module level - use global:
global counter:int = 0;
global config:object = { debug: false };
```

Type Cast Expression

A **type cast** is a compile-time and (for scalar types) runtime type coercion. Syntax: `(type)expr`

```
let arr = (array)Object.Clone(original);
let n   = (int)3.9;           // - 3 (truncates)
let s   = (string)42;         // - "42"
let b   = (bool)0;           // - false
```

Supported Cast Types

Type	Kind	Runtime effect
int	Value	OP_T0_INT - truncates float, converts string/bool/char
float	Value	OP_T0_FLOAT - widens int to double
char	Value	OP_T0_CHAR - converts int/string to a single character value
string	Value	OP_T0_STRING - any value - string representation
bool	Value	Double logical-NOT (! !x) - truthy - true, falsy - false
array	Hint	No bytecode emitted - analyzer type update only
object	Hint	No bytecode emitted - analyzer type update only
class	Hint	No bytecode emitted - analyzer type update only
instance	Hint	No bytecode emitted - analyzer type update only
block	Hint	No bytecode emitted - analyzer type update only
stream	Hint	No bytecode emitted - analyzer type update only

Analyzer Behaviour

- The inferred type of `(T)expr` is always `T`, regardless of what `expr` would otherwise infer to.
- If the cast is the RHS of a `let / var` declaration or assignment, the variable's tracked type is updated to `T` immediately. All subsequent uses of that variable are checked against the new type.
- No warning is emitted for type changes caused by a cast - the programmer's intent is explicit.

Compiler Behaviour

- Value casts (`int`, `float`, `char`, `string`, `bool`) compile the inner expression then emit one conversion opcode.
- Hint casts compile only the inner expression - zero extra bytecode.
- Cast nodes are fully transparent to the optimizer and inliner.

Disambiguation

The parser uses two-token lookahead to distinguish a cast from a grouped expression:

```
(int)x      // cast - type keyword followed immediately by ')'
(int + 1)   // grouped expression - not a cast
(myVar)expr // grouped expression - myVar is not a cast-type keyword
```

Only the eleven listed type names - plus the sized-int aliases `u8` `u16` `u32` `u64` `i8` `i16` `i32` `i64`, which cast exactly like `(int)` - are recognised as cast targets. All other identifiers inside `(...)` are treated as grouped expressions.

`(u8)x` cast vs `u8(x)` built-in - not the same. A sized-int *cast* `(u8)x` converts to an integer without narrowing (it behaves like `(int)`): `(u8)300` is `300`. The sized-int *conversion built-in* `u8(x)` wraps modulo the type width: `u8(300)` is `44`, and `i8(200)` is `-56`. Use the built-in when you want two's-complement wrap-around; use the cast only as a type hint.

Closures

A **closure** is an anonymous function that captures variables from its enclosing scope. Captured variables are snapshotted **by value** at the moment the inner function is created (i.e., when the enclosing function executes the expression `fn(...) { ... }`).

What can be captured

Both parameters and body-local variables of the immediately enclosing function are eligible for capture. Module-level globals are always accessible directly and are not captured - they are read live from the global environment on each access.

Capture semantics

- **Scalars** (`int`, `float`, `bool`, `string`, `char`, `nil`): snapshot copy. Later changes to the outer variable do not affect the captured value.
- **References** (`array`, `object`, `instance`): reference copy. The closure and the outer scope share the same heap object; mutations are visible on both sides.

Example

```
fn makeAdder(x) {
    return fn(n) { return x + n; };
}

let add5 = makeAdder(5);
Console.WriteLine(add5(3)); // 8
```

```
fn makeCounter() {
    let state = [0];
```



```
return fn() {  
    state[0] = state[0] + 1;  
    return state[0];  
};  
}  
  
let c = makeCounter();  
Console.WriteLine(c()); // 1  
Console.WriteLine(c()); // 2
```

Loop capture

Each loop iteration creates a new closure with its own snapshot of the loop variable:

```
fn makeFns() {  
    let funcs = [];  
    iter(i from 0 to 3) {  
        Array.Append(funcs, fn() { return i; });  
    }  
    return funcs;  
}  
  
let fs = makeFns();  
Console.WriteLine(fs[0]()); // 0  
Console.WriteLine(fs[2]()); // 2
```

Runtime representation

A closure is a cloned function object with a private environment table that holds the captured names and values. At call time, the VM sets this table as the active environment, so captured names resolve before the module globals. The `function` type applies to both plain functions and closures; `VM.GetFunctionInfo` returns the same fields for both.

Capturing this

An anonymous function created inside an instance method that references `this` captures the receiver too, so it can read/write fields (`this.field`) after the method has returned. A closure may capture at most **64** distinct outer variables.

Sealed instance fields

Instances are **sealed**: every instance field must be declared with `let` or `var` in the class body. Writing to a field that was never declared is a **compile error** (`Exception.FieldUndeclared`, code 24):

```
class Point {  
    let x = 0;
```

```

let y = 0;
fn Constructor(a, b) {
  this.x = a;      // ok - declared
  this.y = b;      // ok - declared
  this.z = 0;      // COMPILE ERROR: 'this.z' is not declared in class 'Point'
}
}

```

This catches typo'd field names at compile time and lets the VM lay out instances as fixed-size slot arrays. The full field set is the class's own declared fields plus every field inherited from a resolvable base class.

Leniency across module boundaries. If a class extends a base that is not visible in the current compilation unit (an imported base class), its inherited field set cannot be seen, so the seal is relaxed for that subclass: undeclared `this.<field>` writes are allowed and validated by the VM's runtime guard instead of at compile time.

Inheritance and `super`

A subclass reaches its base class through `super`:

- `super(args)` inside `Constructor` chains to the base constructor. It resolves statically on the defining class's base, so it works even when the subclass overrides the constructor.
- `super.method(args)` calls the base class's version of a method statically ("invokespecial"), letting an overriding method reach the one it overrides.

```

class Animal {
  let name = "";
  fn Constructor(n) { this.name = n; }
  fn describe() { return "animal " + this.name; }
}
class Dog : Animal {
  fn Constructor(n) { super(n); }           // chain to Animal.Constructor
  fn describe() { return super.describe() + " (dog)"; }
}

```

`switch` matching

`switch` compares the scrutinee against each `case` label:

- A **value** label matches by equality (`==`). Multiple labels may share a body.
- A **class** label - `case SomeClass:` - matches by type (`is` / `instanceof`, subclass-aware). Imported classes are supported: the most specific matching class wins. An identifier whose type cannot be resolved falls back to value equality (with a compiler warning).
- `default` runs when no case matches. Cases do **not** implicitly fall through to the next case's body; use `break` to leave a case early.

When every case label is a dense range of integer literals, the compiler emits a **jump table** (one $O(1)$ indexed branch) instead of a comparison chain.

```
switch (shape) {
  case Circle:      return "round";      // type match (instanceof)
  case Square: case Rect: return "cornered"; // value/type, shared body
  default:          return "unknown";
}
```

R3 - Built-in Functions

These functions are VM instructions - available everywhere, no namespace required.

Type Information

Function	Description
<code>type(x)</code>	Returns the runtime type as an <code>int</code> bitmask. Compare with <code>Type</code> module constants: <code>type(x) == Type.Int</code> .
<code>is_nil(x)</code>	Returns <code>true</code> if <code>x</code> is <code>nil</code> .
<code>is_array(x)</code>	Returns <code>true</code> if <code>x</code> is an array.
<code>is_object(x)</code>	Returns <code>true</code> if <code>x</code> is an object.

Type Conversion

Function	Converts to	Notes
<code>int(x)</code>	Integer	Truncates floats toward zero
<code>float(x)</code>	Float	Widens int to double
<code>str(x)</code>	String	Converts any value to string
<code>char(x)</code>	Char	32-bit Unicode codepoint

Bitcast / Narrowing

Function	Result type	Wraps on overflow
<code>u8(x)</code>	unsigned 8-bit (0–255)	Yes
<code>i8(x)</code>	signed 8-bit (–128–127)	Yes
<code>u16(x)</code>	unsigned 16-bit	Yes
<code>i16(x)</code>	signed 16-bit	Yes
<code>u32(x)</code>	unsigned 32-bit	Yes
<code>i32(x)</code>	signed 32-bit	Yes

```
let x = u8(300); // - 44 (wraps)
let y = i8(-200); // - 56 (wraps)
```

Length

```
len([1, 2, 3]); // 3 (array element count)
len("hello"); // 5 (byte count, not char count)
len({ a: 1, b: 2 }); // 2 (property count)
```

Guard

```
guard(expr) // Raises Exception.GuardCheck if expr evaluates to nil
```

R4 - The Exception class

`Exception` is a built-in base **class**. Every thrown value - whether from a user `throw` or an internal VM error (index-out-of-bounds, divide-by-zero, ...) - is an `Exception` (or subclass) instance.

Construct / throw:

```
throw new Exception(code, msg); // explicit
throw(code, msg);               // shorthand - desugars to new Exception(code, msg)
class NetworkError : Exception {} // subclass; new NetworkError(503, "down") inherits the
constructor
```

Constructor: `Exception(code:int, msg:string)` - sets `Code`, `Error`, and captures `StackTrace`.

Instance fields:

Field	Type	Description
<code>Code</code>	<code>int</code>	Numeric error code (see the constants below)
<code>Error</code>	<code>string</code>	Human-readable message
<code>StackTrace</code>	<code>array</code>	Call frames at throw time, as strings

Instance methods:

Method	Signature	Description
<code>ToString</code>	<code>ToString() - string</code>	"<TypeName> (code N): message" (uses the subclass name)
<code>Name</code>	<code>Name() - string</code>	The exception's type name (its class name)
<code>StackTraceString</code>	<code>StackTraceString() - string</code>	<code>StackTrace</code> joined into one block of text

Generic class introspection works too (`Class.InstanceOf(e, Exception)`, `Class.Name(e)`, ...). For type dispatch use `e is SomeError` or `switch (e) { case SomeError: ... }`. See the language guide's *Exceptions* section for full semantics.

Error-code constants

Exposed as **static members of the `Exception` class** (`Exception.RuntimeError`, ...). Each has a numeric code (stable ABI - must not be renumbered) used as the `Code` of the raised exception.

Constant	Code	Description
<code>Exception.NullPtr</code>	0	Null-pointer access
<code>Exception.DivByZero</code>	1	Division by zero
<code>Exception.ModByZero</code>	2	Modulo by zero
<code>Exception.InvalidArguments</code>	3	Invalid function arguments
<code>Exception.OutOfBounds</code>	4	Index out of bounds
<code>Exception.IOError</code>	5	I/O error (file, stream, socket)
<code>Exception.RuntimeError</code>	6	General VM runtime error

Constant	Code	Description
Exception.InvalidState	7	VM or object in invalid state
Exception.OutOfMemory	8	Memory allocation failed
Exception.InvalidMemoryAccess	9	Invalid memory access
Exception.SizeLimit	10	Container or allocation size limit exceeded
Exception.GuardCheck	11	<code>guard()</code> failed - value was nil
Exception.StackError	12	Stack overflow / underflow / corruption
(reserved)	13	Intentionally unused
Exception.UnsafeOperation	14	Unsafe operation blocked (requires <code>--unsafe</code>)
Exception.NestingError	15	Excessive recursion or nesting depth
Exception.IllegalInstruction	16	Invalid bytecode instruction
Exception.ExecOutOfMemory	17	Out of memory during execution
Exception.OutOfFibers	18	Fiber limit reached
Exception.ConstAssign	19	Attempt to assign to a constant
Exception.ChecksumError	20	FFI shared-library SHA-256 verification failed (<code>Ffi.Load</code>)
Exception.ClassNonStaticCall	21	Non-static method called without instance
Exception.TypeMismatch	22	Value incompatible with a typed-array element type
Exception.AssertionFailed	23	<code>Debug.Assert</code> / assertion failed
Exception.FieldUndeclared	24	Write to an instance field not declared with <code>let</code> / <code>var</code> (see <i>Sealed instance fields</i>)

Design notes:

- Exceptions 16, 20 indicate corrupted or hostile bytecode - treated as hard failures.
- `OutOfMemory` (8) and `ExecOutOfMemory` (17) are separate to enable more precise diagnostics.
- When an exception fires, the current fiber unwinds to the nearest enclosing error boundary (`try / catch`). If a boundary catches it, the fiber resumes there, other fibers continue, and the VM stays alive. If the exception escapes the top of the fiber with no handler, the VM prints the trace and exits the whole process with the exception's code (see R9).

R5 - Operator Precedence

Operators listed from **highest** (evaluated first) to **lowest** (evaluated last).

Level	Category	Operators	Associativity
1	Primary	literals, identifiers, <code>(...)</code>	N/A
2	Postfix	<code>obj.prop</code> <code>arr[i]</code> <code>fn(...)</code> <code>x++</code> <code>x--</code>	Left
3	Prefix Unary	<code>+x</code> <code>-x</code> <code>!x</code> <code>~x</code> <code>await</code> <code>new</code> <code>guard()</code>	Right
4	Multiplicative & Bitwise	<code>*</code> <code>/</code> <code>%</code> <code>^^</code> <code><<</code> <code>>></code> <code>&</code> <code> </code> <code>^</code>	Left
5	Additive	<code>+</code> <code>-</code>	Left
6	Comparison	<code><</code> <code><=</code> <code>></code> <code>>=</code>	Left
7	Equality	<code>==</code> <code>!=</code> <code>≈</code> <code>in</code> <code>is</code>	Left
8	Logical AND	<code>&&</code> <code>and</code>	Left
9	Logical OR	<code> </code> <code>or</code>	Left
10	Null Coalescing	<code>??</code>	Left
11	Conditional	<code>?</code> <code>:</code>	Right
12	Assignment	<code>=</code> <code>+=</code> <code>-=</code> <code>*=</code> <code>/=</code> <code>%=</code> <code>&=</code> <code> =</code> <code>^=</code> <code><<=</code> <code>>>=</code> <code>??=</code> <code>^^=</code>	Right

Key rules:

- `&&` binds tighter than `||` - `a || b && c` means `a || (b && c)`. `and` / `or` are tokenizer aliases with identical precedence.
- `&&` / `||` short-circuit and yield a **value, not always a bool**: `a && b` evaluates to `b`'s value when `a` is truthy, otherwise `false`; `a || b` evaluates to `true` when `a` is truthy, otherwise `b`'s value. (`nil` and `0` are falsy.)
- Shifts operate on `int64` and mask the shift count to `& 63`, so a count `>= 64` wraps into range. `<<` is a logical left shift; `>>` is an **arithmetic** right shift (the sign bit is replicated), so `-8 >> 1 = -4`.
- `??` (null coalescing) is lower than all arithmetic - `x + y ?? z` means `(x + y) ?? z`.
- Division `/` and power `^^` produce `int` when both operands are `int` (division truncates toward zero; a negative power truncates too, so `2 ^^ -1 = 0`), and `float` when either operand is `float`. Force a float result with a float operand: `a / (b * 1.0)` or `float(a) / b`.
- `^^` is exponentiation (`x ^^ y = x` raised to `y`).
- `+` is **overloaded by operand type** and never a type error (see also **Operator operand checking**):
 - `number + number` - arithmetic (int lane unless an operand is `float`).
 - `char + char` / `char + string` - concatenate into a `string`; `char + int` is arithmetic on the code point (`'a' + 1` is `98`).
 - `array + array` - a **new** array with the elements of both (neither operand is mutated); `array + value` appends `value` to a copy.
 - `object + object` - a **new** merged object (right operand's keys win).
 - anything else - the operands are coerced to their string forms and concatenated (`nil` renders as `"null"`, `true` / `false` as themselves), so `nil + 5` is `"null5"`.
- Comparison (`<` `<=` `>` `>=`) and equality (`==` `!=`) on **containers** are structural and recursive (bounded by the comparison-depth limit below):
 - `array` / `array` and `block` / `block` compare **lexicographically** for ordering (first differing element decides; a prefix is "less"), and element-by-element for equality. `!=` is the negation of `==`.
 - `object` / `object` and `exception` / `exception` compare equal when they hold the same keys mapping to equal values (order-independent).
 - `instance` / `instance` compare **structurally** (same class, field-by-field); `class` / `class` compare by **identity**.
 - Mixed numeric kinds compare by value across `int` / `float` / `char` / `bool` (`'A' == 65` is `true`; `1 == 1.0` is `true`). Two values of otherwise unrelated types are never equal.
- Unary `+x` is an identity operator (yields `x` unchanged); `await expr` is a prefix operator usable in expression position, whereas `yield` is a **statement** only and cannot appear inside a larger expression.
- `≈` (approximate equality) uses `APPROX_REL_EPS = 1e-2` and `APPROX_ABS_EPS = 1e-9`.

R6 - VM Limits and portability

Fibers and Scheduling

Limit	Default	Maximum	Description
Fibers	256	1024	Concurrent fibers tracked by the scheduler
Events	64	64	Event slots (I/O, async, etc.)
Timers	64	64	Active timers
Per-fiber scheduling quantum	10,000	100,000	Scheduling checkpoints (loop back-edges + call/return boundaries) per slice before auto-yield; tune with <code>Fiber.SetQuantum</code>

Call Stack and Execution

Limit	Default	Maximum	Description
Call frames	64	1,024	Call stack depth; override with <code>--frames=</code>
Nested try/catch handlers	24	24	Maximum nested try/catch handlers
Evaluation stack	4,096	65,535	VM evaluation stack size; override with <code>--stack=</code>

Locals and Arguments

Limit	Value	Description
Local variables per function	128	Maximum local variables per function
Function call arguments	16	Maximum arguments in user function calls

Compiler and Parser

Limit	Value	Description
AST nodes per compilation unit	100,000	
Parser nesting depth	256	
Literal nesting depth	64	Max container nesting when building array/object literals (raises <code>Exception.NestingError</code>)
Comparison depth	1,024	Max nesting for structural <code>==</code> /ordering of containers (raises <code>Exception.NestingError</code> ; above the JSON depth cap, so parsed documents always compare)
Clone depth	512	Max nesting for <code>Object.Clone</code> / <code>Array.Clone</code> ; a deeper or cyclic value raises <code>Exception.NestingError</code> instead of overflowing the stack (above the JSON depth cap, so parsed documents always clone)
Print depth	64	Console value printing descends this deep, then elides deeper values with <code>...</code>
<code>VM.Eval()</code> / <code>VM.Compile()</code> source	10 KB	Max source length

Bytecode and Chunks

Limit	Value	Description
Source file size	10 MB	Maximum source file size
Bytecode per function	10 MB	Maximum compiled bytecode per function
Constants per function	16,385	Maximum constant pool entries per function

Strings and Text

Limit	Value	Description
String size	256 MB	Maximum size of a single string
String split parts	100,000	Maximum parts returned by string split
OS command output	16 MB	Maximum captured stdout/stderr from OS commands

Memory Blocks and Buffers

Limit	Value	Description
Single block allocation	2 GiB	Maximum size of a single block allocation
Total block allocations	64 GiB	Total cap across all active block allocations
<code>Memory.Process()</code> input	16 MB	Max bytes per call

Collections

Limit	Value	Description
Array elements / object properties	10,000,000	Max elements in an array or properties in an object
Hash function input	1 GiB	Max input to hashing functions

Classes

Limit	Value	Description
Inheritance chain depth	8	Maximum class inheritance depth
CLI arguments	128	Maximum CLI arguments passed to the VM

Special Methods

Method	Called when	Notes
<code>Constructor(...)</code> <code>new</code> <code>ClassName(...)</code>		Optional. Return value is ignored; <code>new</code> always returns the instance.
<code>Destructor()</code>	Last reference to the instance is released	Optional. <code>this</code> is the dying instance. Exceptions are swallowed. Not inherited. Not run for instances still alive at process exit - teardown tears down the fiber that would execute it, so release anything that must be cleaned up (or do it explicitly) before returning from <code>Main</code> . Do not let <code>this</code> escape (store it in a global, array, or another object's field): the instance is already being freed, so a captured reference dangles - copy out the field values you need instead.

Float Approximation Thresholds

Threshold	Value	Description
Relative epsilon ($\approx$)	$1e-2$	Relative tolerance for the $\approx$ operator
Absolute epsilon ($\approx$)	$1e-9$	Absolute tolerance for the $\approx$ operator

Platform Portability

Flaris targets 64-bit platforms only (x86-64 and ARM64 on Linux, macOS, and Windows).

Compiled Bytecode (.flx)

`.flx` files are fully portable between all supported platforms running the same VM version. The bytecode format uses only fixed-width integer types stored little-endian throughout (header, constant pool, bytecode operands and JIT IR); no pointer-sized fields are stored on disk. Strings used more than once within a file are held once in a shared string pool and referenced by index, so a name mentioned by many functions costs a single copy on disk.

Integer Range

Small integers are stored as tagged pointers (range $\pm 2^{60}$). Integers outside this range are heap-allocated transparently - behavior is identical, only allocation strategy differs. Scripts do not need to account for this.

Pointer-width APIs

Functions that expose raw memory addresses (`Buffer.GetAddress` , `Memory.Alloc`) return 64-bit process-local addresses. Addresses stored in `.flx` constants or serialized to disk are meaningless outside the process that produced them and must never be shared or persisted.

Detection

Use `Os.Arch()` and `Os.Name()` to branch on platform at runtime (`Os.Bits()` is always `64`):

```
if (Os.Arch() == "arm64") {  
    // Apple-silicon / AArch64 path  
}
```

R7 - Performance Model

Understanding how Flaris executes helps write fast, predictable programs.

Execution Model

Flaris compiles source to bytecode, then runs it on a stack-based VM. On ARM64 (`aarch64`) and x86-64, eligible functions are additionally compiled to native machine code at first call - see [R11](#) for details. On other platforms the bytecode interpreter is the only execution tier.

Dynamic Typing and Runtime Checks

Types are checked at runtime. Invalid operations fail fast with no undefined behavior. In hot loops, avoid changing the type of a variable frequently - type stability helps the VM dispatch efficiently.

```
// Good - type-stable loop  
let sum = 0;  
for (let i = 0; i < n; i += 1) {  
    sum += arr[i];  
}
```

Memory Management

Flaris uses **deterministic reference counting**, not a garbage collector.

- Objects are freed immediately when the last reference drops.
- No stop-the-world pauses, and `Destructor` runs at a predictable point.
- Avoid creating short-lived objects inside tight loops.
- Reuse arrays and objects where possible.

Objects come from pre-allocated SLAB pools. A release that reaches zero hands the object to a deferred-free FIFO, which is drained at safe points; this keeps the release path short and bounds cascading frees. Nothing in the VM traces from roots - there is no mark-and-sweep phase and no cycle collector.

Reference cycles are never collected

Pure reference counting cannot reclaim a cycle: each object in the loop is still referenced by another, so no count ever reaches zero. The usual source is a **back-reference** - a child that points at its parent.

```
class Node {
  let _parent = nil;      // holds the parent...
  let _children:array = [];
}
// ...while the parent holds the child. Neither is ever freed.
```

Flaris has no `weak` reference, so the fix is structural. In order of preference:

1. **Do not store the parent.** Most back-references exist to read one or two values. Copy those values into the child at construction instead of the object.
2. **Let the owner pass itself in** as a method argument when the child needs it, so the reference lives only for the duration of the call.
3. **Break the cycle explicitly** before dropping the structure - assign `nil` to the back-reference in a teardown method or `Destructor`.

The same applies to any loop, not just parent/child: two objects referencing each other, a closure stored on the object it captures, or a fiber handle kept in a field of the object whose method the fiber runs.

Finding leaks

Run with `--mem`. It prints live/peak object counts and a leak dump at shutdown, and **exits non-zero when leaks are found**, so it works in CI:

```
flarism --mem myprogram.fl
```

A leak count that is stable across runs but non-zero almost always means a cycle rather than a missing release. `Debug.Refs(value)` reports an individual value's reference count, and `VM.MemoryStats()` returns the same counters programmatically for an in-test check. Note that the project's own suite runs `flarism --jit --mem`, so a program that looks clean without the flag can still fail there.

Locals vs Property Access

Local variables are the fastest storage. Property lookups require a hashmap traversal.

```
// Cache property before a hot loop
let v = obj.value;
for (...) {
  sum += v;    // fast - local
}
```

Arrays vs Objects

Structure	Access	Memory	Best for
array	O(1) index	Compact	Numeric loops, dense data
object	O(1) hash	Flexible	Structured data, named fields

Prefer arrays for anything numeric or iteration-heavy.

Built-ins vs Script Loops

Built-in functions execute closer to the VM and reduce dispatch overhead. Prefer them over equivalent hand-written loops.

```
// Better - built-in
let squares = Array.Select(xs, fn(x){ return x * x; });

// Slower - script loop
let r = [];
for (...) { Array.Append(r, x * x); }
```

Fibers and Scheduling

Fibers are cooperatively scheduled - they never run unless resumed or awaited. No data races by default.

- Batch work inside a fiber before yielding.
- Avoid spawning fibers for tiny tasks.
- Avoid `yield` inside tight numeric loops.
- Default scheduling quantum per fiber: 10,000 checkpoints (tunable via `Fiber.SetQuantum()`, up to 100,000). A checkpoint occurs at each loop back-edge and at call/return boundaries - not at every instruction.

Strings

Strings are immutable. Each concatenation allocates.

- Avoid repeated concatenation in loops.
- Accumulate data first, convert to string once.
- For large text building, collect into an array and use `String.Join()`.

Binary and Numeric Heavy Work

For heavy data processing, use `block` values and `Memory.*` functions. This minimizes object allocation and works similarly to `Span<byte>` in systems languages.

```
let buf = Buffer.Create(1024, 1);
Buffer.Fill(buf, 0);
Memory.Process(Buffer.GetAddress(buf), 1024, 1, fn(x){ return x + 1; });
```

Measure, Don't Guess

Most common bottlenecks are:

- Excessive allocations in hot paths
- Too many small fibers
- Repeated property lookups in loops

- Overuse of dynamic object structures where arrays suffice

Use `--time` (timing) and `--stats` (statistics) flags. Profile real workloads before optimizing.

Mental Model Summary

What	Cost
Local variable read/write	Very cheap
Array index access	Cheap
Property lookup	Moderate (hashmap)
Built-in function call	Cheap
Script function call	Moderate
Object/array allocation	Moderate
String concatenation	Creates new string
Fiber creation	Cheap
Fiber context switch	Cheap

R8 - Standard Library

Flaris ships 38 built-in modules covering everything from math and string processing to networking, concurrency, graphics, and low-level memory access.

Usage

All standard library functions are accessed through their namespace:

```
let sorted = Array.Sort(myArr);
let digest = Hash.Sha256("hello");
let root = Math.Sqrt(2.0);
```

No import statement is needed - all modules are always available. Two modules (`Ffi`, `Memory`) additionally require the `--unsafe` flag at runtime.

Each function table has a `JIT` column stating whether the function can be called directly from JIT-compiled code (see R11 for the legend and the surrounding rules).

Module Index

Category	Module	Description
Data	<code>Array</code>	Map, filter, sort, and aggregate arrays
	<code>Buffer</code>	Low-level fixed-size binary buffers
	<code>Collections</code>	Stack, Queue, HashMap, PriorityQueue, MaxPriorityQueue, LinkedList, Set, and OrderedMap factory functions
	<code>Object</code>	Introspect and manipulate objects and instances
	<code>Class</code>	Type-level reflection and inheritance inspection
	<code>Type</code>	Runtime type constants for use with <code>type()</code>
Strings & Text	<code>String</code>	UTF-8 string operations
	<code>Char</code>	Unicode character classification and conversion
	<code>Regex</code>	POSIX extended regular expression matching
	<code>Json</code>	JSON parse, serialize, and path query

Category	Module	Description
Math & Encoding	Convert	Type conversion across all integer widths
	Math	Full mathematical library
	Random	Seeded, deterministic PCG32 random streams
	Hash	Non-cryptographic and cryptographic hashes
	Crypto	XChaCha20-Poly1305 encryption and CSPRNG
	Compress	zlib/gzip-compatible compression
I/O & Files	Archive	ZIP archive creation and extraction
	Util	Encoding helpers and value comparison
	File	File read, write, and metadata
	Directory	Filesystem directory operations
	Path	Cross-platform path string manipulation
	Stream	Unified I/O for files, sockets, and serial ports
Networking	Console	Terminal I/O, color, and cursor control
	FileWatch	Watch file-states on files and callbacks
	HttpUtil	URL/query/header helpers and a minimal HTTP server; the HTTP <i>client</i> is the <code>Http</code> library (<code>Http.fls</code>)
	Tls	Built-in TLS byte transport (OS-native: Secure Transport / OpenSSL / SChannel)
	Net	DNS, IP conversion, and interface inspection
	Fiber	Green-thread creation, messaging, and lifecycle
Concurrency	Event	One-shot fiber synchronization (max 64 events)
	Scheduler	Low-level fiber scheduling primitives
	Timers	Fiber-based timer scheduling (max 64 timers)
	Gfx	Software rasterizer, up to 32768×32768 canvas
	OS	Process, environment, and system interface
	VM	VM introspection and control
Graphics System	Ffi	Dynamic native library loading ⚠
	Memory	Unsafe raw memory access ⚠
	Time	Unix timestamp manipulation
	Debug	Assertions, introspection, and runtime checks

⚠ = requires `--unsafe` flag (`flarism --unsafe <file>`)

Type Annotation Shorthands

Shorthand	Meaning
any	every type available
number	int or float
numeric	int, float, bool, or char
ptr	int (raw address), string, block, or pointer - anything pointer-like
string block	string or binary block
object instance class	any object-like value

Return type is shown after `-`. Optional arguments are shown as `arg?`.

Array

Namespace: `Array`

Higher-level operations on arrays: mapping, filtering, sorting, aggregation, and capacity management. Most functions create new arrays; some mutate in-place (noted in Description).

Equality: `Contains`, `IndexOf`, and `Distinct` use strict value+type equality — `2 (int)` does not match `2.0 (float)`, unlike the `==` operator and `BinarySearch / Append`, which coerce ints to float when the array is a `[float]` array.

Sort stability: `Sort / SortBy` with the *default* order are not guaranteed stable (they use `qsort`). `Sort` with a user comparator is stable when the comparator is a strict-less predicate (`a < b`); a non-strict (`a <= b`) predicate is also accepted but does not preserve the original order of equal elements.

Function	Signature	Description	JIT
All	<code>All(arr:array, func:fn) - bool</code>	<code>true</code> if all elements satisfy <code>fn(x)</code> . Stops early on first falsy.	—
Any	<code>Any(arr:array, func:fn) - bool</code>	<code>true</code> if any element satisfies <code>fn(x)</code> . Stops early.	—
Append	<code>Append(arr:array, value:any) - arr</code>	Appends <code>value</code> to end of <code>arr</code> . Mutates.	—
Average	<code>Average(arr:array) - float</code>	Returns the arithmetic mean of all elements as <code>float</code> . Returns <code>nil</code> if empty.	✓
BinarySearch	<code>BinarySearch(arr:array, value:any, less?:fn) - int</code>	Binary search over a SORTED array. Returns the index of a match, or <code>-(insertionPoint) - 1</code> when absent (so <code>-(result) - 1</code> is where it would go). Optional <code>less(a,b) - bool</code> must be the same predicate the array was sorted with. $O(\log n)$.	—
Capacity	<code>Capacity(arr:array) - int</code>	Returns current internal capacity (may exceed length).	✓
Clear	<code>Clear(arr:array) - bool</code>	Removes all elements; capacity kept. Mutates.	—
Concat	<code>Concat(a:array, b:array) - array</code>	Returns new array: all of <code>a</code> followed by all of <code>b</code> .	—
Contains	<code>Contains(arr:array, value:any) - bool</code>	<code>true</code> if any element equals <code>value</code> (deep equality).	✓
Create	<code>Create(size:int, type?:int) - array</code>	Creates new array with <code>size</code> pre-allocated slots. Optional <code>type</code> (a <code>Type.* constant</code>) pre-fills all slots with the zero value for that type - <code>Type.Float</code> gives a <code>[float]</code> array of 0.0 values; <code>Type.Int</code> gives a <code>[int]</code> array of 0 values; <code>Type.String</code> gives a <code>[string]</code> array of "" values. Without a type argument all slots are <code>nil</code> .	✓
Distinct	<code>Distinct(arr:array) - array</code>	Returns new array with only first occurrence of each value.	—
Clone	<code>Clone(arr:array) - array</code>	Returns a shallow copy: a new array with the same elements (element refs shared). Preserves a typed array's element type.	—
Count	<code>Count(arr:array, pred:fn) - int</code>	Number of elements for which <code>pred(x)</code> is truthy.	—
Fill	<code>Fill(arr:array, value:any, start?:int, end?:int) - array</code>	Sets elements in <code>[start, end)</code> to <code>value</code> in-place (whole array by default). Negative <code>start / end</code> count from the end. Returns same array. Mutates.	—
Find	<code>Find(arr:array, pred:fn) - any</code>	First element for which <code>pred(x)</code> is truthy, or <code>nil</code> .	—
FindIndex	<code>FindIndex(arr:array, pred:fn) - int</code>	First index for which <code>pred(x)</code> is truthy, or <code>-1</code> .	—
FindLast	<code>FindLast(arr:array, pred:fn) - any</code>	Last element for which <code>pred(x)</code> is truthy, or <code>nil</code> .	—
FindLastIndex	<code>FindLastIndex(arr:array, pred:fn) - int</code>	Last index for which <code>pred(x)</code> is truthy, or <code>-1</code> .	—
ForEach	<code>ForEach(arr:array, func:fn) - array</code>	Calls <code>fn(x)</code> for every element; callback return value is discarded. Returns the original array.	—
First	<code>First(arr:array) - any</code>	Returns first element or <code>nil</code> if empty.	—

Function	Signature	Description	JIT
Flatten	<code>Flatten(arr:array) - array</code>	Returns new array with one level of nested arrays collapsed. Non-array elements are kept as-is.	—
GroupBy	<code>GroupBy(arr:array, func:fn) - object</code>	Groups elements by key. Calls <code>fn(x)</code> for each element; the return value becomes the group key (coerced to string). Returns an object where each key maps to an array of the elements that produced it. Empty array returns <code>{}</code> .	—
Chunk	<code>Chunk(arr:array, size:int) - array</code>	Splits <code>arr</code> into sub-arrays of at most <code>size</code> elements. The last chunk may be smaller. Returns a new array of arrays.	—
CountBy	<code>CountBy(arr:array, func:fn) - object</code>	Like <code>GroupBy</code> but returns <code>{key: count}</code> instead of <code>{key: [elements]}</code> . Calls <code>fn(x)</code> for each element; result is coerced to string as the key.	—
FlatMap	<code>FlatMap(arr:array, func:fn) - array</code>	Calls <code>fn(x)</code> for each element; if the result is an array its elements are spread into the output, otherwise the result itself is appended. One level of flattening.	—
IndexOf	<code>IndexOf(arr:array, value:any) - int</code>	Returns first index of <code>value</code> or <code>-1</code> if not found.	✓
InsertAt	<code>InsertAt(arr:array, index:int, value:any) - arr</code>	Inserts <code>value</code> at <code>index</code> , shifts remaining right; returns the array (<code>nil</code> on out-of-range, or raises on a typed-array element-type mismatch). Mutates.	—
Last	<code>Last(arr:array) - any</code>	Returns last element or <code>nil</code> if empty.	—
Process	<code>Process(arr:array, func:fn) - array</code>	Applies <code>fn(value:any)</code> to every element in-place, replacing each element with the return value. Returns same array. Mutates.	—
ProcessIdx	<code>ProcessIdx(arr:array, func:fn) - array</code>	Applies <code>fn(value:any, index:int)</code> to every element in-place, replacing each element with the return value. Returns same array. Mutates.	—
IsAllInt	<code>IsAllInt(arr:array) - bool</code>	<code>true</code> if every element is an integer (an empty array is <code>true</code>). Use before <code>ProcessCallback</code> / <code>ProcessEvent</code> with an <code>int</code> kernel.	—
IsAllFloat	<code>IsAllFloat(arr:array) - bool</code>	<code>true</code> if every element is a float (an empty array is <code>true</code>). Use before <code>ProcessCallback</code> / <code>ProcessEvent</code> with a <code>float</code> kernel.	—
ProcessCallback	<code>ProcessCallback(fn:function, arr:array, len:int, cb:function) - nil</code>	Processes a scalar array with your worker function <code>fn(arr, len)</code> , then calls <code>cb(arr, result)</code> when finished. Runs on another CPU core when <code>fn</code> is simple enough (only touches <code>arr</code> , plain number math, no calls or allocations) — otherwise runs normally, same result. The array must be all- <code>int</code> , <code>float</code> , <code>char</code> , or <code>bool</code> , and match the kernel's parameter type (<code>fn(arr:[int],...)</code> needs an <code>int</code> array, <code>fn(arr:[float],...)</code> a <code>float</code> array, etc.); anything else (mixed, or holding strings/arrays/objects) returns <code>nil</code> . See <code>Buffer.ProcessCallback</code> .	✓
ProcessEvent	<code>ProcessEvent(fn:function, arr:array, len:int, eventId:int) - nil</code>	Like <code>ProcessCallback</code> , but signals event <code>eventId</code> with the result instead of calling a callback. Wait for several with <code>Event.WaitFor([ids])</code> .	✓
Reduce	<code>Reduce(arr:array, func:fn, initial?:any) - any</code>	Aggregates left-to-right: <code>acc = fn(acc, x)</code> starting from <code>initial</code> . Without <code>initial</code> , seeds with the first element (empty array → <code>nil</code>). <code>acc</code> / result may be any type the callback returns.	—
ReduceRight	<code>ReduceRight(arr:array, func:fn, initial?:any) - any</code>	Like <code>Reduce</code> but folds right-to-left (last element first). Without <code>initial</code> , seeds with the last element (empty array → <code>nil</code>).	—

Function	Signature	Description	JIT
RemoveAt	<code>RemoveAt(arr:array, index:int) - bool</code>	Removes element at <code>index</code> , shifts remaining left. Mutates.	—
Removelf	<code>RemoveIf(arr:array, pred:fn) - int</code>	Removes in place every element for which <code>pred(x)</code> is truthy (compacting survivors) and returns the number removed. Mutates.	—
Reserve	<code>Reserve(arr:array, capacity:int) - bool</code>	Ensures <code>capacity ≥ capacity</code> . Does not affect length. Mutates capacity.	—
Reverse	<code>Reverse(arr:array) - array</code>	Reverses elements in-place. Returns same array. Mutates.	—
Select	<code>Select(arr:array, func:fn) - array</code>	Returns new array by applying <code>fn(x)</code> to every element (map).	—
Shuffle	<code>Shuffle(arr:array) - array</code>	Randomly shuffles elements in-place (Fisher-Yates, CSPRNG). Returns same array. Mutates.	—
ShrinkToFit	<code>ShrinkToFit(arr:array) - bool</code>	Reduces capacity to match current length. Mutates capacity.	—
Skip	<code>Skip(arr:array, count:int) - array</code>	Returns new array with first <code>count</code> elements removed (a non-positive <code>count</code> keeps all).	—
Slice	<code>Slice(arr:array, start:int, end?:int) - array</code>	Returns a new array with elements <code>[start, end)</code> (to the end by default). Negative indices count from the end (Python-style); both are clamped, an empty range yields <code>[]</code> .	—
SequenceEqual	<code>SequenceEqual(a:array, b:array) - bool</code>	<code>true</code> if <code>a</code> and <code>b</code> have the same length and equal elements in order (strict value+type equality: <code>2 ≠ 2.0</code>).	—
MaxBy	<code>MaxBy(arr:array, func:fn) - any</code>	Returns the element for which <code>fn(x)</code> produces the largest key. Returns <code>nil</code> if <code>arr</code> is empty. Builtin functions (e.g. <code>String.Length</code>) use the fast path.	—
MinBy	<code>MinBy(arr:array, func:fn) - any</code>	Returns the element for which <code>fn(x)</code> produces the smallest key. Returns <code>nil</code> if <code>arr</code> is empty. Builtin functions use the fast path.	—
Partition	<code>Partition(arr:array, func:fn) - array</code>	Returns a two-element array <code>[matches, rest]</code> where <code>matches</code> contains elements for which <code>fn(x)</code> is truthy and <code>rest</code> contains the remainder.	—
Sort	<code>Sort(arr:array, fn?:fn) - array</code>	Sorts in-place. Optional comparator <code>fn(a,b)-bool</code> . Mutates.	—
SortedInsert	<code>SortedInsert(arr:array, value:any, less?:fn) - int</code>	Inserts <code>value</code> into a SORTED array keeping it sorted (after any equal run). Returns the insertion index, or <code>nil</code> on failure. Same optional <code>less</code> predicate as <code>Sort / BinarySearch</code> . Mutates.	—
SortBy	<code>SortBy(arr:array, func:fn) - array</code>	Sorts in-place by extracted key: calls <code>fn(x)</code> once per element and compares the results. Builtin functions (e.g. <code>String.Length</code> , <code>Type.Of</code>) use the fast path. Mutates.	—
Tally	<code>Tally(arr:array) - object</code>	Counts occurrences of each element (converted to string). Returns <code>{value: count}</code> .	—
UniqueBy	<code>UniqueBy(arr:array, func:fn) - array</code>	Returns a new array keeping only the first element for each distinct value of <code>fn(x)</code> . Preserves order.	—
Subset	<code>Subset(arr:array, from:int, count:int) - array</code>	Returns new array of <code>count</code> elements starting at index <code>from</code> .	—
Swap	<code>Swap(arr:array, i:int, j:int) - arr</code>	Swaps elements at indices <code>i</code> and <code>j</code> ; returns the array (<code>nil</code> if either index is out of range). Mutates.	—
Take	<code>Take(arr:array, count:int) - array</code>	Returns new array of first <code>count</code> elements.	—
Where	<code>Where(arr:array, func:fn) - array</code>	Returns new array of elements for which <code>fn(x)</code> is truthy (filter).	—
Zip	<code>Zip(a:array, b:array) - array</code>	Returns array of <code>[a[i], b[i]]</code> pairs. Length is <code>min(len(a), len(b))</code> .	—

Array instance method syntax

All `Array` functions where the array is the first argument can be called directly on an array value:

```
var a: array = [3, 1, 4, 1, 5];
a = a.Append(9)           // same as Array.Append(a, 9)
a = a.Sort()              // same as Array.Sort(a)
a.First()                 // 3
a.Last()                  // 9
a.Contains(4)              // true
a.Average()               // arithmetic mean (Sum/Min/Max live in the Math module)
a.Select(fn(x) { return x * 2; })
a.Where(fn(x) { return x > 3; })
a.Reverse()
```

Functions that do not take an array as their first argument (`Create` , and similar factories) are not available as instance methods. Both forms compile to identical bytecode when the variable is typed (`: array` or inferred). Untyped variables fall back to a runtime dispatch.

Buffer

Namespace: `Buffer`

Low-level binary memory operations. A buffer is `count` elements $\times$ `size` bytes per element.

Function	Signature	Description	JIT
ChangeToString	<code>ChangeToString(buf: block) - bool</code>	Reinterprets buffer in-place as <code>string</code> . Destructive - original block handle is invalid after call.	—
Copy	<code>Copy(src: block, srcOff: int, dst: block, dstOff: int, len: int) - bool</code>	Copies <code>len</code> bytes from <code>src</code> to <code>dst</code> . Offsets in bytes. Bounds-checked.	✓
CopyBytesToArray	<code>CopyBytesToArray(buf: block, arr: array, offset: int) - bool</code>	Writes <code>len(buf)</code> raw bytes from <code>buf</code> into <code>arr</code> starting at <code>offset</code> . Existing elements are updated in-place; elements past the current array length are appended. Optional <code>offset</code> defaults to 0.	—
CopyStringAt	<code>CopyStringAt(buf: block, offset: int, s: string, n: int) - bool</code>	Copies <code>n</code> bytes of <code>s</code> into <code>buf</code> at byte <code>offset</code> . <code>n</code> defaults to <code>len(s)</code> . Bounds-checked against the full byte extent (<code>count $\times$ size</code>); returns <code>false</code> if <code>offset + n</code> exceeds it, or on a non-block/non-string.	✓
WriteVarintAt	<code>WriteVarintAt(buf: block, offset: int, value: int) - int</code>	Writes <code>value</code> as an unsigned LEB128 varint (protobuf-style: 7 bits/byte, high bit = continuation, max 10 bytes). Returns the byte count written, or 0 if it would not fit. Zigzag-encode signed values first: <code>(n << 1) ^ (n >> 63)</code> .	✓
ReadVarintAt	<code>ReadVarintAt(buf: block, offset: int) - array</code>	Reads an unsigned LEB128 varint at <code>offset</code> . Returns <code>[value, bytesRead]</code> ; <code>bytesRead</code> is 0 on out-of-range or malformed/truncated input (never reads past the block).	—

Function	Signature	Description	JIT
Create	Create(count:int, size:int) – block	Allocates new buffer: count elements, size bytes each (size in 1–255). nil if count ≤ 0 or size is outside 1–255.	✓ owned ¹
Fill	Fill(buf:block, value:int) – bool	Fills entire buffer with byte value (0–255). Mutates.	✓
ProcessCallback	ProcessCallback(fn:function, buf:block, len:int, blocksize:int, cb:function) – nil	Processes buf with your worker function fn(buf, len, blocksize), then calls cb(buf, result) when finished. Returns immediately. Runs on another CPU core when fn is simple enough (see the note below), otherwise runs normally — same result either way.	✓
ProcessEvent	ProcessEvent(fn:function, buf:block, len:int, blocksize:int, eventId:int) – nil	Like ProcessCallback, but signals event with the result instead of calling a callback. Start several and wait for them all with Event.WaitFor([ids]). See the note below.	✓
FromArray	FromArray(buf:block, arr:array) – bool	Writes integer array into buffer elements. Element size must be 1, 2, 4, or 8. Mutates.	✓
FromString	FromString(s:string) – block	Copies the raw bytes of s into a new block (count = len(s), size = 1). Non-destructive - both string and block remain valid.	✓ owned ¹
GetAddress	GetAddress(buf:block) – int	Returns raw memory address as integer (for FFI/Memory API use). Requires --unsafe (unsafe mode); returns raw pointers, so it is gated like the Memory.* API.	✓
FillRange	FillRange(buf:block, offset:int, length:int, value:int) – bool	Sets length bytes starting at byte offset to byte value (0–255). Bounds-checked against the full backing span (count × size). false on a bad value or out-of-range span.	✓
IndexOf	IndexOf(buf:block, needle:int string, start?:int) – int	Byte offset of the first match at or after start (default 0), or -1. needle is a byte (int 0–255) or a byte substring (string; empty string matches at start).	✓
Compare	Compare(a:block, b:block) – int	memcmp -style ordering (-1 / 0 / 1) over the two backing byte spans; the shorter buffer sorts first when one is a prefix of the other. nil if either argument is not a block.	✓
Equals	Equals(a:block, b:block) – bool	true if both backing spans are the same byte length and byte-for-byte equal.	✓
ReadI8At	ReadI8At(buf:block, offset:int) – int	Reads a sign-extended signed byte at offset. 0 if out of bounds.	✓
ReadI16At	ReadI16At(buf:block, offset:int, bigEndian?:bool) – int	Reads a sign-extended signed 2-byte integer at offset. Optional bigEndian defaults to false. 0 if out of bounds.	✓
ReadI32At	ReadI32At(buf:block, offset:int, bigEndian?:bool) – int	Reads a sign-extended signed 4-byte integer at offset. Optional bigEndian defaults to false. 0 if out of bounds.	✓
ReadFloat32At	ReadFloat32At(buf:block, offset:int, bigEndian?:bool) – float	Reads an IEEE 32-bit float at byte offset. Optional bigEndian defaults to false. 0.0 if out of bounds.	✓
ReadFloat64At	ReadFloat64At(buf:block, offset:int, bigEndian?:bool) – float	Reads an IEEE 64-bit double at byte offset. Optional bigEndian defaults to false. 0.0 if out of bounds.	✓

Function	Signature	Description	JIT
WriteFloat32At	<code>WriteFloat32At(buf:block, offset:int, value:float, bigEndian?:bool) - bool</code>	Stores an IEEE 32-bit float at byte <code>offset</code> . Optional <code>bigEndian</code> defaults to <code>false</code> . <code>false</code> if out of bounds.	✓
WriteFloat64At	<code>WriteFloat64At(buf:block, offset:int, value:float, bigEndian?:bool) - bool</code>	Stores an IEEE 64-bit double at byte <code>offset</code> . Optional <code>bigEndian</code> defaults to <code>false</code> . <code>false</code> if out of bounds.	✓
ReadU8At	<code>ReadU8At(buf:block, offset:int) - int</code>	Reads one byte at <code>offset</code> . Returns 0 if out of bounds.	✓
ReadU16At	<code>ReadU16At(buf:block, offset:int, bigEndian:bool) - int</code>	Reads 2-byte unsigned integer at <code>offset</code> . Optional <code>bigEndian</code> defaults to <code>false</code> (little-endian). Returns 0 if out of bounds.	✓
ReadU32At	<code>ReadU32At(buf:block, offset:int, bigEndian:bool) - int</code>	Reads 4-byte unsigned integer at <code>offset</code> . Optional <code>bigEndian</code> defaults to <code>false</code> . Returns 0 if out of bounds.	✓
ReadU64At	<code>ReadU64At(buf:block, offset:int, bigEndian:bool) - int</code>	Reads 8-byte integer at <code>offset</code> . Optional <code>bigEndian</code> defaults to <code>false</code> . Returns 0 if out of bounds.	✓
Reserve	<code>Reserve(buf:block, capacity:int) - bool</code>	Ensures <code>buf</code> holds at least <code>capacity</code> elements (<code>capacity × size</code> bytes). Grows the buffer in-place via <code>realloc</code> if needed; the block object is mutated and <code>len(buf)</code> reflects the new element count. No-op if already large enough; <code>false</code> on a non-block or zero capacity.	—
Slice	<code>Slice(buf:block, offset:int, count:int) - block</code>	Returns new buffer: <code>count</code> elements starting at element index <code>offset</code> .	✓ owned ¹
StringToArray	<code>StringToArray(s:string, arr:array, offset:int) - int</code>	Writes the raw UTF-8 bytes of <code>s</code> directly into <code>arr</code> at <code>offset</code> without allocating an intermediate block. Returns the number of bytes written (<code>len(s)</code>). Optional <code>offset</code> defaults to 0.	—
ToArray	<code>ToArray(buf:block, offset:int, count:int) - object</code>	Returns object with array of integers from buffer elements. Element size must be 1, 2, 4, or 8.	—
ToString	<code>ToString(buf:block) - string</code>	Copies the raw bytes of <code>buf</code> into a new string. Non-destructive - block remains valid. Total bytes read = <code>count × size</code> .	✓ owned ¹
WriteU8At	<code>WriteU8At(buf:block, offset:int, v:int) - bool</code>	Writes <code>v & 0xFF</code> at byte <code>offset</code> . Returns <code>false</code> if out of bounds.	✓
WriteU16At	<code>WriteU16At(buf:block, offset:int, v:int, bigEndian:bool) - bool</code>	Writes 2-byte integer at <code>offset</code> . Optional <code>bigEndian</code> defaults to <code>false</code> . Returns <code>false</code> if out of bounds.	✓
WriteU32At	<code>WriteU32At(buf:block, offset:int, v:int, bigEndian:bool) - bool</code>	Writes 4-byte integer at <code>offset</code> . Optional <code>bigEndian</code> defaults to <code>false</code> . Returns <code>false</code> if out of bounds.	✓
WriteU64At	<code>WriteU64At(buf:block, offset:int, v:int, bigEndian:bool) - bool</code>	Writes 8-byte integer at <code>offset</code> . Optional <code>bigEndian</code> defaults to <code>false</code> . Returns <code>false</code> if out of bounds.	✓
Fill16	<code>Fill16(buf:block, value:int, bigEndian?:bool) - int</code>	Fills the whole buffer with the 2-byte encoding of <code>value</code> repeated. Returns the number of 2-byte units written; a trailing partial unit is left unchanged. Optional <code>bigEndian</code> defaults to <code>false</code> .	✓
Fill32	<code>Fill32(buf:block, value:int, bigEndian?:bool) - int</code>	Like <code>Fill16</code> with a 4-byte value.	✓

Function	Signature	Description	JIT
Fill64	<code>Fill64(buf: block, value: int, bigEndian?: bool) - int</code>	Like <code>Fill16</code> with an 8-byte value.	✓
FillPattern	<code>FillPattern(buf: block, pattern: block) - int</code>	Tiles <code>pattern</code> 's bytes across the whole buffer extent (the last copy is truncated to fit). Returns the number of bytes written (the buffer extent), or <code>0</code> on a non-block or empty pattern.	✓
Reverse	<code>Reverse(buf: block) - bool</code>	Reverses the buffer's elements in place (each element is <code>size</code> bytes; a byte buffer is reversed byte-for-byte). <code>false</code> on a non-block.	✓
CopyWithin	<code>CopyWithin(buf: block, destOff: int, srcOff: int, length: int) - bool</code>	Moves <code>length</code> bytes within the same buffer from byte <code>srcOff</code> to byte <code>destOff</code> (memmove; the spans may overlap). Bounds-checked against the byte extent; <code>false</code> if out of range.	✓
Concat	<code>Concat(a: block, b: block) - block</code>	Returns a new byte block (element size 1) holding <code>a</code> 's bytes followed by <code>b</code> 's bytes (each buffer's full extent). <code>owned¹</code> <code>nil</code> on a non-block.	—
Append	<code>Append(buf: block, src: block string) - int</code>	Grows <code>buf</code> in place and appends <code>src</code> 's bytes. Returns the new element count, or <code>0</code> on failure (non-block/-src, a byte count that is not a multiple of the element size, or a failed grow — on failure <code>buf</code> is left unchanged).	—

- **ProcessCallback / ProcessEvent** — process a buffer, optionally on another CPU core
 - You give it a **worker function** `fn(buf, len, blocksize)` that reads and writes the buffer, plus a **cb(buf, result)** that runs when the work is finished. `ProcessCallback` returns right away; `cb` is called later. (`ProcessEvent` is the same but signals an event instead of calling `cb` — see below.)
 - `len` and `blocksize` are just two numbers passed straight through to your worker function for it to use as it likes (typically the element count and the element bit-width, 8/16/32/64). The call does not split or reinterpret the buffer for you.
 - Your worker function returns an `int` or a `float`; that value arrives as the second argument of `cb(buf, result)`. If the work hits an error (e.g. an out-of-range index), `result` is `nil` — it never crashes your program.
 - **To make it run on another core (with `--jit`), keep the worker function simple:** it must only read/write the buffer you passed and do plain number math — **no** creating strings/arrays/objects and **no** calling other functions from inside it. Its first parameter must be the buffer and any others must be `int`. A function that follows these rules runs in parallel; any other function still works correctly but runs normally (in line), so you get the same result either way.
 - **cb has no such restriction** — it is an ordinary function. Read the finished buffer, allocate, call anything.
 - **To actually run work in parallel, start several calls at once** over *different* buffers. Each returns immediately, so the jobs run side by side. A single call processes one buffer; it does not split one buffer across cores.
 - `ProcessEvent(fn, buf, len, blocksize, eventId)` delivers the result into an event instead of calling a callback. Start several jobs, each with its own event id, then wait for them all in one line with `Event.WaitFor([ids])` — it blocks until every job is done and returns the results as an array (see the example). Handy when you want to fan out and then continue once everything finishes.

```
// Process 4 buffers in parallel, then wait for all of them in one line.
fn kernel(b:block, len:int, _bs:int): int {
    var sum = 0; var i = 0;
    while (i < len) { b[i] = (b[i] * 2 + 1) % 256; sum = sum + b[i]; i = i + 1; }
    return sum;
}

fn async Main() {
    let bufs = [Buffer.Create(1024,1), Buffer.Create(1024,1),
                Buffer.Create(1024,1), Buffer.Create(1024,1)];
    let ids = [Event.Create(), Event.Create(), Event.Create(), Event.Create()];
    for (let k = 0; k < 4; k++) { Buffer.ProcessEvent(kernel, bufs[k], 1024, 8, ids[k]); }
    let checksums = Event.WaitFor(ids); // blocks until all four events are set
}
```

```
// Process 4 buffers in parallel, using a counter to tell when all are done.
global done = 0;
fn kernel(b:block, len:int, _bs:int): int {
    var sum = 0; var i = 0;
    while (i < len) { b[i] = (b[i] * 2 + 1) % 256; sum = sum + b[i]; i = i + 1; }
    return sum; // delivered to cb as `result`
}

fn onDone(b:block, result:int) { done = done + 1; } // called when a job finishes

fn async Main() {
    let bufs = [Buffer.Create(1024,1), Buffer.Create(1024,1),
                Buffer.Create(1024,1), Buffer.Create(1024,1)];
    for (let k = 0; k < 4; k++) { Buffer.ProcessCallback(kernel, bufs[k], 1024, 8, onDone); }
    while (done < 4) { Fiber.Sleep(1); } // wait for all 4 callbacks
}
```

• FromString / ToString

- Both are non-destructive copies - unlike `ChangeToString`, neither modifies the source object.
- `ToString` reads `count × size` bytes, so a buffer created with `Buffer.Create(n, 4)` will produce a string of `n × 4` bytes. For raw byte buffers intended for string conversion, use `size = 1`.
- Primary use case is passing string data to FFI functions and reading string results back:

```
let buf = Buffer.FromString(http_response.body);
let result = my_ffi_fn(buf);
Console.WriteLine(Buffer.ToString(result));
```

• CopyBytesToArray / StringToArray

- These avoid the round-trip overhead of `ToArray` + a Flaris loop when writing into an existing byte array.
- `StringToArray` reads directly from the string's internal bytes - no block is allocated.

- Both are zero-copy for the destination: existing array slots (0–255 byte values) are updated in-place without heap allocation.

```
// Write UTF-8 bytes of a string into a pre-allocated byte array at offset 10
let n:int = Buffer.StringToArray(s, my_arr, 10);

// Copy raw bytes from a block into an array
let blk:block = Buffer.FromString(s);
Buffer.CopyBytesToArray(blk, my_arr, 0);
```

Class

Namespace: **Class**

Type-level introspection, inheritance inspection, and dynamic reflection on user-defined classes.

Function	Signature	Description	JIT
InstanceOf	<code>InstanceOf(inst:instance, cls:class) - bool</code>	true if instance was created from <code>class</code> or any base class.	✓
SubclassOf	<code>SubclassOf(A:class, B:class) - bool</code>	true if A is B or A descends from B. Siblings (a shared parent) and same-named unrelated classes are not subclasses.	✓
IsClass	<code>IsClass(value:any) - bool</code>	true if value is a class object.	✓
IsInstance	<code>IsInstance(value:any) - bool</code>	true if value is a class instance.	✓
Of	<code>Of(obj:instance class) - class</code>	Runtime class of an instance; identity for a class. The unambiguous "get the type" primitive.	—
GetBase	<code>GetBase(obj:instance class) - class</code>	For a class , its immediate base (or nil). For an instance , its own class — use <code>Of</code> when you want the type unambiguously.	—
GetName	<code>GetName(obj:instance class) - string</code>	Class name string, or nil for other types.	—
Name	<code>Name(obj:instance class) - string</code>	Alias of <code>GetName</code> . Allocation-free (returns the existing class-name string).	—
Fields	<code>Fields(inst:instance) - array</code>	Field name strings on an instance, inherited included.	—
InstanceMethods	<code>InstanceMethods(cls:class) - array</code>	Instance method names declared on the class (not inherited).	—
StaticMethods	<code>StaticMethods(cls:class) - array</code>	Static method names declared on the class (not inherited).	—
AllMethods	<code>AllMethods(obj:instance class) - array</code>	Every resolved method name, inherited included (overrides appear once).	—
HasMethod	<code>HasMethod(target:instance class, name:string) - bool</code>	true if <code>name</code> resolves to a method (inherited included).	—
GetMethod	<code>GetMethod(target:instance class, name:string) - any</code>	Callable handle for method <code>name</code> : a bound method for an instance method, the raw function for a static; nil if absent or cross-context.	—
GetField	<code>GetField(inst:instance, name:string) - any</code>	Value of the declared field <code>name</code> (inherited included), or nil. Methods/consts are not fields.	—
SetField	<code>SetField(inst:instance, name:string, value:any) - bool</code>	Writes a declared field; returns <code>false</code> for an undeclared name or frozen instance. Never fabricates a field.	—

Function	Signature	Description	JIT
Invoke	<code>Invoke(target:instance class, name:string, ...args)</code> – any	Dynamically calls method <code>name</code> on instance or class.	—
Instantiate	<code>Instantiate(name:string)</code> – instance	New instance of the named global class without running its Constructor (fields get declared defaults). Nil for an unknown or non-class name.	—
New	<code>New(target:class string, ...args)</code> – instance	New instance of a class (object or global name) with its Constructor run against <code>args</code> . Nil for an unknown/non-class target.	—

Collections

Namespace: **Collections**

Factory functions for Stack, Queue, HashMap, PriorityQueue, MaxPriorityQueue, LinkedList, Set, and OrderedMap. All factories return `object`.

Stack - `Collections.Stack()` – `object`

Method	Description
<code>s.Push(v)</code>	Push value to top.
<code>s.Pop()</code>	Remove and return top, or <code>nil</code> if empty.
<code>s.Peek()</code>	Return top without removing, or <code>nil</code> if empty.
<code>s.IsEmpty()</code>	<code>true</code> if stack has no elements.
<code>s.Size()</code>	Number of elements.
<code>s.Clear()</code>	Remove all elements.
<code>s.ToArray()</code>	Return current elements as array (bottom-to-top).

Queue - `Collections.Queue()` – `object`

Method	Description
<code>q.Enqueue(v)</code>	Add value to back.
<code>q.Dequeue()</code>	Remove and return front, or <code>nil</code> if empty.
<code>q.Peek()</code>	Return front without removing, or <code>nil</code> if empty.
<code>q.IsEmpty()</code>	<code>true</code> if queue has no elements.
<code>q.Size()</code>	Number of elements.
<code>q.Clear()</code>	Remove all elements.
<code>q.ToArray()</code>	Return current elements as array (front-to-back).

HashMap - `Collections.HashMap(capacity?:int)` – `object`

Keys may be any hashable value (int, char, string, float, ...); lookup matches by hash, which is **by value** for numbers (int, char, bool, and float - equal values, however computed, hit the same entry) and for strings. Other heap objects (arrays, instances, ...) hash by identity, so only the same object matches. Pass an optional `capacity` to pre-size the map when the final key count is known, avoiding intermediate regrows. Iteration order (`Keys` / `Values`) is **not** stable across processes.

Method	Description
<code>m.Set(key, value)</code>	Set key to value.
<code>m.Get(key, default?)</code>	Return value for key, else <code>default</code> (if given) or <code>nil</code> .

Method	Description
<code>m.Has(key)</code>	<code>true</code> if key exists.
<code>m.Delete(key)</code>	Remove key; return <code>true</code> if it existed.
<code>m.Clear()</code>	Remove all entries.
<code>m.Size()</code>	Number of entries.
<code>m.Keys()</code>	Return array of all keys (unspecified order).
<code>m.Values()</code>	Return array of all values (unspecified order).

PriorityQueue - `Collections.PriorityQueue()` - object

Min-heap: `Dequeue()` always returns the smallest value.

Method	Description
<code>q.Enqueue(v, priority)</code>	Insert value with numeric priority.
<code>q.Dequeue()</code>	Remove and return the value with the lowest priority, or <code>nil</code> .
<code>q.Peek()</code>	Return the lowest-priority value without removing, or <code>nil</code> .
<code>q.IsEmpty()</code>	<code>true</code> if queue has no elements.
<code>q.Size()</code>	Number of elements.
<code>q.Clear()</code>	Remove all elements.
<code>q.ToArray()</code>	Return values as array ordered lowest-to-highest priority.

MaxPriorityQueue - `Collections.MaxPriorityQueue()` - object

Max-heap: `Dequeue()` always returns the largest value. Same API as `PriorityQueue`.

LinkedList - `Collections.LinkedList()` - object

Singly-linked list with a cached tail pointer. `PushFront`, `PushBack`, `PopFront`, `PeekFront`, and `PeekBack` are $O(1)$. `PopBack` and `Clear` are $O(n)$.

Method	Description
<code>l.PushFront(v)</code>	Insert at head. Returns <code>self</code> for chaining.
<code>l.PushBack(v)</code>	Insert at tail. Returns <code>self</code> for chaining.
<code>l.PopFront()</code>	Remove and return head value, or <code>nil</code> if empty.
<code>l.PopBack()</code>	Remove and return tail value, or <code>nil</code> if empty. $O(n)$.
<code>l.PeekFront()</code>	Return head value without removing, or <code>nil</code> .
<code>l.PeekBack()</code>	Return tail value without removing, or <code>nil</code> .
<code>l.IsEmpty()</code>	<code>true</code> if list has no elements.
<code>l.Size()</code>	Number of elements.
<code>l.Clear()</code>	Remove all elements. $O(n)$.
<code>l.ToArray()</code>	Return elements as array (front-to-back). $O(n)$.

Set - `Collections.Set(arr?:array)` - object

Hash set: `Add`, `Has`, and `Remove` are one hash lookup each. Optionally seeded with the elements of an array (duplicates collapse). Membership follows the value hash: ints, chars, and strings compare by value; floats and other heap values by identity - use `int/char/string` elements for value semantics.

Method	Description
<code>s.Add(v)</code>	Insert <code>v</code> . Returns <code>true</code> if newly added, <code>false</code> if already present.
<code>s.Has(v)</code>	<code>true</code> if <code>v</code> is a member.

Method	Description
<code>s.Remove(v)</code>	Remove <code>v</code> ; return <code>true</code> if it was present.
<code>s.Union(other)</code>	New Set with every element of both sets.
<code>s.Intersect(other)</code>	New Set with the elements present in both sets.
<code>s.Difference(other)</code>	New Set with the elements of <code>s</code> not in <code>other</code> .
<code>s.IsEmpty()</code>	<code>true</code> if the set has no elements.
<code>s.Size()</code>	Number of elements.
<code>s.Clear()</code>	Remove all elements.
<code>s.ToArray()</code>	Return the elements as an array (hash order, not insertion order).

```
let seen = Collections.Set();
if (seen.Add(id)) {
    // first time we see id
}
let common = Collections.Set(a).Intersect(Collections.Set(b));
```

OrderedMap - Collections.OrderedMap() – object

Insertion-ordered hash map (a LinkedHashMap): $O(1)$ key lookup like `HashMap` plus deterministic insertion-order iteration. Internally a hash index over a doubly-linked node chain.

`Set/Get/Has/Delete/First/Last/PopFirst/PopLast/MoveToEnd/MoveToFront/Size` are $O(1)$;

`Keys/Values/Entries/ToArray/Clear` are $O(n)$. Key equality follows the value hash (ints, chars, and strings compare by value; floats and other heap values by identity). Setting an existing key updates its value in place and leaves its position unchanged; use `MoveToEnd/MoveToFront` to reorder (e.g. to drive an LRU cache).

Method	Description
<code>m.Set(k, v)</code>	Insert or update. Existing keys keep their position. Returns <code>true</code> .
<code>m.Get(k, default?)</code>	Stored value, else <code>default</code> (if given) or <code>nil</code> .
<code>m.Has(k)</code>	<code>true</code> if the key is present.
<code>m.Delete(k)</code>	Remove the entry; <code>true</code> if it existed. $O(1)$.
<code>m.First()</code>	Oldest value (head) without removing, or <code>nil</code> .
<code>m.Last()</code>	Newest value (tail) without removing, or <code>nil</code> .
<code>m.PopFirst()</code>	Remove and return the oldest value (FIFO/LRU eviction), or <code>nil</code> .
<code>m.PopLast()</code>	Remove and return the newest value, or <code>nil</code> .
<code>m.MoveToEnd(k)</code>	Move an existing key to newest position (LRU touch); <code>true</code> if present.
<code>m.MoveToFront(k)</code>	Move an existing key to oldest position; <code>true</code> if present.
<code>m.Keys()</code>	Keys as an array, in insertion order. $O(n)$.
<code>m.Values()</code>	Values as an array, in insertion order. $O(n)$.
<code>m.Entries()</code>	<code>[key, value]</code> pairs as an array, in insertion order. $O(n)$.
<code>m.ToArray()</code>	Values in insertion order (alias of <code>Values</code>). $O(n)$.
<code>m.IsEmpty()</code>	<code>true</code> if the map has no entries.
<code>m.Size()</code>	Number of entries.
<code>m.Clear()</code>	Remove all entries. $O(n)$.

```
// LRU cache: evict the least-recently-used entry when over capacity.
let cache = Collections.OrderedMap();
fn touch(key, value) {
    cache.Set(key, value);
```

```
cache.MoveToEnd(key); // mark as most-recently-used
if (cache.Size() > 100) cache.PopFirst(); // drop the oldest
}
```

Compress

Namespace: **Compress**

Raw zlib compression. Max input 4 GB.

Function	Signature	Description	JIT
Zip	<code>Zip(data:string block, level?:int) - block</code>	Compress data (raw zlib). Optional level 1–9 (default 6). Returns compressed block.	—
Unzip	<code>Unzip(data:block, maxSize?:int) - block</code>	Decompress data (raw zlib/gzip). Optional max output size hint. Returns decompressed block.	—

Archive

Namespace: **Archive**

ZIP archive creation and extraction. Uses a factory pattern: `Open / OpenFile` return reader objects; `Create` returns a writer object.

Factory functions:

Function	Signature	Description	JIT
Open	<code>Open(data:block) - reader</code>	Open a ZIP archive from in-memory bytes. Returns a reader object.	—
OpenFile	<code>OpenFile(path:string) - reader</code>	Open a ZIP archive from disk. Returns a reader object. Requires unsafe mode.	—
Create	<code>Create() - writer</code>	Create a new empty ZIP archive writer.	—

Reader methods (returned by `Open / OpenFile`):

Method	Signature	Description	JIT
List	<code>r.List() - array</code>	List all entries. Each element is <code>{Name:string, Size:int, CompressedSize:int, IsDir:bool}</code> .	—
Read	<code>r.Read(name:string) - block nil</code>	Extract one file by name. Returns raw bytes block, or <code>nil</code> if not found.	—
Extract	<code>r.Extract(name:string, dest:string) - bool</code>	Extract one file to a path on disk. Creates intermediate directories. Returns <code>true</code> on success. Requires unsafe mode.	—
ExtractAll	<code>r.ExtractAll(dir:string) - int</code>	Extract all files into <code>dir</code> . Creates intermediate directories. Returns number of files extracted. Requires unsafe mode.	—

Writer methods (returned by `Create`):

Method	Signature	Description	JIT
Add	<code>w.Add(name:string, data:string block) - nil</code>	Add a file entry from a string or block.	—
AddFile	<code>w.AddFile(name:string, path:string) - nil</code>	Add a file entry by reading from disk. Requires unsafe mode.	—
Build	<code>w.Build() - block</code>	Finalise the archive and return ZIP bytes as a block.	—

Method	Signature	Description	JIT
Save	<code>w.Save(path:string) - bool</code>	Finalise the archive and write it to disk. Returns <code>true</code> on success. Requires unsafe mode.	—

Example:

```
// Create
let w = Archive.Create();
w.Add("readme.txt", "Hello!");
w.Add("data/nums.txt", "1\n2\n3");
let bytes = w.Build();           // returns block
w.Save("/tmp/out.zip");         // or save directly to disk

// Read
let r = Archive.Open(bytes);
let entries = r.List();         // [{Name:"readme.txt", Size:6, ...}, ...]
let raw = r.Read("readme.txt"); // block -> str(raw) == "Hello!"
r.ExtractAll("/tmp/extracted");
```

ConsoleNamespace: **Console**

Terminal I/O, color, and cursor control.

Function	Signature	Description	JIT
Clear	<code>Clear() - nil</code>	Clear the entire terminal screen and move cursor to home (1,1).	—
ClearLine	<code>ClearLine() - nil</code>	Erase the current line and move cursor to column 1.	—
Error	<code>Error(...args:any) - nil</code>	Print to stderr in red.	—
GetCursor	<code>GetCursor() - object</code>	Returns <code>{x, y}</code> cursor position object.	—
IsTTY	<code>IsTTY() - bool</code>	<code>true</code> if stdout is a real terminal (not piped).	—
KeyPressed	<code>KeyPressed() - bool</code>	<code>true</code> if a key is waiting in the input buffer (non-blocking).	—
Ok	<code>Ok(...args:any) - nil</code>	Print to stdout in green.	—
ReadChar	<code>ReadChar() - char</code>	Read a single character from stdin (blocking).	—
ReadFloat	<code>ReadFloat() - float</code>	Read a float from stdin.	—
ReadInt	<code>ReadInt() - int</code>	Read an integer from stdin.	—
ReadKey	<code>ReadKey() - object</code>	Blocking read of one keypress. Returns <code>{Key:string, Char:char, IsSpecial:bool}</code> . Key is "Up", "Down", "Left", "Right", "Home", "End", "PageUp", "PageDown", "Insert", "Delete", "Enter", "Escape", "Backspace", "Tab", "F1" – "F12", or a single printable character. <code>IsSpecial</code> is <code>true</code> for arrows, function keys, and the navigation cluster.	—
ReadLine	<code>ReadLine() - string</code>	Read a line from stdin (blocking, strips newline).	—
ReadPassword	<code>ReadPassword() - string</code>	Read without echo.	—
RestoreCursor	<code>RestoreCursor() - nil</code>	Restore previously saved cursor position.	—
SaveCursor	<code>SaveCursor() - nil</code>	Save current cursor position.	—
SetBackground	<code>SetBackground(color:int) - nil</code>	Set background color. Use <code>ConsoleColor.*</code> constants.	—
SetColor	<code>SetColor(color:int) - nil</code>	Set foreground color. Use <code>ConsoleColor.*</code> constants.	—

Function	Signature	Description	JIT
SetCursor	SetCursor(x:int, y:int) – nil	Move cursor to column <code>x</code> , row <code>y</code> .	—
SetRaw	SetRaw(enable:bool) – nil	Enable or disable terminal raw mode. When <code>true</code> : disables line buffering and echo so individual keypresses are available immediately without pressing Enter. Restore with <code>SetRaw(false)</code> before exit. Has no effect if not running on a TTY.	—
Size	Size() – object	Returns <code>{width, height}</code> of terminal.	—
TryReadChar	TryReadChar() – char	Non-blocking: returns <code>char</code> if key pressed, else <code>\0</code> .	—
Warn	Warn(...args:any) – nil	Print to stderr in yellow.	—
Write	Write(...args:any) – nil	Print to stdout without trailing newline.	—
WriteLines	WriteLines(...args:any) – nil	Print multiple values each on its own line.	—
WriteLine	WriteLine(...args:any) – nil	Print to stdout with trailing newline.	—

ConsoleColor constants: `Default`, `Red`, `Green`, `Blue`, `Yellow`, `Cyan`

Convert

Namespace: **Convert**

Type conversion and parsing with full range of integer widths.

Function	Signature	Description	JIT
ToBool	ToBool(v:int float bool char) – bool	Convert numeric to bool.	✓
ToChar	ToChar(v:int float bool char) – char	Convert numeric to char (Unicode codepoint).	✓
ToFloat	ToFloat(v:int float bool char) – float	Convert numeric to float.	✓
ToInt	ToInt(v:int float bool char) – int	Convert numeric to int.	✓
ToInt8	ToInt8(v:int float bool char) – int	Truncate to signed 8-bit, stored as int.	✓
ToInt16	ToInt16(v:int float bool char) – int	Truncate to signed 16-bit, stored as int.	✓
ToInt32	ToInt32(v:int float bool char) – int	Truncate to signed 32-bit, stored as int.	✓
ToInt64	ToInt64(v:int float bool char) – int	Truncate to signed 64-bit, stored as int.	✓
ToUInt8	ToUInt8(v:int float bool char) – int	Truncate to unsigned 8-bit, stored as int.	✓
ToUInt16	ToUInt16(v:int float bool char) – int	Truncate to unsigned 16-bit, stored as int.	✓
ToUInt32	ToUInt32(v:int float bool char) – int	Truncate to unsigned 32-bit, stored as int.	✓
ToUInt64	ToUInt64(v:int float bool char) – int	Truncate to unsigned 64-bit, stored as int.	✓
ToString	ToString(v:int float bool char string) – string	Precise, round-trippable string form (floats via <code>%.17g</code>). Chars/bools render as their numeric value (<code>ToString('A') → "65"</code>); use <code>str()</code> for the short display form.	✓ owned ¹
ToRadixString	ToRadixString(n:int, base:int) – string nil	<code>n</code> written in an explicit base 2–36 (lowercase digits, <code>-</code> for negatives). <code>nil</code> if base is out of range. Inverse of <code>ParseRadix</code> .	✓ owned ¹
ParseBool	ParseBool(s:string) – bool nil	Parse <code>"true"/"false"</code> (case-insensitive) or <code>"1"/"0"</code> . Returns nil on anything else - never raises.	✓
ParseChar	ParseChar(s:string) – char nil	Parse a one-character string to a char. <code>nil</code> if empty, multi-character, or not valid UTF-8.	✓
ParseFloat	ParseFloat(s:string) – float nil	Parse a decimal/exponent string. Returns nil on failure - never raises. (Accepts <code>inf/nan</code> and hex floats via C <code>strtod</code> .)	✓

Function	Signature	Description	JIT
ParseInt	<code>ParseInt(s:string) - int nil</code>	Parse a signed integer (optional sign, <code>0x/0b/0o</code> prefix, surrounding whitespace). Full <code>Int64Min...Int64Max</code> range. Returns <code>nil</code> on failure - never raises.	✓
ParseUInt	<code>ParseUInt(s:string) - int nil</code>	Parse a non-negative integer (prefixes allowed) that fits the signed 64-bit slot. Returns <code>nil</code> on failure - never raises.	✓
ParseRadix	<code>ParseRadix(s:string, base:int) - int nil</code>	Parse <code>s</code> in an explicit base 2–36; no <code>0x/0b/0o</code> prefix is applied (so <code>"FF"</code> with base 16 is 255). <code>nil</code> on malformed input or out-of-range base.	✓
TryParseInt	<code>TryParseInt(x:any) - int nil</code>	Coerce any value to int: numbers pass through, strings are parsed, anything else is <code>nil</code> . Returns the value or <code>nil</code> .	✓
TryParseUInt	<code>TryParseUInt(x:any) - int nil</code>	As <code>TryParseInt</code> , rejecting negatives. Returns the value or <code>nil</code> .	✓
TryParseFloat	<code>TryParseFloat(x:any) - float nil</code>	Coerce any value to float. Returns the value or <code>nil</code> .	✓
TryParseBool	<code>TryParseBool(s:string) - bool</code>	Success predicate: <code>true</code> if <code>s</code> parses as a bool (see <code>ParseBool</code>), else <code>false</code> . Use <code>ParseBool</code> to recover the value.	✓
TryParseChar	<code>TryParseChar(s:string) - bool</code>	Success predicate: <code>true</code> if <code>s</code> is exactly one valid character, else <code>false</code> . Use <code>ParseChar</code> to recover the value.	✓

Numeric range constants (`Int32Max`, `FloatEpsilon`, ...) live in the `Math` module, alongside `Math.PI` and `Math.NaN`.

Char Module

Namespace: **Char**

Character classification and conversion. All functions operate on the full 32-bit Unicode codepoint stored inside a `char` value. Classification functions are Unicode-aware - they recognise letters, digits, and whitespace across all major scripts, not just ASCII (see per-function notes below). 18 of the 20 functions run at native speed inside JIT functions (see **R11**).

Classification - `char` - `bool`

Function	Signature	Unicode-aware	Description	JIT
IsAlpha	<code>IsAlpha(c:char) - bool</code>	✓	Letter in any major script (Latin, Greek, Cyrillic, Arabic, CJK, Hangul, ...).	✓
IsDigit	<code>IsDigit(c:char) - bool</code>	✓	Decimal digit (Unicode Nd category): ASCII 0–9 and script-native digits (Arabic-Indic, Devanagari, Thai, ...).	✓
IsAlNum	<code>IsAlNum(c:char) - bool</code>	✓	Letter or decimal digit (Unicode-aware for both).	✓
IsUpper	<code>IsUpper(c:char) - bool</code>	✓	Uppercase letter (Latin, Greek, Cyrillic, Armenian, Georgian, Fullwidth, ...).	✓

Function	Signature	Unicode-aware	Description	JIT
IsLower	<code>IsLower(c:char) - bool</code>	✓	Lowercase letter (same script coverage as <code>IsUpper</code>).	✓
IsCased	<code>IsCased(c:char) - bool</code>	✓	Has an upper/lower case distinction - equivalent to <code>IsUpper(c) IsLower(c)</code> .	✓
IsSpace	<code>IsSpace(c:char) - bool</code>	✓	Whitespace: ASCII space/tab/newlines plus Unicode spaces (En Space, Em Space, No-Break Space, ...).	✓
IsNewline	<code>IsNewline(c:char) - bool</code>	✓	Line-ending character: LF <code>\n</code> , CR <code>\r</code> , VT, FF, NEL U+0085, Line Separator U+2028, Paragraph Separator U+2029.	✓
IsAscii	<code>IsAscii(c:char) - bool</code>	-	Codepoint is in the ASCII range (U+0000–U+007F).	✓
IsPunct	<code>IsPunct(c:char) - bool</code>	-	ASCII punctuation only (<code>. , ! ? " ' ; : - () [] { } ...</code>).	✓
IsXDigit	<code>IsXDigit(c:char) - bool</code>	-	Hexadecimal digit: 0–9, A–F, a–f (ASCII only - hex is an ASCII concept).	✓
IsPrint	<code>IsPrint(c:char) - bool</code>	~	Printable: U+0020–U+007E plus codepoints above U+009F.	✓
IsControl	<code>IsControl(c:char) - bool</code>	~	C0 and C1 control characters: codepoint < U+0020 or U+007F–U+009F.	✓

Conversion - char - char

Function	Signature	Unicode-aware	Description	JIT
ToUpper	<code>ToUpper(c:char) - char</code>	✓	Uppercase mapping; non-letters returned unchanged.	✓
ToLower	<code>ToLower(c:char) - char</code>	✓	Lowercase mapping; non-letters returned unchanged.	✓
SwapCase	<code>SwapCase(c:char) - char</code>	✓	Upper-lower, lower-upper, all other characters unchanged.	✓

Numeric and UTF-8 helpers - char - int

Function	Signature	Description	JIT
DigitValue	<code>DigitValue(c:char) - int</code>	Decimal value of a Unicode digit (0–9), or -1 if the character is not a digit. Works for any Nd script block (Arabic-Indic '٣' - 3, Devanagari '३' - 9, ...).	✓
Utf8Len	<code>Utf8Len(c:char) - int</code>	Number of UTF-8 bytes needed to encode the codepoint: 1 (ASCII), 2 (U+0080–U+07FF), 3 (U+0800–U+FFFF), 4 (U+10000+).	✓

Codepoint accessors

Function	Signature	Description	JIT
Code	<code>Code(c:char) - int</code>	Raw Unicode codepoint as an integer ('A' - 65, ' ' - 0x4E2D).	✓
FromInt	<code>FromInt(n:int) - char</code>	Wrap an integer codepoint as a <code>char</code> . No range check performed.	✓

```
// Unicode-aware classification
Char.IsAlpha('ä') // true (Latin Extended)
Char.IsAlpha(' ') // true (CJK)
Char.IsDigit('३') // true (Arabic-Indic digit 3)
Char.IsUpper('Α') // true (Greek capital Alpha)
```

```
// Case conversion
Char.ToUpper('ä') // 'Ä'
Char.ToLower('A') // 'α'
Char.SwapCase('A') // 'a'

// Numeric helpers
Char.DigitValue('r') // 3
Char.DigitValue('a') // -1
Char.Utf8Len('a') // 1
Char.Utf8Len('ä') // 2
Char.Utf8Len(' ') // 3
Char.Utf8Len(Char.FromInt(0x1F600)) // 4

// Char calls run natively inside JIT functions
fn count_letters(s: string, n: int): int {
    let count: int = 0;
    for (let i: int = 0; i < n; i++) {
        if (Char.IsAlpha(s[i])) { count++; }
    }
    return count;
}
```

Char instance method syntax

All `Char` functions can also be called directly on a `char` value:

```
let c = String.CharAt("hello", 0); // 'h'
c.IsAlpha() // true - same as Char.IsAlpha(c)
c.Code() // 104 - same as Char.Code(c)
c.ToUpper() // 'H' - same as Char.ToUpper(c)
```

Both forms compile to identical bytecode when the variable is typed (`: char` or inferred from a builtin return). Untyped variables fall back to a runtime dispatch.

Crypto

Namespace: **Crypto**

Authenticated encryption (XChaCha20-Poly1305), Ed25519 signatures, X25519 key exchange, Argon2id password hashing, HMAC, and a CSPRNG. Public keys and signatures are hex strings.

Function	Signature	Description	JIT
Argon2	<code>Argon2(password:string block, salt:string block, iterations?:int, memKiB?:int) - string</code>	Argon2id password hash (64-hex). <code>salt</code> <code>>=</code> — 8 bytes (16 recommended); <code>iterations</code> default 3; <code>memKiB</code> memory cost in KiB, default 65536 (64 MiB), capped at 1 GiB.	

Function	Signature	Description	JIT
ConstantTimeEquals	<code>ConstantTimeEquals(a:string block, b:string block) - bool</code>	Compare two byte sequences in constant time (no early exit). Use instead of <code>==</code> when verifying MACs, signatures, or tokens so equality checks don't leak timing. Only length inequality returns early.	✓
Decrypt	<code>Decrypt(cipher:string block, cipherLen:int, key:string block, nonce:string block, tag:string) - object</code>	Decrypt <code>cipherLen</code> bytes of XChaCha20-Poly1305 cipher with the 32-byte <code>key</code> , 24-byte <code>nonce</code> , and hex <code>tag</code> returned by <code>Encrypt</code> . Returns <code>{Ok:bool, Plain:block}</code> . On authentication failure <code>Ok</code> is <code>false</code> and <code>Plain</code> is zeroed - always check <code>Ok</code> before using <code>Plain</code> .	—
Ed25519KeyPair	<code>Ed25519KeyPair() - object</code>	Generate an Ed25519 signing key pair. Returns <code>{PublicKey:string(64 hex), SecretKey:string(128 hex)}</code> .	—
Ed25519Sign	<code>Ed25519Sign(secretKey:string, message:string block) - string</code>	Sign <code>message</code> with a hex secret key; returns a 128-hex signature.	✓ owned ¹
Ed25519Verify	<code>Ed25519Verify(signature:string, publicKey:string, message:string block) - bool</code>	Verify a hex signature against a hex public key and message.	✓
Encrypt	<code>Encrypt(data:string block, dataLen:int, key:string block) - object</code>	Encrypt <code>dataLen</code> bytes of <code>data</code> with XChaCha20-Poly1305 under a 32-byte <code>key</code> . A fresh 24-byte nonce is generated per call (never supply your own). Returns <code>{Ok:bool, Cipher:block, Nonce:block, Tag:string(32 hex)}</code> . Max 256 MB.	—
HmacSha1	<code>HmacSha1(key:string block, data:string block) - string</code>	Compute HMAC-SHA1 (40-hex). Legacy/interop: TOTP/HOTP 2FA, OAuth1, AWS SigV2.	✓ owned ¹
HmacSha256	<code>HmacSha256(key:string block, data:string block) - string</code>	Compute HMAC-SHA256 (64-hex).	✓ owned ¹
HmacSha512	<code>HmacSha512(key:string block, data:string block) - string</code>	Compute HMAC-SHA512 (128-hex). E.g. JWT HS512.	✓ owned ¹
RandomBytes	<code>RandomBytes(count:int) - block</code>	Generate <code>count</code> cryptographically random bytes (CSPRNG). Raises if <code>count</code> is not in <code>1..67108864</code> .	—
X25519KeyPair	<code>X25519KeyPair() - object</code>	Generate an X25519 key-exchange key pair. Returns <code>{PublicKey:string(64 hex), SecretKey:string(64 hex)}</code> .	—
X25519Shared	<code>X25519Shared(secretKey:string, peerPublicKey:string) - string</code>	X25519 ECDH; returns the 64-hex raw shared secret. Hash it (e.g. <code>Hash.Blake2b</code>) before using it as a key.	✓ owned ¹

Debug

Namespace: **Debug**

Runtime introspection and assertion tools. Available in all builds; overhead is negligible for non-tracing calls.

Function	Signature	Description	JIT
Assert	Assert(left:any, right:any, ? message:any) - bool	Verify <code>left == right</code> . On mismatch prints both values (plus optional <code>message</code>) and a stack trace, then aborts the VM (<code>EXIT_CODE_ERR_ASSERT</code>) - not a catchable raise. Returns <code>true</code> when it holds.	—
AssertTrue	AssertTrue(condition:any, ? message:any) - bool	Verify <code>condition</code> is truthy. On failure prints it (plus optional <code>message</code>) and a stack trace, then aborts the VM . Returns <code>true</code> when it holds.	—
GuardAddress	GuardAddress(addr:int)	Installs an allocator watchpoint that traps when <code>addr</code> is touched.	—
Here	Here(?label:string)	Prints the current file, line, and function name to stdout (with an optional <code>label</code>).	—
Pool	Pool()	Prints SLAB allocator statistics and a per-slab dump.	—
Refs	Refs(value:any) - int	Returns <code>value</code> 's external reference count (<code>-1</code> for tagged immediates, which have none).	—
Stack	Stack()	Prints the current operand stack to stdout.	—
StackPtr	StackPtr() - int	Returns the current operand-stack depth.	—
Value	Value(v:any)	Prints a detailed internal representation of a value.	—

Directory

Namespace: **Directory**

Filesystem directory operations.

Function	Signature	Description	JIT
Copy	Copy(src:string, dst:string) - bool	Recursively copy directory <code>src</code> into <code>dst</code> (created if missing; existing <code>dst</code> is merged into). Symlinks are copied as links, never dereferenced. <code>false</code> if <code>src</code> is not a directory or any entry could not be copied.	—
Create	Create(path:string) - bool	Create directory at <code>path</code> . Returns <code>false</code> if already exists.	—
CreateTemp	CreateTemp(prefix?:string) - string	Atomically create a uniquely-named directory in the system temp dir and return its path (<code>mkdtemp</code>). Optional name prefix (default <code>"flaris"</code>) must not contain path separators. Not auto-removed - pair with <code>DeleteAll</code> . <code>nil</code> on failure.	—
Delete	Delete(path:string) - bool	Delete directory (must be empty).	—
DeleteAll	DeleteAll(path:string) - bool	Recursively delete a directory and everything under it (<code>rm -rf</code>). Symlinks are removed as links, never followed out of the tree. Refuses an empty path or the filesystem root. <code>true</code> only if every entry was removed.	—
Ensure	Ensure(path:string) - bool	Create directory (and parents) if not exists; no-op if exists.	—
Exists	Exists(path:string) - bool	<code>true</code> if path exists and is a directory.	—
GetModifiedTime	GetModifiedTime(path:string) - int	Returns Unix timestamp of last modification.	—
Glob	Glob(pattern:string, followSymlinks?:bool) - array	Recursive glob: each <code>/</code> -separated segment supports <code>*</code> , <code>?</code> , and <code>[set]</code> ; a bare <code>**</code> segment matches zero or more directory levels (<code>"src/**/*.fls"</code>). Returns matching paths (files and directories) relative to the current directory (or absolute if the pattern is). Case-sensitive; wildcards do not match a leading <code>.</code> (name the dot explicitly). Depth capped at 64. <code>followSymlinks</code> defaults to <code>true</code> ; pass <code>false</code> to keep <code>**</code> from descending through symlinked directories	—

Function	Signature	Description	JIT
		(POSIX). Explicitly named path components are always followed - only <code>**</code> wildcard descent is affected.	
List	<code>List(path:string) - array</code>	Returns array of all entry names (files + dirs) in <code>path</code> .	—
Move	<code>Move(src:string, dst:string) - bool</code>	Rename a directory (atomic within one filesystem; fails across filesystems with no copy+delete fallback).	—
ListDirs	<code>ListDirs(path:string) - array</code>	Returns array of subdirectory names only.	—
ListFiles	<code>ListFiles(path:string, filters?:string) - array</code>	Returns array of entry names, optionally filtered by a comma-separated glob list (<code>"*.txt, *.md"</code> - full <code>* / ? / [set]</code> wildcards, case-insensitive). Empty array if the directory can't be opened (same as <code>List</code>).	—

Event

Namespace: **Event**

Signals fibers can wait on. Max 64 events (a shared, fixed pool).

Function	Signature	Description	JIT
Create	<code>Create() - int</code>	Claims a new event id (<code>0..63</code>), or <code>-1</code> if the 64-id pool is exhausted. Always check for <code>-1</code> .	—
Set	<code>Set(id:int, value?:any)</code>	Signal the event; optionally attach a value. Wakes every fiber waiting on it (broadcast). Raises <code>IndexOutOfBounds</code> for an id outside <code>0..63</code> , or <code>InvalidArgs</code> for an id that was never created.	—
WaitOne	<code>WaitOne(id:int, timeoutMs?:int) - any</code>	Parks the caller until <code>Set</code> is called on this event, then returns the value it was set with. With a positive <code>timeoutMs</code> it resumes with <code>nil</code> once the deadline passes; <code>0</code> /omitted waits forever. Raises <code>IndexOutOfBounds</code> for an out-of-range id.	—
WaitFor	<code>WaitFor(ids:array, timeoutMs?:int) - array</code>	Waits until all of the events in <code>ids</code> have been set, then returns their values as an array. Use it to start several <code>ProcessEvent</code> / <code>Set</code> jobs and continue once they've all finished. (An empty list returns <code>[]</code> right away.) With a positive <code>timeoutMs</code> , if the deadline passes before every event fires it returns <code>nil</code> and stops waiting; <code>0</code> or omitted waits forever. Raises <code>IndexOutOfBounds</code> for an out-of-range id or <code>InvalidArgs</code> for a non-int element.	—
Free	<code>Free(id:int)</code>	Releases an id back to the pool. Use to reclaim an event you created but never delivered (a dropped job) so the fixed pool is not exhausted. Raises <code>IndexOutOfBounds</code> for an out-of-range id.	—
Reset	<code>Reset(id:int)</code>	Clears a claimed event's signalled state and value so it can be <code>Set</code> and waited — on again, without returning the id to the pool. Raises <code>IndexOutOfBounds</code> for an out-of-range id.	—

Lifecycle: an event is reclaimed automatically once the last fiber waiting on it has been served, so a fan-out `WaitFor` join does not leak ids. An event you `Create` but never deliver stays claimed until you `Free` it. A `Set` is broadcast to every fiber already parked on the event; a fiber that starts waiting **after** the last waiter was served and the id reclaimed will block (one-shot events are not retained) — subscribe your waiters before signalling.

Note: `WaitFor` returns the values sorted by event id, which may differ from the order you listed the ids in. If you need to tell them apart, use the values themselves or the buffer each job wrote — not the array position. On timeout the result is `nil` (distinct from the always-non-nil array a successful join returns).

FFI

Namespace: **Ffi**

Dynamic loading of native shared libraries (`.so` , `.dylib`). Requires `--unsafe` flag. FFI calls cross a clean isolation boundary - arguments and return values are marshalled through a structured type system; no VM internals are exposed to native code.

Function	Signature	Description	JIT
Load	<code>Load(path:string, sha256:string nil, flags?:int) - pointer</code>	Load shared library at <code>path</code> . When <code>sha256</code> is a 64-character hex string, the SHA-256 fingerprint of the library is verified before loading - pass <code>nil</code> to skip the check. A digest that is not 64 hex characters is a malformed pin and raises <code>Exception.ChecksumError</code> rather than loading unverified, as does a mismatch. SHA verification is optional security hardening; libraries distributed without a known hash should pass <code>nil</code> . <code>flags</code> may combine <code>Ffi.LAZY</code> (default), <code>Ffi.NOW</code> and <code>Ffi.GLOBAL</code> ; ignored on Windows. The plugin's <code>flaris_abi_version</code> (see <code>FLARIS_PLUGIN_ABI()</code> in <code>ffi_object.h</code>) is verified; an unknown ABI raises <code>Exception.ChecksumError</code> , and a library exporting none is loaded with a warning. Returns a handle, or <code>nil</code> on failure - see <code>Ffi.LastError</code> . Same library is never loaded twice - returns cached handle on repeated calls, and loaded libraries stay mapped for the lifetime of the process.	—
GetFunction	<code>GetFunction(handle:pointer, name:string, sig?:string) - ffi_function</code>	Retrieve a named symbol from a loaded library as a callable value. When <code>sig</code> is provided, argument count, argument types and return type are validated on every call; a mismatched return raises <code>Exception.TypeMismatch</code> , and <code>nil</code> is always accepted as a return value regardless of the declared type. Returns <code>nil</code> if the symbol is missing or <code>sig</code> is malformed - see <code>Ffi.LastError</code> . Omitting <code>sig</code> disables validation (all types accepted).	—
HasFunction	<code>HasFunction(handle:pointer, name:string) - bool</code>	Returns <code>true</code> if the named symbol exists in the library.	—
SetCallback	<code>SetCallback(handle:pointer, name:string, fn:fn) - bool</code>	Does not allocate. Use to guard optional symbols. Register a Flaris function as a named callback. The library must export <code>flaris_set_callback</code> . Callbacks are global - use a lib-specific prefix to avoid name collisions across plugins (e.g. <code>"mylib_ondata"</code>).	—
RemoveCallback	<code>RemoveCallback(handle:pointer, name:string) - bool</code>	Unregister a previously registered callback by name without unloading the library. Useful for one-shot callbacks.	—
GetVersion	<code>GetVersion(handle:pointer) - int</code>	Returns the plugin's version as a <code>uint32</code> in <code>0xMMmmpbb</code> format. Plugin must export <code>uint32_t flaris_version(void)</code> . Returns <code>0</code> if symbol not found.	—
Call	<code>Call(handle:pointer, name:string, ...args) - any</code>	Resolve and call a named symbol in one step without allocating an <code>ffi_function</code> object. Accepts up to 4 extra arguments. Calls <code>dlsym</code> on every invocation - use <code>GetFunction</code> for hot paths.	—
CallAsync	<code>CallAsync(fn:ffi_function, args:array) - any</code>	Runs a pure FFI function (from <code>GetFunction</code>) on the I/O thread pool so the calling fiber suspends instead of blocking the scheduler. Must be <code>await</code> ed. <code>args</code> is the call's argument array; block arguments are refused (they alias VM	—

Function	Signature	Description	JIT
		memory and cannot cross threads). The function must be side-effect-free with respect to the VM: no callbacks, no shared mutable state, thread-safe. A declared return type that does not match resolves to <code>nil</code> .	
LastError	<code>LastError() - string nil</code>	Text of the most recent <code>Load</code> / <code>GetFunction</code> / <code>Call</code> failure, — or <code>nil</code> if the last one succeeded. Those three clear it on entry, so it always describes the most recent call.	
LAZY	<code>LAZY - int</code>	<code>dlopen</code> flag: resolve symbols on first use (the default). —	
NOW	<code>NOW - int</code>	<code>dlopen</code> flag: resolve every symbol at load, so a missing symbol fails <code>Load</code> instead of crashing on first call. —	
GLOBAL	<code>GLOBAL - int</code>	<code>dlopen</code> flag: make the library's symbols available to subsequently loaded libraries. —	

Ffi.Load path resolution. The name resolves in order: (1) the path as given if it exists; (2) the per-user FFI dir - `~/.flaris/ffi` (POSIX) / `%USERPROFILE%\flaris\ffi` (Windows), or `$FLARIS_FFI` - where a **bare name** gets the platform extension appended (`Ffi.Load("myplugin")` → `~/.flaris/ffi/myplugin.{dylib,so,dll}`); (3) `<name><ext>` in the working directory; (4) the system loader (so `"libsqlite3.so"` still works). Prefer a bare name and install the library to `~/.flaris/ffi` for portable, path-free loading. See [ffi.md](#).

FFI Type Mapping

Flaris type	FfiObject	Notes
<code>nil</code>	<code>FFIVAL_NIL</code>	Also what an unsupported or too-deeply-nested value marshals to.
<code>int</code>	<code>FFIVAL_INT (ival)</code>	<code>int64_t</code> .
<code>float</code>	<code>FFIVAL_FLOAT (fval)</code>	<code>double</code> .
<code>bool</code>	<code>FFIVAL_BOOL (ival)</code>	<code>0 / 1</code> .
<code>char</code>	<code>FFIVAL_CHAR (ival)</code>	Unicode code point.
<code>string</code>	<code>FFIVAL_STRING (string.data, string.length)</code>	Arguments borrow the VM's buffer (<code>FFI_FLAG_BORROWED</code>): read it freely during the call, never write it or free it, copy it if you need it afterwards. A string returned to the VM must be <code>malloc'd</code> - the VM takes ownership.
<code>block</code>	<code>FFIVAL_BLOCK (block.memory)</code>	Zero-copy pointer into VM memory, writable by convention. Never free or reallocate. Size is <code>block_size * block_count</code> .
<code>pointer</code>	<code>FFIVAL_POINTER</code> -shaped <code>ptr</code>	Opaque handle (e.g. from <code>Ffi.Load</code>).
<code>array</code>	<code>FFIVAL_ARRAY (array.elements)</code>	Contiguous <code>FfiObject</code> values , not pointers. Deep-copied. Avoid in hot paths.
<code>object</code>	<code>FFIVAL_OBJECT (object.pairs)</code>	Contiguous <code>FfiKV</code> pairs, values embedded. Deep-copied. Argument keys borrow the VM's key strings under the same rules as string arguments. Avoid in hot paths.
<code>array of objects</code>	<code>FFIVAL_TABLE</code> (return only)	Rowset: column names once + row-major cells. Converts to the same array-of-row-objects shape, hashing each column name once per table instead of once per row (~1.5x faster to marshal). Prefer for anything row-shaped.
<code>ffi_function</code>	<code>FFIVAL_FUNC (fn)</code>	A C function pointer handed back to script. It carries no signature, so it is always called fully dynamically.

Prefer `int` and `float` for high-frequency callback arguments. Use `block` for bulk data. `string` is acceptable for low-frequency events (errors, lifecycle). `array` and `object` should be avoided in callbacks called at high rate.

Function Signature Format

Signatures passed to `GetFunction` use the syntax:

```
fn(arg_type, arg_type, ...):return_type
```

Types may be unioned with `|`. Zero-argument functions use `fn():return_type`.

Type name	Matches
<code>any</code>	Any value
<code>nil</code>	Nil
<code>bool</code>	Boolean
<code>int</code>	Integer (<code>u8</code> , <code>u16</code> , <code>u32</code> , <code>u64</code> , <code>i8</code> , <code>i16</code> , <code>i32</code> , <code>i64</code> are aliases)
<code>float</code>	Float
<code>char</code>	Character
<code>string</code>	String
<code>array</code>	Array
<code>object</code>	Object
<code>block</code>	Block (raw memory buffer)
<code>pointer</code> / <code>ptr</code>	Pointer
<code>function</code> / <code>fn</code>	Flaris function (does not match an FFI function - use <code>callable</code> or <code>ffi</code>)
<code>builtin</code>	Builtin function
<code>ffi</code>	FFI function
<code>number</code> / <code>numeric</code>	<code>int float</code>
<code>callable</code>	<code>fn builtin ffi</code>

Examples:

```
let fn = Ffi.GetFunction(lib, "add", "fn(int, int):int");
let fn = Ffi.GetFunction(lib, "process", "fn(string, pointer):bool");
let fn = Ffi.GetFunction(lib, "compute", "fn(number):float");
let fn = Ffi.GetFunction(lib, "reset", "fn():nil");
let fn = Ffi.GetFunction(lib, "anything", "fn(any):any");
```

For a complete plugin development guide including helpers, examples, and ownership rules, see `ffi.md`.

Fiber

Namespace: **Fiber**

Cooperative green-thread creation, communication, and lifecycle management.

Function	Signature	Description	JIT
Cancel	<code>Cancel(f: fiber) - bool</code>	Request cancellation of a running fiber.	—
Cancello	<code>CancelIo(f: fiber) - bool</code>	Abandon <code>f</code> 's pending async I/O without killing the fiber: it wakes from its <code>await</code> with <code>nil</code> and <code>Stream.LastError()</code> reporting 7 (cancelled), then carries on running. Use it to drop a stalled read while still answering the client — <code>Cancel</code> by contrast stops the fiber for good. <code>false</code> if <code>f</code> was not waiting on I/O.	—
FromId	<code>FromId(id: int) - fiber</code>	Returns fiber handle from its integer ID.	—
GetMessage	<code>GetMessage(f: fiber) - any</code>	Consume and return the oldest queued message (FIFO), or <code>nil</code> if the mailbox is empty.	—

Function	Signature	Description	JIT
HasMessage	HasMessage(f: fiber) – bool	true if at least one message is queued for f.	—
Id	Id() – int	Returns current fiber's integer ID.	—
MessageCount	MessageCount(f: fiber) – int	Number of messages currently queued in f's mailbox.	—
IsDone	IsDone(f: fiber) – bool	true if fiber has completed or been cancelled.	—
New	New(func: fn) – fiber	Create a new fiber (not yet started).	—
Resume	Resume(f: fiber nil, detached?: bool) – any	Schedule f to run and yield self. Default (attached): the caller parks until f yields/finishes and that value is routed back. detached=true: the caller just suspends and f runs independently. Passing nil (or a finished f) returns nil. The resume value is delivered later by the scheduler.	—
Run	Run(func: fn) – int	Create and immediately start a detached fiber. Returns its integer ID.	—
SendMessage	SendMessage(f: fiber, msg: any) – bool	Append msg to f's bounded FIFO mailbox (depth 64) and wake it if idle/sleeping. false if f is missing, finished, or its mailbox is full (older messages are kept, not overwritten); true otherwise. Non-blocking.	—
SetQuantum	SetQuantum(f: fiber, quantum: int) – bool	Override the scheduling quantum for f - the number of checkpoints (loop back-edges and call/return boundaries) before yielding to the scheduler. Clamped to [0, 10×FIBER_QUANTUM]; 0 means "use the default".	—
Sleep	Sleep(ms: int) – nil	Suspend current fiber for at least ms milliseconds (other fibers keep running). Negative ms yields immediately.	—
Status	Status(f: fiber) – int	Returns status code. Constants: Fiber.Status.New=0, Running=1, Suspended=2, Finished=3, WaitIO=4, Yield=5, WaitFiber=6.	—

Fiber instance method syntax

All `Fiber` functions where the fiber is the first argument can be called directly on a `fiber` value:

```
let f = Fiber.New(fn() { /* ... */ });
f.Resume()           // same as Fiber.Resume(f)
f.IsDone()           // same as Fiber.IsDone(f)
f.Status()           // same as Fiber.Status(f)
f.Cancel()           // same as Fiber.Cancel(f)
f.HasMessage()       // same as Fiber.HasMessage(f)
f.GetMessage()       // same as Fiber.GetMessage(f)
f.MessageCount()     // same as Fiber.MessageCount(f)
f.SendMessage("hi")  // same as Fiber.SendMessage(f, "hi")
```

Factory functions (`New`, `Run`, `Id`, `Sleep`) are not available as instance methods. Both forms compile to identical bytecode when the variable is typed (`: fiber` or inferred from a builtin return). Untyped variables fall back to a runtime dispatch.

File

Namespace: **File**

File read, write, and metadata operations.

Function	Signature	Description	JIT
AppendLines	<code>AppendLines(path:string, lines:array) - bool</code>	Append array of strings as lines to file.	—
AppendText	<code>AppendText(path:string, text:string) - bool</code>	Append text string to file.	—
Copy	<code>Copy(src:string, dst:string) - bool</code>	Copy file from <code>src</code> to <code>dst</code> .	—
CopyAsync	<code>CopyAsync(src:string, dst:string) - fiber</code>	Non-blocking <code>Copy</code> . Use with <code>await</code> ; resolves to <code>true</code> on success / <code>nil</code> on error. Other fibers run while the file is copied.	—
CreateHardLink	<code>CreateHardLink(target:string, linkPath:string) - bool</code>	Create a hard link at <code>linkPath</code> sharing <code>target</code> 's inode. <code>target</code> must exist and be on the same filesystem.	—
CreateSymlink	<code>CreateSymlink(target:string, linkPath:string) - bool</code>	Create a symbolic link at <code>linkPath</code> pointing to <code>target</code> (need not exist). On Windows requires Developer Mode / admin.	—
Delete	<code>Delete(path:string) - bool</code>	Delete file at <code>path</code> .	—
Exists	<code>Exists(path:string) - bool</code>	<code>true</code> if path exists and is a regular file (symlinks followed). Directories report <code>false</code> - use <code>Directory.Exists</code> .	—
GetModifiedTime	<code>GetModifiedTime(path:string) - int</code>	Returns Unix timestamp of last modification.	—
Lines	<code>Lines(path:string, callback:fn) - int</code>	Stream the file line by line, calling <code>callback(line)</code> for each (trailing CR/LF stripped) without building an array - the memory-cheap alternative to <code>ReadLines</code> for large files. The callback may return <code>false</code> to stop early. Returns the number of lines processed, or <code>nil</code> if the file can't be opened or the callback raises.	—
Move	<code>Move(src:string, dst:string) - bool</code>	Move/rename file.	—
ReadAllBytes	<code>ReadAllBytes(path:string) - block</code>	Read entire file as binary block.	—
ReadAllBytesAsync	<code>ReadAllBytesAsync(path:string) - fiber</code>	Non-blocking <code>ReadAllBytes</code> . Use with <code>await</code> . Other fibers run while the file is read.	—
ReadLines	<code>ReadLines(path:string) - array</code>	Read file into array of strings (one per line).	—
ReadLink	<code>ReadLink(path:string) - string</code>	The target a symbolic link points to (the raw stored target, not a resolved path), or <code>nil</code> if not a symlink / unreadable. POSIX only - <code>nil</code> on Windows.	—
ReadText	<code>ReadText(path:string) - string</code>	Read entire file as UTF-8 string.	—
ReadTextAsync	<code>ReadTextAsync(path:string) - fiber</code>	Non-blocking <code>ReadText</code> . Use with <code>await</code> . Other fibers run while the file is read.	—
Sha256	<code>Sha256(path:string, raw?:bool) - string</code>	Compute SHA-256 of file contents. Optional <code>raw</code> for binary output.	—
Sha256Async	<code>Sha256Async(path:string, raw?:bool) - fiber</code>	Non-blocking <code>Sha256</code> . Use with <code>await</code> ; resolves to the hex string (or 32-byte block if <code>raw</code>), or <code>nil</code> if the file can't be opened.	—
SetMode	<code>SetMode(path:string, mode:int) - bool</code>	Set permission bits (POSIX <code>chmod</code> , low <code>07777</code> of <code>mode</code>). On Windows only the write bit is honored. <code>nil</code> if it can't be applied.	—
SetModifiedTime	<code>SetModifiedTime(path:string, unixSeconds:int) - bool</code>	Set the file's modification time to a Unix timestamp; access time is preserved where the platform allows.	—

Function	Signature	Description	JIT
		<code>nil</code> on failure.	
Size	<code>Size(path:string) - int</code>	File size in bytes, or <code>nil</code> if the file can't be opened.	—
Stat	<code>Stat(path:string) - object</code>	One-syscall metadata: {Size, Modified, Accessed, Created, IsDir, IsFile, IsSymlink, Mode}. Times are Unix timestamps; Mode is the permission bits (octal 7777 mask); Created uses birth time on macOS/BSD, else ctime. Returns <code>nil</code> if the path does not exist. A dangling symlink stats as the link itself.	—
Touch	<code>Touch(path:string) - bool</code>	Create file if missing; set access + modification time to now (like POSIX <code>touch</code>).	—
Truncate	<code>Truncate(path:string, size:int) - bool</code>	Resize the file to exactly <code>size</code> bytes (extension zero-fills). <code>false</code> if <code>size</code> is negative or the file can't be resized.	—
WriteAllBytes	<code>WriteAllBytes(path:string, data:string block) - bool</code>	Write binary data to file (overwrite). <code>false</code> on open/write failure.	—
WriteAtomic	<code>WriteAtomic(path:string, data:string block) - bool</code>	Durable write: stream into a temp file in the same directory, <code>fsync</code> , then <code>rename()</code> over the target - a reader never sees a half-written file, and a crash leaves either the old or new content, never a truncated mix. <code>false</code> on any failure (temp file removed). Atomic only within one filesystem.	—
WriteLines	<code>WriteLines(path:string, lines:array) - bool</code>	Write array of strings as lines (overwrite). <code>nil</code> on open/write failure.	—
WriteText	<code>WriteText(path:string, text:string) - bool</code>	Write text string to file (overwrite). <code>nil</code> on open/write failure.	—
WriteTextAsync	<code>WriteTextAsync(path:string, text:string) - fiber</code>	Non-blocking <code>WriteText</code> . Use with <code>await</code> . Other fibers run while the file is written.	—

FileWatch

Namespace: **FileWatch**

File and directory event monitoring. Up to 64 concurrent watchers. Uses `inotify` on Linux, `kqueue` on macOS/BSD, and `ReadDirectoryChangesW` on Windows. The callback fires on every scheduler tick where an event is pending - no fiber is blocked.

Function	Signature	Description	JIT
Open	<code>Open(path:string, callback:fn) - int</code>	Watch <code>path</code> for changes using the default mode (<code>"write,create,delete"</code>). Returns a slot id (≥ 0) or <code>-1</code> on failure.	—
Open	<code>Open(path:string, mode:string, callback:fn) - int</code>	Watch <code>path</code> with explicit mode.	—
Close	<code>Close(id:int) - bool</code>	Stop watching and release the slot.	—

Mode tokens - comma-separated, any combination:

Token	Description	Linux	macOS	Windows
"write"	File content modified or saved	✓	✓	✓
"create"	New files/dirs created or moved in	✓	✓	✓
"delete"	File deleted or renamed away	✓	✓	✓

Token	Description	Linux	macOS	Windows
"attrib"	Metadata changed (chmod, chown, timestamps)	✓	✓	✓
"open"	File opened for reading or writing	✓	-	-
"all"	All of the above	✓	✓ (except open)	✓ (except open)

Default mode (no mode argument, or no recognized token) is `"write,create,delete"`. Unrecognized tokens are ignored, not errors - `"open"` on a platform without it simply falls back to whatever else the string names (or the default).

Callback signature: `fn(event:object, id:int)`

Field	Type	Description
<code>event.event</code>	string	<code>"modify"</code> , <code>"create"</code> , <code>"delete"</code> , <code>"rename"</code> , <code>"attrib"</code> , <code>"open"</code> , or <code>"unknown"</code>
<code>event.name</code>	string	Linux/Windows: the changed entry's name relative to the watched directory (may be empty). macOS/BSD: the watched path as passed to <code>Open</code> .

Notes:

- The file must exist when `Open` is called. Watching a non-existent path returns `-1`.
- On macOS, editors like `nano` and `vim` save via rename - the file's inode is replaced. After receiving a `"delete"` event, close and re-open the watch to re-attach to the new inode.
- `"open"` mode is Linux-only; elsewhere it is ignored.

Example:

```
var id = FileWatch.Open("./config.json", fn(ev, id) {
    Console.WriteLine(ev.event + " " + ev.name)
})

Fiber.Sleep(30_000)
FileWatch.Close(id)
```

Gfx

Namespace: **Gfx**

Software rasterizer. Colors are `0xAARRGGBB` integers. Max canvas 32768×32768.

Function	Signature	Description	JIT
Create	<code>Create(width:int, height:int) - instance</code>	Create blank canvas.	—
CreateFromFile	<code>CreateFromFile(path:string) - instance</code>	Load PNG from disk into canvas.	—

Static color helpers:

Function	Signature	Description	JIT
RGB	<code>RGB(r:int, g:int, b:int) - int</code>	Pack RGB channels into a fully-opaque <code>0xFFRRGGBB</code> color.	✓
RGBA	<code>RGBA(r:int, g:int, b:int, a:int) - int</code>	Pack RGBA channels into <code>0xAARRGGBB</code> .	✓
HSL	<code>HSL(h:float, s:float, l:float) - int</code>	Convert HSL to <code>0xFFRRGGBB</code> . <code>h</code> is 0–360, <code>s</code> and <code>l</code> are 0.0–1.0.	✓

IGfx instance fields: Width:int, Height:int, Pixels:array

IGfx instance methods:

Method	Signature	Description	JIT
Clear	Clear(color:int)	Fill entire canvas with <code>color</code> .	
SetPixel	SetPixel(x, y, color)	Set single pixel.	
GetPixel	GetPixel(x, y) – int	Return pixel color.	
DrawLine	DrawLine(x1, y1, x2, y2, color)	Bresenham line.	
DrawHLine	DrawHLine(x, y, len, color)	Horizontal line.	
DrawVLine	DrawVLine(x, y, len, color)	Vertical line.	
DrawRect	DrawRect(x, y, w, h, color)	Rectangle outline.	
FillRect	FillRect(x, y, w, h, color)	Filled rectangle.	
DrawCircle	DrawCircle(cx, cy, r, color)	Circle outline (Midpoint algorithm).	
FillPolygon	FillPolygon(points, color)	Fill polygon from array of <code>{x,y}</code> objects.	
StrokePath	StrokePath(points, color)	Stroke open polyline.	
DrawChar	DrawChar(x, y, ch, color, scale?)	Render a single character.	
DrawText	DrawText(x, y, text, color, scale?)	Render a string.	
DrawTextOpts	DrawTextOpts(x, y, text, opts?)	Render a string using an options object. Fields: <code>color</code> (int, default opaque white), <code>scale</code> (int, default 1), <code>spacing</code> (int px between glyphs, default 0), <code>lineSpacing</code> (int px between lines, default 0). Absent/unknown fields use their defaults.	
Blit	Blit(x, y, w, h, pixels)	Copy a flat ARGB array (<code>QRCode.ToPixels</code> , custom bitmaps) onto the canvas at <code>(x, y)</code> . <code>pixels</code> must have at least <code>w × h</code> elements.	
DrawImage	DrawImage(x, y, src)	Blit another <code>IGfx</code> canvas onto this one.	
SaveToFile	SaveToFile(path) – bool	Save canvas as PNG.	
SaveToBlock	SaveToBlock() – block	Return canvas as PNG bytes.	
Draw	Draw(fn)	Batch-draw callback: <code>fn(canvas)</code> .	
CopyRect	CopyRect(x, y, w, h) – instance	Copy sub-rectangle into new canvas.	
Resize	Resize(w, h) – instance	Resize (nearest-neighbor) into new canvas.	
Rescale	Rescale(factor) – instance	Scale by factor into new canvas.	
Filter	Filter(fn, x, y, w, h)	Apply per-pixel callback <code>fn(pixel:int) – int</code> over region. Return <code>nil</code> to leave pixel unchanged. Mutates in place.	
FilterCallback	FilterCallback(fn, x, y, w, h, cb)	Runs your pixel-processing function <code>fn(pixels, x, y, w, h, stride)</code> over the <code>x,y,w,h</code> rectangle of the image, then calls <code>cb(image, result)</code> when finished. <code>pixels</code> is the image's 32-bit ARGB data; reach a pixel with <code>pixels[(y+dy)*stride + (x+dx)]</code> . Runs on another CPU core when <code>fn</code> is simple enough (only touches <code>pixels</code> , plain number math, no calls or allocations) — otherwise runs normally, same result. See <code>Buffer.ProcessCallback</code> .	
FilterEvent	FilterEvent(fn, x, y, w, h, eventId:int)	Like <code>FilterCallback</code> , but signals event <code>eventId</code> with the result instead of calling a callback. Wait for several with	

Method	Signature	Description	JIT
		<code>Event.WaitFor([ids])</code> .	
Grayscale	<code>Grayscale(x, y, w, h)</code>	Convert region to grayscale using BT.601 luminance weights. Alpha channel preserved. Mutates in place.	
Brightness	<code>Brightness(x, y, w, h, factor:float)</code>	Multiply RGB channels by <code>factor</code> (1.0 = no change, 0.5 = half, 2.0 = double). Values clamped to 0–255. Alpha preserved. Mutates in place.	
BlendRect	<code>BlendRect(src:instance, dx:int, dy:int, alpha?:int)</code>	Composite <code>src</code> canvas onto this canvas at <code>(dx, dy)</code> using source-over alpha blending. Optional <code>alpha</code> (0–255) scales the overall opacity of <code>src</code> ; defaults to 255. Mutates in place.	

PNG Compatibility

Reading (`CreateFromFile`, `DrawImage`) - decoded to RGBA8:

- Color types: grayscale, RGB, palette (indexed), grayscale+alpha, and RGBA
- Bit depths: 1/2/4/8/16 as the color type allows (16-bit is reduced to 8-bit; sub-byte depths are unpacked)
- Palette `tRNS` (per-index transparency) is applied; palette images require a `PLTE` chunk
- All five filter types supported (None, Sub, Up, Average, Paeth)
- Not supported: interlaced (Adam7) files, and color-key `tRNS` for grayscale/RGB. These throw `EXCEPTION_IO_ERROR`

Writing (`SaveToFile`, `SaveToBlock`):

- Always writes RGBA 32-bit, 8-bit per channel
- Filter: None (fast, slightly larger files than adaptive filtering)
- Compression: zlib level 6

Color model

- Pixels are 32-bit `0xAARRGGBB` with **straight (non-premultiplied) alpha**.
- Compositing (`Draw`, `BlendRect`, `DrawImage`, semi-transparent primitives) is source-over in **gamma / sRGB space** (not linearized) - the same default as Skia and Cairo. Linear-light blending is not offered; if added it would arrive as an explicit blend-mode option rather than a global change.
- **Bilinear** scaling in `DrawImage` interpolates in **premultiplied alpha**, so a fully transparent source pixel contributes no color and cannot bleed its RGB into scaled edges (no halos). Nearest-neighbour scaling is unaffected.

Hash

Namespace: `Hash`

Non-cryptographic and cryptographic hashes. Input accepts `string` or `block`. Max input 1 GB.

Function	Signature	Description	JIT
Adler32	<code>Adler32(data:string block) - int</code>	Adler-32 checksum.	✓
Blake2b	<code>Blake2b(data:string block, outLen?:int) - string</code>	BLAKE2b hash (hex string). <code>outLen</code> 1-64 bytes, default 32.	✓ owned¹
Blake2s	<code>Blake2s(data:string block) - string</code>	BLAKE2s hash (hex string).	✓ owned¹

Function	Signature	Description	JIT
Crc32	<code>Crc32(data:string block) - int</code>	CRC-32 checksum.	✓
Crc32c	<code>Crc32c(data:string block) - int</code>	CRC-32C (Castagnoli) checksum.	✓
Md5	<code>Md5(data:string block) - string</code>	MD5 (hex string). Legacy/interop checksums only (Content-MD5, ETags); not for security.	✓ owned¹
Sha1	<code>Sha1(data:string block) - string</code>	SHA-1 hash (hex string). Legacy; prefer SHA-256/BLAKE2 for security.	✓ owned¹
Sha256	<code>Sha256(data:string block) - string</code>	SHA-256 hash (hex string).	✓ owned¹
Sha512	<code>Sha512(data:string block) - string</code>	SHA-512 hash (hex string).	✓ owned¹
SipHash24	<code>SipHash24(data:string block, key:string) - int</code>	SipHash-2-4 with 16-byte key.	✓
Xxhash32	<code>Xxhash32(data:string block, seed?:int) - int</code>	xxHash-32. Optional seed.	✓
Xxhash64	<code>Xxhash64(data:string block, seed?:int) - int</code>	xxHash-64. Optional seed.	✓

HttpUtil

Namespace: **HttpUtil**

Small HTTP-utility functions, plus minimal HTTP server. For HTTP-functionality (GET, PUT etc) use the Http-library from Http.fls (shipped as standard-library with source)

```
import { Http, Status, ContentType } from library("Http", "1.0");
```

Function	Signature	Description	JIT
EncodeQuery	<code>EncodeQuery(params:object) - string</code>	Encode object as URL query string (form-urlencoded; + for space). String/int/bool/float values; other value types skipped.	—
UrlEncode	<code>UrlEncode(s:string) - string</code>	Percent-encode a string (RFC 3986, %20 for space). nil for a non-string arg.	—
UrlDecode	<code>UrlDecode(s:string) - string</code>	Decode a percent-encoded string (treats + as space). Malformed %XX is left literal.	—
ParseQuery	<code>ParseQuery(s:string) - object</code>	Parse a query string into an object (percent-decoded, + as space). Keys that decode to empty are skipped.	—
ParseUrl	<code>ParseUrl(url:string) - object</code>	Parse a URL into { scheme, host, port, path, query, fragment }. Handles a missing scheme (defaults to http), no port (80/443 by scheme), IPv6 literals ([::1]:9000), no path (defaults to /), the query, and a #fragment. nil for an empty/ nil arg. Userinfo (user:pass@) is not parsed.	—
ParseHeaders	<code>ParseHeaders(s:string) - object</code>	Parse a CRLF-delimited header block (no status line) into an object. Keys are lowercased for consistent lookup; leading spaces/tabs are trimmed from values. Over-long lines and lines past the internal cap are skipped.	—
DecodeChunked	<code>DecodeChunked(stream:stream) - block</code>	Read HTTP/1.1 chunked transfer-encoding from a stream and return the decoded bytes as a block. Chunk extensions and trailers are ignored. Bounded by the internal response-size cap. nil for a non-stream arg.	—

HTTP Server (requires `--unsafe`):

Function	Signature	Description	JIT
StartServer	StartServer(port:int, handler:fn) – pointer	Start server; handler(request)–response. Returns – server handle.	
HandleRequests	HandleRequests(handle:pointer) – bool	Process one pending request. Call in a loop.	–
WaitReadable	WaitReadable(handle:pointer, timeoutMs?:int) – bool	Park the calling fiber until the listen socket has a pending – connection. Resumes with true on readiness, nil on timeout. timeoutMs <= 0 (default –1) waits without a deadline. Callable from a plain function – no await needed.	
StopServer	StopServer(handle:pointer) – bool	Shutdown server.	–
RouteMatch	RouteMatch(req:object, method:string, pattern:string) – bool	Match req (using its method and path) against the given method + URL pattern. On match, captured path parameters are written into req.params.	–

Response object (returned by the handler): { status:int, body, headers:object }.

- All entries in headers are emitted (e.g. Access-Control-Allow-Origin, Location, Set-Cookie, Cache-Control). Header names and values are validated to block CRLF injection; Content-Type, Content-Length and Connection are managed by the server and skipped if present.
- body may be a **string** or a **block** (bytes). When a block is used, Content-Length is the exact byte count, so binary payloads with embedded NUL bytes (e.g. File.ReadAllBytes for static assets) are sent intact.
- The status line uses the standard reason phrase for the given status (e.g. 404 Not Found), falling back to a per-class generic.
- For a HEAD request the server sends the status line and headers (including the entity's Content-Length) but **no body**.

Server behaviour / limits:

- Each connection is served then closed (Connection: close); there is no keep-alive, HTTP/2, or TLS – for HTTPS put the server behind a TLS-terminating reverse proxy. This is the intended "minimal server" scope; the HTTP client (GET/POST/redirects) lives in Http.fls.
- On POSIX the listen socket sets SO_REUSEADDR, so a restarted server can rebind the port immediately instead of failing during the previous socket's TIME_WAIT window.
- Request intake is bounded (max total size, header count, header/line length) and has both a per-recv timeout and a **total wall-clock deadline**, so a slow client cannot hold the single-threaded server open indefinitely.

Route patterns: :name captures a single segment into params.name. A trailing * is a catch-all matching the remaining path; the matched tail is captured into params["*"] (used by Router.Static).

Json

Namespace: **Json**

JSON parsing, serialization, path-based query and mutation, and NDJSON.

Function	Signature	Description	JIT
Canonicalize	Canonicalize(value:any) – string	Serialize to compact JSON with object (and instance) keys in a deterministic sorted order – for hashing/signing, where equal	–

Function	Signature	Description	JIT
Deserialize	<code>Deserialize(json:string, cls:class) - instance array</code>	values must produce identical bytes (hashmap iteration order is not stable across processes). Keys ordered by UTF-8 byte value. Same cycle/depth guarantees as <code>Stringify</code> . Parse <code>json</code> and map it onto instances of <code>cls</code> by declared field name: an object yields one instance, an array of objects yields an array of instances. Unknown keys are ignored; absent fields keep their class defaults; inherited fields are included. The <code>Constructor</code> is not called - values are written directly to the field slots. Nested objects stay plain objects (fields are not recursively hydrated into their declared class type). Returns <code>nil</code> for malformed JSON, a scalar top level, or an array containing a non-object element.	—
Diff	<code>Diff(a:object, b:object) - object</code>	Produce an RFC 7386 merge-patch such that <code>Merge(copy-of-a, patch)</code> yields <code>b</code> : changed/added keys carry b's value, removed keys map to <code>null</code> , nested objects recurse. <code>nil</code> if either argument is not an object. Cannot represent a key whose value in <code>b</code> is genuinely <code>null</code> (null means delete).	—
Exists	<code>Exists(path:string, value:any) - bool</code>	<code>true</code> if path resolves to a node within <code>value</code> (a parsed object/array, or a class instance for key steps).	✓
FastSelect	<code>FastSelect(path:string, json:string) - any</code>	Fast path extraction from raw JSON text (no full parse). Dot/bracket notation; <code>nil</code> if absent. On duplicate object keys it returns the last (matching <code>Parse</code>).	—
Find	<code>Find(path:string, value:any) - any</code>	Navigate a parsed value by path. Supports wildcards (<code>*</code> , <code>[*]</code>). Returns the matched node or <code>nil</code> . Key steps descend into class instances (by field name); wildcards apply to plain objects/arrays only.	—
IsValid	<code>IsValid(json:string) - bool</code>	<code>true</code> if the string is exactly one well-formed JSON document. Disambiguates the case <code>Parse</code> cannot (it returns <code>nil</code> for both invalid input and a literal <code>null</code>).	✓
Merge	<code>Merge(target:object, patch:object) - bool</code>	Apply an RFC 7386 merge-patch to <code>target</code> in place : a <code>null</code> patch value deletes that key, an object value merges recursively, any other value (including arrays) replaces. <code>false</code> if either argument is not an object. Because <code>null</code> deletes, a patch cannot set a key to <code>null</code> .	—
Parse	<code>Parse(json:string) - any</code>	Parse JSON string into a Flaris value (object/array/string/int/float/bool/nil). <code>nil</code> on malformed input. Duplicate object keys resolve last-wins; a <code>\u0000</code> escape is rejected (Flaris strings cannot hold an embedded NUL). Max depth 256; raises if the document exceeds the internal node cap (<code>MAX_JSON_NODES</code>).	—
ParseLines	<code>ParseLines(text:string) - array</code>	Parse NDJSON / JSON Lines: one JSON value per line. Blank lines skipped; <code>nil</code> if any non-empty line is invalid.	—
Remove	<code>Remove(root:object array, path:string) - bool</code>	Remove the value at <code>path</code> (object key or array index, in place). <code>false</code> on a missing node or type mismatch.	—
Set	<code>Set(root:object array, path:string, value:any) - bool</code>	Set <code>value</code> at <code>path</code> in place, creating missing intermediate objects/arrays. <code>false</code> on a wrong-type intermediate (never clobbered) or an out-of-range array index (only append-at-length allowed).	—
Stringify	<code>Stringify(value:any, pretty?:int) - string</code>	Serialize to JSON. Truthy <code>pretty</code> enables 2-space indenting. Non-finite floats (NaN/Inf) serialize as <code>null</code> . Class instances serialize as objects with every declared field (own + inherited) by name - the inverse of <code>Deserialize</code> . Raises (rather than emit	—

Function	Signature	Description	JIT
		malformed JSON) if <code>value</code> contains a reference cycle or nests past the depth cap.	
StringifyLines	<code>StringifyLines(array:array) - string</code>	Serialize each element as a compact JSON value on its own line (NDJSON); inverse of <code>ParseLines</code> . Raises if any element is cyclic or over-deep.	—

Math

Namespace: **Math**

Full mathematical library. `number` accepts `int` or `float`.

Domain and pole errors (IEEE). Scalar float functions return the IEEE special value on an out-of-domain or pole input rather than `nil`: domain errors yield `NaN` (`Sqrt(-1)`, `Log(-1)`, `Asin(2)`, `Acosh(0.5)`, `Acoth(0.5)`), and poles yield `±Inf` (`Log(0)` → `-Inf`, `Csc(0)` / `Sec` / `Cot` / `Csch` → `±Inf`, `Atanh(1)` → `Inf`, `LogBase(x,1)` → `Inf`). Detect them with `Math.IsNaN` / `Math.IsFinite`.

Function	Signature	Description	JIT
Abs	<code>Abs(x:int float) - int</code>	Absolute value as integer.	✓
AbsF	<code>AbsF(x:int float) - float</code>	Absolute value as float.	✓
Acos	<code>Acos(x:int float) - float</code>	Arc cosine (radians).	✓
Acosh	<code>Acosh(x:int float) - float</code>	Inverse hyperbolic cosine.	✓
Acot	<code>Acot(x:int float) - float</code>	Arc cotangent.	✓
Acoth	<code>Acoth(x:int float) - float</code>	Inverse hyperbolic cotangent.	✓
Asin	<code>Asin(x:int float) - float</code>	Arc sine (radians).	✓
Asinh	<code>Asinh(x:int float) - float</code>	Inverse hyperbolic sine.	✓
Atan	<code>Atan(x:int float) - float</code>	Arc tangent (radians).	✓
Atan2	<code>Atan2(y:int float, x:int float) - float</code>	Two-argument arc tangent.	✓
Atanh	<code>Atanh(x:int float) - float</code>	Inverse hyperbolic tangent.	✓
Cbrt	<code>Cbrt(x:int float) - float</code>	Cube root.	✓
Ceil	<code>Ceil(x:int float) - int</code>	Ceiling (round up).	✓
Clamp	<code>Clamp(x:int float, min:int float, max:int float) - float</code>	Clamp x to [min, max].	✓
CopySign	<code>CopySign(x:int float, y:int float) - float</code>	Returns magnitude of x with sign of y.	✓
Cos	<code>Cos(x:int float) - float</code>	Cosine.	✓
Cosh	<code>Cosh(x:int float) - float</code>	Hyperbolic cosine.	✓
Cot	<code>Cot(x:int float) - float</code>	Cotangent.	✓
Csc	<code>Csc(x:int float) - float</code>	Cosecant.	✓
Csch	<code>Csch(x:int float) - float</code>	Hyperbolic cosecant.	✓
Deg	<code>Deg(x:int float) - float</code>	Convert radians to degrees.	✓
Exp	<code>Exp(x:int float) - float</code>	e^x .	✓
Floor	<code>Floor(x:int float) - int</code>	Floor (round down).	✓
Gcd	<code>Gcd(a:int float, b:int float) - int</code>	Greatest common divisor (≥ 0 ; signs ignored). <code>Gcd(0,0)=0</code> .	✓
Hypot	<code>Hypot(x:int float, y:int float) - float</code>	Euclidean distance $\sqrt{x^2+y^2}$.	✓
IsEven	<code>IsEven(x:int float) - bool</code>	<code>true</code> if x is even.	✓
IsFinite	<code>IsFinite(x:int float) - bool</code>	<code>true</code> if x is finite (not NaN or Inf). Note: return type is float in registry; treat as bool.	✓
IsInteger	<code>IsInteger(x:int float) - bool</code>	<code>true</code> if x has no fractional part. Note: return type is float in registry; treat as bool.	✓

Function	Signature	Description	JIT
IsNaN	<code>IsNaN(x:int float) - bool</code>	true if x is NaN.	✓
IsNegative	<code>IsNegative(x:int float) - bool</code>	true if x < 0.	✓
IsOdd	<code>IsOdd(x:int float) - bool</code>	true if x is odd.	✓
IsPositive	<code>IsPositive(x:int float) - bool</code>	true if x > 0.	✓
Lcm	<code>Lcm(a:int float, b:int float) - int</code>	Least common multiple (≥ 0). <code>Lcm(x,0)=0</code> . May overflow <code>int64</code> for large inputs.	✓
Log	<code>Log(x:int float) - float</code>	Natural logarithm.	✓
Log10	<code>Log10(x:int float) - float</code>	Base-10 logarithm.	✓
LogBase	<code>LogBase(x:int float, base:int float) - float</code>	Logarithm with arbitrary base.	✓
Max	<code>Max(a:int float, b:int float) - int</code>	Larger of a, b as integer.	✓
MaxF	<code>MaxF(a:int float, b:int float) - float</code>	Larger of a, b as float.	✓
MaxOf	<code>MaxOf(arr:array) - int float</code>	Largest element of a numeric array (type preserved). <code>nil</code> if empty or non-numeric.	—
Mean	<code>Mean(arr:array) - float</code>	Arithmetic mean of array elements.	✓
Median	<code>Median(arr:array) - float</code>	Median of array elements.	✓
Min	<code>Min(a:int float, b:int float) - int</code>	Smaller of a, b as integer.	✓
MinF	<code>MinF(a:int float, b:int float) - float</code>	Smaller of a, b as float.	✓
MinOf	<code>MinOf(arr:array) - int float</code>	Smallest element of a numeric array (type preserved). <code>nil</code> if empty or non-numeric.	—
Mod	<code>Mod(a:int float, b:int float) - int float</code>	Euclidean/floored modulo: result takes sign of b, so <code>Mod(-1,5)=4</code> . <code>nil</code> if <code>b=0</code> .	—
Mode	<code>Mode(arr:array) - float</code>	Mode of array elements.	✓
Pow	<code>Pow(base:int float, exp:int float) - float</code>	base^{exp} .	✓
Product	<code>Product(arr:array) - int float</code>	Product of array elements (int if all-int, else float). <code>1</code> if empty; <code>nil</code> if non-numeric.	—
Rad	<code>Rad(x:int float) - float</code>	Convert degrees to radians.	✓
RandBool	<code>RandBool() - bool</code>	Random boolean.	—
RandChoice	<code>RandChoice(arr:array) - any</code>	Random element from array.	—
RandFloat	<code>RandFloat(min:int float, max:int float) - float</code>	Random float in <code>[min, max]</code> .	✓
Random	<code>Random() - float</code>	Random float in <code>[0.0, 1.0)</code> .	—
RandomInt	<code>RandomInt(min:int float, max:int float) - int</code>	Random integer in <code>[min, max]</code> .	✓
RandomNormal	<code>RandomNormal(mean:int float, stddev:int float) - float</code>	Normal (Gaussian) distribution sample.	✓
Round	<code>Round(x:int float) - int</code>	Round to nearest integer (half-up).	✓
RoundTo	<code>RoundTo(x:int float, decimals:int) - float</code>	Round to <code>decimals</code> fractional digits (clamped <code>0..15</code>).	✓
Sec	<code>Sec(x:int float) - float</code>	Secant.	✓
Sech	<code>Sech(x:int float) - float</code>	Hyperbolic secant.	✓
Sign	<code>Sign(x:int float) - int</code>	Returns -1, 0, or 1.	✓
Sin	<code>Sin(x:int float) - float</code>	Sine.	✓
Sinh	<code>Sinh(x:int float) - float</code>	Hyperbolic sine.	✓
Sqrt	<code>Sqrt(x:int float) - float</code>	Square root.	✓
Stddev	<code>Stddev(arr:array) - float</code>	Standard deviation of array elements.	✓
Sum	<code>Sum(arr:array) - int float</code>	Sum of array elements (int if all-int, else float). <code>0</code> if empty; <code>nil</code> if non-numeric.	—
Tan	<code>Tan(x:int float) - float</code>	Tangent.	✓
Tanh	<code>Tanh(x:int float) - float</code>	Hyperbolic tangent.	✓
Trunc	<code>Trunc(x:int float) - float</code>	Truncate toward zero.	✓

Function	Signature	Description	JIT
Variance	Variance(arr:array) – float	Variance of array elements.	✓
Wrap	Wrap(x:int float, min:int float, max:int float) – float	Wrap x into [min, max] range.	✓

Additional scalar functions (less common):

Function	Signature	Description	JIT
Clamp01	Clamp01(x:int float) – float	Clamp to [0.0, 1.0].	✓
Combination	Combination(n:int float, r:int float) – int	Binomial coefficient C(n,r).	✓
Distance	Distance(x1:int float, y1:int float, x2:int float, y2:int float) – float	2D Euclidean distance.	✓
Expm1	Expm1(x:int float) – float	e^x – 1 (accurate near zero).	✓
Factorial	Factorial(n:int float) – int	n!	✓
Fract	Fract(x:int float) – float	Fractional part of x.	✓
IEEERemainder	IEEERemainder(x:int float, y:int float) – float	IEEE 754 remainder.	✓
IsClose	IsClose(a:int float, b:int float) – bool	Approximate equality check.	✓
Lerp	Lerp(a:int float, b:int float, t:int float) – float	Linear interpolation.	✓
Log1p	Log1p(x:int float) – float	ln(1 + x) (accurate near zero).	✓
Log2	Log2(x:int float) – float	Base-2 logarithm.	✓
Permutation	Permutation(n:int float, r:int float) – int	P(n,r) ordered permutations.	✓

Special functions (Gamma, LnGamma, Erf, Erfc run at native speed inside JIT functions):

Function	Signature	Description	JIT
Gamma	Gamma(x:int float) – float	Gamma function $\Gamma(x)$.	✓
LnGamma	LnGamma(x:int float) – float	Natural log of $\Gamma(x)$ (no overflow for large x).	✓
Digamma	Digamma(x:int float) – float	Digamma $\psi(x) = \Gamma'(x)/\Gamma(x)$.	✓
Beta	Beta(a:int float, b:int float) – float	Beta function $B(a,b) = \Gamma(a)\Gamma(b)/\Gamma(a+b)$.	✓
LnBeta	LnBeta(a:int float, b:int float) – float	Natural log of B(a,b).	✓
Erf	Erf(x:int float) – float	Error function.	✓
Erfc	Erfc(x:int float) – float	Complementary error function 1–erf(x).	✓
GammaP	GammaP(a:int float, x:int float) – float	Regularized lower incomplete gamma $P(a,x) \in [0,1]$. χ^2/gamma CDFs.	✓
GammaQ	GammaQ(a:int float, x:int float) – float	Regularized upper incomplete gamma $Q(a,x) = 1-P(a,x)$.	✓
BetaInc	BetaInc(a:int float, b:int float, x:int float) – float	Regularized incomplete beta $I_x(a,b) \in [0,1]$. Student-t/F/beta CDFs.	✓

Bit intrinsics (integer operands; all four run at native speed inside JIT functions):

Function	Signature	Description	JIT
BitCount	BitCount(n:int) – int	Number of set bits in the 64-bit two's-complement representation of n.	✓
BitLength	BitLength(n:int) – int	Minimum bits to represent n (0 for n=0). Equivalent to Python <code>int.bit_length()</code> .	✓
LeadingZeros	LeadingZeros(n:int) – int	Number of leading zero bits in the 64-bit representation (64 when n=0).	✓

Function	Signature	Description	JIT
TrailingZeros	TrailingZeros(n:int) - int	Number of trailing zero bits in the 64-bit representation (64 when n=0).	✓

Matrix operations (operate on 2D arrays of numbers):

Function	Signature	Description	JIT
MatAdd	MatAdd(a:array, b:array) - array	Element-wise matrix addition.	—
MatApply	MatApply(mat:array, vec:array) - array	Multiply matrix by vector.	—
MatDet	MatDet(mat:array) - float	Matrix determinant.	✓
MatIdentity	MatIdentity(n:int float) - array	Create nxn identity matrix.	—
MatInverse	MatInverse(mat:array) - array	Matrix inverse.	—
MatMul	MatMul(a:array, b:array) - array	Matrix multiplication.	—
MatRotate2D	MatRotate2D(angle:int float) - array	2D rotation matrix for angle radians.	—
MatScale2D	MatScale2D(sx:int float, sy:int float) - array	2D scale matrix.	—
MatSub	MatSub(a:array, b:array) - array	Element-wise matrix subtraction.	—
MatTranslate2D	MatTranslate2D(tx:int float, ty:int float) - array	2D translation matrix.	—
MatTranspose	MatTranspose(mat:array) - array	Transpose matrix.	—

Vector operations (operate on flat arrays of numbers):

Function	Signature	Description	JIT
VecAdd	VecAdd(a:array, b:array) - array	Element-wise addition.	—
VecDistance	VecDistance(a:array, b:array) - float	Euclidean distance between two vectors.	✓
VecDot	VecDot(a:array, b:array) - float	Dot product.	✓
VecLength	VecLength(v:array) - float	Vector magnitude.	✓
VecNormalize	VecNormalize(v:array) - array	Unit vector.	—
VecScale	VecScale(v:array, s:int float) - array	Scale vector by scalar.	—
VecSub	VecSub(a:array, b:array) - array	Element-wise subtraction.	—

Constants:

Constant	Value
Math.PI	3.14159265358979323846
Math.TAU	6.28318530717958647692
Math.E	2.71828182845904523536
Math.Infinity	+Inf
Math.NaN	NaN

Range constants: Int8Min/Max, Int16Min/Max, Int32Min/Max, Int64Min/Max, IntMin/Max, UInt8Min/Max, UInt16Min/Max, UInt32Min/Max, UInt64Min/Max, Float32Min/Max/Epsilon, Float64Min/Max/Epsilon, FloatMin/Max/Epsilon, Float32Lowest, Float64Lowest, FloatLowest.

Notes. UInt64Max is capped at Int64Max (9223372036854775807) because constants live in a signed 64-bit slot. For floats, *Min is the smallest *positive normal* (like C's FLT_MIN / DBL_MIN), **not** the most-negative value - use *Lowest (= - *Max) for that.

Random

Namespace: **Random**

Seeded, deterministic pseudo-random streams (PCG32). The same seed always yields the same sequence - use for reproducible tests, simulations, and procedural generation. Each generator is an independent stream; `Math.Random` and friends stay on the global CSPRNG. **Not cryptographic** - use `Crypto.RandomBytes` for keys and tokens.

Function	Signature	Description	JIT
New	<code>New(seed?:int) - object</code>	Create a generator. Without a seed the stream is seeded from the OS CSPRNG — (fast but non-reproducible).	

Generator methods:

Method	Description
<code>r.Next()</code>	Uniform int in <code>[0, 2³²)</code> - one raw PCG32 output.
<code>r.NextInt(min, max)</code>	Uniform int in <code>[min, max]</code> (inclusive, like <code>Math.RandomInt</code>). Unbiased (rejection sampling). <code>nil</code> if <code>min > max</code> .
<code>r.NextFloat()</code>	Uniform float in <code>[0, 1)</code> with 53-bit precision.
<code>r.NextNormal(mean?, stddev?)</code>	Normally distributed float (Box-Muller). Defaults: mean 0, stddev 1.

```
let r = Random.New(42);
r.NextInt(1, 6);      // same die roll on every run
r.NextFloat();        // deterministic [0,1)
```

Memory

Namespace: **Memory**

Unsafe raw memory access. Requires `--unsafe` flag. `ptr = int|string|block|pointer` (any pointer-like value - see legend).

Region requirement. Every access is guarded: the `[ptr, ptr+width)` span must lie inside a **registered** region, otherwise the call returns `nil`/no-op (or raises `Forbidden/Invalid memory-access`). Registered regions come from `Memory.Alloc`, `Memory.Pin`, or a `block`'s backing store (`Buffer.Create/GetAddress`). A plain `string` address is **not** registered, so passing a bare string is rejected — copy it into a `block` (`Buffer.FromString`) or `Pin` its address first.

Function	Signature	Description	JIT
Alloc	<code>Alloc(size:int) - int</code>	Allocate <code>size</code> bytes; returns the raw address. <code>nil</code> if <code>size <= 0</code> or exceeds the 2 GiB per-allocation cap. Caller must <code>Free</code> .	—
BitClear	<code>BitClear(ptr:int string block pointer, bit:int) - nil</code>	Clear bit <code>bit</code> at address.	✓ <code>--unsafe</code>
BitSet	<code>BitSet(ptr:int string block pointer, bit:int) - nil</code>	Set bit <code>bit</code> at address.	✓ <code>--unsafe</code>

Function	Signature	Description	JIT
BitTest	BitTest(ptr:int string block pointer, bit:int) - int	Return 0 or 1 for bit at address.	✓ -- unsafe
BitToggle	BitToggle(ptr:int string block pointer, bit:int) - nil	Toggle bit at address.	✓ -- unsafe
Compare	Compare(a:int string block pointer, b:int string block pointer, len:int) - int	memcmp of len bytes. Returns <0, 0, >0.	✓ -- unsafe
Copy	Copy(dst:int string block pointer, src:int string block pointer, len:int) - nil	memcpy len bytes — the spans must not overlap (use Move if they might). Region-guarded.	✓ -- unsafe
Move	Move(dst:int string block pointer, src:int string block pointer, len:int) - nil	memmove len bytes; the dst and src spans may overlap. Region-guarded.	✓ -- unsafe
Fill	Fill(addr:int string block pointer, value:int, len:int) - nil	Set len bytes at addr to byte value (0–255). Region-guarded; nil/no-op out of region or on a bad value.	✓ -- unsafe
Free	Free(addr:int) - nil	Free previously allocated memory.	—
IsLittleEndian	IsLittleEndian() - bool	true if host is little-endian.	—
Pin	Pin(addr:int, size:int) - int	Register an externally-owned region (e.g. an FFI buffer) so the peek/poke guards accept it; returns addr. Does not allocate or take ownership. Must be paired with Unpin — a region left pinned at exit is reported as a leaked block region (and fails under --mem).	—
Process	Process(addr:int, size:int, chunkSize:int, func:fn) - bool	Iterate memory in chunks calling fn(ptr, len).	—
ProcessCallback	ProcessCallback(fn, address:int, start:int, len:int, cb)	Processes a raw memory range with your worker function fn(address, start, len) (which reads/writes it via Memory.Read*/Write*), then calls cb(address, result) when finished. Runs on another CPU core when fn is simple enough (only Memory.Read*/Write* on that address, plain number math, no calls or allocations) — otherwise runs normally, same result. Don't allocate or free memory while the job is running.	✓ -- unsafe
ProcessEvent	ProcessEvent(fn, address:int, start:int, len:int, eventId:int)	Like ProcessCallback, but signals event eventId with the result instead of calling a callback. Wait for several with Event.WaitFor([ids]).	✓ -- unsafe
Read8	Read8(ptr:int string block pointer, offset?:int) - int	Read 1 byte at ptr+offset.	✓ -- unsafe
Read16	Read16(ptr:int string block pointer, offset?:int) - int	Read 2 bytes (little-endian).	✓ -- unsafe
Read32	Read32(ptr:int string block pointer, offset?:int) - int	Read 4 bytes (little-endian).	✓ -- unsafe
Read64	Read64(ptr:int string block pointer, offset?:int) - int	Read 8 bytes (little-endian).	✓ -- unsafe
ReadBytes	ReadBytes(ptr:int string block pointer, offset?:int) - block	Read a NUL-terminated byte run at ptr (+ offset) into a new block	✓ owned ¹

Function	Signature	Description	JIT
		(excluding the terminator). The scan is bounded to the registered region. <code>nil</code> out of region or with no terminator inside it.	-- unsafe
ReadFloat32	<code>ReadFloat32(ptr:int string block pointer, offset?:int) - float</code>	Read IEEE-754 single.	✓ -- unsafe
ReadFloat64	<code>ReadFloat64(ptr:int string block pointer, offset?:int) - float</code>	Read IEEE-754 double.	✓ -- unsafe
ReadString	<code>ReadString(ptr:int string block pointer, offset?:int) - string</code>	Read a NUL-terminated string from <code>ptr (+ offset)</code> . The scan is bounded to the registered region (never runs off the allocation). <code>nil</code> out of region or with no terminator inside it.	✓ owned¹ -- unsafe
Swap16	<code>Swap16(v:int) - int</code>	Byte-swap 16-bit value.	✓ -- unsafe
Swap32	<code>Swap32(v:int) - int</code>	Byte-swap 32-bit value.	✓ -- unsafe
Swap64	<code>Swap64(v:int) - int</code>	Byte-swap 64-bit value.	✓ -- unsafe
Unpin	<code>Unpin(addr:int) - bool</code>	Remove a region previously registered with <code>Pin</code> ; returns <code>true</code> if it was tracked. Does not free the underlying memory.	--
Write8	<code>Write8(ptr:int string block pointer, value:int, offset?:int) - nil</code>	Write 1 byte.	✓ -- unsafe
Write16	<code>Write16(ptr:int string block pointer, value:int, offset?:int) - nil</code>	Write 2 bytes (little-endian).	✓ -- unsafe
Write32	<code>Write32(ptr:int string block pointer, value:int, offset?:int) - nil</code>	Write 4 bytes (little-endian).	✓ -- unsafe
Write64	<code>Write64(ptr:int string block pointer, value:int, offset?:int) - nil</code>	Write 8 bytes (little-endian).	✓ -- unsafe
WriteBytes	<code>WriteBytes(ptr:int string block pointer, s:string, offset?:int) - nil</code>	Write <code>s</code> 's bytes (WITHOUT a NUL terminator) at <code>ptr (+ offset)</code> . The byte count comes from <code>strlen</code> , so an embedded NUL truncates. Region-guarded.	✓ -- unsafe
WriteFloat32	<code>WriteFloat32(ptr:int string block pointer, value:float, offset?:int) - nil</code>	Write IEEE-754 single.	✓ -- unsafe
WriteFloat64	<code>WriteFloat64(ptr:int string block pointer, value:float, offset?:int) - nil</code>	Write IEEE-754 double.	✓ -- unsafe
WriteString	<code>WriteString(ptr:int string block pointer, s:string, offset?:int) - nil</code>	Write <code>s</code> 's bytes plus a NUL terminator at <code>ptr (+ offset)</code> . Region-guarded.	✓ -- unsafe

Net

Namespace: **Net**

Network utility functions: DNS resolution, IP parsing/conversion, address classification, CIDR math, and connectivity checks. Available on all platforms. Synchronous resolvers (`ResolveDNS` , `ResolveIPv4` , `ResolveIPv6` , `ReverseDNS`) block the single-threaded VM for the duration of the lookup — prefer the `*Async` forms on hot/UI paths.

Function	Signature	Description	JIT
BytesToIP	<code>BytesToIP(bytes:array) - string nil</code>	Convert a 4- (IPv4) or 16-element (IPv6) byte array — to an IP string. <code>nil</code> for any other length.	—
GetHostname	<code>GetHostname() - string nil</code>	Returns machine hostname, or <code>nil</code> on error.	—
GetLocalIP	<code>GetLocalIP() - string nil</code>	Returns primary local IP address (IPv4 preferred; falls back to IPv6). <code>nil</code> if fully offline.	—
IpInRange	<code>IpInRange(ip:string, cidr:string) - bool</code>	<code>true</code> if <code>ip</code> falls inside <code>cidr</code> (e.g. "10.1.2.3", "10.0.0.0/8"). A v4 address is never inside a v6 range. <code>false</code> on malformed input.	—
IpToBytes	<code>IpToBytes(ip:string) - array nil</code>	Convert an IP string to its byte array (4 or 16 ints). <code>nil</code> if not a valid literal.	—
IPVersion	<code>IPVersion(ip:string) - int</code>	4, 6, or 0 if <code>ip</code> is not a valid literal.	—
IsLinkLocal	<code>IsLinkLocal(ip:string) - bool</code>	<code>true</code> for 169.254.0.0/16 or fe80::/10.	—
IsLoopback	<code>IsLoopback(ip:string) - bool</code>	<code>true</code> for 127.0.0.0/8 or ::1.	—
IsMulticast	<code>IsMulticast(ip:string) - bool</code>	<code>true</code> for 224.0.0.0/4 or ff00::/8.	—
IsPrivate	<code>IsPrivate(ip:string) - bool</code>	<code>true</code> for RFC1918 (10/8, 172.16/12, 192.168/16) or IPv6 ULA fc00::/7. IPv4-mapped IPv6 is unwrapped first.	—
IsReachable	<code>IsReachable(host:string, port:int, timeoutMs?:int) - bool</code>	TCP reachability check. Resolves both IPv4 and IPv6 and tries each under one total deadline. <code>timeoutMs</code> defaults to 1000 (clamped to a sane range).	—
IsReachableAsync	<code>IsReachableAsync(host:string, port:int, timeoutMs?:int) - fiber</code>	Non-blocking <code>IsReachable</code> . Use with <code>await</code> ; resolves to a <code>bool</code> .	—
IsValidIP	<code>IsValidIP(s:string) - bool</code>	<code>true</code> if string is a valid IPv4 or IPv6 literal.	—
IsValidIPv4	<code>IsValidIPv4(s:string) - bool</code>	<code>true</code> only for a valid IPv4 literal.	—
IsValidIPv6	<code>IsValidIPv6(s:string) - bool</code>	<code>true</code> only for a valid IPv6 literal.	—
ParseCIDR	<code>ParseCIDR(cidr:string) - object nil</code>	Parse <code>addr/prefix</code> into { <code>network</code> , <code>prefix</code> , <code>version</code> } (plus <code>netmask</code> , <code>broadcast</code> for IPv4). <code>nil</code> on malformed input or out-of-range prefix.	—
ResolveDNS	<code>ResolveDNS(host:string) - array</code>	Resolve hostname to array of IP strings (IPv4 and IPv6), capped at 32. Empty array on failure.	—
ResolveDNSAsync	<code>ResolveDNSAsync(host:string) - fiber</code>	Non-blocking <code>ResolveDNS</code> . Use with <code>await</code> . Other fibers run during the DNS lookup.	—
ResolveIPv4	<code>ResolveIPv4(host:string) - array</code>	Resolve to array of IPv4 address strings.	—
ResolveIPv4Async	<code>ResolveIPv4Async(host:string) - fiber</code>	Non-blocking <code>ResolveIPv4</code> . Use with <code>await</code> . Other fibers run during the DNS lookup.	—
ResolveIPv6	<code>ResolveIPv6(host:string) - array</code>	Resolve to array of IPv6 address strings.	—
ResolveIPv6Async	<code>ResolveIPv6Async(host:string) - fiber</code>	Non-blocking <code>ResolveIPv6</code> . Use with <code>await</code> . Other fibers run during the DNS lookup.	—
ReverseDNS	<code>ReverseDNS(ip:string) - string nil</code>	Reverse lookup: IP → hostname. <code>nil</code> on a bad literal / no PTR record.	—
ReverseDNSAsync	<code>ReverseDNSAsync(ip:string) - fiber</code>	Non-blocking <code>ReverseDNS</code> . Use with <code>await</code> ; resolves to the hostname string, or <code>nil</code> on a bad literal / no PTR record.	—

Object

Namespace: **Object**

Introspection and manipulation of `object` values, plus read-only reflection over `instance` and `class` values.

`Keys`, `Values`, `Count` and `Entries` accept an `instance` or `class` and reflect its declared **fields** (base-first, in declaration order): the field names, an instance's current field values (a class's field defaults), and the field count. The mutating and higher-order helpers (`Merge`, `Delete`, `Clear`, `Select`, `Where`, `Pick`, `MapKeys`, `Invert`, `Reduce`, `Any`, `All`) are **object-only** — a class's schema is sealed.

Property iteration order follows the hashmap's bucket order, which is salted per process (hash-flood defense) and therefore **not stable across runs**; use `Collections.OrderedMap` when insertion order matters. Field reflection order is the class's declaration order.

Function	Signature	Description	JIT
Clear	<code>Clear(obj:object) - bool</code>	Remove all properties from <code>obj</code> . Mutates. Raises if <code>obj</code> is — frozen.	
Clone	<code>Clone(value:any) - any</code>	Deep clone: scalars, strings, arrays (including flat <code>[float]</code>), objects, classes and instances are copied recursively; functions, builtins, modules and exceptions are shared (ref-counted, not copied). A cyclic value (or nesting past the clone-depth limit) raises <code>Exception.NestingError</code> rather than looping forever.	—
Count	<code>Count(obj:object instance class) - int</code>	Number of properties (object) or declared fields (instance/class).	✓
Delete	<code>Delete(obj:object, key:string) - bool</code>	Delete property <code>key</code> . Returns <code>true</code> if existed. Raises if <code>obj</code> is frozen.	—
Entries	<code>Entries(obj:object instance class) - array</code>	Array of <code>[key, value]</code> pairs (object properties, or field name → value).	—
Freeze	<code>Freeze(obj:object) - bool</code>	Makes <code>obj</code> read-only. Any subsequent write throws — including via <code>Merge</code> / <code>Delete</code> / <code>Clear</code> . Shallow only. Returns <code>false</code> for a non-object. Query with <code>IsFrozen</code> .	—
FromEntries	<code>FromEntries(pairs:array) - object</code>	Inverse of <code>Entries</code> : builds an object from an array of <code>[key, value]</code> pairs (string keys). A later duplicate key overwrites the earlier. <code>nil</code> for a malformed entry.	—
FromKeys	<code>FromKeys(arr:array, value:any) - object</code>	Creates new object from string key array, all set to <code>value</code> .	—
All	<code>All(obj:object, func:fn) - bool</code>	Returns <code>true</code> if <code>fn(key, value)</code> is truthy for every entry. Returns <code>true</code> for an empty object.	—
Any	<code>Any(obj:object, func:fn) - bool</code>	Returns <code>true</code> if <code>fn(key, value)</code> is truthy for at least one entry. Returns <code>false</code> for an empty object.	—
HasKey	<code>HasKey(obj:object, key:string) - bool</code>	<code>true</code> if property <code>key</code> exists.	✓
Invert	<code>Invert(obj:object) - object</code>	Returns a new object with keys and values swapped. Values are coerced to string to become keys.	—
IsFrozen	<code>IsFrozen(obj:object) - bool</code>	<code>true</code> if <code>obj</code> was frozen by <code>Freeze</code> . <code>false</code> for an unfrozen object or a non-object.	✓
IsNil	<code>IsNil(val:any) - bool</code>	<code>true</code> if <code>val</code> is <code>nil</code> . Single-arg; usable as a fast-path callback to <code>Array.Where</code> , <code>Array.Any</code> , <code>Array.GroupBy</code> , etc.	✓
Keys	<code>Keys(obj:object instance class) - array</code>	Array of property names (object) or field names (instance/class).	—
MapKeys	<code>MapKeys(obj:object, func:fn) - object</code>	Returns a new object with keys transformed by <code>fn(key)</code> . Values are kept; new keys are coerced to string. Builtin functions (e.g. <code>String.ToUpper</code>) use the fast path.	—
Merge	<code>Merge(dst:object, src:object) - bool</code>	Copy all <code>src</code> properties into <code>dst</code> . Mutates <code>dst</code> . Shallow. Raises if <code>dst</code> is frozen.	—

Function	Signature	Description	JIT
NotNil	<code>NotNil(val:any) - bool</code>	true if <code>val</code> is not <code>nil</code> . Single-arg; usable as a fast-path ✓ callback to <code>Array.Where</code> , <code>Array.Any</code> , <code>Array.GroupBy</code> , etc.	
Pick	<code>Pick(obj:object, keys:array) - object</code>	Returns new object with only the listed keys.	—
Reduce	<code>Reduce(obj:object, func:fn, initial:any) - any</code>	Folds over entries: <code>acc = fn(acc, key, value)</code> starting from <code>initial</code> .	—
Select	<code>Select(obj:object, func:fn) - object</code>	Returns new object with values transformed by <code>fn(key, value)</code> .	—
TryGet	<code>TryGet(obj:object, key:string, default:any) - any</code>	Returns <code>obj[key]</code> if present, otherwise <code>default</code> .	—
Values	<code>Values(obj:object instance class) - array</code>	Array of property values (object) or field values (instance/class).	—
Where	<code>Where(obj:object, func:fn) - object</code>	Returns new object containing only keys where <code>fn(key, value)</code> is truthy.	—

Object instance method syntax

All `Object` functions where the object is the first argument can be called directly on an `object` value:

```
var o = { name: "Alice", age: 30 };
o.Keys()           // ["name", "age"] - same as Object.Keys(o)
o.Values()         // ["Alice", 30]
o.HasKey("age")    // true
o.Count()          // 2
o.Delete("age")    // same as Object.Delete(o, "age")
o.Merge({ city: "Stockholm" }) // same as Object.Merge(o, ...)
o.Where(fn(k, v) { return v != nil; })
o.Select(fn(k, v) { return String.ToUpper(v); })
```

Functions that do not take an object as first argument (`FromKeys`, `FromEntries`) are not available as instance methods. Both forms compile to identical bytecode when the variable is typed (`: object` or inferred).

Untyped variables fall back to a runtime dispatch.

Type

Namespace: **Type**

Integer constants representing every runtime type. Used with the `type()` built-in, which returns the type of a value as an `int` bitmask.

```
type(42)          == Type.Int;           // true
type("hello")    == Type.String;        // true
(type(x) & Type.Callable) != 0; // true for functions, builtins, bound methods, FFI
(type(x) & Type.Number) != 0;  // true for int or float

Type.Is(x, Type.Callable);           // true for any function type
```

```

Type.Is(x, Type.Int | Type.Float);          // ad-hoc union without a named composite

Type.Name(42);                             // "int"
Type.Name([]);                             // "array"

Type.Assert(x, Type.Int);                  // throws if x is not int
Type.Assert(x, Type.Number, "expected a number");

```

Functions

Function	Signature	Description	JIT
Assert	Assert(x:any, mask:int, msg?:string) - bool	Throws InvalidArgs if type(x) does not match mask. Optional custom message. Returns true on success.	✓
Is	Is(x:any, mask:int) - bool	Returns true if type(x) matches any bit in mask. Use with composite constants.	✓
IsArray	IsArray(x:any) - bool	true if x is an array. Single-arg; usable as a fast-path callback to Array.Where, Array.Select, Array.GroupBy, etc.	✓
IsBlock	IsBlock(x:any) - bool	true if x is a memory block/typed buffer.	✓
IsBool	IsBool(x:any) - bool	true if x is a bool.	✓
IsCallable	IsCallable(x:any) - bool	true if x is any callable: function, builtin, bound method , or FFI. Superset of IsFunction.	✓
IsChar	IsChar(x:any) - bool	true if x is a character.	✓
IsClass	IsClass(x:any) - bool	true if x is a class definition. For "is x an instance of class C" use Class.InstanceOf.	✓
IsException	IsException(x:any) - bool	true if x is an exception object.	✓
IsFiber	IsFiber(x:any) - bool	true if x is a fiber.	✓
IsFloat	IsFloat(x:any) - bool	true if x is a float.	✓
IsFunction	IsFunction(x:any) - bool	true if x is a plain function, builtin, or FFI function. Does not include bound methods - use IsCallable for those.	✓
IsInstance	IsInstance(x:any) - bool	true if x is a class instance (of any class). For a specific class use Class.InstanceOf.	✓
IsInt	IsInt(x:any) - bool	true if x is an integer.	✓
IsModule	IsModule(x:any) - bool	true if x is a built-in module.	✓
IsNil	IsNil(x:any) - bool	true if x is nil.	✓
IsNumber	IsNumber(x:any) - bool	true if x is an integer or float.	✓
IsObject	IsObject(x:any) - bool	true if x is a plain object.	✓
IsStream	IsStream(x:any) - bool	true if x is a file/network stream.	✓
IsString	IsString(x:any) - bool	true if x is a string.	✓
Name	Name(x:any) - string	Returns the type as a human-readable string: "int", "array", "fiber", etc. Single-arg; usable as a callback.	✓ owned¹
NameOf	NameOf(code:int) - string	Returns the readable name of a Type.* code or a composite mask, e.g. NameOf(Type.Number) - "int float". The inverse of Of.	✓ owned¹
Of	Of(x:any) - int	Returns the type bitmask of x - the referenceable equivalent of the type() built-in. Single-arg; usable as a callback (e.g. Array.GroupBy(items, Type.Of)).	✓

The numeric values below are implementation-defined bit flags. Always compare against the Type.* constants (type(x) == Type.Int, Type.Is(x, Type.Number)), never against hardcoded hex - the underlying bits may change between releases.

Primitive Types

Constant	Value	Matches
Type.Nil	0x000001	nil
Type.Bool	0x000080	true, false
Type.Char	0x000002	character literals
Type.Int	0x000008	integers
Type.Float	0x000004	floating-point numbers
Type.String	0x000010	strings

Collection Types

Constant	Value	Matches
Type.Array	0x000040	arrays
Type.Object	0x000020	plain objects

Callable Types

Constant	Value	Matches
Type.Function	0x0010000	user-defined functions
Type.BuiltinFunction	0x0020000	built-in functions
Type.FFIFunction	0x0800000	FFI-loaded native functions
Type.BoundMethod	0x0100000	bound instance methods

Object-Oriented Types

Constant	Value	Matches
Type.Class	0x1000000	class definitions
Type.Instance	0x2000000	class instances
Type.Module	0x0040000	built-in modules

Runtime Types

Constant	Value	Matches
Type.Fiber	0x0400000	fibers
Type.Exception	0x0080000	exception objects
Type.Stream	0x4000000	file/network streams
Type.Block	0x0000100	memory blocks
Type.Pointer	0x0200000	raw pointers (internal)

Composite Constants

These are bitmask unions - use `&` not `==`.

Constant	Matches
Type.Number	int or float
Type.Numeric	int, float, bool, or char
Type.Callable	any callable: fn, builtin, bound method, FFI
Type.ObjectLike	object, instance, class, or module

Os

Namespace: **Os**

Process, environment, and system interface. Functions marked **unsafe** require the `--unsafe` flag.

Function	Signature	Description	JIT
Arch	<code>Arch()</code> – string	Returns the CPU architecture: "x86_64", "x86", "mips", "arm64", "arm", or "unknown".	—
Args	<code>Args()</code> – array	Returns CLI argument strings as passed to the script.	—
Bits	<code>Bits()</code> – int	Returns the pointer width of the platform in bits. Flaris ships 64-bit-only, so this currently always returns 64.	—
Chdir	<code>Chdir(path:string)</code> – bool	Change working directory. Process-global - affects every fiber.	—
GetCwd	<code>GetCwd()</code> – string	Returns the current working directory, or nil if it cannot be read (e.g. the directory was removed, or its path exceeds 4096 bytes). Pairs with <code>Chdir</code> .	—
Exec	<code>Exec(path:string, args?:array)</code> – nil	Replace the current process image with <code>path</code> via <code>execvp</code> . Only returns on error. Does not require <code>--unsafe</code> .	—
Execute	<code>Execute(cmd:string)</code> – string	Run a shell command via <code>popen</code> ; returns stdout+stderr merged as a string. Returns nil on <code>popen</code> failure, output exceeding 16 MB, or a <code>pclose</code> error; the command's own non-zero exit status does not map to nil (the captured output is returned regardless). Does not require <code>--unsafe</code> . The command string is passed directly to the shell with no escaping or sanitization - never pass untrusted or externally-supplied input (command-injection risk). To run a program with separate arguments and no shell, use <code>Os.Spawn / Os.RunEx</code> (which capture output via <code>execvp</code>) or <code>Os.Exec</code> (which replaces the image).	—
ExecuteAsync	<code>ExecuteAsync(cmd:string)</code> – fiber	Non-blocking <code>Execute</code> . Use with <code>await</code> . Other fibers run while the command runs. Same security warning as <code>Execute</code> applies.	—
Getenv	<code>Getenv(name:string)</code> – string	Read environment variable. Returns nil if not set.	—
GetArgValue	<code>GetArgValue(name:string, default?:string)</code> – string	Extract value from CLI arg matching <code>name=value</code> . Returns <code>default</code> (or nil) if not found.	—
Gid	<code>Gid()</code> – int	Returns the real group ID of the process.	—
IsRoot	<code>IsRoot()</code> – bool	true if the process is running as root (UID 0).	—
Kill	<code>Kill(pid:int, signal:int)</code> – bool	Send a signal to a process. Use standard signal numbers (e.g. 15 for SIGTERM, 9 for SIGKILL). Requires <code>--unsafe</code> .	—
Name	<code>Name()</code> – string	Returns OS name: "Windows", "Linux" or "macOS".	—
Pid	<code>Pid()</code> – int	Returns the current process ID.	—
Ppid	<code>Ppid()</code> – int	Returns the parent process ID.	—
Setenv	<code>Setenv(name:string, value:string)</code> – bool	Set environment variable.	—
Sleep	<code>Sleep(ms:int)</code> – bool	Sleep current OS thread for <code>ms</code> milliseconds. Blocks all fibers - prefer <code>Fiber.Sleep</code> .	—
TempDir	<code>TempDir()</code> – string	Returns path to system temp directory.	—
Uid	<code>Uid()</code> – int	Returns the real user ID of the process.	—
Wait	<code>Wait(pid:int)</code> – int	Wait (blocking) for child process to exit. Returns its exit code, the negative signal number if killed by a signal, or nil on <code>waitpid</code> error.	—
Spawn	<code>Spawn(cmd:string, args?:array)</code> – object	Fork and exec <code>cmd</code> with separate stdin/stdout/stderr pipes. Returns <code>{Pid:int, Stdin:int, Stdout:int, Stderr:int}</code> where the int values are raw file descriptors for use with	—

Function	Signature	Description	JIT
		ReadPipe/WritePipe/ClosePipe. Returns <code>nil</code> on error. Stdout and stderr fds are non-blocking.	
ReadPipe	<code>ReadPipe(fd:int) - string</code>	Non-blocking read from a pipe fd returned by <code>Spawn</code> . Returns <code>nil</code> when no data is available or the pipe is closed.	—
WritePipe	<code>WritePipe(fd:int, data:string) - int</code>	Write data to a pipe fd (stdin of a spawned process). Returns bytes written, or <code>-1</code> on error.	—
ClosePipe	<code>ClosePipe(fd:int) - bool</code>	Close a pipe fd. Call on stdin to signal EOF to the child process.	—
IsAlive	<code>IsAlive(pid:int) - bool</code>	Returns <code>true</code> if the process with <code>pid</code> is still running. Non-blocking. Does not reap the child - safe to call before <code>TryWait/Wait</code> .	—
TryWait	<code>TryWait(pid:int) - object</code>	Non-blocking wait. Returns <code>{Alive:bool, Exit:int}</code> . When <code>Alive</code> is <code>false</code> the child has been reaped and <code>Exit</code> holds the exit code. Prefer over polling <code>IsAlive + Wait</code> to avoid double-reap.	—
KillChild	<code>KillChild(pid:int, signal:int) - bool</code>	Send signal to a process previously spawned via <code>Os.Spawn</code> . Does not require <code>--unsafe</code> . Returns <code>false</code> if <code>pid</code> was not spawned by this VM instance. For sending signals to arbitrary PIDs use <code>Os.Kill</code> (requires <code>--unsafe</code>).	—
RunEx	<code>RunEx(cmd:string, args?:array) - object</code>	Run <code>cmd</code> and wait for it to finish. Returns <code>{Exit:int, Stdout:string, Stderr:string}</code> with stdout and stderr captured separately. Returns <code>nil</code> on fork/spawn failure.	—
GetEnvAll	<code>GetEnvAll() - object</code>	Returns all environment variables as an object <code>{NAME: "value", ...}</code> .	—

Path

Namespace: **Path**

Cross-platform path string manipulation.

Function	Signature	Description	JIT
ChangeExtension	<code>ChangeExtension(path:string, ext:string) - string</code>	Replace file extension.	—
Combine	<code>Combine(part1:string, part2:string, ...more) - string</code>	Join path segments with separator. Up to 15 arguments.	—
GetCurrentDir	<code>GetCurrentDir() - string</code>	Returns current working directory.	—
GetDirectoryName	<code>GetDirectoryName(path:string) - string</code>	Returns directory portion of path.	—
GetExtension	<code>GetExtension(path:string) - string</code>	Returns extension including dot (e.g. <code>".txt"</code>).	—
GetFileName	<code>GetFileName(path:string) - string</code>	Returns filename with extension.	—
GetFileNameWithoutExtension	<code>GetFileNameWithoutExtension(path:string) - string</code>	Returns filename without extension.	—
GetFullPath	<code>GetFullPath(path:string) - string</code>	Resolve to absolute path (lexical <code>./</code> tidy; does not resolve <code>..</code> - use <code>Resolve</code>).	—
GetRelativePath	<code>GetRelativePath(from:string, to:string) - string</code>	Relative path that reaches to when walked from directory from (using <code>../</code> to climb). Both are resolved to absolute canonical form first. <code>""</code> when equal. Pure lexical; <code>nil</code> if <code>getcwd</code> fails.	—
GetRoot	<code>GetRoot(path:string) - string</code>	Returns <code>"/"</code> for a POSIX-absolute path, <code>""</code> otherwise	—

Function	Signature	Description	JIT
HasExtension	<code>HasExtension(path:string) - bool</code>	(Windows drive roots are not recognized). <code>true</code> if path has an extension. —	
IsAbsolute	<code>IsAbsolute(path:string) - bool</code>	<code>true</code> for a POSIX root (/), a Windows drive root (C:\ / C:/), a leading backslash, or a UNC path (\\server). Pure lexical test - no filesystem access. —	
Normalize	<code>Normalize(path:string) - string</code>	Convert \ to / , collapse // runs and drop ./ segments. Does not resolve .. or touch the filesystem. —	
Resolve	<code>Resolve(path:string) - string</code>	Absolute path with . and .. resolved (relative inputs anchored to the current directory). Unlike Normalize/GetFullPath this collapses .. , giving a canonical form for containment checks. Purely lexical - does not resolve symlinks or require the path to exist. <code>nil</code> if <code>getcwd</code> fails. —	
TempFile	<code>TempFile(prefix?:string) - string</code>	Create a unique empty file in the system temp dir (<code>mkstemp</code>) and return its path. Optional name prefix (default "flaris") must not contain path separators. The file is not auto-deleted. <code>nil</code> on failure. —	

Regex

Namespace: **Regex**

Regular expression matching, search, replace, and split. The engine is a lightweight custom implementation supporting the syntax documented below. Patterns are matched against UTF-8 strings, but character-level operations are byte-based.

Functions

Function	Signature	Description	JIT
IsMatch	<code>IsMatch(input:string, pattern:string) - bool</code>	<code>true</code> if pattern matches anywhere in input.	✓
Match	<code>Match(input:string, pattern:string) - string</code>	Returns first match substring, or <code>nil</code> .	✓ owned ¹
Matches	<code>Matches(input:string, pattern:string) - array</code>	Returns all non-overlapping match substrings.	—
Capture	<code>Capture(input:string, pattern:string) - array</code>	Returns <code>[fullMatch, group1, group2, ...]</code> for the first match, or <code>nil</code> if no match. Element 0 is the whole match; subsequent elements are the capture groups in order of their opening (. A group that did not participate in the match (e.g. an unmatched optional group) is <code>nil</code> .	—

Function	Signature	Description	JIT
Replace	Replace(input:string, pattern:string, replacement:string) - string	Replace all matches with replacement string.	✓ owned ¹
ReplaceFn	ReplaceFn(input:string, pattern:string, fn:fn) - string	Replace each match with the result of <code>fn(match)</code> . The callback — receives the matched substring and must return a string (<code>nil</code> deletes the match). If the callback raises, or returns a value that is neither a string nor <code>nil</code> , the exception propagates to the caller (wrap the call in <code>try/catch</code> to handle it).	—
Split	Split(input:string, pattern:string) - array	Split input at pattern boundaries. Always includes remainder as last element.	—

Supported syntax

Anchors

Syntax	Description
<code>^</code>	Match start of string
<code>\$</code>	Match end of string

Wildcards and quantifiers

Syntax	Description
<code>.</code>	Any character except <code>\n</code> and <code>\r</code>
<code>*</code>	Zero or more (greedy)
<code>+</code>	One or more (greedy)
<code>?</code>	Zero or one
<code>{n}</code>	Exactly <code>n</code> repetitions
<code>{n,}</code>	At least <code>n</code> repetitions
<code>{n,m}</code>	Between <code>n</code> and <code>m</code> repetitions (inclusive)

Character classes

Syntax	Description
<code>[abc]</code>	Match any of <code>a</code> , <code>b</code> , <code>c</code>
<code>[^abc]</code>	Match anything not in set
<code>[a-z]</code>	Character range
<code>[a-zA-Z0-9]</code>	Multiple ranges
<code>[[:space:]]</code>	POSIX named class (see below)

POSIX named classes (for use inside `[...]`)

Class	Equivalent	Description
<code>[:alpha:]</code>	<code>[a-zA-Z]</code>	Letters
<code>[:digit:]</code>	<code>[0-9]</code>	Decimal digits
<code>[:alnum:]</code>	<code>[a-zA-Z0-9_]</code>	Letters, digits, underscore
<code>[:space:]</code>	<code>[\t\n\r\f\v]</code>	Whitespace
<code>[:blank:]</code>	<code>[\t]</code>	Space and tab only
<code>[:upper:]</code>	<code>[A-Z]</code>	Uppercase letters
<code>[:lower:]</code>	<code>[a-z]</code>	Lowercase letters
<code>[:punct:]</code>		Punctuation characters
<code>[:print:]</code>		Printable characters (including space)

Class	Equivalent	Description
<code>[:graph:]</code>		Printable characters (excluding space)
<code>[:cntrl:]</code>		Control characters
<code>[:xdigit:]</code>	<code>[0-9a-fA-F]</code>	Hexadecimal digits

Escape sequences

Syntax	Description
<code>\d</code>	Digit <code>[0-9]</code>
<code>\D</code>	Non-digit
<code>\w</code>	Word character <code>[a-zA-Z0-9_]</code>
<code>\W</code>	Non-word character
<code>\s</code>	Whitespace
<code>\S</code>	Non-whitespace
<code>\t</code>	Tab character
<code>\n</code>	Newline character
<code>\r</code>	Carriage return
<code>\b</code>	Word boundary (zero-width)
<code>\B</code>	Non-word boundary (zero-width)
<code>\.</code> <code>*</code> etc.	Escaped literal character

Alternation

Syntax	Description
<code>a b</code>	Match <code>a</code> or <code>b</code> . Left branch is tried first. Multiple alternatives supported: <code>a b c</code> .

Grouping and capture

Syntax	Description
<code>(...)</code>	Capturing group. Scopes alternation/quantifiers (e.g. <code>(ab)+</code> , <code>(cat dog)s</code>) and records the matched substring for <code>Regex.Capture</code> . Groups are numbered 1, 2, ... by the position of their opening <code>(</code> .
<code>(?:...)</code>	Non-capturing group. Groups for alternation/quantifiers without producing a capture.

Groups may be nested and quantified. With `Regex.Capture`, a quantified group reports the substring of its **last** iteration (e.g. `(ab)+` over `"abab"` captures `"ab"`).

Flags

Syntax	Description
<code>(?i)</code>	Case-insensitive matching. Must appear at the start of the pattern.

Examples

```
// Quantifiers
Regex.IsMatch("year 2024", "[0-9]{4}") // true
Regex.IsMatch("123", "^([0-9]{2,4})$") // true

// Alternation
Regex.IsMatch("my dog", "cat|dog") // true
Regex.Match("I have a fish", "cat|dog|fish") // "fish"
```

```
// Capture groups
Regex.Capture("2024-12-31", "(\\d+)-(\\d+)-(\\d+)") // ["2024-12-31", "2024", "12", "31"]
Regex.Capture("name: John", "(\\w+): (\\w+)") // ["name: John", "name", "John"]
Regex.Capture("foo", "(a)?(foo)") // ["foo", nil, "foo"]
Regex.Capture("nope", "(\\d+)") // nil

// POSIX named classes
Regex.Split("a b c", "[[:space:]]+") // ["a", "b", "c"]
Regex.Replace("hello!", "[[:punct:]]+", "") // "hello"

// Escape sequences
Regex.IsMatch("a\\tb", "a\\tb") // true
Regex.Split("line1\\nline2\\nline3", "\\n") // ["line1", "line2", "line3"]

// Word boundaries
Regex.IsMatch("concatenate", "\\bcat\\b") // false
Regex.IsMatch("my cat", "\\bcat\\b") // true
Regex.Matches("x1 22 333", "\\b\\d+\\b") // ["1", "22", "333"]

// Case-insensitive
Regex.IsMatch("HELLO WORLD", "(?i)hello") // true
Regex.Match("HELLO", "(?i)[a-z]+") // "HELLO"
Regex.Replace("F00 bar foo", "(?i)foo", "baz") // "baz bar baz"
```

Limitations

- No backreferences (`\1`, `\2`, etc.)
- No lookahead or lookbehind
- No non-greedy quantifiers (`*?`, `+?`) - all repetitions are greedy
- No Unicode-aware matching - character classes and ranges operate on bytes
- `(?i)` must be at the start of the pattern; inline flags like `foo(?i)bar` are not supported
- Up to 31 capture groups per pattern; pattern length is limited to 511 bytes

Matching semantics

The engine finds the **leftmost** match. Among matches starting at the same position, alternation is leftmost-first (`a|b` prefers `a`) and quantifiers are greedy. For multi-alternative patterns, the match starting at the earliest position in the input wins regardless of branch order - e.g. `Regex.Match("dog cat", "cat|dog")` returns `"dog"`. Matching is linear in the length of the input (no catastrophic backtracking).

Scheduler

Namespace: `Scheduler`

Low-level fiber scheduling primitives.

Function	Signature	Description	JIT
HasReadyFibers	<code>HasReadyFibers() - bool</code>	<code>true</code> if any fiber is ready to run.	—
Schedule	<code>Schedule(f: fiber) - nil</code>	Enqueue fiber <code>f</code> for next scheduler tick.	—

Stream

Namespace: **Stream**

Unified I/O for files, network sockets, pipes, and serial ports. All functions operate on a `stream` value backed by a file descriptor.

Function	Signature	Description	JIT
Accept	<code>Accept(s: stream) - stream</code>	Accept an incoming connection on a listening socket. Returns <code>nil</code> if no connection is pending.	—
Close	<code>Close(s: stream) - bool</code>	Close the stream. Flushes first.	—
Connect	<code>Connect(proto: string, host: string, port: int) - stream</code>	Open socket. <code>proto</code> : <code>"tcp"</code> or <code>"udp"</code> . <code>host</code> : hostname or IP (v4/v6). Returns <code>nil</code> on failure.	—
ConnectAsync	<code>ConnectAsync(proto: string, host: string, port: int) - fiber</code>	Non-blocking <code>Connect</code> : the DNS lookup and TCP handshake are offloaded to the io pool, so the fiber suspends instead of stalling the VM. Use with <code>await</code> ; resolves to a socket stream, or <code>nil</code> on failure. Raises on a bad <code>proto</code> / <code>port</code> .	—
Copy	<code>Copy(src: stream, dst: stream, limit?: int) - bool</code>	Copy data from <code>src</code> to <code>dst</code> . Optional byte limit.	—
Flush	<code>Flush(s: stream) - bool</code>	Flush write buffer (<code>fsync</code> for files, <code>tcdrain</code> for serial, no-op for sockets).	—
IsOpen	<code>IsOpen(s: stream) - bool</code>	<code>true</code> if stream is still open.	—
LastError	<code>LastError() - int</code>	Why the calling fiber's most recent async operation (<code>ReadAsync</code> , <code>WriteAsync</code> , <code>WaitReadable</code>) resolved to <code>nil</code> . <code>0</code> = it succeeded, so a <code>nil</code> with <code>LastError() == 0</code> means a clean end-of-file, not a failure. Per-fiber, so concurrent fibers never overwrite each other's reason. Codes: <code>1</code> timeout, <code>2</code> stream closed/invalid, <code>3</code> connection reset, <code>4</code> read exceeded the maximum block size, <code>5</code> descriptor cannot be multiplexed, <code>6</code> other OS error, <code>7</code> cancelled by <code>Fiber.CancelIo</code> .	—
Listen	<code>Listen(proto: string, port: int, backlog?: int) - stream</code>	Create a listening socket. <code>proto</code> : <code>"tcp"</code> or <code>"udp"</code> . Binds a dual-stack IPv6 socket (<code>IPV6_V6ONLY=0</code> , so IPv4 clients connect too) and falls back to IPv4-only if v6 is unavailable. Sets <code>SO_REUSEADDR</code> and <code>SO_REUSEPORT</code> . <code>backlog</code> defaults to <code>128</code> .	—
LocalAddr	<code>LocalAddr(s: stream) - string</code>	Local endpoint of a socket as <code>"ip:port"</code> (<code>"[ip]:port"</code> for IPv6). <code>nil</code> for a non-socket or on error.	—
Open	<code>Open(path: string, mode: string) - stream</code>	Open file stream. Mode: <code>"r"</code> (read-only), <code>"w"</code> (write/create/truncate), or <code>"rw"</code> (read-write/create). Returns <code>nil</code> on failure.	—
OpenSerial	<code>OpenSerial(port: string, baud: int) - stream</code>	Open serial port in raw mode. Valid baud rates: <code>9600</code> , <code>19200</code> , <code>38400</code> , <code>57600</code> , <code>115200</code> .	—
Peek	<code>Peek(s: stream) - int</code>	Return the next byte (0–255) without consuming it. Sockets use <code>MSG_PEEK</code> ; files read one byte and rewind. Returns <code>nil</code> at EOF, on error, or for a pipe/serial stream (which have no non-destructive read).	—

Function	Signature	Description	JIT
PeerAddr	PeerAddr(s:stream) – string	Remote endpoint of a connected socket as "ip:port" ("[ip]:port" for IPv6). nil for a non-socket or on error.	—
Pipe	Pipe() – array	Create a pipe. Returns [readStream, writeStream].	—
ReadAll	ReadAll(s:stream, limit?:int, timeout?:int) – block	Read until EOF or limit bytes and return a block. On a file a short read ends the read; on a socket/pipe it reads until the peer closes (EOF) or, for a socket, the timeout fires (seconds, default 3) — bytes already read are returned, not discarded. Returns nil only on a hard I/O error.	—
ReadAsync	ReadAsync(s:stream, count?:int, timeout?:int, cb?:fn) – fiber	Async read. count > 0 completes once exactly count bytes have arrived; count omitted or 0 reads until EOF. await resolves to a block, or nil. A block can be shorter than count when EOF arrives first or when the deadline fires after some bytes have arrived — a timeout hands over what it already has rather than discarding it, so always check len(), and call Stream.LastError() (1 = timed out) to tell a short block from a complete one. nil means nothing was received at all, or a hard error. timeout in ms; <= 0 (default) = no deadline. With cb, cb(block) is invoked on completion and the awaited result is nil.	—
ReadByte	ReadByte(s:stream) – int	Read one byte as integer.	—
ReadBytes	ReadBytes(s:stream, count:int) – block	Read exactly count bytes into block.	—
ReadLine	ReadLine(s:stream) – string	Read until \n, stripping a preceding \r (so \r\n and \n both work). A lone \r is kept as data — no speculative read-ahead, so no byte is lost on a socket/pipe. Returns "" at EOF.	—
ReadString	ReadString(s:stream, count:int) – string	Read count bytes as string.	—
Seek	Seek(s:stream, pos:int, whence?:int) – bool	Seek a file stream. whence: 0 from start (default), 1 from the current position, 2 from the end (pos may be negative). Returns true on success.	—
SendFile	SendFile(s:stream, path:string, offset?:int, count?:int) – int	Zero-copy file-to-socket transfer. Uses sendfile on Linux and macOS/BSD; falls back to read/write loop on other targets (e.g. OpenWrt). Returns bytes sent.	—
SetKeepAlive	SetKeepAlive(s:stream, on:bool) – bool	Enable/disable TCP keep-alive probes (SO_KEEPALIVE) on a socket. Returns true if set.	—
SetNoDelay	SetNoDelay(s:stream, on:bool) – bool	Enable/disable TCP_NODELAY (disable Nagle) so small writes go out immediately. Returns true if set.	—
SetTimeout	SetTimeout(s:stream, readMs:int, writeMs?:int) – bool	Set read and write timeouts in milliseconds on a socket stream. If writeMs is omitted, both directions use readMs.	—
Shutdown	Shutdown(s:stream, how:string) – bool	Half-close a socket. how: "r" stops reads, "w" stops writes (sends EOF to the peer while the read side stays open), anything else shuts both. Returns true on success.	—
Size	Size(s:stream) – int	Byte length of the underlying file (fstat). nil for a socket/pipe or on error.	—
Stderr	Stderr() – stream	A dup() of the process stderr as a writable stream. Closing it (or its GC) leaves the real stderr open.	—
Stdin	Stdin() – stream	A dup() of the process stdin as a readable stream. Closing it leaves the real stdin open.	—
Stdout	Stdout() – stream	A dup() of the process stdout as a writable stream. Closing it leaves the real stdout open.	—
Tell	Tell(s:stream) – int	Return current byte offset.	—

Function	Signature	Description	JIT
Truncate	Truncate(s:stream, size:int) – bool	Grow or shrink the file to exactly size bytes (ftruncate). Returns true on success.	—
WaitReadable	WaitReadable(s:stream, timeoutMs?:int) – bool	Park the calling fiber until s (typically a listen socket from Listen) becomes readable. Resumes with true on readiness, nil on timeout. timeoutMs <= 0 (default -1) waits without a deadline. Suspends without await — callable from a plain function, so a server loop is while (Stream.WaitReadable(srv, -1)) { let c = Stream.Accept(srv); ... }.	—
WriteAll	WriteAll(s:stream, data:string block) – bool	Write all bytes, retrying on partial writes.	—
WriteAsync	WriteAsync(s:stream, data:string block, timeout?:int) – fiber	Async write of all of data . await resolves to the bytes written (int), or nil on timeout/error. timeout in ms; <= 0 (default) = no deadline. data is retained while the write is in flight. Use Stream.LastError() to tell a timeout from an I/O error.	—
WriteByte	WriteByte(s:stream, value:int) – bool	Write single byte. Returns true on success, false on a write error.	—
WriteBytes	WriteBytes(s:stream, buf:block, count:int) – int	Write count bytes from block; returns bytes written.	—
WriteString	WriteString(s:stream, s:string) – bool	Write string bytes.	—
ReadU8	ReadU8(s:stream) – int	Read one unsigned byte (0–255). Returns nil on EOF.	—
ReadU16	ReadU16(s:stream, bigEndian?:bool) – int	Read 2 bytes as unsigned 16-bit integer. Default little-endian.	—
ReadU32	ReadU32(s:stream, bigEndian?:bool) – int	Read 4 bytes as unsigned 32-bit integer. Default little-endian.	—
ReadU64	ReadU64(s:stream, bigEndian?:bool) – int	Read 8 bytes as 64-bit integer. Default little-endian.	—
WriteU8	WriteU8(s:stream, value:int) – bool	Write one byte. Returns true on success.	—
WriteU16	WriteU16(s:stream, value:int, bigEndian?:bool) – bool	Write 2 bytes. Default little-endian. Returns true on success.	—
WriteU32	WriteU32(s:stream, value:int, bigEndian?:bool) – bool	Write 4 bytes. Default little-endian. Returns true on success.	—
WriteU64	WriteU64(s:stream, value:int, bigEndian?:bool) – bool	Write 8 bytes. Default little-endian. Returns true on success.	—
ReadI8	ReadI8(s:stream) – int	Read one signed byte (–128...127). nil on EOF.	—
ReadI16	ReadI16(s:stream, bigEndian?:bool) – int	Read a signed 16-bit integer. Default little-endian. nil on EOF.	—
ReadI32	ReadI32(s:stream, bigEndian?:bool) – int	Read a signed 32-bit integer. Default little-endian. nil on EOF.	—
ReadI64	ReadI64(s:stream, bigEndian?:bool) – int	Read a signed 64-bit integer. Default little-endian. nil on EOF.	—
ReadF32	ReadF32(s:stream, bigEndian?:bool) – float	Read a 32-bit IEEE-754 float. Default little-endian. nil on EOF.	—
ReadF64	ReadF64(s:stream, bigEndian?:bool) – float	Read a 64-bit IEEE-754 double. Default little-endian. nil on EOF.	—
WriteF32	WriteF32(s:stream, value:float,	Write a 32-bit IEEE-754 float. Default little-endian.	—

Function	Signature	Description	JIT
WriteF64	<code>bigEndian?:bool) - bool</code> <code>WriteF64(s:stream, value:float, bigEndian?:bool) - bool</code>	Write a 64-bit IEEE-754 double. Default little-endian.	—
ReadVarint	<code>ReadVarint(s:stream) - int</code>	Decode an unsigned LEB128 varint (1–10 bytes). <code>nil</code> on EOF/truncation or an over-long encoding.	—
WriteVarint	<code>WriteVarint(s:stream, value:int) - bool</code>	Encode <code>value</code> as an unsigned LEB128 varint. Returns <code>true</code> on success.	—

Stream instance method syntax

All `Stream` functions where the stream is the first argument can be called directly on a `stream` value:

```
let s = Stream.Open("data.bin", "r");
s.ReadLine()           // same as Stream.ReadLine(s)
s.ReadByte()           // same as Stream.ReadByte(s)
s.WriteByte(0x0A)      // same as Stream.WriteByte(s, 0x0A)
s.Seek(0)              // same as Stream.Seek(s, 0)
s.Tell()               // same as Stream.Tell(s)
s.IsOpen()             // same as Stream.IsOpen(s)
s.Flush()              // same as Stream.Flush(s)
s.Close()              // same as Stream.Close(s)
```

Factory functions (`Open`, `Connect`, `Listen`, `Pipe`, `OpenSerial`) are not available as instance methods. Both forms compile to identical bytecode when the variable is typed (`: stream` or inferred from a builtin return). Untyped variables fall back to a runtime dispatch.

String

Namespace: **String**

UTF-8 string operations. `Length` returns **byte count**; use `Utf8Length` for codepoint count.

All `String` functions can also be called directly on a string value:

```
let s = "Hello, World!";
s.Length()             // 13    - same as String.Length(s)
s.ToLower()            // "hello, world!"
s.Contains("World")    // true
s.Substr(7, 5)         // "World"
s.Replace("o", "0")    // "Hell0, W0rld!"
s.Split(",")           // ["Hello", " World!"]
" hi ".Trim()          // "hi"
```

Both forms compile to identical bytecode when the variable is typed (`: string` or inferred from a builtin return). Untyped variables fall back to a runtime dispatch.

Building strings in a loop. `s += piece` is linear: when `s` is the only reference to its value, the piece is appended into the existing buffer (which grows geometrically) instead of copying the whole accumulated string each time. Building 80,000 pieces costs about 2 ms.

The optimization is invisible to program semantics - strings remain values. As soon as the value is shared, the append falls back to producing a fresh string, so an alias never observes a later mutation:

```
let s = "hello";
let t = s;      // t shares the value
s += " world"; // s becomes a new string; t is still "hello"
```

`String.Join(parts, sep)` is still the better choice when the pieces are already in an array, since it sizes the result once.

A `nil` argument in a string position never crashes and follows one policy: transforms and formatters return `nil`, predicates return `false`, index searches return `-1`, `Count` returns `0`; `Length` and `CharAt` treat `nil` as the empty string.

Function	Signature	Description	JIT
CharAt	<code>CharAt(s:string, pos:int) - char</code>	Character at byte position <code>pos</code> as a <code>char</code> immediate. Returns <code>'\0'</code> if <code>pos</code> is out of range. Zero allocation - prefer over <code>Substr(s, pos, 1)</code> in tight loops.	✓
Contains	<code>Contains(s:string, sub:string) - bool</code>	<code>true</code> if <code>sub</code> appears anywhere in <code>s</code> .	✓
Count	<code>Count(s:string, sub:string) - int</code>	Returns the number of non-overlapping occurrences of <code>sub</code> in <code>s</code> . Returns <code>0</code> if <code>sub</code> is empty.	✓
Create	<code>Create(n:int, fill?:int char) - string</code>	Create a mutable string of length <code>n</code> filled with <code>fill</code> (default space). Returns <code>""</code> for <code>n ≤ 0</code> , <code>nil</code> for <code>n</code> exceeding the string size limit or <code>fill</code> outside 1–255 (a NUL fill would embed <code>'\0'</code> bytes and corrupt length bookkeeping). Intended for use with JIT functions that write characters in-place via <code>s[i] = ch</code> .	✓ owned ¹
Empty	<code>Empty - string</code>	Property: the empty string constant <code>""</code> . Not a function.	—
EndsWith	<code>EndsWith(s:string, suffix:string) - bool</code>	<code>true</code> if <code>s</code> ends with <code>suffix</code> .	✓
EqualsIgnoreCase	<code>EqualsIgnoreCase(a:string, b:string) - bool</code>	Case-insensitive equality. ASCII only (<code>A–Z</code> folds to <code>a–z</code>) and locale-independent, matching <code>ToLower</code> / <code>ToUpper</code> ; non-ASCII bytes compare exactly, so <code>"Å"</code> and <code>"å"</code> are not equal.	✓
Format	<code>Format(fmt:string, ...args) - string</code>	C#-style positional formatting: <code>{0}</code> , <code>{1,width}</code> (negative — width = left-align), <code>{0:X4}</code> / <code>{0:x}</code> hex, <code>{0:D3}</code> zero-padded decimal, <code>{0:F2}</code> or <code>{0:.2f}</code> fixed-point (ints included: <code>{0:F2}</code> of 5 is <code>"5.00"</code>). <code>{ {</code> and <code> } }</code> emit literal braces. The template is used literally - a source literal's escapes are already expanded by the compiler, so <code>fmt</code> is never unescaped again and a <code>\t</code> in e.g. a Windows path survives. <code>:X / :x / :D</code> require an int and yield <code>nil</code> otherwise; other specs fall back to the value's default rendering. Returns <code>nil</code> on a malformed template. At least 1 arg required.	

Function	Signature	Description	JIT
FormatArray	<code>FormatArray(fmt:string, values:array) - string</code>	Like <code>Format</code> , but placeholder indices <code>{0}</code> , <code>{1}</code> , ... refer to elements of <code>values</code> . Identical grammar and rendering - owned ¹ both share one engine. Returns <code>nil</code> on a malformed template or out-of-range index.	✓
IndexOf	<code>IndexOf(s:string, sub:string) - int</code>	Byte index of first occurrence of <code>sub</code> , or <code>-1</code> .	✓
IndexOfAnyFrom	<code>IndexOfAnyFrom(s:string, charset:string, start:int) - int</code>	Byte index of the first character in <code>s</code> (starting from <code>start</code>) that appears anywhere in <code>charset</code> , or <code>-1</code> . Uses a single <code>strcspn</code> scan - efficient for finding the first of several possible delimiter characters.	✓
IndexOfFrom	<code>IndexOfFrom(s:string, sub:string, start:int) - int</code>	Byte index of first occurrence of <code>sub</code> at or after byte offset <code>start</code> , or <code>-1</code> . Avoids allocating a substring; use instead of <code>IndexOf(Substr(s, start))</code> in parsing loops.	✓
LastIndexOf	<code>LastIndexOf(s:string, sub:string) - int</code>	Byte index of the last occurrence of <code>sub</code> , or <code>-1</code> . An empty <code>sub</code> returns <code>Length(s)</code> (as <code>IndexOf</code> returns <code>0</code>).	✓
LastIndexOfAny	<code>LastIndexOfAny(s:string, charset:string) - int</code>	Byte index of the last character of <code>s</code> that appears anywhere in <code>charset</code> , or <code>-1</code> . Use to split on the last of several delimiters, e.g. <code>LastIndexOfAny(path, "/\\")</code> .	✓
Insert	<code>Insert(s:string, pos:int, sub:string) - string</code>	Insert <code>sub</code> at byte position <code>pos</code> .	✓ owned ¹
IsAlnum	<code>IsAlnum(s:string) - bool</code>	<code>true</code> if all characters are alphanumeric. <code>false</code> for <code>""</code> .	✓
IsAlpha	<code>IsAlpha(s:string) - bool</code>	<code>true</code> if all characters are alphabetic. <code>false</code> for <code>""</code> .	✓
IsUpper	<code>IsUpper(s:string) - bool</code>	<code>true</code> if all alphabetic characters are uppercase and at least one alphabetic character is present, so <code>"A1"</code> is <code>true</code> but <code>"123"</code> and <code>""</code> are <code>false</code> .	✓
IsLower	<code>IsLower(s:string) - bool</code>	<code>true</code> if all alphabetic characters are lowercase and at least one alphabetic character is present, so <code>"a1"</code> is <code>true</code> but <code>"123"</code> and <code>""</code> are <code>false</code> .	✓
IsNumeric	<code>IsNumeric(s:string) - bool</code>	<code>true</code> if all characters are numeric digits. <code>false</code> for <code>""</code> .	✓
IsWhitespace	<code>IsWhitespace(s:string) - bool</code>	<code>true</code> if all characters are whitespace. <code>false</code> for <code>""</code> .	✓
Join	<code>Join(arr:array, sep:str char) - string</code>	Concatenate array elements with separator.	✓ owned ¹
JsonPretty	<code>JsonPretty(json:string) - string</code>	Pretty-print a JSON string.	✓ owned ¹
Left	<code>Left(s:string, n:int) - string</code>	Return first <code>n</code> bytes.	✓ owned ¹
Length	<code>Length(s:string) - int</code>	Byte length of string.	✓
Levenshtein	<code>Levenshtein(a:string, b:string) - int</code>	Edit distance (insert/delete/substitute) between the strings. Byte-based (equals character distance for ASCII). $O(n \cdot m)$ in C - use for "did you mean", fuzzy matching, dedup. Raises if <code>n*m</code> exceeds <code>MAX_LEVENSHTEIN_CELLS</code> (100M) rather than stalling the VM on huge inputs.	✓
NaturalCompare	<code>NaturalCompare(a:string, b:string) - int</code>	Three-way compare with digit runs compared numerically, so <code>"file2" < "file10"</code> . ASCII case-insensitive with a byte-wise tiebreak (total, deterministic order). Sort naturally: <code>Array.Sort(arr, fn(x,y) { return String.NaturalCompare(x,y) < 0; })</code> .	✓
ProcessCallback	<code>ProcessCallback(fn:function, s:string, len:int, cb:function) - nil</code>	Processes a string with your worker function <code>fn(s, len)</code> , then calls <code>cb(s, result)</code> when finished. A read-only kernel (reads <code>s[i]</code> , returns an <code>int</code> / <code>float</code> — a hash, checksum, count, or scan) runs on another CPU core when simple enough; a kernel that writes to the string runs	✓

Function	Signature	Description	JIT
		normally (inline) so the string stays correct. Returns immediately. See <code>Buffer.ProcessCallback</code> .	
ProcessEvent	<code>ProcessEvent(fn:function, s:string, len:int, eventId:int) - nil</code>	Like <code>ProcessCallback</code> , but signals event <code>eventId</code> with the result instead of calling a callback. Wait for several with <code>Event.WaitFor([ids])</code> .	✓
PadLeft	<code>PadLeft(s:string, width:int, fill?:str char) - string</code>	Left-pad to <code>width</code> with <code>fill</code> (default space).	✓ owned ¹
PadRight	<code>PadRight(s:string, width:int, fill?:str char) - string</code>	Right-pad to <code>width</code> with <code>fill</code> (default space).	✓ owned ¹
Repeat	<code>Repeat(s:string, n:int) - string</code>	Concatenate <code>s</code> with itself <code>n</code> times. Returns "" for <code>n ≤ 0</code> and <code>nil</code> if the result would exceed the string size limit.	✓ owned ¹
Replace	<code>Replace(s:string, old:string, new:string, count?:int) - string</code>	Replace occurrences of <code>old</code> with <code>new</code> , left to right. Without <code>count</code> (or with a negative one) every occurrence is replaced; <code>0</code> replaces none. Matching is non-overlapping, so <code>Replace("aaa", "aa", "b")</code> is <code>"ba"</code> .	✓ owned ¹
Reverse	<code>Reverse(s:string) - string</code>	Reverse bytes of string.	✓ owned ¹
Right	<code>Right(s:string, n:int) - string</code>	Return last <code>n</code> bytes.	✓ owned ¹
Sanitize	<code>Sanitize(s:string) - string</code>	Keep only ASCII letters, digits and <code>_</code> ; every other character (spaces, punctuation, non-ASCII) is removed. Useful for identifiers/filenames.	✓ owned ¹
Fields	<code>Fields(s:string) - array</code>	Split around runs of whitespace, discarding empty parts, so — leading/trailing/repeated spaces are absorbed: <code>Fields("a b ")</code> is <code>["a","b"]</code> where <code>Split(s, " ")</code> would yield empty parts. The usual way to tokenise a line.	—
Split	<code>Split(s:string, sep:str char, limit?:int) - array</code>	Split by separator; returns array of strings. With a positive <code>limit</code> the result holds at most that many parts and the last keeps the unsplit remainder (<code>Split("a,b,c", ",", 2)</code> is <code>["a","b,c"]</code>); <code>limit ≤ 0</code> means no limit.	—
SplitLines	<code>SplitLines(s:string) - array</code>	Split by newlines; returns array of strings.	—
StartsWith	<code>StartsWith(s:string, prefix:string) - bool</code>	<code>true</code> if <code>s</code> starts with <code>prefix</code> .	✓
Substr	<code>Substr(s:string, start:int, len?:int) - string</code>	Byte substring starting at <code>start</code> , optional <code>len</code> .	✓ owned ¹
ToAscii	<code>ToAscii(s:string, replacement?:str char int) - string</code>	Strip/replace non-ASCII characters.	✓ owned ¹
ToLower	<code>ToLower(s:string) - string</code>	Lowercase ASCII characters. Non-ASCII is left untouched - use <code>Utf8ToLower</code> for Unicode.	✓ owned ¹
ToUpper	<code>ToUpper(s:string) - string</code>	Uppercase ASCII characters. Non-ASCII is left untouched - use <code>Utf8ToUpper</code> for Unicode.	✓ owned ¹
Trim	<code>Trim(s:string, cutset?:str char) - string</code>	Remove leading and trailing whitespace, or every leading/trailing character present in <code>cutset</code> when given. An empty <code>cutset</code> trims nothing.	✓ owned ¹
TrimLeft	<code>TrimLeft(s:string, cutset?:str char) - string</code>	Remove leading whitespace, or leading <code>cutset</code> characters.	✓ owned ¹
TrimRight	<code>TrimRight(s:string, cutset?:str char) - string</code>	Remove trailing whitespace, or trailing <code>cutset</code> characters.	✓ owned ¹
Truncate	<code>Truncate(s:string, maxlen:int) - string</code>	Truncate to <code>maxLen</code> bytes.	✓ owned ¹

UTF-8 helpers:

Function	Signature	Description	JIT
Utf8ByteIndexOf	<code>Utf8ByteIndexOf(s:string, cpIndex:int) - int</code>	Byte offset of codepoint at index <code>cpIndex</code> .	✓
Utf8Concat	<code>Utf8Concat(...parts:str char int) - string</code>	Concatenate strings, chars, or codepoints. 1–5 args.	✓ owned ¹
Utf8CpAt	<code>Utf8CpAt(s:string, bytePos:int) - char</code>	Codepoint at byte position <code>bytePos</code> .	✓
Utf8CpNext	<code>Utf8CpNext(s:string, bytePos:int) - int</code>	Byte offset of next codepoint after <code>bytePos</code> .	✓
Utf8GetCharAt	<code>Utf8GetCharAt(s:string, cpIndex:int) - char</code>	Character at codepoint index <code>cpIndex</code> .	✓
Utf8Length	<code>Utf8Length(s:string) - int</code>	Number of Unicode codepoints.	✓
Utf8Substr	<code>Utf8Substr(s:string, cpStart:int, cpLen:int) - string</code>	Substring by codepoint indices.	✓ owned ¹
Utf8ToLower	<code>Utf8ToLower(s:string) - string</code>	Unicode-aware lowercase: <code>Utf8ToLower("ÄÄÖ")</code> is "ääö", where the ASCII-only <code>ToLower</code> leaves it unchanged. Uses the same tables as <code>Char.ToLower</code> ; uncased codepoints pass through. Returns <code>nil</code> on invalid UTF-8.	✓ owned ¹
Utf8ToUpper	<code>Utf8ToUpper(s:string) - string</code>	Unicode-aware uppercase: <code>Utf8ToUpper("ääö")</code> is "ÄÄÖ". Returns <code>nil</code> on invalid UTF-8.	✓ owned ¹

Tls

Namespace: **Tls**

Built-in TLS byte transport - the encrypted counterpart to the TCP side of `Stream`, and the transport beneath the `Https` library (most code should use `Https`). TLS comes from the OS stack (macOS Secure Transport, Linux OpenSSL via `dlopen`, Windows SChannel) and the system trust store. Certificates and hostnames are verified by default (TLS $\geq$ 1.2); connections are opaque `int` handles from a generation-tagged table, so a stale or forged handle is rejected rather than dereferenced.

The synchronous calls are **blocking**: because the VM is cooperative single-threaded, a read or write on a slow peer parks every fiber until data arrives or the timeout fires. A `recv` timeout is reported through `LastError` (the read returns what it has, or `nil`), so on a stalled connection check `LastError` rather than trusting an early EOF. Buffered reads are bounded (`ReadLine` $\leq$ 1 MB per line, `ReadAll` $\leq$ 256 MB) so a hostile peer cannot exhaust memory.

Up to 512 socket/stream operations may be in flight at once (one per parked fiber); past that a further `*Async` call raises `OutOfMemory`. Descriptor numbers are unrestricted — the pump uses `select()` while every descriptor fits its table and switches to `poll()` automatically when one does not, so a process holding thousands of files can still do async I/O.

For concurrency each call has an `*Async` twin (`ConnectAsync`, `WriteAsync`, `ReadLineAsync`, `ReadNAsync`, `ReadAllAsync`) that offloads the blocking work to the I/O thread pool and suspends only the calling fiber - other fibers keep running. The pool queues up to 256 jobs; beyond that a submission runs inline on the VM thread, which still produces the right answer but stalls the scheduler for the duration. `--stats` reports how often that happened, and `--io-threads=<n>` raises the worker count. `await` the result. Only one async op may be in flight per handle, and a handle must not be `Close`d or reused until its async op resolves.

Certificate pinning (trust-on-first-use): after `Connect`, compare the peer fingerprint against a known value and drop the connection on mismatch.

```
let h = Tls.Connect("api.example.com", 443);
if (h != 0 && Tls.PeerCert(h).sha256 != PINNED_SHA256) { Tls.Close(h); h = 0; }
```

Function	Signature	Description	JIT
Available	<code>Available()</code> – <code>bool</code>	<code>true</code> if a TLS backend is compiled in and usable (on Linux, also that <code>libssl</code> loaded at runtime).	—
Connect	<code>Connect(host:string, port:int, insecure?:bool, timeoutMs?:int, options?:object)</code> – <code>int</code>	Open a TLS connection; returns a handle, or <code>0</code> on failure. <code>insecure=true</code> skips certificate/hostname verification (trusted hosts only). <code>timeoutMs</code> (default 30000; <code>0</code> /negative is treated as the default) bounds the TCP connect; <code>options.readTimeoutMs</code> overrides the per-read/write timeout. <code>options.sni</code> overrides the SNI name sent and verified against the certificate (default: <code>host</code>) - useful when connecting to an IP but validating a hostname.	—
PeerCert	<code>PeerCert(h:int)</code> – <code>object nil</code>	The peer's leaf certificate as <code>{ sha256: string, subject: string }</code> (<code>sha256</code> = lowercase hex fingerprint of the DER cert). Use for pinning / trust-on-first-use. <code>nil</code> if unavailable.	—
Write	<code>Write(h:int, data:string)</code> – <code>int</code>	Encrypt and send all bytes; returns bytes written, or <code>-1</code> on error.	—
ReadLine	<code>ReadLine(h:int)</code> – <code>string nil</code>	Read one CRLF/LF-terminated line, terminator stripped; <code>nil</code> at EOF. A single line is capped at 1 MB - a longer unterminated line returns <code>nil</code> with the reason in <code>LastError</code> .	—
ReadN	<code>ReadN(h:int, n:int)</code> – <code>string</code>	Read exactly <code>n</code> bytes (fewer only at EOF). Binary-safe.	—
ReadAll	<code>ReadAll(h:int)</code> – <code>string</code>	Read until the peer closes the connection. Capped at 256 MB - a larger stream returns <code>nil</code> with the reason in <code>LastError</code> .	—
Close	<code>Close(h:int)</code> – <code>bool</code>	Close the connection and free the handle.	—
LastError	<code>LastError(h:int)</code> – <code>string nil</code>	Last error message for the handle, or <code>nil</code> .	—
ConnectAsync	<code>ConnectAsync(host, port, insecure?, timeoutMs?)</code> – <code>fiber</code>	Async <code>Connect</code> ; <code>await</code> resolves to the handle (or <code>0</code>). Suspends only the caller.	—
WriteAsync	<code>WriteAsync(h:int, data:string)</code> – <code>fiber</code>	Async <code>Write</code> ; <code>await</code> resolves to bytes written (or <code>-1</code>).	—
ReadLineAsync	<code>ReadLineAsync(h:int)</code> – <code>fiber</code>	Async <code>ReadLine</code> ; <code>await</code> resolves to the line (or <code>nil</code>).	—
ReadNAsync	<code>ReadNAsync(h:int, n:int)</code> – <code>fiber</code>	Async <code>ReadN</code> ; <code>await</code> resolves to the bytes (or <code>nil</code>).	—
ReadAllAsync	<code>ReadAllAsync(h:int)</code> – <code>fiber</code>	Async <code>ReadAll</code> ; <code>await</code> resolves to the body (or <code>nil</code>).	—

Time

Namespace: **Time**

Unix timestamp manipulation (seconds since epoch). `timestamp` is always `int`. Wall-clock time (`Now`, `NowMillis`) is Unix time; monotonic time (`Millis`, `Ticks`) has an arbitrary epoch and is only meaningful as a delta. All broken-down fields and formatting are UTC. The richer calendar/timezone API (local zones, leap years, durations, relative "ago") lives in the `Datetime` library.

Function	Signature	Description	JIT
AddDays	AddDays(t:int, n:number) – int	Add n days to timestamp. Fractional values allowed; nil if the result would overflow.	✓
AddHours	AddHours(t:int, n:number) – int	Add n hours to timestamp. Fractional values allowed; nil if the result would overflow.	✓
AddMinutes	AddMinutes(t:int, n:number) – int	Add n minutes to timestamp. Fractional values allowed; nil if the result would overflow.	✓
AddSeconds	AddSeconds(t:int, n:number) – int	Add n seconds to timestamp. Fractional values allowed; nil if the result would overflow.	✓
AddWeeks	AddWeeks(t:int, n:number) – int	Add n weeks to timestamp. Fractional values allowed; nil if the result would overflow.	✓
Day	Day(t:int) – int	Day of month (1–31).	✓
DiffDays	DiffDays(a:int, b:int) – int	Whole days between two timestamps (a – b), truncated toward zero.	✓
Format	Format(t:int, fmt?:string) – string	Format timestamp as string (UTC). Named aliases: "short" (YYYY-MM-DD), "long" / "log" (YYYY-MM-DD HH:MM:SS), "time" (HH:MM:SS), "iso" (YYYY-MM-DDTHH:MM:SSZ). Default: "long". Custom strftime patterns also accepted (capped at 128 chars); an empty pattern yields "" . nil on a bad time or format failure.	✓ owned ¹
FromMillis	FromMillis(ms:int) – int	Convert epoch milliseconds to seconds (truncate). Inverse of ToMillis .	✓
Hour	Hour(t:int) – int	Hour (0–23).	✓
Millis	Millis() – int	Monotonic millisecond timestamp for timing/intervals. The epoch is arbitrary (not Unix time); use only for deltas. For wall-clock milliseconds use NowMillis .	—
Minute	Minute(t:int) – int	Minute (0–59).	✓
Month	Month(t:int) – int	Month (1–12).	✓
Now	Now() – int	Current Unix timestamp in seconds.	—
NowMillis	NowMillis() – int	Current Unix time in whole milliseconds (wall clock). FromMillis(NowMillis()) == Now() .	—
Parse	Parse(s:string, fmt?:string) – int	Parse timestamp string (UTC). Named aliases: "ISO" (default, %Y-%m-%dT%H:%M:%S), "RFC" (%a, %d %b %Y %H:%M:%S), "DATE_ONLY" (%Y-%m-%d). Custom strftime patterns also accepted. Returns nil on failure.	✓
Second	Second(t:int) – int	Second (0–59).	✓
Ticks	Ticks() – int	Monotonic nanosecond counter. The epoch is undefined; use only for deltas with TicksToMs / TicksToUs .	—
TicksToMs	TicksToMs(ticks:int) – float	Convert a nanosecond tick delta to milliseconds. Example: Time.TicksToMs(Time.Ticks() – t0) - 1.234	✓
TicksToUs	TicksToUs(ticks:int) – float	Convert a nanosecond tick delta to microseconds. Example: Time.TicksToUs(Time.Ticks() – t0) - 1234.5	✓
TimeZoneOffsetMins	TimeZoneOffsetMins(t?:int) – int	Local timezone offset in minutes from UTC, DST-aware. Uses the current time, or the given timestamp.	—
ToMillis	ToMillis(seconds:int) – int	Convert whole seconds to milliseconds. Inverse of FromMillis .	✓
Weekday	Weekday(t:int) – int	Day of week: 0=Sunday ... 6=Saturday.	✓
Year	Year(t:int) – int	Four-digit year.	✓

Timers

Namespace: **Timers**

Fiber-based timer scheduling. Max 64 simultaneous timers. Timer resolution is approx 1-3ms.

Function	Signature	Description	JIT
Cancel	<code>Cancel(id:int float) - bool</code>	Cancel a timer by its handle.	—
IsActive	<code>IsActive(id:int float) - bool</code>	<code>true</code> if <code>id</code> refers to a live timer (a scheduled one-shot not yet fired, or a still-repeating interval).	—
SetAtTime	<code>SetAtTime(timestamp:int float, func:fn) - int nil</code>	Fire <code>fn</code> once at the given Unix timestamp (seconds since epoch); a time already in the past fires as soon as possible. Returns a handle, or <code>nil</code> on a NaN timestamp or when the timer table (max 64) is full.	—
SetImmediate	<code>SetImmediate(func:fn) - int nil</code>	Fire <code>fn</code> on the next scheduler pump (a zero-delay one-shot), akin to Node's <code>setImmediate</code> . Returns a handle, or <code>nil</code> when the timer table is full.	—
SetInterval	<code>SetInterval(ms:int float, func:fn) - int nil</code>	Fire <code>fn</code> every <code>ms</code> milliseconds. Returns a handle, or <code>nil</code> if <code>ms</code> is NaN or <code><= 0</code> , or the timer table is full.	—
SetTimeout	<code>SetTimeout(ms:int float, func:fn) - int nil</code>	Fire <code>fn</code> once after <code>ms</code> milliseconds (a negative delay fires as soon as possible). Returns a handle, or <code>nil</code> on a NaN delay or when the timer table is full.	—

Util

Namespace: **Util**

Encoding, hashing utilities, and value comparison.

Function	Signature	Description	JIT
AscToHex	<code>AscToHex(s:string, upper?:bool) - string</code>	Convert ASCII string to hex representation.	✓ owned ¹
Base64Encode	<code>Base64Encode(data:string block) - string</code>	Encode string or block as standard Base64 (padded, <code>+/</code> alphabet).	✓ owned ¹
Base64UrlEncode	<code>Base64UrlEncode(data:string block) - string</code>	Encode string or block as Base64url (no padding, <code>-_</code> alphabet). Suitable for JWT.	✓ owned ¹
Base64DecodeString	<code>Base64DecodeString(s:string, strict?:bool) - string</code>	Decode Base64 or Base64url string to UTF-8 string. Accepts unpadded input and <code>-_</code> alphabet. Returns <code>nil</code> on invalid input. With <code>strict=true</code> , also rejects non-canonical encodings (interior <code>=</code> , an impossible length, or non-zero trailing bits).	✓ owned ¹
Base64DecodeBlock	<code>Base64DecodeBlock(s:string, strict?:bool) - block</code>	Decode Base64 or Base64url string to binary block. Accepts unpadded input and <code>-_</code> alphabet. Returns <code>nil</code> on invalid input. With <code>strict=true</code> , also rejects non-canonical encodings.	✓ owned ¹
Base32Encode	<code>Base32Encode(data:string block) - string</code>	Encode as RFC 4648 Base32 (A-Z2-7, <code>=</code> padded). The encoding used for TOTP/2FA secrets.	✓ owned ¹
Base32DecodeString	<code>Base32DecodeString(s:string, strict?:bool) - string</code>	Decode Base32 to a string. Accepts lowercase and missing padding. Returns <code>nil</code> on invalid input. With <code>strict=true</code> , also rejects non-canonical encodings (interior <code>=</code> or non-zero trailing bits).	✓ owned ¹
Base32DecodeBlock	<code>Base32DecodeBlock(s:string, strict?:bool) - block</code>	Decode Base32 to a binary block. Accepts lowercase and missing padding. Returns <code>nil</code> on invalid input.	✓ owned ¹

Function	Signature	Description	JIT
		invalid input. With <code>strict=true</code> , also rejects non-canonical encodings.	
BinToHex	<code>BinToHex(buf:block) - string</code>	Convert binary block to hex string.	✓ owned ¹
Coalesce	<code>Coalesce(args...:any) - any</code>	Return first non-nil argument. Also accepts a single array and scans it.	—
CoalesceEmpty	<code>CoalesceEmpty(args...:any) - any</code>	Return first non-nil and non-empty argument. Empty means <code>nil</code> , <code>""</code> , <code>0/0.0</code> , <code>{}</code> or <code>[]</code> (but not <code>false</code>). Also accepts a single array.	—
Compare	<code>Compare(a:any, b:any) - int</code>	Three-way comparison usable as a sort comparator: returns -1, 0, or 1 (matches the language's <code><</code> ordering).	✓
CompareVersions	<code>CompareVersions(a:string, b:string) - int</code>	Semantic-version compare: returns -1, 0, or 1. <code>MAJOR.MINOR.PATCH</code> numerically (<code>"1.2.3" < "1.10.0"</code>), prerelease per semver (<code>"1.0.0-rc.1" < "1.0.0"</code> ; numeric identifiers before alphanumeric), build metadata (<code>+. . .</code>) ignored. Tolerates a leading <code>v</code> ; missing components count as 0.	✓
VersionSatisfies	<code>VersionSatisfies(version:string, range:string) - bool</code>	<code>true</code> if <code>version</code> falls within an npm-style semver <code>range</code> . Supports <code>^/~/>/>=/</<=/=</code> , x-ranges (<code>1.2.x</code> , <code>1.x</code> , <code>*</code>), hyphen ranges (<code>1.2.3 - 2.3.4</code>), space-separated <code>AND</code> , and <code> </code> <code>OR</code> . A prerelease version matches only a set that itself names a prerelease with the same <code>MAJOR.MINOR.PATCH</code> .	✓
VersionMax	<code>VersionMax(versions:array, range:string) - string</code>	Highest version string in <code>versions</code> that satisfies— <code>range</code> (see <code>VersionSatisfies</code>), or <code>nil</code> if none. Non-string elements are skipped. The dependency-resolver primitive.	—
HexToAsc	<code>HexToAsc(s:string) - string</code>	Convert hex string back to ASCII.	✓ owned ¹
HexToBin	<code>HexToBin(s:string) - block</code>	Convert hex string to binary block.	✓ owned ¹
Nanoid	<code>Nanoid(size?:int) - string</code>	Generate a URL-safe random ID. Optional size.	✓ owned ¹
Uuid	<code>Uuid() - string</code>	Generate a UUID v4 string.	✓ owned ¹
Uuid7	<code>Uuid7() - string</code>	Generate an RFC 9562 UUID v7: a 48-bit millisecond timestamp plus random bits, so ids sort lexicographically by creation time (millisecond granularity). Ideal for database keys.	✓ owned ¹
UuidFromBytes	<code>UuidFromBytes(b:block) - string</code>	Format a 16-byte block as a canonical lowercase UUID string. <code>nil</code> unless the block is exactly 16 bytes. Inverse of <code>UuidToBytes</code> .	✓ owned ¹
UuidIsValid	<code>UuidIsValid(s:string) - bool</code>	<code>true</code> if <code>s</code> is a canonical 8-4-4-4-12 hex UUID (any version, either case).	✓
UuidToBytes	<code>UuidToBytes(s:string) - block</code>	Parse a canonical UUID string into its 16 raw bytes, or <code>nil</code> if invalid. Inverse of <code>UuidFromBytes</code> .	✓ owned ¹

VM

Namespace: **VM**

Virtual machine introspection and control.

Function	Signature	Description	JIT
Compile	<code>Compile(src:string, arity?:int) - function nil</code>	Compile a STATEMENT body into a callable function. <code>src</code> becomes the body of a function taking <code>arity</code> parameters named <code>arg0..argN</code> (use <code>return</code> for the result). Max 10 KB; returns <code>nil</code> on compile failure.	—
CompactMemory	<code>CompactMemory()</code>	Run memory compaction pass.	—
CurrentAllocations	<code>CurrentAllocations() - int</code>	Current live allocation count.	—
Eval	<code>Eval(src:string, ...args) - any</code>	Compile and immediately evaluate <code>src</code> . Tried first as an EXPRESSION (<code>Eval("40 + 2") → 42</code>); if that fails to compile it is retried as a statement body, where an explicit <code>return</code> provides the result (<code>nil</code> otherwise). Extra call arguments are bound to <code>arg0..argN</code> : <code>Eval("arg0 * arg1", 6, 7) → 42</code> . Max 10 KB; returns <code>nil</code> on compile failure (the process is never terminated by a bad <code>src</code>). Each <code>Eval</code> compiles as its own unit: it cannot reference host globals by name, but a <code>global</code> declaration in it can rebind an existing host global (warns on stderr). Eval'd code runs with full VM authority - never pass it untrusted input.	—
Exit	<code>Exit(code:int) - nil</code>	Terminate VM with exit code.	—
GetFunctionInfo	<code>GetFunctionInfo(fn:callable) - object</code>	Return an object describing a function. Accepts <code>function</code> , <code>builtin</code> , <code>ffi</code> , and bound methods. See field table below.	—
GetModuleInfo	<code>GetModuleInfo(name:string) - object nil</code>	Return an object describing a built-in module. Members <code>array</code> contains <code>{Name, Signature}</code> per function; <code>Constants</code> array contains <code>{Name, Type, Value}</code> per constant. Returns <code>nil</code> if the module name is not found.	—
GetStartTimeMs	<code>GetStartTimeMs() - int</code>	VM start time in milliseconds since epoch.	—
Import	<code>VM.Import(name:string, version:string, fingerprint?:string)</code>	Load or return cached module by name and version requirement. <code>version</code> uses the same syntax as <code>library(): "1.0", ">=1.2 <2.0"</code> , etc. Optional <code>fingerprint</code> is the whole-file SHA-256 (<code>== sha256sum</code> ; a <code>sha256</code> : prefix is accepted) - if provided and mismatched, the VM halts regardless of <code>--no-verify</code> . <code>let m = VM.Import("jwt", "1.0");</code> . See version syntax table in Guide §6.	—
MemoryStats	<code>MemoryStats() - object</code>	Snapshot of allocator counters: <code>Current</code> (live objects), <code>Peak</code> (high-water live objects), <code>HeapBytes</code> (cumulative bytes requested from malloc/calloc/realloc), <code>BlockRegions / BlockBytes</code> (live block allocations and their size), <code>SlabCacheFree</code> (objects held in the free-list cache). A superset of <code>CurrentAllocations / PeakAllocations</code> for tooling and leak checks.	—
OnSignal	<code>OnSignal(signum:int, handler:fn) - bool</code>	Register a signal handler and return <code>true</code> . Only signals the VM installs an OS handler for are accepted: <code>SIGHUP</code> , <code>SIGINT</code> , <code>SIGUSR1</code> , <code>SIGUSR2</code> , <code>SIGTERM</code> (numbers are platform-specific). Any other signal - including the uncatchable	—

Function	Signature	Description	JIT
PatchFunction	<code>PatchFunction(original:fn, replacement:fn) - bool</code>	SIGKILL/SIGSTOP, which could never be delivered - or a non-function handler returns <code>false</code> . Replace function at runtime. Requires running <code>unsafe-mode --unsafe</code> . The replacement may be a Flaris-function, a builtin-function or a FFI-function. Pass <code>nil</code> as replacement to remove the patch. Caveats: call sites the compiler inlined (trivial single- <code>return</code> functions) and calls made from inside JIT-compiled functions bypass the patch; self-recursive calls inside the original body also keep calling the original.	—
PeakAllocations	<code>PeakAllocations() - int</code>	Peak allocation count since VM start.	—
RaiseSignal	<code>RaiseSignal(signum:int) - bool</code>	Send signal <code>signum</code> to the current process (<code>raise</code>). A handler registered via <code>OnSignal</code> runs on the next signal pump; an uncatchable signal takes the OS default effect. Returns <code>false</code> on an out-of-range number.	—
SignalName	<code>SignalName(signum:int) - string nil</code>	Canonical name (e.g. "SIGINT") for a signal number, or <code>nil</code> if unknown on this platform.	—
SignalNumber	<code>SignalNumber(name:string) - int nil</code>	Platform signal number for a name like "SIGINT" (the "SIG" prefix is optional), or <code>nil</code> if unknown. Use this instead of hardcoding numbers, which are platform-specific (<code>SIGUSR1</code> is 10 on Linux, 30 on macOS).	—

GetFunctionInfo - returned object fields

Field	Type	Description
<code>Kind</code>	<code>string</code>	"function", "builtin", "ffi", or "unknown"
<code>Name</code>	<code>string nil</code>	Function name. For builtins resolved as <code>"Module.Name"</code> . Nil if anonymous.
<code>Arity</code>	<code>int</code>	Number of required parameters.
<code>Optional</code>	<code>int</code>	Number of optional parameters (builtins only; 0 for scripted and FFI).
<code>Returns</code>	<code>string</code>	Return type as a human-readable string (e.g. "string", "int nil").
<code>ReturnType</code>	<code>int</code>	Return type as a raw type-mask integer.
<code>ArgTypes</code>	<code>array<int></code>	Per-argument type masks in declaration order.
<code>ArgTypeNames</code>	<code>array<string></code>	Per-argument type names in declaration order.
<code>IsStatic</code>	<code>bool</code>	<code>true</code> if the function is a static method.
<code>IsAsync</code>	<code>bool</code>	<code>true</code> if declared <code>async</code> .
<code>IsMethod</code>	<code>bool</code>	<code>true</code> if the function is an instance method (thiscall).
<code>Owner</code>	<code>string nil</code>	Class name for methods; nil for free functions.
<code>Signature</code>	<code>string</code>	Rendered signature, e.g. <code>"fn Module.Name(int, string?) - bool"</code> .

Bound methods are automatically unwrapped - passing a bound method returns info about the underlying function.

R9 - Debug Reference

Debug Opcodes

The compiler emits these opcodes when debug symbols are enabled (default; disabled by `--strip` flag):

Opcode	Arguments	Purpose	When emitted
<code>OP_DBG_LINE</code>	<code>u16 line</code>	Mark source line number	Before each statement
<code>OP_DBG_FUNC_NAME</code>	<code>u16 const_idx</code>	Mark function name	At function entry

Opcode	Arguments	Purpose	When emitted
OP_DBG_FILE_NAME	u16 const_idx	Mark source file	Once per compilation unit
OP_DBG_BREAK	u16 code	Breakpoint	At breakpoint statement

Bytecode size overhead with debug symbols: ~15%. **Execution overhead:** ~2–3%.

Bytecode Example

```
fn compute(x) { return x * 2; }
```

With debug symbols:

```
[0000] OP_DBG_FILE_NAME    "main.fls"
[0003] OP_DBG_FUNC_NAME    "compute"
[0006] OP_DBG_LINE         line=2
[0009] OP_GET_LOCAL        x
[0011] OP_IMM8              2
[0013] OP_MULTIPLY
[0014] OP_RETURN
```

With `--strip` (stripped):

```
[0000] OP_GET_LOCAL
[0002] OP_IMM8              2
[0004] OP_MULTIPLY
[0005] OP_RETURN
```

Quick Reference

CLI Flags for Debugging

Flag	Effect	When to use
(default)	Debug symbols ON	Development, testing
<code>--strip</code>	Debug symbols OFF	Production, embedded
<code>--verbose</code>	Verbose output	High-level debugging
<code>--trace</code>	Very verbose	Opcode-level tracing
<code>--time</code>	Timing report	Performance analysis
<code>--stats</code>	Statistics	Memory debugging
<code>--debug</code>	Stop at <code>Main</code> , open the prompt	Interactive debugging

Debug Module Quick Summary

Function	Purpose
<code>Debug.Stack()</code>	Dump operand stack to stdout
<code>Debug.Pool()</code>	Memory slab statistics (requires <code>--unsafe</code>)
<code>Debug.Value(v)</code>	Inspect value (type, ref-count, address) (requires <code>--unsafe</code>)
<code>Debug.Assert(a, b)</code>	Equality check - aborts the VM on failure

Function	Purpose
<code>Debug.AssertTrue(c)</code>	Truthiness check - aborts the VM on failure
<code>Debug.GuardAddress(p)</code>	Watch memory address - traps on access (requires <code>--unsafe</code>)
<code>Debug.StackPtr()</code>	Current stack depth as integer
<code>Debug.Refs(o)</code>	External reference count of a value (requires <code>--unsafe</code>)
<code>Debug.Here(msg?)</code>	Print location marker (file/line/function)

Breakpoint Statement

```
breakpoint 100;           // Unconditional - always triggers
if (error) breakpoint 200; // Conditional
```

Codes are user-defined integers. Use a namespace convention (e.g. 1000 = startup, 2000 = after validation) to make them meaningful.

A `breakpoint` statement is not the only way in. `--debug` stops at the start of `Main` and opens the prompt there, so a session can begin without editing the source at all - set breakpoints by function or by line and continue:

```
$ flarism --debug app.fl
[Entry] Main (app.fl:40) [fiber 1]
(fdb) b inspectOrder
breakpoint 1 at app.fl:21 (inspectOrder)
(fdb) b Order.Describe
breakpoint 2 at app.fl:18 (Order.Describe)
(fdb) c

[Breakpoint 1] inspectOrder (app.fl:21) [fiber 1]
(fdb) n
[Step] inspectOrder (app.fl:22) [fiber 1]
(fdb)          <- an empty line repeats the step
```

The stop waits for `Main` rather than the first line executed: the module prologue runs first and is what defines the program's classes and functions, so stopping before it finished would leave nothing to set a breakpoint on.

Reaching a breakpoint opens an interactive prompt that pauses the whole VM and lets you walk the call stack and inspect values:

```
[Breakpoint:42] inspectOrder (app.fl:26) [fiber 1]
(fdb) bt
-> #0  inspectOrder (app.fl:26)
     #1  handleRequest (app.fl:33)
     #2  Main (app.fl:39)
(fdb) locals
[0] order      = (instance) Instance: Order, {1017, 249.5, espresso}
[1] retries    = (int) 2
[2] items      = (array) [10, 20, 30]
```

```
(fdb) p order.label
order.label = (string) espresso
(fdb) c
```

Command	Purpose
s / n	Step into / step over one source line
finish	Run until the selected frame returns
(empty line)	Repeat the last step
b [<file>:]<line>	Set a line breakpoint (bare line = the selected frame's file)
b <function>	Break at a function's first line; <code>Class.Method</code> works too
delete [<id>]	Delete one breakpoint, or all of them
info	List breakpoints with hit counts
bt	Backtrace, innermost frame as <code>#0</code>
frame <n>	Select frame <code><n></code>
up / down [n]	Move <code>n</code> frames outward / inward (default 1)
locals	Locals of the selected frame, named from debug info
p <path>	Print a value: <code>name</code> , <code>name.field</code> , <code>name[0]</code> , or a chain of those
list	Source lines around the selected frame
dis [n]	Bytecode around the instruction the selected frame is stopped at, <code>n</code> either side (default 6)
dis ir [fn]	JIT IR of the selected frame, or of a named function
stack	Operand stack contents
fibers	Every live fiber, its state, position and what it waits on
c	Continue execution
q	Quit (exit code 5)
help	Command list

Inspection is read-only: the prompt never allocates a value, changes a reference count or raises, so looking at a program cannot change how it behaves.

Editor Debugging (DAP)

`flarism --dap=<port>` serves the Debug Adapter Protocol, so an editor can drive the same debugger: breakpoints in the gutter, stepping, a variables pane and a call stack. The VSCode extension in `vscode-flaris/` contributes a `flaris` debug type that starts the VM and connects for you; a minimal `launch.json` is:

```
{
  "type": "flaris",
  "request": "launch",
  "name": "Debug current Flaris file",
  "program": "${file}",
  "cwd": "${workspaceFolder}"
}
```

The transport is a TCP socket on loopback, never stdio: the debugged program owns stdout, and a protocol sharing that stream would be corrupted by the first `Console.WriteLine`. The adapter binds `127.0.0.1` only, so a session is never reachable off the machine.

A DAP thread is a fiber. The mapping is exact - fibers have ids, their own call stacks and their own scheduling state - so the editor's thread picker becomes a fiber picker.

Supported requests: `initialize`, `launch`, `setBreakpoints`, `configurationDone`, `threads`, `stackTrace`, `scopes`, `variables`, `evaluate`, `disassemble`, `continue`, `next`, `stepIn`, `stepOut`, `pause`, `disconnect`.

`evaluate` resolves the same paths the prompt's `p` does (`order.items[0]`), not arbitrary expressions, and reports a failure rather than inventing a value. Conditional breakpoints and `setVariable` are not implemented, and the adapter says so during `initialize` so the editor does not offer them.

Pausing and editing breakpoints while running. Requests are normally served while the program is stopped, but `pause` and `setBreakpoints` also have to be noticed mid-run. The adapter checks its socket every 128 line events - often enough to feel immediate, rare enough that it is not paying a syscall per source line. A `pause` takes effect at the next line boundary, which is where every other stop lands too; there is no way to suspend mid-instruction. Requests that would read a paused state (`stackTrace`, `variables`, stepping) are refused while running rather than answered with stale frames.

Disassembly. The editor's Disassembly View is served from the same disassembler `dis` uses at the prompt. A bytecode VM has no machine address to report, but the client parses memory references and instruction addresses *numerically*, so a position is dressed as one: the frame level occupies the high byte and the bytecode offset the low 24 bits (`0x0000003A` is offset `0x3A` in frame 0). A non-numeric reference makes the view report "Disassembly not available". Only real instructions are returned - padding rows would need invented addresses, and the client indexes rows by address.

Open it by right-clicking a frame in the Call Stack and choosing **Open Disassembly View**. Its step buttons step by source line: instruction-granularity stepping is not offered, because the debugger's hook is on line events.

Because `--dap` implies `--debug`, the JIT is disabled for the session and the program stops at the start of `Main` before the first line runs.

Uncaught Exceptions

An exception that reaches the top of a fiber with no handler is reported with its own code and message, and the process exits with that code:

```
Unhandled exception (code 42): boom
Stack trace (depth: 2):
  at Main (app.flr) [ip=34, line=9]
  at level3 (app.flr) [ip=31, line=3]
```

Under `--debug` the prompt opens at the throw site instead, with every frame still standing, so the values that produced the failure can be inspected:

```
[Unhandled exception: code 42] level3 (app.flr:3) [fiber 1]
(fdb) locals
  [0] x                = (int) 11
  [1] doomed            = (int) 22
(fdb) up
-> #1 Main (app.flr:9)
```

The program cannot resume from there - `c` and `q` both exit, with the exception's code.

Inspecting Fibers

`fibers` lists every live fiber with its state, current position and what it is parked on. Scheduling is cooperative and single-threaded, so the whole program is frozen while the prompt is open and the listing is a consistent snapshot - which is what makes it useful for a program that has stopped making progress:

```
(fdb) fibers
-> #1   running      frames=2   Main (app.fl:18)
      #2   yielded    frames=1   producer (app.fl:3)
      #3   wait-fiber frames=1   consumer (app.fl:9) waiting on fiber 2
```

Stepping and line breakpoints are decided at `OP_DBG_LINE`, the opcode that marks each source line. Nothing is examined unless a step is pending or a breakpoint exists, so an ordinary run costs one predicted branch, and a `--strip` build never executes the opcode at all.

Notes:

- A step belongs to the fiber that issued it. Other fibers run between two line events of the stepped one, and their lines never satisfy the step.
- `--debug` disables the JIT. A JIT-compiled function runs natively and never executes a line event, so breakpoints and steps could not reach inside one; turning the JIT off is what makes stopping reliable rather than silent.
- The prompt opens when stdin is a terminal. Otherwise the breakpoint prints its location and execution continues, so a breakpoint left in committed code cannot hang a piped or CI run. Pass `--debug` to open the prompt anyway, which is also how a session is driven from a script.
- `--strip` removes variable names, not breakpoints: slots and their values are still listed, as `<unnamed>`.
- `bt` shows the frames that actually exist. A tail call reuses its caller's frame, so a function that ends in `return f(...)` does not appear above `f`.
- Names resolve against the selected frame's locals first, then module and global scope.

Real-World Debugging Examples

Breakpoint with Bytecode Context

`dis` at the prompt shows the instruction stream around the instruction the selected frame is stopped at, marked `>>`. `--verbose` prints the same window automatically when the breakpoint fires.

```
[Breakpoint:100] main (app.fl:1649) [fiber 1]
(fdb) dis
bytecode at ip=4957 in main:
  [4929] LESS_EQUAL
  [4930] CALL1_BUILTIN      mod-id:12 func-id:1 -> Console.WriteLine
  [4933] POP
>> [4957] BREAKPOINT      code=100
  [4960] LINE              line=1651
```


Selecting an outer frame with `up` marks the call that frame is waiting inside, so `dis` always points at where that frame actually is.

The window always starts on a real instruction boundary. Because instructions are variable length, it is found by decoding forward from the start of the chunk rather than by stepping back a fixed number of bytes.

For a JIT-compiled function, `dis ir <fn>` shows the typed register IR it was compiled from. A function containing a `breakpoint` is never JIT-eligible, so name the kernel you want rather than expecting IR for the frame you stopped in:

```
(fdb) dis ir kernel
JIT IR for kernel:
— JIT IR (regs=5 ir_bytes=54) —
0022 ARITH_IMM8    r3, MUL, r2, 3
0027 ARITH_II      r1, ADD, r1, r3
0031 LOOP          -40  -000c
0034 RET           r1
```

Memory Guard Trigger

```
Debug.GuardAddress(critical_data);
run_entire_program(critical_data);
// Any code that touches this object will stop here:
// [Mem-guard] R/W-operation on memory address 0xC98CDC1D0 in function main, line 1491
```

Tracking Stack Leaks

```
let before = Debug.StackPtr();
risky_operation(data);
let after = Debug.StackPtr();
if (after != before) {
    Console.WriteLine("Stack leak: delta =", after - before);
    Debug.Stack();
}
```

Full Debug Session Pattern

```
fn investigate(data) {
    Debug.Here("start");
    let alloc_before = VM.CurrentAllocations();
    let stack_before = Debug.StackPtr();
    Debug.Pool();

    Debug.GuardAddress(data);
    breakpoint 100;
```

```
let result = risky_operation(data);

breakpoint 200;

if (Debug.StackPtr() != stack_before) Console.WriteLine("STACK LEAK!");
Debug.Value(result);
Debug.Pool();
Console.WriteLine("Alloc delta:", VM.CurrentAllocations() - alloc_before);
Debug.Here("end");
}
```

For FFI plugin development, see `ffi.md`.

R10 - Writing Fast Flaris Code

This chapter translates the execution model from R7 into concrete coding patterns. Each recommendation is grounded in how many VM dispatch cycles a pattern consumes per iteration.

Dispatch count is the most direct proxy for loop speed.

Counting loops: use `iter` for known integer ranges

`iter` is the fastest loop construct in Flaris. A single fused instruction both increments the counter and checks the bound - one dispatch for what `while` needs three to do.

```
// Best - 2 dispatches/iteration overhead
iter (i from 0 to n) {
    sum += arr[i];
}
```

Equivalent `while` costs at least 3 dispatches overhead (compare-jump + increment + loop-back). Use `iter` whenever the range is a fixed integer expression known at the loop head.

`iter` allocates no iterator object. `i` is a local slot that starts at `from` and stops just before `to`.

Comparing locals to constants: `while` fuses to a single instruction

When the loop condition is `local < constant` (or `<=`, `>`, `>=`), the compiler fuses the compare and branch into a single instruction instead of the four-instruction sequence a generic condition would require. A literal bound up to ± 127 encodes in 1 byte; larger bounds encode in 4 bytes - both still one dispatch.

```
// Fused - 1 dispatch for the entire condition
while (i < 1_000_000) { ... } // large literal bound (4-byte encoding)
while (i < 100)          { ... } // small literal bound (1-byte encoding)
```

```
// Also fused – local vs local
while (i < len) { ... }
```

All comparison operators (`<`, `<=`, `>`, `>=`, `==`, `!=`) are eligible. The bound must be a literal integer or a local variable resolved at compile time. If it is an expression (e.g. `arr.count + 1`), cache it into a local before the loop:

```
let limit = arr.count + 1;      // evaluated once
while (i < limit) { ... }      // fused, 1 dispatch
```

Incrementing counters: prefer `++` over `+= 1` over `= x + 1`

Three ways to increment a counter produce meaningfully different instruction sequences:

Form	Encoding	Dispatches
<code>i++</code>	2 bytes	1
<code>i += 1</code>	4 bytes	1
<code>i = i + 1</code>	6 bytes (padded)	3

`i = i + 1` is peephole-optimized in-place, but the padding bytes that fill the freed space still execute as individual dispatches.

`i++` is the most compact form - a 2-byte instruction with no arithmetic subcode. It is both the smallest encoding and the clearest signal to the compiler.

```
// Best – 2 bytes, 1 dispatch
while (i < n) {
    process(arr[i]);
    i++;
}

// Good – 4 bytes, 1 dispatch
while (i < n) {
    process(arr[i]);
    i += 1;
}

// Avoid – 6 bytes padded, 3 dispatches
while (i < n) {
    process(arr[i]);
    i = i + 1;
}
```

The same rule applies to `--` and any compound-assign: `n += x` is always cheaper than `n = n + x`.

Compound assignment beats explicit assignment for all arithmetic

`n += x`, `n -= x`, `n *= x` etc. compile to a compact fused instruction (zero stack, 1 dispatch). Writing `n = n + x` instead forces the optimizer to patch the result in-place, leaving padding bytes that still dispatch individually.

```
// Good - zero-stack, 1 dispatch
sum += arr[i];
n    *= factor;

// Slower - fused + 1-3 padding dispatches depending on operand sizes
sum = sum + arr[i];
n    = n * factor;
```

This matters most inside loops. The operand encoding determines how much padding remains after optimization.

Cache repeated property and array accesses into locals

Every `obj.prop` or `arr[expr]` inside a loop re-executes a hashmap lookup or bounds-checked index. Pull stable reads out of the loop body.

```
// Without caching - hashmap + bounds check every iteration
while (i < data.count) {
    process(data.items[i]);
    i++;
}

// With caching - local reads only inside the loop
let items = data.items;
let count = data.count;
while (i < count) {
    process(items[i]);
    i++;
}
```

This combines with the fused-compare rule: `while (i < count)` fuses to a single instruction because `count` is now a local.

For array element reads inside loops, the compiler uses a zero-stack fast path when both the array and the index are locals, avoiding the generic dispatch path entirely.

The same principle applies to method calls. Storing a method reference in a local before the loop avoids repeating the hash lookup on every iteration:

```
// Without caching - hash lookup on every call
while (i < count) {
```

```
    obj.process(data[i]);
    i++;
}

// With caching - direct call, no lookup
let fn = obj.process;
while (i < count) {
    fn(data[i]);
    i++;
}
```

Keep variables type-stable inside loops

The VM fast-paths for arithmetic and comparison check the type tag on every operation. A variable that starts as an integer and becomes a float mid-loop forces the slow path on every subsequent instruction.

```
// Good - tag-int fast path throughout
let sum: int = 0;
iter (i from 0 to n) {
    sum += values[i]; // stays int the whole way
}

// Risky - if any values[i] is a float, sum flips type mid-loop
let sum = 0;
iter (i from 0 to n) {
    sum += values[i];
}
```

Type annotations on locals (`let sum: int = 0`) are enforced by the analyzer and allow the compiler to take the fastest arithmetic path.

Type your function parameters for faster calls

Flaris supports two styles: quick untyped scripts and fully typed code. The style you choose affects call performance.

When a function is called, the VM validates each argument's type against the declared parameter type. For untyped parameters the check still runs - it just accepts anything. For performance-critical functions that are called many times, this overhead adds up.

```
// Untyped - runtime validates each argument on every call
fn fib(n) {
    if (n < 2) return n;
    return fib(n - 1) + fib(n - 2);
}
```

```
// Typed - compiler verifies types statically; runtime skips the check entirely
fn fib(n: int): int {
  if (n < 2) return n;
  return fib(n - 1) + fib(n - 2);
}
```

When the compiler can prove at the call site that every argument's type matches the declared parameter type, it emits `CALL_TYPED` instead of `CALL`. The VM fast path then skips argument validation completely - saving roughly 10 cycles per call.

Style	Validation	Bytecode emitted
Untyped params	At runtime, every call	<code>CALL</code>
Typed params + typed arguments	Skipped - done at compile time	<code>CALL_TYPED</code>

The gain is proportional to call frequency. A function called 100 million times saves ~300 ms on M-series hardware when typed. A function called once saves nothing measurable.

Guideline: write untyped for glue code, one-off scripts, and prototypes. Add type annotations to any function that sits on a hot call path.

Inline small functions to remove call overhead

A function marked `inline` (or a trivial expression-bodied function the compiler auto-inlines) has its body expanded into each call site, removing the call frame entirely. Arguments are evaluated once, in order, and bound to fresh temporaries; name collisions with caller locals, over-deep recursive inlining (past 8 levels), and arity mismatches fall back to a normal call, so inlining never changes behavior.

```
fn inline sq(x) { return x * x; } // expands at each call site
fn dist2(a, b) { return sq(a) + sq(b); }
```

Inline aggressively only for small, hot helpers - a large inlined body is duplicated at every call site and can bloat the bytecode.

Tail calls run in constant stack space

When a `return` directly returns a call - `return f(x);` - the compiler emits a **tail call** that reuses the current frame instead of pushing a new one. A self-recursive tail call (`return self(...)`) is lowered even more tightly (`OP_TAIL_SELF`, reading the current frame's function directly). This lets recursion-as-iteration run in $O(1)$ stack:

```
fn sum(n, acc) {
  if (n == 0) { return acc; }
  return sum(n - 1, acc + n); // tail call - constant stack depth
}
```

Note: a self tail call always re-enters the original function body; it is not affected by a later `VM.PatchFunction` override of that name.

Strip debug symbols for production

By default, the compiler injects debug tracking instructions for line numbers and symbol names. These execute as individual dispatches inside every function, including loop bodies. With `--strip` the compiler omits them.

Mode	Extra dispatches per loop iteration
Default (debug on)	1–3 per statement
<code>--strip</code> (debug off)	0

For a 10 million-iteration loop with two body statements, debug symbols add roughly 20–30 ms on M-series hardware. Use `--strip` for any compiled production binary.

```
# Development – full debug symbols, optimizations on
flarismv myapp.flc

# Production – strip debug symbols (optimizations are on by default)
flarismv --strip myapp.flc

# Compile for distribution
flarismv --compile myapp.flc myapp.flx --strip
```

Dispatch budget summary for common loop patterns

All figures measured with `--strip` on M-series hardware. "Overhead" is the number of VM dispatches per iteration that are not the loop body.

Pattern	Overhead dispatches	Notes
<code>iter (i from 0 to N) { ... }</code>	2	fused increment+check + loop-back
<code>while (i < N) { i++; }</code>	3	fused compare + compact increment + loop-back
<code>while (i < N) { i += 1; }</code>	3	fused compare + compound assign + loop-back
<code>while (i < N) { i = i + 1; }</code>	5	same + 2 padding dispatches
<code>while (i < N) { ... }</code> without <code>--strip</code>	+1–3	per body statement, debug tracking
<code>obj.method()</code> per iteration	-	hash lookup each call
<code>let fn = obj.method; fn()</code> per iteration	-	direct call; saves the lookup

The gap between the best and worst row above is 2.5x. On a 10 million-iteration loop that is the difference between 70 ms and 175 ms for a single counter pattern alone.

Worked example: arithmetic benchmark

```
// Baseline – explicit assignments, no fusion (slowest)
fn slow_sum(n) {
  let sum = 0;
  let i = 0;
  while (i < n) {
    sum = sum + i; // padded, 3 dispatches
    i = i + 1;     // padded, 3 dispatches
  }
}
```



```

    }
    return sum;
}

// Optimized – compound assign + iter (fastest)
fn fast_sum(n) {
    let sum = 0;
    iter (i from 0 to n) {
        sum += i;          // fused, 1 dispatch
    }
    return sum;
}

```

`fast_sum` inner loop (with `--strip`): 3 dispatches per iteration for 2 meaningful operations - fused increment+check, compound assign, loop-back.

R11 - JIT Compilation

On ARM64 and x86-64, Flaris includes a second execution tier: a function-level JIT compiler that translates eligible functions to native machine code. JIT-compiled functions run much faster than the bytecode interpreter for numeric workloads.

JIT is controlled by a single flag, `--jit`, which gates the pipeline at both ends:

Step	Command	Effect
Compile	<code>flarism --compile --jit app.flx</code>	Generates JIT IR and embeds it in the <code>.flx</code> alongside standard bytecode
Run	<code>flarism --jit --exec app.flx</code>	Compiles the embedded JIT IR to native code on first function load

Without `--jit` at compile time, no JIT IR is written - the `.flx` is standard bytecode only. Without `--jit` at run time, any JIT IR in the file is loaded but not compiled; functions run on the bytecode interpreter.

The JIT backend is automatically selected at run time based on the CPU: ARM64 on Apple Silicon / Raspberry Pi / etc., and x86-64 on Linux/macOS/Windows PCs. On other platforms the bytecode interpreter is the only execution tier.

The JIT boundary - why it exists

The JIT compiler works on **unboxed, statically-typed values**. Inside a JIT-compiled function, integers are bare 64-bit integers, floats are bare 64-bit doubles, and strings, blocks, and typed arrays are raw pointers to their objects. There is no dynamic dispatch, and heap allocation is limited to a few tracked patterns.

What runs natively inside a JIT function:

Allowed in a JIT function	Constraint
Integer/float/bool/char arithmetic and comparisons	-
Typed-array and block element access (<code>a[i]</code> , <code>Buffer.*At</code>)	Array/block must arrive as a typed parameter or a tracked local

Allowed in a JIT function	Constraint
String comparison (<code>==</code> , <code>!=</code> , <code><</code> , <code><=</code> , <code>></code> , <code>>=</code>)	Both operands strings (variables, literals, or string fields of <code>this</code>)
String concatenation <code>s + t</code>	Only as a local initializer, a return value, or a builtin-call argument
<code>Buffer.Create</code> / <code>Array.Create</code> and other allocating builtins	Only as a local initializer or a return value (the frame owns and releases the result)
Homogeneous <code>[int]</code> / <code>[float]</code> array literals with computed elements (<code>[x, y, x*2]</code>)	Only as a local initializer (allocated + filled as an owned typed array). Fully-constant literals are already const-folded
<code>this.<field></code> reads (scalar unboxed; string/array/block/object/float borrowed) and stores (int/bool/char inline, <code>float</code> copy-on-write, object fields borrowed)	Declared fields of the method's own class
<code><recv>.<field></code> reads and stores, where <code>recv</code> is a class-typed parameter (<code>a: Vec2</code>) or a new -pinned local	The receiver's class is known at compile time; int/bool/char store inline, <code>float</code> fields store copy-on-write (no per-store allocation), object fields take a borrowed value
<code><recv>.method(...)</code> on a class-typed parameter or new -pinned local	The method must be a non-overridden instance method that is itself JIT-eligible
<code>let p = new C()</code> for a constructor-less class	The frame owns and releases the instance; <code>p</code> must not be reassigned
Calls to other JIT-compiled script functions	Callee must be JIT-eligible and return a scalar (object returns stay at the interpreter boundary)
Calls to stdlib builtins marked in the <code>JIT</code> column of R8	See the column legend below
<code>param == nil</code> tests	Folds to a constant - a native call always has all arguments present and typed

What still requires the interpreter:

Not allowed in a JIT function	Why
Object literals <code>{a: 1}</code> ; <code>[char]</code> / <code>[bool]</code> or heterogeneous array literals	Untracked heap allocation (homogeneous <code>[int]</code> / <code>[float]</code> literals in a declaration are lowered)
<code>new Foo(...)</code> with a constructor, or with arguments	Constructor dispatch — only argument-less <code>let p = new C()</code> on a constructor-less class is lowered
<code>obj.field</code> on an object of unknown class	Dynamic property lookup — resolved only for <code>this</code> , a class-typed parameter, or a new -pinned local
<code>try/catch</code> , <code>yield</code> , <code>await</code> , closures	Interpreter machinery
Global variable reads/writes	Globals live in the interpreter environment

A function that uses an unsupported construct is simply not JIT-compiled - it runs on the bytecode interpreter with identical semantics. `flarism --check file.fl` prints the per-function verdicts and the exact reason a function stays interpreted.

The `JIT` column in the R8 stdlib tables

Cell	Meaning
✓	Callable directly from JIT-compiled code (scalar result, arguments passed borrowed)
✓ owned ¹	Callable; the fresh object result must be consumed where the frame can track it: a local declaration's initializer, a return value, or a string-concatenation operand
✓ -- unsafe	Callable, but the builtin itself requires unsafe mode (<code>--unsafe</code>) at runtime
—	Not callable natively - a call makes the surrounding function fall back to the interpreter

Callable builtins require every argument to be a JIT value (typed scalar, string, block, or typed array); optional trailing arguments may be omitted exactly as in interpreted code. The core builtins `len(x)` , `int(x)` , and

`float(x)` are also JIT-lowered.

The practical rule remains: **assemble complex data structures in normal interpreter code; pass typed arrays, blocks, and scalars into JIT functions that do the heavy computation.**

Structuring code for JIT

The recommended split is:

```
// Interpreter layer: set up data, call the compute function, use results
fn process_signal(raw: array) : array {
  let n: int = len(raw);
  let buf: [float] = Array.Create(n, Type.Float);
  // copy raw values into typed float array
  for (let i: int = 0; i < n; i++) { buf[i] = float(raw[i]); }

  // hand off to JIT - all heavy work happens here
  normalize(buf, n);

  return buf;
}

// JIT layer: typed arrays in, typed scalars/arrays out - no allocation
fn normalize(a: [float], n: int) : int {
  let max: float = 0.0;
  let i: int = 0;
  while (i < n) {
    if (a[i] > max) { max = a[i]; }
    i++;
  }
  if (max == 0.0) { return 0; }
  i = 0;
  while (i < n) { a[i] = a[i] / max; i++; }
  return 1;
}
```

Key points:

- The interpreter function creates the `[float]` array and fills it - that allocation lives outside the compute-heavy path
- The compute function receives the array as a parameter and works on it in-place - no allocation, no boxing
- Multiple JIT functions can call each other directly without re-entering the interpreter (see [JIT fast-call for callbacks](#))

Small objects can also flow through JIT code when their class is statically known - a class-typed parameter, or a local constructed in place. Field reads and writes and method calls on such a receiver are native; `float` fields are stored copy-on-write, so a tight update loop allocates nothing:

```

class Vec2 {
  let x: float = 0.0;
  let y: float = 0.0;
  fn len2() : float { return this.x * this.x + this.y * this.y; }
}

// Class-typed parameter: `v`'s class is verified once at entry, then
// `v.x`, `v.y` and `v.len2()` are resolved to direct slot access / calls.
fn dist2(v: Vec2) : float { return v.len2(); }

// `new`-pinned local: construct in place, fill fields, loop - all native.
fn sweep(n: int) : float {
  let p = new Vec2();           // constructor-less class, frame-owned
  p.x = 3.0;
  p.y = 4.0;
  let acc: float = 0.0;
  for (let i: int = 0; i < n; i++) { acc = acc + p.len2(); }
  return acc;
}

```

Eligibility (automatic)

There is no keyword to mark a function for JIT. When `--jit` is active, **every eligible function is compiled automatically** - just write ordinary functions with typed parameters. A function is **JIT-eligible** when all of the following hold:

- All parameters have explicit type annotations
- The return type is explicitly annotated, **or** the analyzer can infer it: when every `return` expression provably has the same single concrete scalar type, the function is typed automatically (chains of unannotated functions calling each other resolve too). Mixed returns (`return 1; / return 2.5;`) keep the function dynamic.
- Functions with no value returns (procedures) are also eligible - they produce `nil` exactly like the interpreter
- Parameter and return types are: `int`, `float`, `bool`, `char`, `string`, or `block`; typed arrays `[int]`, `[float]`, `[char]`, `[bool]`; or a **user-defined class** (`a: Vec2`) - the receiver's class is verified at the native entry, so field and method access on it is resolved statically. Optional parameters (`k?: int`) are fine - a native call always has every argument present and typed, and other call shapes take the interpreter path
- The function body contains **none** of: `try/catch`, `yield`, `await`, member-access on an object of unknown class (`obj.field` - resolved for `this`, a class-typed parameter, or a `new`-pinned local), object literals (`{k: v}`), global variable access, or calls to non-eligible script functions. `new` is allowed only as `let p = new C()` for a constructor-less class, and array literals only as a `let a = [...]` initializer of a homogeneous `[int]` / `[float]` array
- String operations are native: comparisons (`==`, `!=`, ordering) anywhere, concatenation (`+`) as a local initializer, return value, or builtin-call argument, and string-literal returns
- `foreach` over a typed-array or `string` parameter/local is allowed; `foreach (v, i in arr)` with an index variable is also supported

Allowed exceptions:

- Calls to every stdlib builtin marked ✓ in the R8 JIT column run at native speed; ✓ owned¹ builtins (`Buffer.Create`, `Array.Create`, `String.Replace`, ...) additionally require their result to be consumed as a local initializer, a return value, or a concat operand
- Class methods are compiled too: `this.<field>` access (scalar/borrowed reads and `int/bool/char/float/object` stores), `this.method()`, static and `super.method()` calls all stay native when the callee is eligible and returns a scalar (`int`, `float`, `bool`, or `char`)
- **Objects flow through JIT code** when their class is statically known - a class-typed parameter (`a: Vec2`), or a local pinned by `let p = new C()` (constructor-less class, not reassigned). On such a receiver, `p.field` reads and stores (`int/bool/char` inline, `float` copy-on-write, object fields borrowed) and `p.method(...)` calls are all native. This lets a self-contained numeric kernel construct a small object, fill its fields, and loop over them without leaving native code
- Calls to other eligible functions become direct JIT-to-JIT calls when the callee returns a scalar; object-returning script callees keep the caller on the interpreter

Fibers and long JIT loops: a JIT-compiled function runs as one native call and cannot be preempted mid-loop. Every loop back-edge decrements a scheduling counter; on expiry (every 10,000 iterations) the JIT pumps pending IO (timers, sockets, completions) and flags the fiber so it yields to other fibers as soon as the JIT call returns. Async work therefore keeps flowing during long numeric kernels, but a single very long JIT call still occupies the CPU until it finishes.

```
fn sum_to(n: int) : int {
  let s: int = 0;
  for (let i: int = 1; i <= n; i++) { s = s + i; }
  return s;
}
```

A function that fails the eligibility check simply runs on the bytecode interpreter - no error, no warning. Use `--verbose` to see which functions qualified:

```
flarism --verbose --jit myfile.fl
# [JIT] signature-eligible: sum_to (line 20)
# [JIT] signature-eligible: normalize (line 35)
```

If a function you expected to be JIT-compiled does not appear, look for one of the disqualifiers above - untyped parameters and member access on arbitrary objects (`obj.field`) are the most common surprises; `--check` prints the exact reason per function.

Checking eligibility with `--check`

`flarism --check <file.fl>` compiles the file without running it and prints a per-function eligibility report - the quickest way to see what runs native and, for anything that doesn't, exactly why:

```
flarism --check mymath.fl
```

JIT eligibility (✓ native with `--jit` · interpreted):

```

✓ sum_to                (line 12)
✓ dot                   (line 20)
· scale_all             (line 28) – expression has no JIT lowering
· label                 (line 34) – a parameter is untyped, or not a JIT type
(int/float/bool/char/string/block, or a typed array)
· Point.dist            (line 51) – member access (only this.<field> is JIT-lowered)

2 of 5 functions are JIT-eligible.
```

The reason on each · line names the first disqualifier and its line, so you can fix or restructure without guessing. It also works on library files, e.g. `flarism --check libs/Signals.flx --libs='./libs'`, to see which functions of a module are numeric kernels.

JIT IR in compiled .flx files

`--jit` controls the JIT at both ends. When you compile a .flx with `--jit`, the JIT code for every eligible function is embedded alongside the regular bytecode; without `--jit` at compile time, none is written:

```

flarism --compile --jit myfile.flx myfile.flx # compile – embeds JIT code
flarism --jit --exec myfile.flx              # execute – runs functions natively
flarism    --exec myfile.flx                  # execute – runs bytecode only (JIT code present but
inactive)
```

The embedded JIT code is architecture-independent - the same .flx runs natively on any platform with a JIT backend, and stays fully portable: it is only activated when `--jit` is passed at run time; otherwise functions run on the bytecode interpreter. (Self-contained binaries built with `--embed` always run with the JIT active.)

Supported operations

Category	What is JIT-compiled
Int arithmetic	<code>+</code> <code>-</code> <code>*</code> <code>/</code> <code>%</code> on <code>int</code> locals
Float arithmetic	<code>+</code> <code>-</code> <code>*</code> <code>/</code> on <code>float</code> locals; unary negation
Bitwise / shifts	<code>&</code> <code> </code> <code>^</code> <code><<</code> <code>>></code> on <code>int</code>
Comparisons	<code><</code> <code><=</code> <code>></code> <code>>=</code> <code>==</code> <code>!=</code> (int and float)
Logical	<code>&&</code> and <code> </code> with short-circuit evaluation
Bool operations	<code>bool</code> params/return; <code>true</code> / <code>false</code> literals; <code>bool</code> in <code>if</code>
Control flow	<code>if/else</code> , <code>while</code> , <code>for</code> , <code>iter</code> , <code>foreach</code> , <code>return</code>
Switch	<code>switch</code> on <code>int</code> , <code>bool</code> , or <code>char</code> expression with literal case values; <code>break</code> and <code>fallthrough</code> supported
Counter ops	<code>i++</code> <code>i--</code> <code>i += n</code> <code>i -= n</code> (also <code>*=</code> , <code>/=</code> , <code>%=</code> , <code>&=</code> , <code> =</code> , <code>^=</code> , <code><<=</code> , <code>>>=</code>)
Type conversions	<code>float(n)</code> widens <code>int</code> ; <code>int(x)</code> truncates <code>float</code>
Int array reads	<code>a[i]</code> and <code>a[constant]</code> where <code>a</code> : <code>[int]</code>
Int array writes	<code>a[i] = expr</code> and <code>a[constant] = expr</code> where <code>a</code> : <code>[int]</code>
Float array reads	<code>a[i]</code> and <code>a[constant]</code> where <code>a</code> : <code>[float]</code> (returns unboxed double)
Float array writes	<code>a[i] = expr</code> and <code>a[constant] = expr</code> where <code>a</code> : <code>[float]</code>
Char array reads	<code>a[i]</code> and <code>a[constant]</code> where <code>a</code> : <code>[char]</code> (returns unboxed codepoint)
Char array writes	<code>a[i] = expr</code> and <code>a[constant] = expr</code> where <code>a</code> : <code>[char]</code>

Category	What is JIT-compiled
Bool array reads	<code>a[i]</code> and <code>a[constant]</code> where <code>a: [bool]</code> (returns unboxed 0/1)
Bool array writes	<code>a[i] = expr</code> and <code>a[constant] = expr</code> where <code>a: [bool]</code>
Array returns	Functions may return a typed array (<code>[int]</code> , <code>[float]</code> , etc.); the array object is passed through unboxed
Logical NOT	<code>!expr</code> where <code>expr</code> is a bool or bool-typed local
String char reads	<code>s[i]</code> and <code>s[constant]</code> where <code>s: string</code> (returns codepoint)
String char writes	<code>s[i] = expr</code> and <code>s[constant] = expr</code> where <code>s: string</code>
Array length	<code>len(a)</code> where <code>a</code> is any array or string
Foreach - int/char/bool array	<code>foreach (v in a)</code> and <code>foreach (v, i in a)</code> where <code>a: [int] / [char] / [bool]</code> ; element is unboxed
Foreach - float array	<code>foreach (v in a)</code> where <code>a: [float]</code> ; element is unboxed double
Foreach - string	<code>foreach (c in s)</code> where <code>s: string</code> ; yields unboxed codepoints
Math builtins	<code>Math.Sin</code> , <code>Math.Sqrt</code> , etc. - see table below
Char builtins	<code>Char.IsAlpha</code> , <code>Char.ToUpper</code> , etc. - see table below
Math int builtins	<code>Math.BitCount</code> , <code>Math.LeadingZeros</code> , <code>Math.TrailingZeros</code> , <code>Math.BitLength</code>
String comparison	<code>==</code> <code>!=</code> <code><</code> <code><=</code> <code>></code> <code>>=</code> on strings (variables, literals, string fields of <code>this</code>) - exact interpreter semantics
String concatenation	<code>a + b + ...</code> on strings, as a local initializer, return value, or builtin-call argument
String literal returns	<code>return "text"</code> in string-returning functions
Block (buffer) access	<code>b[i]</code> reads/writes on <code>block</code> params/locals; <code>Buffer.ReadU8At</code> -family via descriptor rows
Buffer/array creation	<code>let b = Buffer.Create(n, w) / Array.Create(n, Type.X)</code> as local initializer; the result may be returned
Stdlib builtin calls	Every builtin marked ✓ in the R8 JIT column (pure or JIT-safe, scalar or owned-object result)
Class fields	<code>this.<field></code> reads/writes for scalar fields of the method's own class (declared <code>let x: int = ...</code>)
Method calls	<code>this.method(...)</code> , <code>ClassName.Static(...)</code> , <code>super.method(...)</code> to eligible scalar-returning callees
Optional parameters	<code>k?: int</code> params; <code>k == nil</code> folds to a constant on the native path
Void procedures	Functions with no value returns compile and produce <code>nil</code> at the boundary

Out-of-bounds array or string accesses in JIT-compiled functions return `nil` safely.

switch in JIT functions - requirements and limitations:

- The switch expression must have a static type of `int` , `bool` , or `char`
- All `case` values must be integer, bool, or char literals - computed case expressions disqualify the function
- Multiple values per case use stacked `case` labels: `case 1: case 2: body` (not `case 1, 2:`)
- `break` inside a case body works correctly
- Fallthrough (no `break` between cases) is supported
- Nested `break` (e.g., `break` inside an `if` inside a `case`) is **not** supported in the JIT - it will be emitted but will not break out of the enclosing switch

Fiber scheduling: a JIT-compiled function runs as one native call and cannot be preempted mid-call. Every loop back-edge decrements a scheduling counter; on expiry (every 10,000 iterations) the JIT pumps pending IO and flags the fiber so it yields as soon as the native call returns. Async work keeps flowing during long numeric kernels, but a single very long JIT call still occupies the CPU until it finishes.

Operator precedence note: In Flaris, bitwise `&` has higher precedence than comparison operators (`>=`, `<=`, etc.). Use `&&` (logical AND) to combine comparison results: `x >= lo && x <= hi`.

Math builtins in JIT functions

JIT functions may call the following `Math.*` functions directly, at native speed (no interpreter overhead). All take and return `float`:

Call	Meaning
<code>Math.Sqrt(x)</code>	Square root
<code>Math.Cbrt(x)</code>	Cube root
<code>Math.Sin(x)</code> <code>Math.Cos(x)</code> <code>Math.Tan(x)</code>	Trigonometric (radians)
<code>Math.Asin(x)</code> <code>Math.Acos(x)</code> <code>Math.Atan(x)</code>	Inverse trigonometric
<code>Math.Atan2(y, x)</code>	Angle of the point <code>(x, y)</code>
<code>Math.Sinh(x)</code> <code>Math.Cosh(x)</code> <code>Math.Tanh(x)</code>	Hyperbolic
<code>Math.Asinh(x)</code> <code>Math.Acosh(x)</code> <code>Math.Atanh(x)</code>	Inverse hyperbolic
<code>Math.Exp(x)</code>	e to the power x
<code>Math.Expm1(x)</code>	$\text{Exp}(x) - 1$, accurate for small x
<code>Math.Pow(x, y)</code>	x to the power y
<code>Math.Log(x)</code>	Natural logarithm
<code>Math.Log10(x)</code> <code>Math.Log2(x)</code>	Base-10 / base-2 logarithm
<code>Math.Log1p(x)</code>	$\text{Log}(1 + x)$, accurate for small x
<code>Math.LogBase(x, base)</code>	Logarithm of x in the given base
<code>Math.AbsF(x)</code>	Absolute value
<code>Math.Floor(x)</code> <code>Math.Ceil(x)</code> <code>Math.Round(x)</code> <code>Math.Trunc(x)</code>	Rounding
<code>Math.Hypot(x, y)</code>	$\text{Sqrt}(x^2 + y^2)$ without overflow
<code>Math.MinF(x, y)</code> <code>Math.MaxF(x, y)</code>	Smaller / larger of two floats
<code>Math.IEEERemainder(x, y)</code>	IEEE floating-point remainder
<code>Math.CopySign(x, y)</code>	Magnitude of x with the sign of y
<code>Math.Deg(x)</code> <code>Math.Rad(x)</code>	Radians→degrees / degrees→radians
<code>Math.Fract(x)</code>	Fractional part, $x - \text{Floor}(x)$
<code>Math.Clamp01(x)</code>	Clamp to the range <code>[0, 1]</code>
<code>Math.Sech(x)</code> <code>Math.Csch(x)</code>	$1/\text{Cosh}(x)$ / $1/\text{Sinh}(x)$
<code>Math.Cot(x)</code> <code>Math.Csc(x)</code> <code>Math.Sec(x)</code>	$1/\text{Tan}(x)$ / $1/\text{Sin}(x)$ / $1/\text{Cos}(x)$
<code>Math.Acot(x)</code> <code>Math.Acoth(x)</code>	Inverse cotangent / inverse hyperbolic cotangent
<code>Math.Gamma(x)</code> <code>Math.LnGamma(x)</code>	Gamma $\Gamma(x)$ / log-gamma
<code>Math.Erf(x)</code> <code>Math.Erfc(x)</code>	Error function / complementary error function

These always return a `float` and follow IEEE semantics (see the *Math* module below): domain errors return NaN, poles return `±Inf`. The fused form `a * b + c` is also recognised and optimised.

```
fn gaussian(x: float, mu: float, sigma: float): float {
    let d: float = (x - mu) / sigma;
    return Math.Exp(-0.5 * d * d) / (sigma * Math.Sqrt(6.283185307));
}
```

Math integer intrinsics in JIT functions

The following `Math.*` functions take and return `int` and run at native speed:

Call	Description
<code>Math.BitCount(n)</code>	Number of set bits (popcount)
<code>Math.LeadingZeros(n)</code>	Count of leading zero bits (64-bit)
<code>Math.TrailingZeros(n)</code>	Count of trailing zero bits (64-bit)
<code>Math.BitLength(n)</code>	Bit length: $\text{floor}(\log_2(\text{abs}(n))) + 1$, or 0 for 0
<code>Math.Abs(n)</code>	Integer absolute value (int arg only)
<code>Math.Sign(n)</code>	-1 / 0 / 1 (int arg only)
<code>Math.Min(a, b)</code>	Integer minimum
<code>Math.Max(a, b)</code>	Integer maximum

`Math.Abs` / `Sign` / `Min` / `Max` are `int|float -> int` in general; the native path applies only when **all** arguments are `int`. A call with a float argument runs on the interpreter. `Math.Lerp(a, b, t)` (float) also runs natively, computed as `a + (b - a) * t` with the same rounding as the interpreter.

Char builtins in JIT functions

JIT functions may call the following `Char.*` functions directly, at native speed:

Call	Description	Returns
<code>Char.IsAlpha(c)</code>	True if <code>c</code> is an alphabetic character	<code>bool</code>
<code>Char.IsDigit(c)</code>	True if <code>c</code> is a decimal digit	<code>bool</code>
<code>Char.IsAlNum(c)</code>	True if <code>c</code> is alphanumeric	<code>bool</code>
<code>Char.IsUpper(c)</code>	True if <code>c</code> is uppercase	<code>bool</code>
<code>Char.IsLower(c)</code>	True if <code>c</code> is lowercase	<code>bool</code>
<code>Char.IsSpace(c)</code>	True if <code>c</code> is whitespace	<code>bool</code>
<code>Char.IsPunct(c)</code>	True if <code>c</code> is punctuation (ASCII)	<code>bool</code>
<code>Char.IsXDigit(c)</code>	True if <code>c</code> is a hex digit (ASCII)	<code>bool</code>
<code>Char.IsPrint(c)</code>	True if <code>c</code> is printable	<code>bool</code>
<code>Char.IsControl(c)</code>	True if <code>c</code> is a control character	<code>bool</code>
<code>Char.IsAscii(c)</code>	True if <code>c</code> < 128	<code>bool</code>
<code>Char.IsCased(c)</code>	True if <code>c</code> has case (upper/lower)	<code>bool</code>
<code>Char.IsNewLine(c)</code>	True if <code>c</code> is a line-ending character	<code>bool</code>
<code>Char.ToUpper(c)</code>	Uppercase mapping of <code>c</code>	<code>char</code>
<code>Char.ToLower(c)</code>	Lowercase mapping of <code>c</code>	<code>char</code>
<code>Char.SwapCase(c)</code>	Toggle case of <code>c</code>	<code>char</code>
<code>Char.Utf8Len(c)</code>	Number of UTF-8 bytes needed to encode <code>c</code>	<code>int</code>
<code>Char.DigitValue(c)</code>	Numeric digit value of <code>c</code> (0–9, -1 if not a digit)	<code>int</code>
<code>Char.Code(c)</code>	Unicode codepoint of <code>c</code>	<code>int</code>
<code>Char.FromInt(n)</code>	Char for codepoint <code>n</code> (<code>n</code> masked to 32 bits)	<code>char</code>

Parameters and return values are unboxed `char` or `int` inside the JIT. The function signature must use `char` or `int` as appropriate.

Block indexing and memory access in JIT functions

JIT functions may index a `block` parameter and call a small set of memory/buffer functions directly, with the same bounds checks and results as the interpreter:

Construct	Meaning
<code>b[i]</code> (read)	Signed element read of block <code>b</code> (honours the block's element size)
<code>b[i] = v</code> (write)	Truncated element write

Construct	Meaning
<code>Memory.Read8/16/32/64(addr[, off])</code>	Unsigned integer at a raw integer address, native byte order (requires <code>--unsafe</code>)
<code>Memory.Write8/16/32/64(addr, val[, off])</code>	Store an integer at a raw integer address (requires <code>--unsafe</code>)
<code>Memory.Swap16/32/64(v)</code>	Byte-order (endianness) swap of a plain integer
<code>Memory.Compare(a, b, len)</code>	Compare <code>len</code> bytes at two raw integer addresses, like <code>memcmp</code> (requires <code>--unsafe</code>)
<code>Memory.ReadFloat32/64(addr[, off])</code>	Read an IEEE float/double at a raw integer address (requires <code>--unsafe</code>)
<code>Memory.WriteFloat32/64(addr, val[, off])</code>	Store an IEEE float/double at a raw integer address (requires <code>--unsafe</code>)
<code>Buffer.ReadU8/16/32/64At(buf, off)</code>	Little-endian unsigned read from a block; <code>0</code> if out of range
<code>Buffer.WriteU8/16/32/64At(buf, off, v)</code>	Little-endian bounded store into a block; returns <code>bool</code>

Two constraints apply: `Memory.*` runs natively only when the address argument is a plain `int` (passing a `block/string/pointer` uses the interpreter), and the `Buffer.Read/Write` calls run natively in their little-endian form (the default) — passing the optional big-endian flag falls back to the interpreter.

```
fn count_upper(s: string, n: int): int {
    let count: int = 0;
    let i: int = 0;
    while (i < n) {
        if (Char.IsUpper(s[i])) { count = count + 1; }
        i++;
    }
    return count;
}
```

JIT fast-call for callbacks

When a JIT-compiled function is passed as a callback to one of the builtins below, the builtin runs it natively per element instead of through the bytecode interpreter. This eliminates the per-element interpreter overhead for high-throughput loops.

Array methods - callback receives the element (and index where noted):

Method	Callback signature
<code>Array.ForEach(arr, fn)</code>	<code>fn(elem: T): any</code>
<code>Array.Select(arr, fn)</code>	<code>fn(elem: T): U</code>
<code>Array.Where(arr, fn)</code>	<code>fn(elem: T): bool</code>
<code>Array.Reduce(arr, fn, init)</code>	<code>fn(acc: T, elem: T): T</code>
<code>Array.ReduceRight(arr, fn, init)</code>	<code>fn(acc: T, elem: T): T</code>
<code>Array.Any(arr, fn)</code>	<code>fn(elem: T): bool</code>
<code>Array.All(arr, fn)</code>	<code>fn(elem: T): bool</code>
<code>Array.Find(arr, fn) / FindIndex / FindLast / FindLastIndex</code>	<code>fn(elem: T): bool</code>
<code>Array.Count(arr, fn)</code>	<code>fn(elem: T): bool</code>
<code>Array.RemoveIf(arr, fn)</code>	<code>fn(elem: T): bool</code>
<code>Array.Process(arr, fn)</code>	<code>fn(elem: T): T</code> - mutates in place
<code>Array.ProcessIdx(arr, fn)</code>	<code>fn(elem: T, idx: int): T</code> - mutates in place

Method	Callback signature
<code>Array.GroupBy(arr, fn)</code>	<code>fn(elem: T): key</code>
<code>Array.SortBy(arr, fn)</code>	<code>fn(elem: T): key</code>
<code>Array.FlatMap(arr, fn)</code>	<code>fn(elem: T): array T</code>
<code>Array.CountBy(arr, fn)</code>	<code>fn(elem: T): key</code>
<code>Array.Partition(arr, fn)</code>	<code>fn(elem: T): bool</code>
<code>Array.MinBy(arr, fn)</code>	<code>fn(elem: T): key</code>
<code>Array.MaxBy(arr, fn)</code>	<code>fn(elem: T): key</code>
<code>Array.UniqueBy(arr, fn)</code>	<code>fn(elem: T): key</code>
<code>Array.Sort(arr, fn)</code>	<code>fn(a: T, b: T): bool</code> - comparator; return true if <code>a < b</code>

Other builtins:

Method	Callback signature
<code>Memory.Process(addr, count, size, fn)</code>	<code>fn(val: int, idx: int): int</code> - rewrites each element
<code>Gfx.Filter(x, y, w, h, fn)</code>	<code>fn(pixel: int): int</code> - rewrites each ARGB pixel

The fast path is selected automatically at runtime when `--jit` is active and the passed function is JIT-eligible. Otherwise (e.g. `--jit` is absent, or the function isn't eligible) the builtin falls back to the normal interpreter path transparently.

```
fn double_pixel(px: int): int {
    let r: int = (px >> 16) & 0xFF;
    let g: int = (px >> 8) & 0xFF;
    let b: int = px & 0xFF;
    let a: int = (px >> 24) & 0xFF;
    // Brighten RGB channels, clamp to 255
    if (r * 2 > 255) { r = 255; } else { r = r * 2; }
    if (g * 2 > 255) { g = 255; } else { g = g * 2; }
    if (b * 2 > 255) { b = 255; } else { b = b * 2; }
    return (a << 24) | (r << 16) | (g << 8) | b;
}

let img = Gfx.Create(800, 600);
// ...draw something...
img.Filter(0, 0, 800, 600, double_pixel); // native ARM64 per-pixel loop
```

Calling convention and memory

JIT-compiled functions share the same call sites as interpreter functions - callers use the normal path. The compiled function pointer is stored and invoked directly once available.

Machine code is allocated with `mmap/mprotect` (no entitlement required on macOS).

Full workflow example

```
# 1. Compile with JIT IR embedded
flarism --compile --jit myapp.flx myapp.flx
```

```
# 2. Verify which functions are eligible (structural check, --jit not required here)
flarism --verbose myapp.flx

# 3. Run with JIT active
flarism --jit --exec myapp.flx

# 4. Run without JIT (fallback to bytecode interpreter, even if JIT IR present)
flarism --exec myapp.flx

# 5. Source file directly: compile+run in one step with JIT
flarism --jit myapp.flx
```

Benchmark reference

Workload	JIT	Interpreter	Speedup
Gaussian sum to 10 M	~9.5 ms	~63 ms	6.6×
sum_range(10 000) × 10 000 reps	~54 ms	~606 ms	11×
count_down(10 000) × 10 000 reps	~53 ms	~954 ms	18×

Results are wall-clock time on a single fiber. Workloads with tight integer loops show the largest gains; workloads dominated by string or object operations are unaffected.

R12 - Static Analyzer

Between the parser and the optimizer, every compile runs a ten-pass static analyzer over the AST. It resolves names and types, reports problems, and decorates the tree with the facts the optimizer, code generator and JIT depend on (`resolved_fn_decl` drives inlining, purity stamps drive LICM, JIT eligibility drives native compilation).

The analyzer runs on every compile - `flarism file.flx`, `--compile`, `VM.Eval`, and module loading. It is not a separate lint step and cannot be skipped.

Diagnostic format

```
✖ Error: file.flx:12:5: 1000: Operator '-' does not accept string as its left operand.
⚠ Warning: file.flx:19:9: 2000: Local 'tmp' is never read (prefix with '_' to suppress).
```

`file:line:column: code: message` - all diagnostics go to **stderr**. The summary lines (`Errors detected: N`) go to stdout.

Diagnostics are collected during analysis and emitted once, sorted by source position. Output is therefore **reproducible**: the same input produces byte-identical output on every run, which makes it safe to diff build logs or assert on compiler output in tests.

A compile stops after **100 errors** and stops recording after **200 warnings**; the remainder are counted and summarised. Compilation fails if any error was reported. Warnings never fail a compile.

Text interpolated into a diagnostic (identifiers, string-literal contents) has its control bytes escaped as `\xNN`, so source containing terminal escape sequences cannot repaint the terminal or forge compiler output.

Error codes (1000+)

Errors cannot be suppressed; the **Name** column is what appears in JSON diagnostics.

Code	Name	Symbol	Meaning
1000	type-error	TYPE_ERROR	Operand, assignment or argument type mismatch
1001	arg-count	ARG_COUNT	Wrong number of arguments to a function, method or constructor
1002	invalid-return	INVALID_RETURN	A path can finish without returning a declared value, or <code>return</code> ; where a value is required
1003	undefined-symbol	UNDEFINED_SYMBOL	Name does not resolve - undefined variable, function, export or base class
1004	redeclaration	REDECLARATION	Name already declared in this scope
1005	const-assign	CONST_ASSIGN	Write to a <code>const</code> binding, or to one of its properties or elements
1006	invalid-control-flow	INVALID_CONTROL_FLOW	<code>break</code> / <code>continue</code> outside a loop or switch, <code>return</code> outside a function, or control leaving a <code>finally</code>
1007	invalid-throw	INVALID_THROW	Thrown value is not an <code>Exception</code> (or subclass) instance
1008	—	INVALID_CAST	<i>Reserved</i> - a cast is an assertion and is not checked
1009	—	IMPORT_RESOLUTION	<i>Reserved</i> - see Reserved codes
1010	class-this-usage	CLASS_THIS_USAGE	<code>this</code> used outside an instance method
1011	async-await-misuse	ASYNC_AWAIT_MISUSE	<code>await</code> outside an <code>async</code> function, or applied to a non-fiber
1012	missing-entry	MISSING_ENTRY	The program has neither <code>fn Main</code> nor any <code>export</code>
1013	invalid-assign	INVALID_ASSIGN	Assignment used where a value is required (e.g. <code>if (x = 1)</code>)
1014	undeclared-field	UNDECLARED_FIELD	Member is not declared on a sealed instance - unknown field or method
1015	circular-inheritance	CIRCULAR_INHERITANCE	A class extends itself through its base chain

Warning codes (2000+)

The **Name** column is the exact spelling accepted by `-wno-<name>` and `// flaris-ignore[<name>]`, and the one emitted in JSON diagnostics. It is not always a mechanical kebab-casing of the symbol - `2002`, `2003` and `2010` differ - so copy it from this column rather than deriving it.

Code	Name	Symbol	Meaning
2000	unused-symbol	UNUSED_SYMBOL	Declared but never read - local, parameter, global or import
2001	possible-nil-flow	POSSIBLE_NIL_FLOW	A <code>guard</code> expression may be nil
2002	float-eq	FLOAT_EQ_SUGGEST_APPROX	Exact <code>==</code> / <code>!=</code> on two floats; use the approximate operator <code>≈</code>
2003	oob-index	OOB_INDEX_POSSIBLE	Constant index outside the bounds of an array literal
2004	infinite-loop	INFINITE_LOOP	Constant-true loop with no <code>break</code> , <code>return</code> or <code>throw</code>
2005	—	RECURSION_NO_BASE	<i>Reserved</i> - no recursion analysis is implemented
2006	unreachable-code	UNREACHABLE_CODE	Statement after a <code>return</code> , <code>throw</code> , <code>break</code> or <code>continue</code>
2007	—	STYLE	<i>Reserved</i> - no style rules are implemented
2008	shadowed-variable	SHADOWED_VARIABLE	Declaration hides one in an enclosing scope
2009	empty-block	EMPTY_BLOCK	Empty function, <code>if</code> , <code>else</code> , loop, <code>try</code> , <code>catch</code> or <code>finally</code> body
2010	type-change	TYPE_ERROR	Assignment changes a variable's declared type

Code	Name	Symbol	Meaning
2011	no-effect	NO_EFFECT	Expression statement with no effect (e.g. a bare literal)
2012	class-reserved-name	CLASS_RESERVED_NAME	Method shadows a built-in instance property and can never be called
2013	duplicate-key	DUPLICATE_KEY	Repeated key in an object literal - the first value is kept
2014	duplicate-case	DUPLICATE_CASE	Repeated case value - only the first match runs
2015	override-mismatch	OVERRIDE_MISMATCH	An override takes a different number of parameters than the method it replaces
2016	class-shadows-module	CLASS_SHADOWS_MODULE	A non-static method is named after a built-in namespace <i>and</i> the class calls that namespace - the call resolves to the method and yields <code>nil</code>

Every warning above (the *Reserved* codes excepted) is suppressible; **no error is**. A raw number works anywhere a name does, e.g. `-Wno-2000`.

Reserved codes

Codes marked *Reserved* are never emitted. Their numbers stay retired rather than being reused, so tooling that matches on a code can rely on it never changing meaning.

`1009 IMPORT_RESOLUTION` is reserved rather than implemented because telling "this module does not exist" apart from "this module exports nothing" needs a resolver with no side effects, and the runtime's module resolver performs network fetches and version-pin enforcement - neither is acceptable during a compile. An unresolvable import is reported by the runtime loader instead.

Suppressing an unused-symbol warning

Prefix the name with `_`. This works for **locals, parameters and globals**:

```
fn handler(_event, payload) { // _event is intentionally unused
  let _debugOnly = payload; // kept for inspection, never read
  return payload;
}
```

Loop variables and `catch` variables are exempt automatically - iterating or catching without using the value is normal.

Suppressing a specific warning

`// flaris-ignore[<name>]` silences the listed warnings. A pragma trailing a line of code applies to that line; one on a line of its own applies to the next line, so it can sit above what it refers to:

```
global buildStamp = 1; // flaris-ignore[unused-symbol]

// flaris-ignore[shadowed-variable, unused-symbol]
let value = compute();
```

With no bracket list (`// flaris-ignore`) it silences every warning on the target line. Names are the **Name** column of the warning table above; a raw number (`2000`) also works.

Warnings only. There is deliberately no way to suppress an error: an error means the program is invalid and would not survive code generation.

Warning flags

Flag	Effect
<code>-Wno-<name></code>	Silence one warning across the whole compile, e.g. <code>-Wno-unused-symbol</code> . <code><name></code> is the Name column of the warning table above (or its raw number). An unknown or non-suppressible name is a hard error, so a typo cannot silently do nothing.
<code>-Werror</code>	Any warning that survives suppression fails the compile
<code>-w</code>	Silence all warnings

`-Wno-` rejects an unknown name, and refuses error codes, rather than silently doing nothing. All three flags also apply to `VM.Eval` and `VM.Compile`.

Machine-readable diagnostics

`--diagnostics json` writes a single JSON array to **stdout** instead of the human-readable form, for editors and CI:

```
flarism --compile app.flx app.flx --diagnostics json
```

```
[
  {"severity":"warning","code":2000,"name":"unused-
symbol","file":"app.flx","line":1,"column":1,"message":"Global 'unusedOne' is never read."},
  {"severity":"error","code":1000,"name":"type-
error","file":"app.flx","line":4,"column":29,"message":"'Operator '-' does not accept string as its
left operand."}
]
```

In this mode stdout carries nothing but the array - the usual summary and fingerprint lines are suppressed so the output can be piped straight into a parser. Diagnostics keep their sorted order.

Suggestions

When a name does not resolve, the analyzer offers the closest declared name within a short edit distance:

```
1003: 'totl' is not defined. Did you mean 'total'?
1014: Class 'Counter' has no member 'cont'. Did you mean 'count'?
```

Nothing is suggested when no candidate is close enough. Ties are broken by name, so the suggestion is reproducible.

Control flow rules

- `break` and `continue` require an enclosing loop; `break` also accepts an enclosing `switch`.
- `return` requires an enclosing function.
- Control may **not** leave a `finally` body via `return`, `break` or `continue` (matching C# CS0157). A `throw` from a `finally` is allowed, and a `break` / `continue` targeting a loop opened *inside* the `finally`

is fine.

- A function with a declared return type must return on every path. Falling off the end would yield nil, which a typed caller cannot detect.

These rules apply inside **function literals** exactly as they do in named functions - a lambda body is its own control-flow scope.

Operator operand checking

Arithmetic and bitwise operators are checked against the operand types the VM actually accepts:

Operators	Accepted operands
<code>- * / % ^^</code>	<code>int, float, bool, char</code>
<code>& ^ << >></code>	<code>int, char, bool</code> - a <code>float</code> operand is an error, not a widening
<code>~</code>	<code>int</code>
<code>+</code>	anything - see below

`+` is never a type error: it is overloaded by operand type (numeric add, char and container concatenation/merge), and when no numeric, array or object rule applies it falls back to string concatenation, so any operand pair produces a value (`nil + 5` is `"null5"`). The full per-type rules are listed under [R5 - Operator Precedence](#).

A diagnostic is only reported when an operand's type is fully known and shares nothing with the operator's domain. Dynamic values - unannotated parameters, `any`, results of dynamic calls - are never diagnosed.

Bitwise and shift operators always produce an `int`, whatever the operands infer to, because the VM refuses to run them on the float lane at all.

Class checks

Instances are sealed: the member set is fixed by the class body (see [Sealed instance fields](#)). Where the analyzer can prove which class a value belongs to, it checks member access against that class - own members and inherited ones:

```
class Counter {
  let count = 0;
  fn Constructor(start) { this.count = start; }
  fn Bump(step) { this.count = this.count + step; }
}

let c = new Counter(1, 2, 3); // 1001 - Constructor expects 0-1 arguments
c.Nope();                    // 1014 - Counter has no method 'Nope'
Console.WriteLine(c.missing); // 1014 - Counter has no member 'missing'
```

A receiver's class is known when it is `this`, a local initialised by `new C()`, or a parameter annotated with a class type (`fn f(v: Vec2)`). Any other receiver is dynamic and is not checked.

Base classes may be declared in any order - `class Dog : Animal` compiles with `Animal` declared later in the file, just as calling a function declared later does.

A class whose base is imported from a compiled module is **lenient**: its full member set is not visible, so no membership check fires for it and the VM's runtime guard applies instead.

Limitations

- **Argument minimums are a lower bound.** An unannotated parameter is indistinguishable from an optional one once parsed, so a missing argument on an unannotated parameter is not reported. Passing too many is always caught.
- **Typed arrays and unions share bits.** `[string]` and the union `array|string` have the same internal representation; an ambiguous type is read as a typed array.
- **Member checks need a provable receiver.** Chained access (`a.b.Method()`) and values from dynamic sources are not checked.
- **No nil-narrowing.** `if (x != nil)` does not narrow `x` in the branch. Expression types are memoized per AST node, so a variable has one type throughout a function; flow-sensitive narrowing would need that layer replaced.
- **No exhaustiveness checking for `switch` over an enum.**

Definite-assignment analysis is not needed: `let x;` without an initializer is a syntax error, so every variable is assigned where it is declared.